



OpenShift Container Platform 4.18

Hosted control planes

Using hosted control planes with OpenShift Container Platform

OpenShift Container Platform 4.18 Hosted control planes

Using hosted control planes with OpenShift Container Platform

Legal Notice

Copyright © 2025 Red Hat, Inc.

The text of and illustrations in this document are licensed by Red Hat under a Creative Commons Attribution–Share Alike 3.0 Unported license ("CC-BY-SA"). An explanation of CC-BY-SA is available at

<http://creativecommons.org/licenses/by-sa/3.0/>

. In accordance with CC-BY-SA, if you distribute this document or an adaptation of it, you must provide the URL for the original version.

Red Hat, as the licensor of this document, waives the right to enforce, and agrees not to assert, Section 4d of CC-BY-SA to the fullest extent permitted by applicable law.

Red Hat, Red Hat Enterprise Linux, the Shadowman logo, the Red Hat logo, JBoss, OpenShift, Fedora, the Infinity logo, and RHCE are trademarks of Red Hat, Inc., registered in the United States and other countries.

Linux® is the registered trademark of Linus Torvalds in the United States and other countries.

Java® is a registered trademark of Oracle and/or its affiliates.

XFS® is a trademark of Silicon Graphics International Corp. or its subsidiaries in the United States and/or other countries.

MySQL® is a registered trademark of MySQL AB in the United States, the European Union and other countries.

Node.js® is an official trademark of Joyent. Red Hat is not formally related to or endorsed by the official Joyent Node.js open source or commercial project.

The OpenStack® Word Mark and OpenStack logo are either registered trademarks/service marks or trademarks/service marks of the OpenStack Foundation, in the United States and other countries and are used with the OpenStack Foundation's permission. We are not affiliated with, endorsed or sponsored by the OpenStack Foundation, or the OpenStack community.

All other trademarks are the property of their respective owners.

Abstract

This document provides instructions for managing hosted control planes for OpenShift Container Platform. With hosted control planes, you create control planes as pods on a hosting cluster without the need for dedicated physical or virtual machines for each control plane.

Table of Contents

CHAPTER 1. HOSTED CONTROL PLANES RELEASE NOTES	10
1.1. HOSTED CONTROL PLANES RELEASE NOTES FOR OPENSHIFT CONTAINER PLATFORM 4.18	10
1.1.1. New features and enhancements	10
1.1.1.1. Comparing hosted control planes with OpenShift Container Platform	10
1.1.1.2. Proxy configuration for hosted control planes	10
1.1.1.3. Default container runtime for worker nodes is crun	10
1.1.2. Bug fixes	10
1.1.3. Known issues	11
1.1.4. General Availability and Technology Preview features	13
CHAPTER 2. HOSTED CONTROL PLANES OVERVIEW	14
2.1. INTRODUCTION TO HOSTED CONTROL PLANES	14
2.1.1. Architecture of hosted control planes	14
2.1.2. Benefits of hosted control planes	15
2.2. DIFFERENCES BETWEEN HOSTED CONTROL PLANES AND OPENSHIFT CONTAINER PLATFORM	16
2.2.1. Cluster creation and lifecycle	16
2.2.2. Cluster configuration	16
2.2.3. etcd encryption	16
2.2.4. Operators and control plane	16
2.2.5. Updates	17
2.2.6. Machine configuration and management	17
2.2.7. Networking	18
2.2.8. Web console	18
2.3. RELATIONSHIP BETWEEN HOSTED CONTROL PLANES, MULTICLUSTER ENGINE OPERATOR, AND RHACM	19
2.3.1. Discovering multicluster engine Operator hosted clusters in RHACM	20
2.4. VERSIONING FOR HOSTED CONTROL PLANES	20
2.4.1. Management cluster	20
2.4.2. HyperShift Operator	21
2.4.3. hosted control planes CLI	21
2.4.4. hypershift.openshift.io API	21
2.4.5. Control Plane Operator	21
2.5. GLOSSARY OF COMMON CONCEPTS AND PERSONAS FOR HOSTED CONTROL PLANES	22
2.5.1. Concepts	22
2.5.2. Personas	22
CHAPTER 3. PREPARING TO DEPLOY HOSTED CONTROL PLANES	24
3.1. REQUIREMENTS FOR HOSTED CONTROL PLANES	24
3.1.1. Support matrix for hosted control planes	24
3.1.1.1. Management cluster support	24
3.1.1.2. Hosted cluster support	25
3.1.1.3. Hosted cluster platform support	26
3.1.1.4. Multi-architecture support	26
3.1.1.5. Updates of multicluster engine Operator	28
3.1.1.6. Technology Preview features	29
3.1.2. FIPS-enabled hosted clusters	29
3.1.3. CIDR ranges for hosted control planes	29
3.2. SIZING GUIDANCE FOR HOSTED CONTROL PLANES	30
3.2.1. Pod limits	30
3.2.2. Request-based resource limit	31
3.2.3. Load-based limit	31
3.2.4. Sizing calculation example	32

3.2.5. Shared infrastructure between hosted and standalone control planes	34
3.3. OVERRIDING RESOURCE UTILIZATION MEASUREMENTS	34
3.3.1. Overriding resource utilization measurements for a hosted cluster	34
3.3.2. Disabling the metric service monitoring	35
3.4. INSTALLING THE HOSTED CONTROL PLANES COMMAND-LINE INTERFACE	36
3.4.1. Installing the hosted control planes command-line interface from the terminal	36
3.4.2. Installing the hosted control planes command-line interface by using the web console	37
3.4.3. Installing the hosted control planes command-line interface by using the content gateway	38
3.5. DISTRIBUTING HOSTED CLUSTER WORKLOADS	39
3.5.1. Labeling management cluster nodes	40
3.5.2. Priority classes	41
3.5.3. Custom taints and tolerations	41
3.5.4. Control plane isolation	42
3.5.4.1. Network policy isolation	42
3.5.4.2. Control plane pod isolation	42
3.6. ENABLING OR DISABLING THE HOSTED CONTROL PLANES FEATURE	43
3.6.1. Manually enabling the hosted control planes feature	43
3.6.1.1. Manually enabling the hypershift-addon managed cluster add-on for local-cluster	44
3.6.2. Disabling the hosted control planes feature	44
3.6.2.1. Uninstalling the HyperShift Operator	44
3.6.2.2. Disabling the hosted control planes feature	45
CHAPTER 4. DEPLOYING HOSTED CONTROL PLANES	47
4.1. DEPLOYING HOSTED CONTROL PLANES ON AWS	47
4.1.1. Preparing to deploy hosted control planes on AWS	47
4.1.1.1. Prerequisites to configure a management cluster	47
4.1.1.2. Creating the Amazon Web Services S3 bucket and S3 OIDC secret	48
4.1.1.3. Creating a routable public zone for hosted clusters	49
4.1.1.4. Creating an AWS IAM role and STS credentials	50
4.1.1.5. Enabling AWS PrivateLink for hosted control planes	54
4.1.1.6. Enabling external DNS for hosted control planes on AWS	54
4.1.1.6.1. Prerequisites	55
4.1.1.6.2. Setting up external DNS for hosted control planes	56
4.1.1.6.3. Creating the public DNS hosted zone	57
4.1.1.6.4. Creating a hosted cluster by using the external DNS on AWS	58
4.1.1.7. Creating a hosted cluster on AWS	60
4.1.1.7.1. Accessing a hosted cluster on AWS by using the kubeadm credentials	62
4.1.1.7.2. Accessing a hosted cluster on AWS by using the hcp CLI	62
4.1.1.8. Creating a hosted cluster in multiple zones on AWS	63
4.1.1.8.1. Creating a hosted cluster by providing AWS STS credentials	64
4.1.1.9. Running hosted clusters on an ARM64 architecture	65
4.1.1.9.1. Creating a hosted cluster on an ARM64 OpenShift Container Platform cluster	65
4.1.1.9.2. Creating an ARM or AMD NodePool object on AWS hosted clusters	66
4.1.1.10. Creating a private hosted cluster on AWS	67
4.1.1.10.1. Accessing a private management cluster on AWS	68
4.2. DEPLOYING HOSTED CONTROL PLANES ON BARE METAL	69
4.2.1. Preparing to deploy hosted control planes on bare metal	69
4.2.1.1. Prerequisites to configure a management cluster	69
4.2.1.2. Bare metal firewall, port, and service requirements	70
4.2.1.3. Bare metal infrastructure requirements	71
4.2.2. DNS configurations on bare metal	72
4.2.3. Creating a hosted cluster on bare metal	73
4.2.3.1. Creating a hosted cluster on bare metal by using the console	74

4.2.3.2. Creating a hosted cluster on bare metal by using a mirror registry	76
4.2.4. Verifying hosted cluster creation	77
4.3. DEPLOYING HOSTED CONTROL PLANES ON OPENSHIFT VIRTUALIZATION	78
4.3.1. Requirements to deploy hosted control planes on OpenShift Virtualization	78
4.3.1.1. Prerequisites	79
4.3.1.2. Firewall and port requirements	80
4.3.2. Live migration for compute nodes	81
4.3.3. Creating a hosted cluster with the KubeVirt platform	82
4.3.3.1. Creating a hosted cluster with the KubeVirt platform by using the CLI	82
4.3.3.2. Creating a hosted cluster with the KubeVirt platform by using external infrastructure	83
4.3.3.3. Creating a hosted cluster by using the console	84
4.3.4. Configuring the default ingress and DNS for hosted control planes on OpenShift Virtualization	85
4.3.5. Customizing ingress and DNS behavior	86
4.3.5.1. Deploying a hosted cluster that specifies the base domain	86
4.3.5.2. Setting up the load balancer	88
4.3.5.3. Setting up a wildcard DNS	89
4.3.6. Configuring MetalLB	89
4.3.7. Configuring additional networks, guaranteed CPUs, and VM scheduling for node pools	91
4.3.7.1. Adding multiple networks to a node pool	91
4.3.7.1.1. Using an additional network as default	92
4.3.7.2. Requesting guaranteed CPU resources	92
4.3.7.3. Scheduling KubeVirt VMs on a set of nodes	93
4.3.8. Scaling a node pool	93
4.3.8.1. Adding node pools	94
4.3.9. Verifying hosted cluster creation on OpenShift Virtualization	95
4.4. DEPLOYING HOSTED CONTROL PLANES ON NON-BARE-METAL AGENT MACHINES	96
4.4.1. Preparing to deploy hosted control planes on non-bare-metal agent machines	97
4.4.1.1. Prerequisites for deploying hosted control planes on non-bare-metal agent machines	98
4.4.1.2. Firewall, port, and service requirements for non-bare-metal agent machines	98
4.4.1.3. Infrastructure requirements for non-bare-metal agent machines	99
4.4.2. Configuring DNS on non-bare-metal agent machines	100
4.4.3. Creating a hosted cluster on non-bare-metal agent machines by using the CLI	101
4.4.3.1. Creating a hosted cluster on non-bare-metal agent machines by using the web console	102
4.4.3.2. Creating a hosted cluster on bare metal by using a mirror registry	103
4.4.4. Verifying hosted cluster creation on non-bare-metal agent machines	105
4.5. DEPLOYING HOSTED CONTROL PLANES ON IBM Z	106
4.5.1. Prerequisites to configure hosted control planes on IBM Z	106
4.5.2. IBM Z infrastructure requirements	107
4.5.3. DNS configuration for hosted control planes on IBM Z	107
4.5.4. Creating a hosted cluster on bare metal	108
4.5.5. Creating an InfraEnv resource for hosted control planes on IBM Z	110
4.5.6. Adding IBM Z agents to the InfraEnv resource	110
4.5.6.1. Adding IBM Z KVM as agents	110
4.5.6.2. Adding IBM Z LPAR as agents	112
4.5.6.3. Adding IBM z/VM as agents	113
4.5.7. Scaling the NodePool object for a hosted cluster on IBM Z	114
4.6. DEPLOYING HOSTED CONTROL PLANES ON IBM POWER	116
4.6.1. Prerequisites to configure hosted control planes on IBM Power	117
4.6.2. IBM Power infrastructure requirements	118
4.6.3. DNS configuration for hosted control planes on IBM Power	118
4.6.4. Creating a hosted cluster on bare metal	119
CHAPTER 5. MANAGING HOSTED CONTROL PLANES	121

5.1. MANAGING HOSTED CONTROL PLANES ON AWS	121
5.1.1. Prerequisites to manage AWS infrastructure and IAM permissions	121
5.1.1.1. Infrastructure requirements for AWS	121
5.1.1.2. Unmanaged infrastructure for the HyperShift Operator in an AWS account	121
5.1.1.3. Unmanaged infrastructure requirements for a management AWS account	122
5.1.1.4. Infrastructure requirements for a management AWS account	122
5.1.1.5. Infrastructure requirements for an AWS account in a hosted cluster	123
5.1.1.6. Kubernetes-managed infrastructure in a hosted cluster AWS account	123
5.1.2. Identity and Access Management (IAM) permissions	123
5.1.3. Creating AWS infrastructure and IAM resources separate	130
5.1.3.1. Creating the AWS infrastructure separately	130
5.1.3.2. Creating the AWS IAM resources	133
5.1.3.3. Creating a hosted cluster separately	134
5.1.4. Transitioning a hosted cluster from single-architecture to multi-architecture	135
5.1.5. Creating node pools on the multi-architecture hosted cluster	138
5.2. MANAGING HOSTED CONTROL PLANES ON BARE METAL	139
5.2.1. Accessing the hosted cluster	139
5.2.2. Scaling the NodePool object for a hosted cluster	140
5.2.2.1. Adding node pools	142
5.2.2.2. Enabling node auto-scaling for the hosted cluster	143
5.2.2.3. Disabling node auto-scaling for the hosted cluster	145
5.2.3. Handling ingress in a hosted cluster on bare metal	145
5.2.4. Enabling machine health checks on bare metal	148
5.2.5. Disabling machine health checks on bare metal	149
5.3. MANAGING HOSTED CONTROL PLANES ON OPENSHIFT VIRTUALIZATION	150
5.3.1. Accessing the hosted cluster	150
5.3.2. Enabling node auto-scaling for the hosted cluster	151
5.3.3. Configuring storage for hosted control planes on OpenShift Virtualization	152
5.3.3.1. Mapping KubeVirt CSI storage classes	153
5.3.3.2. Mapping a single KubeVirt CSI volume snapshot class	155
5.3.3.3. Mapping multiple KubeVirt CSI volume snapshot classes	156
5.3.3.4. Configuring KubeVirt VM root volume	157
5.3.3.5. Enabling KubeVirt VM image caching	157
5.3.3.6. KubeVirt CSI storage security and isolation	158
5.3.3.7. Configuring etcd storage	158
5.3.4. Attaching NVIDIA GPU devices by using the hcp CLI	159
5.3.5. Attaching NVIDIA GPU devices by using the NodePool resource	160
5.3.6. Spreading node pool VMs by using topologySpreadConstraint	162
5.4. MANAGING HOSTED CONTROL PLANES ON NON-BARE-METAL AGENT MACHINES	163
5.4.1. Accessing the hosted cluster	163
5.4.2. Scaling the NodePool object for a hosted cluster	164
5.4.2.1. Adding node pools	166
5.4.2.2. Enabling node auto-scaling for the hosted cluster	167
5.4.2.3. Disabling node auto-scaling for the hosted cluster	169
5.4.3. Handling ingress in a hosted cluster on non-bare-metal agent machines	169
5.4.4. Enabling machine health checks on non-bare-metal agent machines	173
5.4.5. Disabling machine health checks on non-bare-metal agent machines	174
5.5. MANAGING HOSTED CONTROL PLANES ON IBM POWER	174
5.5.1. Creating an InfraEnv resource for hosted control planes on IBM Power	174
5.5.2. Adding IBM Power agents to the InfraEnv resource	175
5.5.3. Scaling the NodePool object for a hosted cluster on IBM Power	176
CHAPTER 6. DEPLOYING HOSTED CONTROL PLANES IN A DISCONNECTED ENVIRONMENT	179

6.1. INTRODUCTION TO HOSTED CONTROL PLANES IN A DISCONNECTED ENVIRONMENT	179
6.2. DEPLOYING HOSTED CONTROL PLANES ON OPENSHIFT VIRTUALIZATION IN A DISCONNECTED ENVIRONMENT	179
6.2.1. Prerequisites	179
6.2.2. Configuring image mirroring for hosted control planes in a disconnected environment	180
6.2.3. Applying objects in the management cluster	182
6.2.4. Deploying multicloud engine Operator for a disconnected installation of hosted control planes	183
6.2.5. Configuring TLS certificates for a disconnected installation of hosted control planes	183
6.2.5.1. Adding the registry CA to the management cluster	183
6.2.5.2. Adding the registry CA to the worker nodes for the hosted cluster	184
6.2.6. Creating a hosted cluster on OpenShift Virtualization	185
6.2.6.1. Requirements to deploy hosted control planes on OpenShift Virtualization	185
6.2.6.2. Creating a hosted cluster with the KubeVirt platform by using the CLI	185
6.2.6.3. Configuring the default ingress and DNS for hosted control planes on OpenShift Virtualization	187
6.2.6.4. Customizing ingress and DNS behavior	188
6.2.6.4.1. Deploying a hosted cluster that specifies the base domain	188
6.2.6.4.2. Setting up the load balancer	189
6.2.6.4.3. Setting up a wildcard DNS	190
6.2.7. Finishing the deployment	191
6.2.7.1. Monitoring the control plane	191
6.2.7.2. Monitoring the data plane	192
6.3. DEPLOYING HOSTED CONTROL PLANES ON BARE METAL IN A DISCONNECTED ENVIRONMENT	192
6.3.1. Disconnected environment architecture for bare metal	192
6.3.2. Requirements to deploy hosted control planes on bare metal in a disconnected environment	195
6.3.3. Extracting the release image digest	196
6.3.4. DNS configurations on bare metal	196
6.3.5. Deploying a registry for hosted control planes in a disconnected environment	197
6.3.6. Setting up a management cluster for hosted control planes in a disconnected environment	199
6.3.7. Configuring the web server for hosted control planes in a disconnected environment	200
6.3.8. Configuring image mirroring for hosted control planes in a disconnected environment	201
6.3.9. Applying objects in the management cluster	203
6.3.10. Deploying AgentServiceConfig resources	205
6.3.11. Configuring TLS certificates for a disconnected installation of hosted control planes	207
6.3.11.1. Adding the registry CA to the management cluster	207
6.3.11.2. Adding the registry CA to the worker nodes for the hosted cluster	208
6.3.12. Creating a hosted cluster on bare metal	209
6.3.12.1. Deploying hosted cluster objects	209
6.3.12.2. Creating a NodePool object for the hosted cluster	214
6.3.12.3. Creating an InfraEnv resource for the hosted cluster	215
6.3.12.4. Creating bare metal hosts for the hosted cluster	216
6.3.12.5. Scaling up the node pool	219
6.4. DEPLOYING HOSTED CONTROL PLANES ON IBM Z IN A DISCONNECTED ENVIRONMENT	220
6.4.1. Prerequisites to deploy hosted control planes on IBM Z in a disconnected environment	220
6.4.2. Adding credentials and the registry certificate authority to the management cluster	221
6.4.3. Update the registry certificate authority in the AgentServiceConfig resource with the mirror registry	222
6.4.4. Adding the registry certificate authority to the hosted cluster	223
6.5. MONITORING USER WORKLOAD IN A DISCONNECTED ENVIRONMENT	224
6.5.1. Resolving user workload monitoring issues	224
6.5.2. Verifying the status of the hosted control plane feature	225
6.5.3. Configuring the hypershift-addon managed cluster add-on to run on an infrastructure node	226
CHAPTER 7. UPDATING HOSTED CONTROL PLANES	227
7.1. REQUIREMENTS TO UPGRADE HOSTED CONTROL PLANES	227

7.2. SETTING CHANNELS IN A HOSTED CLUSTER	227
7.3. UPDATING THE OPENSHIFT CONTAINER PLATFORM VERSION IN A HOSTED CLUSTER	230
7.3.1. The multicloud engine Operator hub management cluster	230
7.3.2. Supported OpenShift Container Platform versions in a hosted cluster	231
7.4. UPDATES FOR THE HOSTED CLUSTER	232
7.5. UPDATES FOR NODE POOLS	232
7.5.1. Replace updates for node pools	233
7.5.2. In place updates for node pools	233
7.6. UPDATING NODE POOLS IN A HOSTED CLUSTER	233
7.7. UPDATING A CONTROL PLANE IN A HOSTED CLUSTER	234
7.8. UPDATING A HOSTED CLUSTER BY USING THE MULTICLOUD ENGINE OPERATOR CONSOLE	235
CHAPTER 8. HIGH AVAILABILITY FOR HOSTED CONTROL PLANES	237
8.1. ABOUT HIGH AVAILABILITY FOR HOSTED CONTROL PLANES	237
8.1.1. Impact of the failed management cluster component	237
8.2. RECOVERING AN UNHEALTHY ETCD CLUSTER	237
8.2.1. Checking the status of an etcd cluster	237
8.2.2. Recovering a failing etcd pod	238
8.3. BACKING UP AND RESTORING ETCD IN AN ON-PREMISE ENVIRONMENT	239
8.3.1. Backing up and restoring etcd on a hosted cluster in an on-premise environment	239
8.4. BACKING UP AND RESTORING ETCD ON AWS	242
8.4.1. Taking a snapshot of etcd for a hosted cluster	243
8.4.2. Restoring an etcd snapshot on a hosted cluster	244
8.5. BACKING UP AND RESTORING A HOSTED CLUSTER ON OPENSHIFT VIRTUALIZATION	245
8.5.1. Backing up a hosted cluster on OpenShift Virtualization	245
8.5.2. Restoring a hosted cluster on OpenShift Virtualization	247
8.6. DISASTER RECOVERY FOR A HOSTED CLUSTER IN AWS	248
8.6.1. Overview of the backup and restore process	249
8.6.2. Backing up a hosted cluster on AWS	254
8.6.3. Restoring a hosted cluster	260
8.6.4. Deleting a hosted cluster from your source management cluster	262
8.7. DISASTER RECOVERY FOR A HOSTED CLUSTER BY USING OADP	265
8.7.1. Prerequisites	265
8.7.2. Preparing AWS to use OADP	265
8.7.3. Preparing bare metal to use OADP	266
8.7.4. Backing up the data plane workload	266
8.7.5. Backing up the control plane workload	266
8.7.6. Restoring a hosted cluster by using OADP	269
8.7.7. Observing the backup and restore process	272
8.7.8. Using the velero CLI to describe the Backup and Restore resources	272
CHAPTER 9. AUTHENTICATION AND AUTHORIZATION FOR HOSTED CONTROL PLANES	274
9.1. CONFIGURING THE OAUTH SERVER FOR A HOSTED CLUSTER BY USING THE CLI	274
9.2. CONFIGURING THE OAUTH SERVER FOR A HOSTED CLUSTER BY USING THE WEB CONSOLE	276
9.3. ASSIGNING COMPONENTS IAM ROLES BY USING THE CCO IN A HOSTED CLUSTER ON AWS	277
9.4. VERIFYING THE CCO INSTALLATION IN A HOSTED CLUSTER ON AWS	278
9.5. ENABLING OPERATORS TO SUPPORT CCO-BASED WORKFLOWS WITH AWS STS	278
CHAPTER 10. HANDLING MACHINE CONFIGURATION FOR HOSTED CONTROL PLANES	285
10.1. CONFIGURING NODE POOLS FOR HOSTED CONTROL PLANES	285
10.2. REFERENCING THE KUBELET CONFIGURATION IN NODE POOLS	286
10.3. CONFIGURING NODE TUNING IN A HOSTED CLUSTER	287
10.4. DEPLOYING THE SR-IOV OPERATOR FOR HOSTED CONTROL PLANES	289
10.5. CONFIGURING THE NTP SERVER FOR HOSTED CLUSTERS	291

CHAPTER 11. USING FEATURE GATES IN A HOSTED CLUSTER	295
11.1. ENABLING FEATURE SETS BY USING FEATURE GATES	295
CHAPTER 12. OBSERVABILITY FOR HOSTED CONTROL PLANES	297
12.1. CONFIGURING METRICS SETS FOR HOSTED CONTROL PLANES	297
12.1.1. Configuring the SRE metrics set	297
12.2. ENABLING MONITORING DASHBOARDS IN A HOSTED CLUSTER	299
12.2.1. Dashboard customization	300
CHAPTER 13. NETWORKING FOR HOSTED CONTROL PLANES	302
13.1. CONTROL PLANE WORKLOADS THAT NEED TO ACCESS EXTERNAL SERVICES	302
13.2. COMPUTE NODES THAT NEED TO ACCESS AN IGNITION ENDPOINT	302
13.3. COMPUTE NODES THAT NEED TO ACCESS THE API SERVER	302
13.4. MANAGEMENT CLUSTERS THAT NEED EXTERNAL ACCESS	303
13.5. MANAGEMENT CLUSTER THAT USES A PROXY AND A HOSTED CLUSTER WITH A SECONDARY NETWORK AND NO DEFAULT POD NETWORK	303
13.6. ADDITIONAL RESOURCES	303
CHAPTER 14. TROUBLESHOOTING HOSTED CONTROL PLANES	304
14.1. GATHERING INFORMATION TO TROUBLESHOOT HOSTED CONTROL PLANES	304
14.2. GATHERING OPENSHIFT CONTAINER PLATFORM DATA FOR A HOSTED CLUSTER	305
14.2.1. Gathering data for a hosted cluster by using the CLI	305
14.2.2. Gathering data for a hosted cluster by using the web console	306
14.3. ENTERING THE MUST-GATHER COMMAND IN A DISCONNECTED ENVIRONMENT	306
14.4. TROUBLESHOOTING HOSTED CLUSTERS ON OPENSHIFT VIRTUALIZATION	307
14.4.1. HostedCluster resource is stuck in a partial state	307
14.4.2. No worker nodes are registered	307
14.4.3. Worker nodes are stuck in the NotReady state	308
14.4.4. Ingress and console cluster operators are not coming online	308
14.4.5. Load balancer services for the hosted cluster are not available	309
14.4.6. Hosted cluster PVCs are not available	309
14.4.7. VM nodes are not correctly joining the cluster	310
14.4.8. RHCOS image mirroring fails	310
14.4.9. Return non-bare-metal clusters to the late binding pool	311
14.5. TROUBLESHOOTING HOSTED CLUSTERS ON BARE METAL	312
14.5.1. Nodes fail to be added to hosted control planes on bare metal	312
14.6. RESTARTING HOSTED CONTROL PLANE COMPONENTS	312
14.7. PAUSING THE RECONCILIATION OF A HOSTED CLUSTER AND HOSTED CONTROL PLANE	314
14.8. SCALING DOWN THE DATA PLANE TO ZERO	315
CHAPTER 15. DESTROYING A HOSTED CLUSTER	317
15.1. DESTROYING A HOSTED CLUSTER ON AWS	317
15.1.1. Destroying a hosted cluster on AWS by using the CLI	317
15.2. DESTROYING A HOSTED CLUSTER ON BARE METAL	318
15.2.1. Destroying a hosted cluster on bare metal by using the CLI	318
15.2.2. Destroying a hosted cluster on bare metal by using the web console	318
15.3. DESTROYING A HOSTED CLUSTER ON OPENSHIFT VIRTUALIZATION	318
15.3.1. Destroying a hosted cluster on OpenShift Virtualization by using the CLI	318
15.4. DESTROYING A HOSTED CLUSTER ON IBM Z	319
15.4.1. Destroying a hosted cluster on x86 bare metal with IBM Z compute nodes	319
15.5. DESTROYING A HOSTED CLUSTER ON IBM POWER	320
15.5.1. Destroying a hosted cluster on IBM Power by using the CLI	320
15.6. DESTROYING A HOSTED CLUSTER ON NON-BARE-METAL AGENT MACHINES	320
15.6.1. Destroying a hosted cluster on non-bare-metal agent machines	320

15.6.2. Destroying a hosted cluster on non-bare-metal agent machines by using the web console	320
---	-----

CHAPTER 16. MANUALLY IMPORTING A HOSTED CLUSTER	322
16.1. LIMITATIONS OF MANAGING IMPORTED HOSTED CLUSTERS	322
16.2. ADDITIONAL RESOURCES	322
16.3. MANUALLY IMPORTING HOSTED CLUSTERS	322
16.4. MANUALLY IMPORTING A HOSTED CLUSTER ON AWS	323
16.5. DISABLING THE AUTOMATIC IMPORT OF HOSTED CLUSTERS INTO MULTICLUSTER ENGINE OPERATOR	324

CHAPTER 1. HOSTED CONTROL PLANES RELEASE NOTES

Release notes contain information about new and deprecated features, changes, and known issues.

1.1. HOSTED CONTROL PLANES RELEASE NOTES FOR OPENSHIFT CONTAINER PLATFORM 4.18

With this release, hosted control planes for OpenShift Container Platform 4.18 is available. Hosted control planes for OpenShift Container Platform 4.18 supports multicluster engine for Kubernetes Operator version 2.8.

1.1.1. New features and enhancements

This release adds improvements related to the following concepts:

1.1.1.1. Comparing hosted control planes with OpenShift Container Platform

The OpenShift Container Platform documentation now highlights the differences between hosted control planes and standalone OpenShift Container Platform. For more information, see [Differences between hosted control planes and OpenShift Container Platform](#).

1.1.1.2. Proxy configuration for hosted control planes

Configuring proxy support for hosted control planes has a few differences from configuring proxy support for standalone OpenShift Container Platform. For more information, see [Networking for hosted control planes](#).

1.1.1.3. Default container runtime for worker nodes is crun

In hosted control planes for OpenShift Container Platform 4.18 or later, the default container runtime for worker nodes is changed from runC to crun.

1.1.2. Bug fixes

- Previously, you could not create node pools that use ARM64 architecture on a hosted control planes cluster with the platform parameter set to either **agentBareMetal** or **none**. This issue did not exist for a hosted control planes cluster that runs on an AWS or an Microsoft Azure platform. With this release, a fix ensures that you can now create node pools that use ARM64 architecture in a hosted control planes cluster that has **platform** set to either **agentBareMetal** or **none**. ([OCPBUGS-48579](#))
- Previously, when you attempted to use the hosted control planes CLI to create a cluster in a disconnected environment, the installation command failed. An issue existed with the registry that hosts the command. With this release, a fix to the command registry means that you can use the hosted control planes CLI to create a cluster in a disconnected environment. ([OCPBUGS-48170](#))
- Previously, incorrect addresses were being passed to the Kubernetes EndpointSlice on a cluster, and this issue prevented the installation of the MetalLB Operator on an Agent-based cluster in an IPv6 disconnected environment. With this release, a fix modifies the address evaluation method. Red Hat Marketplace pods can now successfully connect to the cluster API server, so that the installation of MetalLB Operator and handling of ingress traffic in IPv6 disconnected environments can occur. ([OCPBUGS-46665](#))

- Previously, in hosted control planes on the Agent platform, ARM64 architecture was not allowed in the **NodePool** API. As a consequence, heterogeneous clusters could not be deployed on the Agent platform. In this release, the API now allows ARM64 architecture node pools on the Agent platform. ([OCPBUGS-4673](#))
- Previously, the default **node-monitor-grace-period** value was 50 seconds. As a consequence, nodes did not stay ready for the duration of time that Kubernetes components needed to reconnect, coordinate, and complete their requests. With this release, the default **node-monitor-grace-period** value is 55 seconds. As a result, the issue is resolved and deployments have enough time to be completed. ([OCPBUGS-46008](#))
- Previously, the provider ID for the **IBMPowerVSMachine** object was not populated properly, due to the improper retrieval of IBM Cloud Workspace ID. As a consequence, any certificate signing requests (CSRs) were pending in hosted clusters. With this release, the provider ID is properly populated for the **IBMPowerVSMachine** object. As a result, no CSRs are pending, and all COs are moved to the available state. ([OCPBUGS-44944](#))
- Previously, when you created a hosted cluster by using a shared VPC where the private DNS hosted zones existed in the cluster creator account, the private link controller failed to create the route53 DNS records in the local zone. With this release, the ingress shared role adds records to the private link controller. The VPC endpoint is used to share the role to create the VPC endpoint in the VPC owner account. A hosted cluster is created in a shared VPC configuration, where the private hosted zones exist in the cluster creator account. ([OCPBUGS-44630](#))
- Previously, when the hosted cluster **controllerAvailabilityPolicy** parameter was set to **SingleReplica**, the **podAntiAffinity** rules on networking components blocked the availability of the components. With this release, the issue is resolved. ([OCPBUGS-39313](#))
- Previously, periodic conformance jobs in hosted control planes failed because of changes to the core operating system. These failed jobs caused the OpenShift API deployment to fail. With this release, an update recursively copies individual trusted certificate authority (CA) certificates instead of copying a single file, so that the periodic conformance jobs succeed and the OpenShift API runs as expected. ([OCPBUGS-38943](#))
- Previously, you could not create **NodePool** resources with ARM64 architecture on non-AWS or Azure platforms. This bug resulted in validation errors that prevented the addition of bare-metal compute nodes and caused CEL validation blocks when creating a **NodePool** resource. The fix modifies the **NodePool** spec validation rules by adding "**self.platform.type == 'None'**". As a result, you can now create **NodePool** resources with ARM64 architecture specifications on non-AWS or Azure bare-metal platforms, expanding functionality. ([OCPBUGS-46440](#))
- Previously, when you created a hosted cluster in a shared VPC, the private link controller sometimes failed to assume the shared VPC role to manage the VPC endpoints in the shared VPC. With this release, a client is created for every reconciliation in the private link controller so that you can recover from invalid clients. As a result, the hosted cluster endpoints and the hosted cluster are created successfully. ([OCPBUGS-45184](#))

1.1.3. Known issues

- If the annotation and the **ManagedCluster** resource name do not match, the multicluster engine for Kubernetes Operator console displays the cluster as **Pending import**. The cluster cannot be used by the multicluster engine Operator. The same issue happens when there is no annotation and the **ManagedCluster** name does not match the **Infra-ID** value of the **HostedCluster** resource.

- When you use the multicluster engine for Kubernetes Operator console to add a new node pool to an existing hosted cluster, the same version of OpenShift Container Platform might appear more than once in the list of options. You can select any instance in the list for the version that you want.
- When a node pool is scaled down to 0 workers, the list of hosts in the console still shows nodes in a **Ready** state. You can verify the number of nodes in two ways:
 - In the console, go to the node pool and verify that it has 0 nodes.
 - On the command-line interface, run the following commands:
 - Verify that 0 nodes are in the node pool by running the following command:

```
$ oc get nodepool -A
```
 - Verify that 0 nodes are in the cluster by running the following command:

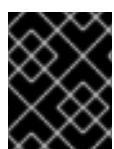
```
$ oc get nodes --kubeconfig
```
 - Verify that 0 agents are reported as bound to the cluster by running the following command:

```
$ oc get agents -A
```
- When you create a hosted cluster in an environment that uses the dual-stack network, you might encounter the following DNS-related issues:
 - **CrashLoopBackOff** state in the **service-ca-operator** pod: When the pod tries to reach the Kubernetes API server through the hosted control plane, the pod cannot reach the server because the data plane proxy in the **kube-system** namespace cannot resolve the request. This issue occurs because in the HAProxy setup, the front end uses an IP address and the back end uses a DNS name that the pod cannot resolve.
 - Pods stuck in the **ContainerCreating** state: This issue occurs because the **openshift-service-ca-operator** resource cannot generate the **metrics-tls** secret that the DNS pods need for DNS resolution. As a result, the pods cannot resolve the Kubernetes API server. To resolve these issues, configure the DNS server settings for a dual stack network.
- On the Agent platform, the hosted control planes feature periodically rotates the token that the Agent uses to pull ignition. As a result, if you have an Agent resource that was created some time ago, it might fail to pull ignition. As a workaround, in the Agent specification, delete the secret of the **IgnitionEndpointTokenReference** property then add or modify any label on the Agent resource. The system re-creates the secret with the new token.
- If you created a hosted cluster in the same namespace as its managed cluster, detaching the managed hosted cluster deletes everything in the managed cluster namespace including the hosted cluster. The following situations can create a hosted cluster in the same namespace as its managed cluster:
 - You created a hosted cluster on the Agent platform through the multicluster engine for Kubernetes Operator console by using the default hosted cluster cluster namespace.
 - You created a hosted cluster through the command-line interface or API by specifying the hosted cluster namespace to be the same as the hosted cluster name.

- When you use the console or API to specify an IPv6 address for the **spec.services.servicePublishingStrategy.nodePort.address** field of a hosted cluster, a full IPv6 address with 8 hexets is required. For example, instead of specifying **2620:52:0:1306::30**, you need to specify **2620:52:0:1306:0:0:0:30**.

1.1.4. General Availability and Technology Preview features

Some features in this release are currently in Technology Preview. These experimental features are not intended for production use. For more information about the scope of support for these features, see [Technology Preview Features Support Scope](#) on the Red Hat Customer Portal.



IMPORTANT

For IBM Power and IBM Z, you must run the control plane on machine types based on 64-bit x86 architecture, and node pools on IBM Power or IBM Z.

Table 1.1. Hosted control planes GA and TP tracker

Feature	4.16	4.17	4.18
Hosted control planes for OpenShift Container Platform on Amazon Web Services (AWS)	General Availability	General Availability	General Availability
Hosted control planes for an ARM64 OpenShift Container Platform cluster on Amazon Web Services	Technology Preview	General Availability	General Availability
Hosted control planes for OpenShift Container Platform on bare metal	General Availability	General Availability	General Availability
Hosted control planes for OpenShift Container Platform on OpenShift Virtualization	General Availability	General Availability	General Availability
Hosted control planes for OpenShift Container Platform using non-bare-metal agent machines	Technology Preview	Technology Preview	Technology Preview
Hosted control planes for OpenShift Container Platform on IBM Power	Technology Preview	General Availability	General Availability
Hosted control planes for OpenShift Container Platform on IBM Z	Technology Preview	General Availability	General Availability
Hosted control planes for OpenShift Container Platform on RHOSP	Not Available	Developer Preview	Developer Preview

CHAPTER 2. HOSTED CONTROL PLANES OVERVIEW

You can deploy OpenShift Container Platform clusters by using two different control plane configurations: standalone or hosted control planes. The standalone configuration uses dedicated virtual machines or physical machines to host the control plane. With hosted control planes for OpenShift Container Platform, you create control planes as pods on a management cluster without the need for dedicated virtual or physical machines for each control plane.

2.1. INTRODUCTION TO HOSTED CONTROL PLANES

Hosted control planes is available by using a [supported version of multicluster engine for Kubernetes Operator](#) on the following platforms:

- Bare metal by using the Agent provider
- Non-bare-metal Agent machines, as a Technology Preview feature
- OpenShift Virtualization
- Amazon Web Services (AWS)
- IBM Z
- IBM Power

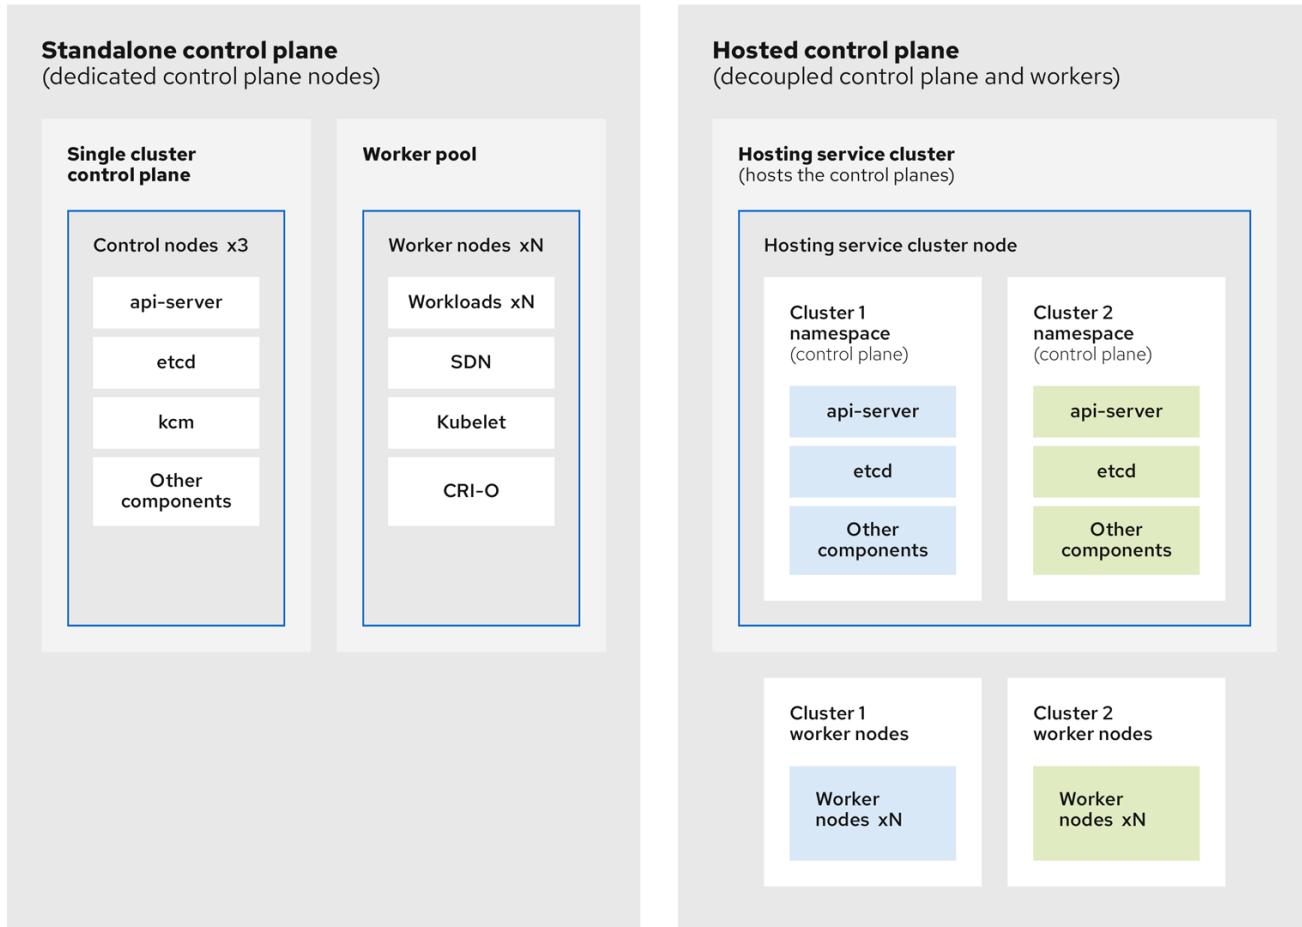
The hosted control planes feature is enabled by default.

2.1.1. Architecture of hosted control planes

OpenShift Container Platform is often deployed in a coupled, or standalone, model, where a cluster consists of a control plane and a data plane. The control plane includes an API endpoint, a storage endpoint, a workload scheduler, and an actuator that ensures state. The data plane includes compute, storage, and networking where workloads and applications run.

The standalone control plane is hosted by a dedicated group of nodes, which can be physical or virtual, with a minimum number to ensure quorum. The network stack is shared. Administrator access to a cluster offers visibility into the cluster's control plane, machine management APIs, and other components that contribute to the state of a cluster.

Although the standalone model works well, some situations require an architecture where the control plane and data plane are decoupled. In those cases, the data plane is on a separate network domain with a dedicated physical hosting environment. The control plane is hosted by using high-level primitives such as deployments and stateful sets that are native to Kubernetes. The control plane is treated as any other workload.



272_OpenShift_1122

2.1.2. Benefits of hosted control planes

With hosted control planes, you can pave the way for a true hybrid-cloud approach and enjoy several other benefits.

- The security boundaries between management and workloads are stronger because the control plane is decoupled and hosted on a dedicated hosting service cluster. As a result, you are less likely to leak credentials for clusters to other users. Because infrastructure secret account management is also decoupled, cluster infrastructure administrators cannot accidentally delete control plane infrastructure.
- With hosted control planes, you can run many control planes on fewer nodes. As a result, clusters are more affordable.
- Because the control planes consist of pods that are launched on OpenShift Container Platform, control planes start quickly. The same principles apply to control planes and workloads, such as monitoring, logging, and auto-scaling.
- From an infrastructure perspective, you can push registries, HAProxy, cluster monitoring, storage nodes, and other infrastructure components to the tenant's cloud provider account, isolating usage to the tenant.
- From an operational perspective, multicluster management is more centralized, which results in fewer external factors that affect the cluster status and consistency. Site reliability engineers have a central place to debug issues and navigate to the cluster data plane, which can lead to shorter Time to Resolution (TTR) and greater productivity.

2.2. DIFFERENCES BETWEEN HOSTED CONTROL PLANES AND OPENSHIFT CONTAINER PLATFORM

Hosted control planes is a form factor of OpenShift Container Platform. Hosted clusters and the stand-alone OpenShift Container Platform clusters are configured and managed differently. See the following tables to understand the differences between OpenShift Container Platform and hosted control planes:

2.2.1. Cluster creation and lifecycle

OpenShift Container Platform	Hosted control planes
You install a standalone OpenShift Container Platform cluster by using the openshift-install binary or the Assisted Installer.	You install a hosted cluster by using the hypershift.openshift.io API resources such as HostedCluster and NodePool , on an existing OpenShift Container Platform cluster.

2.2.2. Cluster configuration

OpenShift Container Platform	Hosted control planes
You configure cluster-scoped resources such as authentication, API server, and proxy by using the config.openshift.io API group.	You configure resources that impact the control plane in the HostedCluster resource.

2.2.3. etcd encryption

OpenShift Container Platform	Hosted control planes
You configure etcd encryption by using the APIServer resource with AES-GCM or AES-CBC. For more information, see "Enabling etcd encryption".	You configure etcd encryption by using the HostedCluster resource in the SecretEncryption field with AES-CBC or KMS for Amazon Web Services.

2.2.4. Operators and control plane

OpenShift Container Platform	Hosted control planes
A standalone OpenShift Container Platform cluster contains separate Operators for each control plane component.	A hosted cluster contains a single Operator named Control Plane Operator that runs in the hosted control plane namespace on the management cluster.
etcd uses storage that is mounted on the control plane nodes. The etcd cluster Operator manages etcd.	etcd uses a persistent volume claim for storage and is managed by the Control Plane Operator.

OpenShift Container Platform	Hosted control planes
The Ingress Operator, network related Operators, and Operator Lifecycle Manager (OLM) run on the cluster.	The Ingress Operator, network related Operators, and Operator Lifecycle Manager (OLM) run in the hosted control plane namespace on the management cluster.
The OAuth server runs inside the cluster and is exposed through a route in the cluster.	The OAuth server runs inside the control plane and is exposed through a route, node port, or load balancer on the management cluster.

2.2.5. Updates

OpenShift Container Platform	Hosted control planes
The Cluster Version Operator (CVO) orchestrates the update process and monitors the ClusterVersion resource. Administrators and OpenShift components can interact with the CVO through the ClusterVersion resource. The oc adm upgrade command results in a change to the ClusterVersion.Spec.DesiredUpdate field in the ClusterVersion resource.	The hosted control planes update results in a change to the .spec.release.image field in the HostedCluster and NodePool resources. Any changes to the ClusterVersion resource are ignored.
After you update an OpenShift Container Platform cluster, both the control plane and compute machines are updated.	After you update the hosted cluster, only the control plane is updated. You perform node pool updates separately.

2.2.6. Machine configuration and management

OpenShift Container Platform	Hosted control planes
The MachineSets resource manages machines in the openshift-machine-api namespace.	The NodePool resource manages machines on the management cluster.
A set of control plane machines are available.	A set of control plane machines do not exist.
You enable a machine health check by using the MachineHealthCheck resource.	You enable a machine health check through the .spec.management.autoRepair field in the NodePool resource.
You enable autoscaling by using the ClusterAutoscaler and MachineAutoscaler resources.	You enable autoscaling through the spec.autoScaling field in the NodePool resource.

OpenShift Container Platform	Hosted control planes
Machines and machine sets are exposed in the cluster.	Machines, machine sets, and machine deployments from upstream Cluster CAPI Operator are used to manage machines but are not exposed to the user.
All machine sets are upgraded automatically when you update the cluster.	You update your node pools independently from the hosted cluster updates.
Only an in-place upgrade is supported in the cluster.	Both replace and in-place upgrades are supported in the hosted cluster.
The Machine Config Operator manages configurations for machines.	The Machine Config Operator does not exist in hosted control planes.
You configure machine Ignition by using the MachineConfig , KubeletConfig , and ContainerRuntimeConfig resources that are selected from a MachineConfigPool selector.	You configure the MachineConfig , KubeletConfig , and ContainerRuntimeConfig resources through the config map referenced in the spec.config field of the NodePool resource.
The Machine Config Daemon (MCD) manages configuration changes and updates on each of the nodes.	For an in-place upgrade, the node pool controller creates a run-once pod that updates a machine based on your configuration.
You can modify the machine configuration resources such as the SR-IOV Operator.	You cannot modify the machine configuration resources.

2.2.7. Networking

OpenShift Container Platform	Hosted control planes
The Kube API server communicates with nodes directly, because the Kube API server and nodes exist in the same Virtual Private Cloud (VPC).	The Kube API server communicates with nodes through Konnectivity. The Kube API server and nodes exist in a different Virtual Private Cloud (VPC).
Nodes communicate with the Kube API server through the internal load balancer.	Nodes communicate with the Kube API server through an external load balancer or a node port.

2.2.8. Web console

OpenShift Container Platform	Hosted control planes
The web console shows the status of a control plane.	The web console does not show the status of a control plane.
You can update your cluster by using the web console.	You cannot update the hosted cluster by using the web console.

OpenShift Container Platform	Hosted control planes
The web console displays the infrastructure resources such as machines.	The web console does not display the infrastructure resources.
You can configure machines through the MachineConfig resource by using the web console.	You cannot configure machines by using the web console.

Additional resources

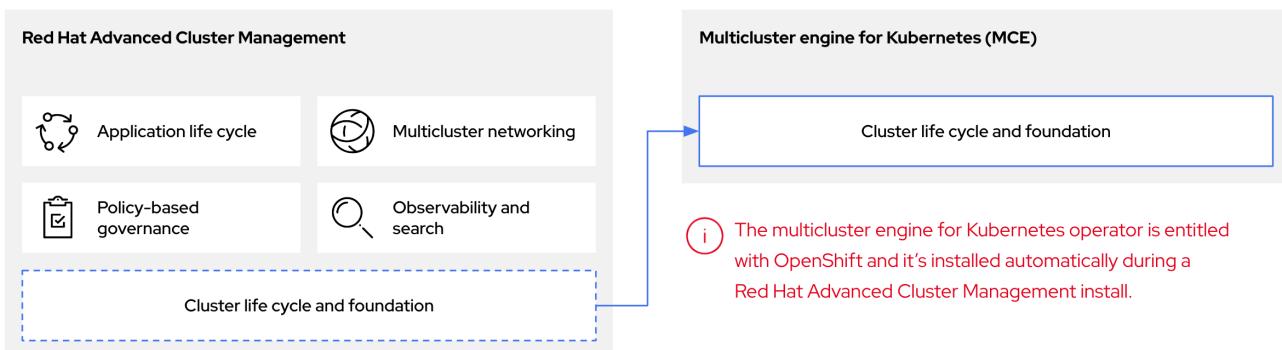
- [Enabling etcd encryption](#)

2.3. RELATIONSHIP BETWEEN HOSTED CONTROL PLANES, MULTICLUSTER ENGINE OPERATOR, AND RHACM

You can configure hosted control planes by using the multicluster engine for Kubernetes Operator. The multicluster engine is an integral part of Red Hat Advanced Cluster Management (RHACM) and is enabled by default with RHACM. The multicluster engine Operator cluster lifecycle defines the process of creating, importing, managing, and destroying Kubernetes clusters across various infrastructure cloud providers, private clouds, and on-premises data centers.

The multicluster engine Operator is the cluster lifecycle Operator that provides cluster management capabilities for OpenShift Container Platform and RHACM hub clusters. The multicluster engine Operator enhances cluster fleet management and supports OpenShift Container Platform cluster lifecycle management across clouds and data centers.

Figure 2.1. Cluster life cycle and foundation



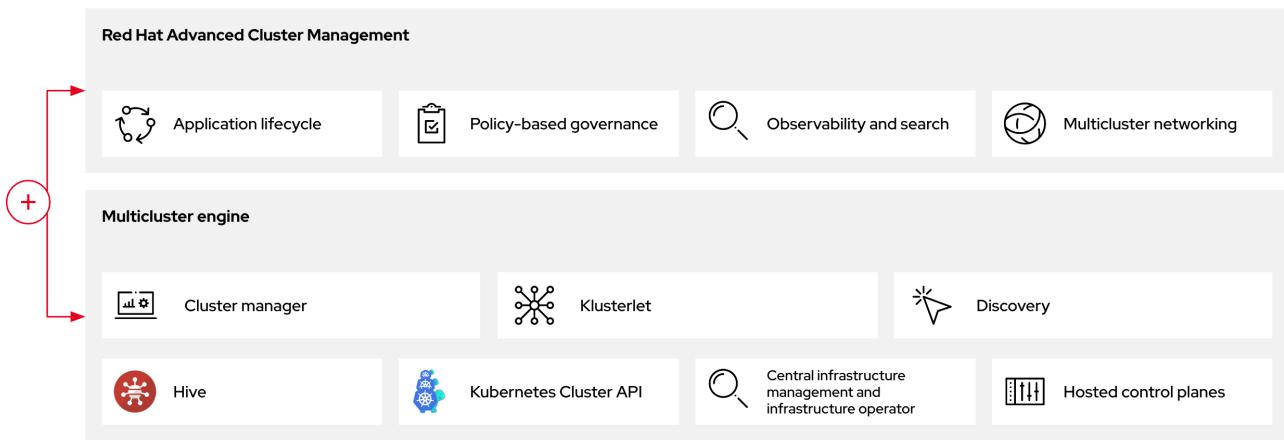
You can use the multicluster engine Operator with OpenShift Container Platform as a standalone cluster manager or as part of a RHACM hub cluster.

TIP

A management cluster is also known as the hosting cluster.

You can deploy OpenShift Container Platform clusters by using two different control plane configurations: standalone or hosted control planes. The standalone configuration uses dedicated virtual machines or physical machines to host the control plane. With hosted control planes for OpenShift Container Platform, you create control planes as pods on a management cluster without the need for dedicated virtual or physical machines for each control plane.

Figure 2.2. RHACM and the multicluster engine Operator introduction diagram



2.3.1. Discovering multicluster engine Operator hosted clusters in RHACM

If you want to bring hosted clusters to a Red Hat Advanced Cluster Management (RHACM) hub cluster to manage them with RHACM management components, see the instructions in the [Red Hat Advanced Cluster Management official documentation](#).

2.4. VERSIONING FOR HOSTED CONTROL PLANES

The hosted control planes feature includes the following components, which might require independent versioning and support levels:

- Management cluster
- HyperShift Operator
- Hosted control planes (**hcp**) command-line interface (CLI)
- **hypershift.openshift.io** API
- Control Plane Operator

2.4.1. Management cluster

In management clusters for production use, you need multicluster engine for Kubernetes Operator, which is available through OperatorHub. The multicluster engine Operator bundles a supported build of the HyperShift Operator. For your management clusters to remain supported, you must use the version of OpenShift Container Platform that multicluster engine Operator runs on. In general, a new release of multicluster engine Operator runs on the following versions of OpenShift Container Platform:

- The latest General Availability version of OpenShift Container Platform
- Two versions before the latest General Availability version of OpenShift Container Platform

The full list of OpenShift Container Platform versions that you can install through the HyperShift Operator on a management cluster depends on the version of your HyperShift Operator. However, the list always includes at least the same OpenShift Container Platform version as the management cluster and two previous minor versions relative to the management cluster. For example, if the management cluster is running 4.17 and a supported version of multicluster engine Operator, the HyperShift Operator can install 4.17, 4.16, 4.15, and 4.14 hosted clusters.

With each major, minor, or patch version release of OpenShift Container Platform, two components of hosted control planes are released:

- The HyperShift Operator
- The **hcp** command-line interface (CLI)

2.4.2. HyperShift Operator

The HyperShift Operator manages the lifecycle of hosted clusters that are represented by the **HostedCluster** API resources. The HyperShift Operator is released with each OpenShift Container Platform release. The HyperShift Operator creates the **supported-versions** config map in the **hypershift** namespace. The config map contains the supported hosted cluster versions.

You can host different versions of control planes on the same management cluster.

Example **supported-versions** config map object

```
apiVersion: v1
data:
  supported-versions: '{"versions":["4.18"]}'
kind: ConfigMap
metadata:
  labels:
    hypershift.openshift.io/supported-versions: "true"
  name: supported-versions
  namespace: hypershift
```

2.4.3. hosted control planes CLI

You can use the **hcp** CLI to create hosted clusters. You can download the CLI from multicluster engine Operator. When you run the **hcp version** command, the output shows the latest OpenShift Container Platform that the CLI supports against your **kubeconfig** file.

2.4.4. hypershift.openshift.io API

You can use the **hypershift.openshift.io** API resources, such as, **HostedCluster** and **NodePool**, to create and manage OpenShift Container Platform clusters at scale. A **HostedCluster** resource contains the control plane and common data plane configuration. When you create a **HostedCluster** resource, you have a fully functional control plane with no attached nodes. A **NodePool** resource is a scalable set of worker nodes that is attached to a **HostedCluster** resource.

The API version policy generally aligns with the policy for [Kubernetes API versioning](#).

Updates for hosted control planes involve updating the hosted cluster and the node pools. For more information, see "Updates for hosted control planes".

2.4.5. Control Plane Operator

The Control Plane Operator is released as part of each OpenShift Container Platform payload release image for the following architectures:

- amd64
- arm64

- multi-arch

Additional resources

- [AMD64 release images](#)
- [ARM64 release images](#)
- [Multi-arch release images](#)

2.5. GLOSSARY OF COMMON CONCEPTS AND PERSONAS FOR HOSTED CONTROL PLANES

When you use hosted control planes for OpenShift Container Platform, it is important to understand its key concepts and the personas that are involved.

2.5.1. Concepts

hosted cluster

An OpenShift Container Platform cluster with its control plane and API endpoint hosted on a management cluster. The hosted cluster includes the control plane and its corresponding data plane.

hosted cluster infrastructure

Network, compute, and storage resources that exist in the tenant or end-user cloud account.

hosted control plane

An OpenShift Container Platform control plane that runs on the management cluster, which is exposed by the API endpoint of a hosted cluster. The components of a control plane include etcd, the Kubernetes API server, the Kubernetes controller manager, and a VPN.

hosting cluster

See *management cluster*.

managed cluster

A cluster that the hub cluster manages. This term is specific to the cluster lifecycle that the multicloud engine for Kubernetes Operator manages in Red Hat Advanced Cluster Management. A managed cluster is not the same thing as a *management cluster*. For more information, see [Managed cluster](#).

management cluster

An OpenShift Container Platform cluster where the HyperShift Operator is deployed and where the control planes for hosted clusters are hosted. The management cluster is synonymous with the *hosting cluster*.

management cluster infrastructure

Network, compute, and storage resources of the management cluster.

node pool

A resource that manages a set of compute nodes that are associated with a hosted cluster. The compute nodes run applications and workloads within the hosted cluster.

2.5.2. Personas

cluster instance administrator

Users who assume this role are the equivalent of administrators in standalone OpenShift Container Platform. This user has the **cluster-admin** role in the provisioned cluster, but might not have power

over when or how the cluster is updated or configured. This user might have read-only access to see some configuration projected into the cluster.

cluster instance user

Users who assume this role are the equivalent of developers in standalone OpenShift Container Platform. This user does not have a view into OperatorHub or machines.

cluster service consumer

Users who assume this role can request control planes and worker nodes, drive updates, or modify externalized configurations. Typically, this user does not manage or access cloud credentials or infrastructure encryption keys. The cluster service consumer persona can request hosted clusters and interact with node pools. Users who assume this role have RBAC to create, read, update, or delete hosted clusters and node pools within a logical boundary.

cluster service provider

Users who assume this role typically have the **cluster-admin** role on the management cluster and have RBAC to monitor and own the availability of the HyperShift Operator as well as the control planes for the tenant's hosted clusters. The cluster service provider persona is responsible for several activities, including the following examples:

- Owning service-level objects for control plane availability, uptime, and stability
- Configuring the cloud account for the management cluster to host control planes
- Configuring the user-provisioned infrastructure, which includes the host awareness of available compute resources

CHAPTER 3. PREPARING TO DEPLOY HOSTED CONTROL PLANES

3.1. REQUIREMENTS FOR HOSTED CONTROL PLANES

In the context of hosted control planes, a *management cluster* is an OpenShift Container Platform cluster where the HyperShift Operator is deployed and where the control planes for hosted clusters are hosted.

The control plane is associated with a hosted cluster and runs as pods in a single namespace. When the cluster service consumer creates a hosted cluster, it creates a worker node that is independent of the control plane.

The following requirements apply to hosted control planes:

- In order to run the HyperShift Operator, your management cluster needs at least three worker nodes.
- You can run both the management cluster and the worker nodes on-premise, such as on a bare-metal platform or on OpenShift Virtualization. In addition, you can run both the management cluster and the worker nodes on cloud infrastructure, such as Amazon Web Services (AWS).
- If you use a mixed infrastructure, such as running the management cluster on AWS and your worker nodes on-premise, or running your worker nodes on AWS and your management cluster on-premise, you must use the **PublicAndPrivate** publishing strategy and follow the latency requirements in the support matrix.
- In Bare Metal Host (BMH) deployments, where the Bare Metal Operator starts machines, the hosted control plane must be able to reach baseboard management controllers (BMCs). If your security profile does not permit the Cluster Baremetal Operator to access the network where the BMHs have their BMCs in order to enable Redfish automation, you can use BYO ISO support. However, in BYO mode, OpenShift Container Platform cannot automate the powering on of BMHs.

3.1.1. Support matrix for hosted control planes

Because multicluster engine for Kubernetes Operator includes the HyperShift Operator, releases of hosted control planes align with releases of multicluster engine Operator. For more information, see [OpenShift Operator Life Cycles](#).

3.1.1.1. Management cluster support

Any supported standalone OpenShift Container Platform cluster can be a management cluster.



NOTE

A single-node OpenShift Container Platform cluster is not supported as a management cluster. If you have resource constraints, you can share infrastructure between a standalone OpenShift Container Platform control plane and hosted control planes. For more information, see "Shared infrastructure between hosted and standalone control planes".

The following table maps multicluster engine Operator versions to the management cluster versions that support them:

Table 3.1. Supported multicluster engine Operator versions for OpenShift Container Platform management clusters

Management cluster version	Supported multicluster engine Operator version
4.14 – 4.15	2.4
4.14 – 4.16	2.5
4.14 – 4.17	2.6
4.15 – 4.17	2.7
4.16 – 4.18	2.8

3.1.1.2. Hosted cluster support

For hosted clusters, no direct relationship exists between the management cluster version and the hosted cluster version. The hosted cluster version depends on the HyperShift Operator that is included with your multicluster engine Operator version.



NOTE

Ensure a maximum latency of 200 ms between the management cluster and hosted clusters. This requirement is especially important for mixed infrastructure deployments, such as when your management cluster is on AWS and your worker nodes are on-premise.

The following table maps multicluster engine Operator versions to the hosted cluster versions that you can create by using the HyperShift Operator that is associated with that version of multicluster engine Operator:



NOTE

Although the HyperShift Operator supports the hosted cluster versions in the following table, multicluster engine Operator supports only as far back as 2 versions earlier than the current version. For example, if the current hosted cluster version is 4.18, multicluster engine Operator supports as far back as version 4.16. If you want to use a hosted cluster version that is earlier than one of the versions that multicluster engine Operator supports, you can detach your hosted clusters from multicluster engine Operator to be unmanaged, or you can use an earlier version of multicluster engine Operator. For more information, see [The multicluster engine for Kubernetes operator 2.8 Support Matrix](#).

Table 3.2. Hosted cluster versions that can be created by multicluster engine Operator versions

Hosted cluster version	multicluster engine Operator 2.4	multicluster engine Operator 2.5	multicluster engine Operator 2.6	multicluster engine Operator 2.7	multicluster engine Operator 2.8
4.14	Yes	Yes	Yes	Yes	Yes
4.15	No	Yes	Yes	Yes	Yes
4.16	No	No	Yes	Yes	Yes
4.17	No	No	No	Yes	Yes
4.18	No	No	No	No	Yes

3.1.1.3. Hosted cluster platform support

A hosted cluster supports only one infrastructure platform. For example, you cannot create multiple node pools on different infrastructure platforms.

The following table indicates which OpenShift Container Platform versions are supported for each platform of hosted control planes.



IMPORTANT

For IBM Power and IBM Z, you must run the control plane on machine types based on 64-bit x86 architecture, and node pools on IBM Power or IBM Z.

In the following table, the management cluster version is the OpenShift Container Platform version where the multicluster engine Operator is enabled:

Table 3.3. Required OpenShift Container Platform versions for platforms

Hosted cluster platform	Management cluster version	Hosted cluster version
Amazon Web Services	4.16 - 4.18	4.16 - 4.18
IBM Power	4.17 - 4.18	4.17 - 4.18
IBM Z	4.17 - 4.18	4.17 - 4.18
OpenShift Virtualization	4.14 - 4.18	4.14 - 4.18
Bare metal	4.14 - 4.18	4.14 - 4.18
Non-bare-metal agent machines (Technology Preview)	4.16 - 4.18	4.16 - 4.18

3.1.1.4. Multi-architecture support

The following tables indicate the support status for hosted control planes on multiple architectures, organized by platform.

Table 3.4. Multi-architecture support for hosted control planes on AWS

OpenShift Container Platform version	ARM64 control planes	ARM64 worker nodes
4.18	General Availability	General Availability
4.17	General Availability	General Availability
4.16	Technology Preview	General Availability
4.15	Technology Preview	General Availability
4.14	Technology Preview	General Availability

Table 3.5. Multi-architecture support for hosted control planes on bare metal

OpenShift Container Platform version	ARM64 control planes	ARM64 worker nodes
4.18	Not available	Technology Preview
4.17	Not available	Technology Preview
4.16	Not available	Technology Preview
4.15	Not available	Technology Preview
4.14	Not available	Technology Preview

Table 3.6. Multi-architecture support for hosted control planes on non-bare-metal agent machines

OpenShift Container Platform version	ARM64 control planes	ARM64 worker nodes
4.18	Not available	Not available
4.17	Not available	Not available

Table 3.7. Multi-architecture support for hosted control planes on IBM Power

OpenShift Container Platform version	ARM64 control planes	ARM64 worker nodes
4.18	Not available	Not available
4.17	Not available	Not available

Table 3.8. Multi-architecture support for hosted control planes on IBM Z

OpenShift Container Platform version	ARM64 control planes	ARM64 worker nodes
4.18	Not available	Not available
4.17	Not available	Not available

Table 3.9. Multi-architecture support for hosted control planes on OpenShift Virtualization

OpenShift Container Platform version	ARM64 control planes	ARM64 worker nodes
4.18	Not available	Not available
4.17	Not available	Not available
4.16	Not available	Not available
4.15	Not available	Not available
4.14	Not available	Not available

3.1.1.5. Updates of multicluster engine Operator

When you update to another version of the multicluster engine Operator, your hosted cluster can continue to run if the HyperShift Operator that is included in the version of multicluster engine Operator supports the hosted cluster version. The following table shows which hosted cluster versions are supported on which updated multicluster engine Operator versions.



NOTE

Although the HyperShift Operator supports the hosted cluster versions in the following table, multicluster engine Operator supports only as far back as 2 versions earlier than the current version. For example, if the current hosted cluster version is 4.18, multicluster engine Operator supports as far back as version 4.16. If you want to use a hosted cluster version that is earlier than one of the versions that multicluster engine Operator supports, you can detach your hosted clusters from multicluster engine Operator to be unmanaged, or you can use an earlier version of multicluster engine Operator. For more information, see [The multicluster engine for Kubernetes operator 2.8 Support Matrix](#).

Table 3.10. Updated multicluster engine Operator version support for hosted clusters

Updated multicluster engine Operator version	Supported hosted cluster version
Updated from 2.4 to 2.5	OpenShift Container Platform 4.14
Updated from 2.5 to 2.6	OpenShift Container Platform 4.14 – 4.15
Updated from 2.6 to 2.7	OpenShift Container Platform 4.14 – 4.16
Updated from 2.7 to 2.8	OpenShift Container Platform 4.14 – 4.17

For example, if you have an OpenShift Container Platform 4.14 hosted cluster on the management cluster and you update from multicluster engine Operator 2.4 to 2.5, the hosted cluster can continue to run.

3.1.1.6. Technology Preview features

The following list indicates Technology Preview features for this release:

- Custom taints and tolerations for hosted control planes
- NVIDIA GPU devices on hosted control planes for OpenShift Virtualization

3.1.2. FIPS-enabled hosted clusters

The binaries for hosted control planes are FIPs-compliant, with the exception of the hosted control planes command-line interface, **hcp**.

If you want to deploy a FIPS-enabled hosted cluster, you must use a FIPS-enabled management cluster. To enable FIPS mode for your management cluster, you must run the installation program from a Red Hat Enterprise Linux (RHEL) computer configured to operate in FIPS mode. For more information about configuring FIPS mode on RHEL, see [Switching RHEL to FIPS mode](#).

When running RHEL or Red Hat Enterprise Linux CoreOS (RHCOS) booted in FIPS mode, OpenShift Container Platform core components use the RHEL cryptographic libraries that have been submitted to NIST for FIPS 140-2/140-3 Validation on only the x86_64, ppc64le, and s390x architectures.

After you set up your management cluster in FIPS mode, the hosted cluster creation process runs on that management cluster.

Additional resources

- [The multicluster engine for Kubernetes Operator 2.8 support matrix](#)
- [Red Hat OpenShift Container Platform Operator Update Information Checker](#)
- [Shared infrastructure between hosted and standalone control planes](#)

3.1.3. CIDR ranges for hosted control planes

For deploying hosted control planes on OpenShift Container Platform, use the following required Classless Inter-Domain Routing (CIDR) subnet ranges:

- **v4InternalSubnet**: 100.65.0.0/16 (OVN-Kubernetes)
- **clusterNetwork**: 10.132.0.0/14 (pod network)
- **serviceNetwork**: 172.31.0.0/16

For more information about OpenShift Container Platform CIDR range definitions, see "CIDR range definitions".

Additional resources

- [CIDR range definitions](#)

3.2. SIZING GUIDANCE FOR HOSTED CONTROL PLANES

Many factors, including hosted cluster workload and worker node count, affect how many hosted control planes can fit within a certain number of worker nodes. Use this sizing guide to help with hosted cluster capacity planning. This guidance assumes a highly available hosted control planes topology. The load-based sizing examples were measured on a bare-metal cluster. Cloud-based instances might have different limiting factors, such as memory size.

You can override the following resource utilization sizing measurements and disable the metric service monitoring.

See the following highly available hosted control planes requirements, which were tested with OpenShift Container Platform version 4.12.9 and later:

- 78 pods
- Three 8 GiB PVs for etcd
- Minimum vCPU: approximately 5.5 cores
- Minimum memory: approximately 19 GiB

Additional resources

- [Overriding resource utilization measurements](#)
- [Distributing hosted cluster workloads](#)

3.2.1. Pod limits

The **maxPods** setting for each node affects how many hosted clusters can fit in a control-plane node. It is important to note the **maxPods** value on all control-plane nodes. Plan for about 75 pods for each highly available hosted control plane.

For bare-metal nodes, the default **maxPods** setting of 250 is likely to be a limiting factor because roughly three hosted control planes fit for each node given the pod requirements, even if the machine has plenty of resources to spare. Setting the **maxPods** value to 500 by configuring the **KubeletConfig** value allows for greater hosted control plane density, which can help you take advantage of additional compute resources.

Additional resources

- Configuring the maximum number of pods per node

3.2.2. Request-based resource limit

The maximum number of hosted control planes that the cluster can host is calculated based on the hosted control plane CPU and memory requests from the pods.

A highly available hosted control plane consists of 78 pods that request 5 vCPUs and 18 GB memory. These baseline numbers are compared to the cluster worker node resource capacities to estimate the maximum number of hosted control planes.

3.2.3. Load-based limit

The maximum number of hosted control planes that the cluster can host is calculated based on the hosted control plane pods CPU and memory utilizations when some workload is put on the hosted control plane Kubernetes API server.

The following method is used to measure the hosted control plane resource utilizations as the workload increases:

- A hosted cluster with 9 workers that use 8 vCPU and 32 GiB each, while using the KubeVirt platform
- The workload test profile that is configured to focus on API control-plane stress, based on the following definition:
 - Created objects for each namespace, scaling up to 100 namespaces total
 - Additional API stress with continuous object deletion and creation
 - Workload queries-per-second (QPS) and Burst settings set high to remove any client-side throttling

As the load increases by 1000 QPS, the hosted control plane resource utilization increases by 9 vCPUs and 2.5 GB memory.

For general sizing purposes, consider the 1000 QPS API rate that is a *medium* hosted cluster load, and a 2000 QPS API that is a *heavy* hosted cluster load.



NOTE

This test provides an estimation factor to increase the compute resource utilization based on the expected API load. Exact utilization rates can vary based on the type and pace of the cluster workload.

The following example shows hosted control plane resource scaling for the workload and API rate definitions:

Table 3.11. API rate table

QPS (API rate)	vCPU usage	Memory usage (GiB)
Low load (Less than 50 QPS)	2.9	11.1

QPS (API rate)	vCPU usage	Memory usage (GiB)
Medium load (1000 QPS)	11.9	13.6
High load (2000 QPS)	20.9	16.1

The hosted control plane sizing is about control-plane load and workloads that cause heavy API activity, etcd activity, or both. Hosted pod workloads that focus on data-plane loads, such as running a database, might not result in high API rates.

3.2.4. Sizing calculation example

This example provides sizing guidance for the following scenario:

- Three bare-metal workers that are labeled as **hypershift.openshift.io/control-plane** nodes
- **maxPods** value set to 500
- The expected API rate is medium or about 1000, according to the load-based limits

Table 3.12. Limit inputs

Limit description	Server 1	Server 2
Number of vCPUs on worker node	64	128
Memory on worker node (GiB)	128	256
Maximum pods per worker	500	500
Number of workers used to host control planes	3	3
Maximum QPS target rate (API requests per second)	1000	1000

Table 3.13. Sizing calculation example

Calculated values based on worker node size and API rate	Server 1	Server 2	Calculation notes
Maximum hosted control planes per worker based on vCPU requests	12.8	25.6	Number of worker vCPUs ÷ 5 total vCPU requests per hosted control plane

Maximum hosted control planes per worker based on vCPU usage	5.4	10.7	Number of vCPUs ÷ (2.9 measured idle vCPU usage + (QPS target rate ÷ 1000) × 9.0 measured vCPU usage per 1000 QPS increase)
Maximum hosted control planes per worker based on memory requests	7.1	14.2	Worker memory GiB ÷ 18 GiB total memory request per hosted control plane
Maximum hosted control planes per worker based on memory usage	9.4	18.8	Worker memory GiB ÷ (11.1 measured idle memory usage + (QPS target rate ÷ 1000) × 2.5 measured memory usage per 1000 QPS increase)
Maximum hosted control planes per worker based on per node pod limit	6.7	6.7	500 maxPods ÷ 75 pods per hosted control plane
Minimum of previously mentioned maximums	5.4	6.7	
	vCPU limiting factor	maxPods limiting factor	
Maximum number of hosted control planes within a management cluster	16	20	Minimum of previously mentioned maximums × 3 control-plane workers

Table 3.14. Hosted control planes capacity metrics

Name	Description
mce_hs_addon_request_based_hcp_capacity_gauge	Estimated maximum number of hosted control planes the cluster can host based on a highly available hosted control planes resource request.
mce_hs_addon_low_qps_based_hcp_capacity_gauge	Estimated maximum number of hosted control planes the cluster can host if all hosted control planes make around 50 QPS to the clusters Kube API server.
mce_hs_addon_medium_qps_based_hcp_capacity_gauge	Estimated maximum number of hosted control planes the cluster can host if all hosted control planes make around 1000 QPS to the clusters Kube API server.

mce_hs_addon_high_qps_based_hcp_capacity_gauge	Estimated maximum number of hosted control planes the cluster can host if all hosted control planes make around 2000 QPS to the clusters Kube API server.
mce_hs_addon_average_qps_based_hcp_capacity_gauge	Estimated maximum number of hosted control planes the cluster can host based on the existing average QPS of hosted control planes. If you do not have an active hosted control planes, you can expect low QPS.

3.2.5. Shared infrastructure between hosted and standalone control planes

As a service provider, you can more effectively use your resources by sharing infrastructure between a standalone OpenShift Container Platform control plane and hosted control planes. A 3-node OpenShift Container Platform cluster can be a management cluster for a hosted cluster.

Sharing infrastructure can be beneficial in constrained environments, such as in small-scale deployments where you need resource efficiency.

Before you share infrastructure, ensure that your infrastructure has enough resources to support hosted control planes. On the OpenShift Container Platform management cluster, nothing else can be deployed except hosted control planes. Ensure that the management cluster has enough CPU, memory, storage, and network resources to handle the combined load of the hosted clusters. Workload must not be demanding, and it must fall within a low queries-per-second (QPS) profile. For more information about resources and workload, see "Sizing guidance for hosted control planes".

Additional resources

- [Sizing guidance for hosted control planes](#)

3.3. OVERRIDING RESOURCE UTILIZATION MEASUREMENTS

The set of baseline measurements for resource utilization can vary in each hosted cluster.

3.3.1. Overriding resource utilization measurements for a hosted cluster

You can override resource utilization measurements based on the type and pace of your cluster workload.

Procedure

1. Create the **ConfigMap** resource by running the following command:

```
$ oc create -f <your-config-map-file.yaml>
```

Replace `<your-config-map-file.yaml>` with the name of your YAML file that contains your **hcp-sizing-baseline** config map.

- Create the **hcp-sizing-baseline** config map in the **local-cluster** namespace to specify the measurements you want to override. Your config map might resemble the following YAML file:

```
kind: ConfigMap
apiVersion: v1
metadata:
  name: hcp-sizing-baseline
  namespace: local-cluster
data:
  incrementalCPUUsagePer1KQPS: "9.0"
  memoryRequestPerHCP: "18"
  minimumQPSPerHCP: "50.0"
```

- Delete the **hypershift-addon-agent** deployment to restart the **hypershift-addon-agent** pod by running the following command:

```
$ oc delete deployment hypershift-addon-agent \
-n open-cluster-management-agent-addon
```

Verification

- Observe the **hypershift-addon-agent** pod logs. Verify that the overridden measurements are updated in the config map by running the following command:

```
$ oc logs hypershift-addon-agent -n open-cluster-management-agent-addon
```

Your logs might resemble the following output:

Example output

```
2024-01-05T19:41:05.392Z INFO agent.agent-reconciler agent/agent.go:793 setting
cpuRequestPerHCP to 5
2024-01-05T19:41:05.392Z INFO agent.agent-reconciler agent/agent.go:802 setting
memoryRequestPerHCP to 18
2024-01-05T19:53:54.070Z INFO agent.agent-reconciler
agent/hcp_capacity_calculation.go:141 The worker nodes have 12.000000 vCPUs
2024-01-05T19:53:54.070Z INFO agent.agent-reconciler
agent/hcp_capacity_calculation.go:142 The worker nodes have 49.173369 GB memory
```

If the overridden measurements are not updated properly in the **hcp-sizing-baseline** config map, you might see the following error message in the **hypershift-addon-agent** pod logs:

Example error

```
2024-01-05T19:53:54.052Z ERROR agent.agent-reconciler agent/agent.go:788 failed to get
configmap from the hub. Setting the HCP sizing baseline with default values. {"error": "
"configmaps \\"hcp-sizing-baseline\\" not found"}
```

3.3.2. Disabling the metric service monitoring

After you enable the **hypershift-addon** managed cluster add-on, metric service monitoring is configured by default so that OpenShift Container Platform monitoring can gather metrics from **hypershift-addon**.

Procedure

You can disable metric service monitoring by completing the following steps:

1. Log in to your hub cluster by running the following command:

```
$ oc login
```

2. Edit the **hypershift-addon-deploy-config** add-on deployment configuration specification by running the following command:

```
$ oc edit addondeploymentconfig hypershift-addon-deploy-config \
-n multicluster-engine
```

3. Add the **disableMetrics=true** customized variable to the specification, as shown in the following example:

```
apiVersion: addon.open-cluster-management.io/v1alpha1
kind: AddOnDeploymentConfig
metadata:
  name: hypershift-addon-deploy-config
  namespace: multicluster-engine
spec:
  customizedVariables:
    - name: hcMaxNumber
      value: "80"
    - name: hcThresholdNumber
      value: "60"
    - name: disableMetrics ①
      value: "true"
```

- 1 The **disableMetrics=true** customized variable disables metric service monitoring for both new and existing **hypershift-addon** managed cluster add-ons.

4. Apply the changes to the configuration specification by running the following command:

```
$ oc apply -f <filename>.yaml
```

3.4. INSTALLING THE HOSTED CONTROL PLANES COMMAND-LINE INTERFACE

The hosted control planes command-line interface, **hcp**, is a tool that you can use to get started with hosted control planes. For Day 2 operations, such as management and configuration, use GitOps or your own automation tool.

3.4.1. Installing the hosted control planes command-line interface from the terminal

You can install the hosted control planes command-line interface (CLI), **hcp**, from the terminal.

Prerequisites

- On an OpenShift Container Platform cluster, you have installed multicluster engine for Kubernetes Operator 2.5 or later. The multicluster engine Operator is automatically installed when you install Red Hat Advanced Cluster Management. You can also install multicluster engine Operator without Red Hat Advanced Management as an Operator from OpenShift Container Platform OperatorHub.

Procedure

- Get the URL to download the **hcp** binary by running the following command:

```
$ oc get ConsoleCLIDownload hcp-cli-download -o json | jq -r ".spec"
```

- Download the **hcp** binary by running the following command:

```
$ wget <hcp_cli_download_url> ①
```

① Replace **hcp_cli_download_url** with the URL that you obtained from the previous step.

- Unpack the downloaded archive by running the following command:

```
$ tar xvzf hcp.tar.gz
```

- Make the **hcp** binary file executable by running the following command:

```
$ chmod +x hcp
```

- Move the **hcp** binary file to a directory in your path by running the following command:

```
$ sudo mv hcp /usr/local/bin/.
```



NOTE

If you download the CLI on a Mac computer, you might see a warning about the **hcp** binary file. You need to adjust your security settings to allow the binary file to be run.

Verification

- Verify that you see the list of available parameters by running the following command:

```
$ hcp create cluster <platform> --help ①
```

① You can use the **hcp create cluster** command to create and manage hosted clusters. The supported platforms are **aws**, **agent**, and **kubevirt**.

3.4.2. Installing the hosted control planes command-line interface by using the web console

You can install the hosted control planes command-line interface (CLI), **hcp**, by using the OpenShift Container Platform web console.

Prerequisites

- On an OpenShift Container Platform cluster, you have installed multicluster engine for Kubernetes Operator 2.5 or later. The multicluster engine Operator is automatically installed when you install Red Hat Advanced Cluster Management. You can also install multicluster engine Operator without Red Hat Advanced Management as an Operator from OpenShift Container Platform OperatorHub.

Procedure

1. From the OpenShift Container Platform web console, click the **Help icon → Command Line Tools**.

2. Click **Download hcp CLI** for your platform.

3. Unpack the downloaded archive by running the following command:

```
$ tar xvzf hcp.tar.gz
```

4. Run the following command to make the binary file executable:

```
$ chmod +x hcp
```

5. Run the following command to move the binary file to a directory in your path:

```
$ sudo mv hcp /usr/local/bin/.
```



NOTE

If you download the CLI on a Mac computer, you might see a warning about the **hcp** binary file. You need to adjust your security settings to allow the binary file to be run.

Verification

- Verify that you see the list of available parameters by running the following command:

```
$ hcp create cluster <platform> --help ①
```

- ① You can use the **hcp create cluster** command to create and manage hosted clusters. The supported platforms are **aws**, **agent**, and **kubevirt**.

3.4.3. Installing the hosted control planes command-line interface by using the content gateway

You can install the hosted control planes command-line interface (CLI), **hcp**, by using the content gateway.

Prerequisites

- On an OpenShift Container Platform cluster, you have installed multicluster engine for Kubernetes Operator 2.5 or later. The multicluster engine Operator is automatically installed when you install Red Hat Advanced Cluster Management. You can also install multicluster

engine Operator without Red Hat Advanced Management as an Operator from OpenShift Container Platform OperatorHub.

Procedure

1. Navigate to the [content gateway](#) and download the **hcp** binary.
2. Unpack the downloaded archive by running the following command:

```
$ tar xvzf hcp.tar.gz
```

3. Make the **hcp** binary file executable by running the following command:

```
$ chmod +x hcp
```

4. Move the **hcp** binary file to a directory in your path by running the following command:

```
$ sudo mv hcp /usr/local/bin/.
```



NOTE

If you download the CLI on a Mac computer, you might see a warning about the **hcp** binary file. You need to adjust your security settings to allow the binary file to be run.

Verification

- Verify that you see the list of available parameters by running the following command:

```
$ hcp create cluster <platform> --help ①
```

- ① You can use the **hcp create cluster** command to create and manage hosted clusters. The supported platforms are **aws**, **agent**, and **kubevirt**.

3.5. DISTRIBUTING HOSTED CLUSTER WORKLOADS

Before you get started with hosted control planes for OpenShift Container Platform, you must properly label nodes so that the pods of hosted clusters can be scheduled into infrastructure nodes. Node labeling is also important for the following reasons:

- To ensure high availability and proper workload deployment. For example, to avoid having the control plane workload count toward your OpenShift Container Platform subscription, you can set the **node-role.kubernetes.io/infra** label.
- To ensure that control plane workloads are separate from other workloads in the management cluster.
- To ensure that control plane workloads are configured at the correct multi-tenancy distribution level for your deployment. The distribution levels are as follows:
 - Everything shared: Control planes for hosted clusters can run on any node that is designated for control planes.
 - Request serving isolation: Serving pods are requested in their own dedicated nodes.

- Nothing shared: Every control plane has its own dedicated nodes.

For more information about dedicating a node to a single hosted cluster, see "Labeling management cluster nodes".



IMPORTANT

Do not use the management cluster for your workload. Workloads must not run on nodes where control planes run.

3.5.1. Labeling management cluster nodes

Proper node labeling is a prerequisite to deploying hosted control planes.

As a management cluster administrator, you use the following labels and taints in management cluster nodes to schedule a control plane workload:

- hypershift.openshift.io/control-plane: true**: Use this label and taint to dedicate a node to running hosted control plane workloads. By setting a value of **true**, you avoid sharing the control plane nodes with other components, for example, the infrastructure components of the management cluster or any other mistakenly deployed workload.
- hypershift.openshift.io/cluster: \${HostedControlPlane Namespace}**: Use this label and taint when you want to dedicate a node to a single hosted cluster.

Apply the following labels on the nodes that host control-plane pods:

- node-role.kubernetes.io/infra**: Use this label to avoid having the control-plane workload count toward your subscription.
- topology.kubernetes.io/zone**: Use this label on the management cluster nodes to deploy highly available clusters across failure domains. The zone might be a location, rack name, or the hostname of the node where the zone is set. For example, a management cluster has the following nodes: **worker-1a**, **worker-1b**, **worker-2a**, and **worker-2b**. The **worker-1a** and **worker-1b** nodes are in **rack1**, and the **worker-2a** and **worker-2b** nodes are in **rack2**. To use each rack as an availability zone, enter the following commands:

```
$ oc label node/worker-1a node/worker-1b topology.kubernetes.io/zone=rack1
```

```
$ oc label node/worker-2a node/worker-2b topology.kubernetes.io/zone=rack2
```

Pods for a hosted cluster have tolerations, and the scheduler uses affinity rules to schedule them. Pods tolerate taints for **control-plane** and the **cluster** for the pods. The scheduler prioritizes the scheduling of pods into nodes that are labeled with **hypershift.openshift.io/control-plane** and **hypershift.openshift.io/cluster: \${HostedControlPlane Namespace}**.

For the **ControllerAvailabilityPolicy** option, use **HighlyAvailable**, which is the default value that the hosted control planes command-line interface, **hcp**, deploys. When you use that option, you can schedule pods for each deployment within a hosted cluster across different failure domains by setting **topology.kubernetes.io/zone** as the topology key. Scheduling pods for a deployment within a hosted cluster across different failure domains is available only for highly available control planes.

Procedure

To enable a hosted cluster to require its pods to be scheduled into infrastructure nodes, set **HostedCluster.spec.nodeSelector**, as shown in the following example:

```
spec:
  nodeSelector:
    node-role.kubernetes.io/infra: ""
```

This way, hosted control planes for each hosted cluster are eligible infrastructure node workloads, and you do not need to entitle the underlying OpenShift Container Platform nodes.

3.5.2. Priority classes

Four built-in priority classes influence the priority and preemption of the hosted cluster pods. You can create the pods in the management cluster in the following order from highest to lowest:

- **hypershift-operator**: HyperShift Operator pods.
- **hypershift-etcd**: Pods for etcd.
- **hypershift-api-critical**: Pods that are required for API calls and resource admission to succeed. These pods include pods such as **kube-apiserver**, aggregated API servers, and web hooks.
- **hypershift-control-plane**: Pods in the control plane that are not API-critical but still need elevated priority, such as the cluster version Operator.

3.5.3. Custom taints and tolerations

By default, pods for a hosted cluster tolerate the **control-plane** and **cluster** taints. However, you can also use custom taints on nodes so that hosted clusters can tolerate those taints on a per-hosted-cluster basis by setting **HostedCluster.spec.tolerations**.



IMPORTANT

Passing tolerations for a hosted cluster is a Technology Preview feature only. Technology Preview features are not supported with Red Hat production service level agreements (SLAs) and might not be functionally complete. Red Hat does not recommend using them in production. These features provide early access to upcoming product features, enabling customers to test functionality and provide feedback during the development process.

For more information about the support scope of Red Hat Technology Preview features, see [Technology Preview Features Support Scope](#).

Example configuration

```
spec:
  tolerations:
    - effect: NoSchedule
      key: kubernetes.io/custom
      operator: Exists
```

You can also set tolerations on the hosted cluster while you create a cluster by using the **--tolerations** hcp CLI argument.

Example CLI argument

```
--toleration="key=kubernetes.io/custom,operator=Exists,effect=NoSchedule"
```

For fine granular control of hosted cluster pod placement on a per-hosted-cluster basis, use custom tolerations with **nodeSelectors**. You can co-locate groups of hosted clusters and isolate them from other hosted clusters. You can also place hosted clusters in infra and control plane nodes.

Tolerations on the hosted cluster spread only to the pods of the control plane. To configure other pods that run on the management cluster and infrastructure-related pods, such as the pods to run virtual machines, you need to use a different process.

3.5.4. Control plane isolation

You can configure hosted control planes to isolate network traffic or control plane pods.

3.5.4.1. Network policy isolation

Each hosted control plane is assigned to run in a dedicated Kubernetes namespace. By default, the Kubernetes namespace denies all network traffic.

The following network traffic is allowed through the network policy that is enforced by the Kubernetes Container Network Interface (CNI):

- Ingress pod-to-pod communication in the same namespace (intra-tenant)
- Ingress on port 6443 to the hosted **kube-apiserver** pod for the tenant
- Metric scraping from the management cluster Kubernetes namespace with the **network.openshift.io/policy-group: monitoring** label is allowed for monitoring

3.5.4.2. Control plane pod isolation

In addition to network policies, each hosted control plane pod is run with the **restricted** security context constraint. This policy denies access to all host features and requires pods to be run with a UID and with SELinux context that is allocated uniquely to each namespace that hosts a customer control plane.

The policy ensures the following constraints:

- Pods cannot run as privileged.
- Pods cannot mount host directory volumes.
- Pods must run as a user in a pre-allocated range of UIDs.
- Pods must run with a pre-allocated MCS label.
- Pods cannot access the host network namespace.
- Pods cannot expose host network ports.
- Pods cannot access the host PID namespace.
- By default, pods drop the following Linux capabilities: **KILL**, **MKNOD**, **SETUID**, and **SETGID**.

The management components, such as **kubelet** and **crio**, on each management cluster worker node are protected by an SELinux label that is not accessible to the SELinux context for pods that support hosted control planes.

The following SELinux labels are used for key processes and sockets:

- **kubelet**: **system_u:system_r:unconfined_service_t:s0**
- **crio**: **system_u:system_r:container_runtime_t:s0**
- **crio.sock**: **system_u:object_r:container_var_run_t:s0**
- <example user container processes> **system_u:system_r:container_t:s0:c14,c24**

3.6. ENABLING OR DISABLING THE HOSTED CONTROL PLANES FEATURE

The hosted control planes feature, as well as the **hypershift-addon** managed cluster add-on, are enabled by default. If you want to disable the feature, or if you disabled it and want to manually enable it, see the following procedures.

3.6.1. Manually enabling the hosted control planes feature

If you need to manually enable hosted control planes, complete the following steps.

Procedure

1. Run the following command to enable the feature:

```
$ oc patch mce multiclusterengine --type=merge -p \
'{"spec": {"overrides": {"components": [{"name": "hypershift", "enabled": true}]} }}'
```

- 1** The default **MultiClusterEngine** resource instance name is **multiclusterengine**, but you can get the **MultiClusterEngine** name from your cluster by running the following command: **\$ oc get mce**.

2. Run the following command to verify that the **hypershift** and **hypershift-local-hosting** features are enabled in the **MultiClusterEngine** custom resource:

```
$ oc get mce multiclusterengine -o yaml
```

- 1** The default **MultiClusterEngine** resource instance name is **multiclusterengine**, but you can get the **MultiClusterEngine** name from your cluster by running the following command: **\$ oc get mce**.

Example output

```
apiVersion: multicluster.openshift.io/v1
kind: MultiClusterEngine
metadata:
  name: multiclusterengine
spec:
```

```

overrides:
  components:
    - name: hypershift
      enabled: true
    - name: hypershift-local-hosting
      enabled: true

```

3.6.1.1. Manually enabling the hypershift-addon managed cluster add-on for local-cluster

Enabling the hosted control planes feature automatically enables the **hypershift-addon** managed cluster add-on. If you need to enable the **hypershift-addon** managed cluster add-on manually, complete the following steps to use the **hypershift-addon** to install the HyperShift Operator on **local-cluster**.

Procedure

1. Create the **ManagedClusterAddon** add-on named **hypershift-addon** by creating a file that resembles the following example:

```

apiVersion: addon.open-cluster-management.io/v1alpha1
kind: ManagedClusterAddOn
metadata:
  name: hypershift-addon
  namespace: local-cluster
spec:
  installNamespace: open-cluster-management-agent-addon

```

2. Apply the file by running the following command:

```
$ oc apply -f <filename>
```

Replace **filename** with the name of the file that you created.

3. Confirm that the **hypershift-addon** managed cluster add-on is installed by running the following command:

```
$ oc get managedclusteraddons -n local-cluster hypershift-addon
```

If the add-on is installed, the output resembles the following example:

NAME	AVAILABLE	DEGRADED	PROGRESSING
hypershift-addon	True		

Your **hypershift-addon** managed cluster add-on is installed and the hosting cluster is available to create and manage hosted clusters.

3.6.2. Disabling the hosted control planes feature

You can uninstall the HyperShift Operator and disable the hosted control planes feature. When you disable the hosted control planes feature, you must destroy the hosted cluster and the managed cluster resource on multicluster engine Operator, as described in the *Managing hosted clusters* topics.

3.6.2.1. Uninstalling the HyperShift Operator

To uninstall the HyperShift Operator and disable the **hypershift-addon** from the **local-cluster**, complete the following steps:

Procedure

- Run the following command to ensure that there is no hosted cluster running:

```
$ oc get hostedcluster -A
```

IMPORTANT

If a hosted cluster is running, the HyperShift Operator does not uninstall, even if the **hypershift-addon** is disabled.

- Disable the **hypershift-addon** by running the following command:

```
$ oc patch mce multiclusterengine --type=merge -p \①
'{"spec": {"overrides": {"components": [{"name": "hypershift-local-hosting", "enabled": false}]}}}'
```

- The default **MultiClusterEngine** resource instance name is **multiclusterengine**, but you can get the **MultiClusterEngine** name from your cluster by running the following command: **\$ oc get mce**.

NOTE

You can also disable the **hypershift-addon** for the **local-cluster** from the multicloud engine Operator console after disabling the **hypershift-addon**.

3.6.2.2. Disabling the hosted control planes feature

To disable the hosted control planes feature, complete the following steps.

Prerequisites

- You uninstalled the HyperShift Operator. For more information, see "Uninstalling the HyperShift Operator".

Procedure

- Run the following command to disable the hosted control planes feature:

```
$ oc patch mce multiclusterengine --type=merge -p \①
'{"spec": {"overrides": {"components": [{"name": "hypershift", "enabled": false}]}}}'
```

- The default **MultiClusterEngine** resource instance name is **multiclusterengine**, but you can get the **MultiClusterEngine** name from your cluster by running the following command: **\$ oc get mce**.

- You can verify that the **hypershift** and **hypershift-local-hosting** features are disabled in the **MultiClusterEngine** custom resource by running the following command:

```
$ oc get mce multiclusterengine -o yaml ①
```

- 1 The default **MultiClusterEngine** resource instance name is **multiclusterengine**, but you can get the **MultiClusterEngine** name from your cluster by running the following command: `$ oc get mce`.

See the following example where **hypershift** and **hypershift-local-hosting** have their **enabled:** flags set to **false**:

```
apiVersion: multicluster.openshift.io/v1
kind: MultiClusterEngine
metadata:
  name: multiclusterengine
spec:
  overrides:
    components:
      - name: hypershift
        enabled: false
      - name: hypershift-local-hosting
        enabled: false
```

CHAPTER 4. DEPLOYING HOSTED CONTROL PLANES

4.1. DEPLOYING HOSTED CONTROL PLANES ON AWS

A *hosted cluster* is an OpenShift Container Platform cluster with its API endpoint and control plane that are hosted on the management cluster. The hosted cluster includes the control plane and its corresponding data plane. To configure hosted control planes on premises, you must install multicluster engine for Kubernetes Operator in a management cluster. By deploying the HyperShift Operator on an existing managed cluster by using the **hypershift-addon** managed cluster add-on, you can enable that cluster as a management cluster and start to create the hosted cluster. The **hypershift-addon** managed cluster add-on is enabled by default for the **local-cluster** managed cluster.

You can use the multicluster engine Operator console or the hosted control plane command-line interface (CLI), **hcp**, to create a hosted cluster. The hosted cluster is automatically imported as a managed cluster. However, you can [disable this automatic import feature into multicluster engine Operator](#).

4.1.1. Preparing to deploy hosted control planes on AWS

As you prepare to deploy hosted control planes on Amazon Web Services (AWS), consider the following information:

- Each hosted cluster must have a cluster-wide unique name. A hosted cluster name cannot be the same as any existing managed cluster in order for multicluster engine Operator to manage it.
- Do not use **clusters** as a hosted cluster name.
- Run the management cluster and workers on the same platform for hosted control planes.
- A hosted cluster cannot be created in the namespace of a multicluster engine Operator managed cluster.

4.1.1.1. Prerequisites to configure a management cluster

You must have the following prerequisites to configure the management cluster:

- You have installed the multicluster engine for Kubernetes Operator 2.5 and later on an OpenShift Container Platform cluster. The multicluster engine Operator is automatically installed when you install Red Hat Advanced Cluster Management (RHACM). The multicluster engine Operator can also be installed without RHACM as an Operator from the OpenShift Container Platform OperatorHub.
- You have at least one managed OpenShift Container Platform cluster for the multicluster engine Operator. The **local-cluster** is automatically imported in the multicluster engine Operator version 2.5 and later. You can check the status of your hub cluster by running the following command:

```
$ oc get managedclusters local-cluster
```

- You have installed the [aws command-line interface \(CLI\)](#).
- You have installed the hosted control plane CLI, **hcp**.

Additional resources

- [Configuring Ansible Automation Platform jobs to run on hosted clusters](#)
- [Advanced configuration](#)
- [Enabling the central infrastructure management service](#)
- [Manually enabling the hosted control planes feature](#)
- [Disabling the hosted control planes feature](#)
- [Deploying the SR-IOV Operator for hosted control planes](#)

4.1.2. Creating the Amazon Web Services S3 bucket and S3 OIDC secret

Before you can create and manage hosted clusters on Amazon Web Services (AWS), you must create the S3 bucket and S3 OIDC secret.

Procedure

- 1 Create an S3 bucket that has public access to host OIDC discovery documents for your clusters by running the following commands:

```
$ aws s3api create-bucket --bucket <bucket_name> \ ①  
  --create-bucket-configuration LocationConstraint=<region> \ ②  
  --region <region> ③
```

① Replace **<bucket_name>** with the name of the S3 bucket you are creating.

②③ To create the bucket in a region other than the **us-east-1** region, include this line and replace **<region>** with the region you want to use. To create a bucket in the **us-east-1** region, omit this line.

```
$ aws s3api delete-public-access-block --bucket <bucket_name> ①
```

① Replace **<bucket_name>** with the name of the S3 bucket you are creating.

```
$ echo '{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Principal": "*",  
      "Action": "s3:GetObject",  
      "Resource": "arn:aws:s3:::<bucket_name>/*" ①  
    }  
  ]  
' | envsubst > policy.json
```

① Replace **<bucket_name>** with the name of the S3 bucket you are creating.

```
$ aws s3api put-bucket-policy --bucket <bucket_name> \
--policy file://policy.json
```

- 1 Replace <bucket_name> with the name of the S3 bucket you are creating.



NOTE

If you are using a Mac computer, you must export the bucket name in order for the policy to work.

- 2 Create an OIDC S3 secret named **hypershift-operator-oidc-provider-s3-credentials** for the HyperShift Operator.
- 3 Save the secret in the **local-cluster** namespace.
- 4 See the following table to verify that the secret contains the following fields:

Table 4.1. Required fields for the AWS secret

Field name	Description
bucket	Contains an S3 bucket with public access to host OIDC discovery documents for your hosted clusters.
credentials	A reference to a file that contains the credentials of the default profile that can access the bucket. By default, HyperShift only uses the default profile to operate the bucket .
region	Specifies the region of the S3 bucket.

- 5 To create an AWS secret, run the following command:

```
$ oc create secret generic <secret_name> \
--from-file=credentials=<path>/aws/credentials \
--from-literal=bucket=<s3_bucket> \
--from-literal=region=<region> \
-n local-cluster
```



NOTE

Disaster recovery backup for the secret is not automatically enabled. To add the label that enables the **hypershift-operator-oidc-provider-s3-credentials** secret to be backed up for disaster recovery, run the following command:

```
$ oc label secret hypershift-operator-oidc-provider-s3-credentials \
-n local-cluster cluster.open-cluster-management.io/backup=true
```

4.1.3. Creating a routable public zone for hosted clusters

To access applications in your hosted clusters, you must configure the routable public zone. If the public zone exists, skip this step. Otherwise, the public zone affects the existing functions.

Procedure

- To create a routable public zone for DNS records, enter the following command:

```
$ aws route53 create-hosted-zone \  
  --name <basedomain> \ ①  
  --caller-reference $(whoami)-$(date --rfc-3339=date)
```

- 1 Replace **<basedomain>** with your base domain, for example, **www.example.com**.

4.1.4. Creating an AWS IAM role and STS credentials

Before creating a hosted cluster on Amazon Web Services (AWS), you must create an AWS IAM role and STS credentials.

Procedure

- 1 Get the Amazon Resource Name (ARN) of your user by running the following command:

```
$ aws sts get-caller-identity --query "Arn" --output text
```

Example output

```
arn:aws:iam::1234567890:user/<aws_username>
```

Use this output as the value for **<arn>** in the next step.

- 2 Create a JSON file that contains the trust relationship configuration for your role. See the following example:

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Principal": {  
        "AWS": "<arn>" ①  
      },  
      "Action": "sts:AssumeRole"  
    }  
  ]  
}
```

- 1 Replace **<arn>** with the ARN of your user that you noted in the previous step.

- 3 Create the Identity and Access Management (IAM) role by running the following command:

```
$ aws iam create-role \  
 \
```

```
--role-name <name> \1
--assume-role-policy-document file://<file_name>.json \2
--query "Role.Arn"
```

- 1 Replace <name> with the role name, for example, **hcp-cli-role**.
- 2 Replace <file_name> with the name of the JSON file you created in the previous step.

Example output

```
arn:aws:iam::820196288204:role/myrole
```

4. Create a JSON file named **policy.json** that contains the following permission policies for your role:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "EC2",
      "Effect": "Allow",
      "Action": [
        "ec2:CreateDhcpOptions",
        "ec2:DeleteSubnet",
        "ec2:ReplaceRouteTableAssociation",
        "ec2:DescribeAddresses",
        "ec2:DescribeInstances",
        "ec2:DeleteVpcEndpoints",
        "ec2:CreateNatGateway",
        "ec2:CreateVpc",
        "ec2:DescribeDhcpOptions",
        "ec2:AttachInternetGateway",
        "ec2:DeleteVpcEndpointServiceConfigurations",
        "ec2:DeleteRouteTable",
        "ec2:AssociateRouteTable",
        "ec2:DescribeInternetGateways",
        "ec2:DescribeAvailabilityZones",
        "ec2:CreateRoute",
        "ec2:CreateInternetGateway",
        "ec2:RevokeSecurityGroupEgress",
        "ec2:ModifyVpcAttribute",
        "ec2:DeleteInternetGateway",
        "ec2:DescribeVpcEndpointConnections",
        "ec2:RejectVpcEndpointConnections",
        "ec2:DescribeRouteTables",
        "ec2:ReleaseAddress",
        "ec2:AssociateDhcpOptions",
        "ec2:TerminateInstances",
        "ec2:CreateTags",
        "ec2:DeleteRoute",
        "ec2:CreateRouteTable",
        "ec2:DetachInternetGateway",
        "ec2:DescribeVpcEndpointServiceConfigurations",
        "ec2:DescribeNatGateways",
```

```
        "ec2:DisassociateRouteTable",
        "ec2:AllocateAddress",
        "ec2:DescribeSecurityGroups",
        "ec2:RevokeSecurityGroupIngress",
        "ec2>CreateVpcEndpoint",
        "ec2:DescribeVpcs",
        "ec2:DeleteSecurityGroup",
        "ec2:DeleteDhcpOptions",
        "ec2:DeleteNatGateway",
        "ec2:DescribeVpcEndpoints",
        "ec2:DeleteVpc",
        "ec2:CreateSubnet",
        "ec2:DescribeSubnets"
    ],
    "Resource": "*"
},
{
    "Sid": "ELB",
    "Effect": "Allow",
    "Action": [
        "elasticloadbalancing>DeleteLoadBalancer",
        "elasticloadbalancing>DescribeLoadBalancers",
        "elasticloadbalancing>DescribeTargetGroups",
        "elasticloadbalancing>DeleteTargetGroup"
    ],
    "Resource": "*"
},
{
    "Sid": "IAMPassRole",
    "Effect": "Allow",
    "Action": "iam:PassRole",
    "Resource": "arn:*.iam::*:role/*-worker-role",
    "Condition": {
        "ForAnyValue:StringEqualsIfExists": {
            "iam:PassedToService": "ec2.amazonaws.com"
        }
    }
},
{
    "Sid": "IAM",
    "Effect": "Allow",
    "Action": [
        "iam>CreateInstanceProfile",
        "iam>DeleteInstanceProfile",
        "iam:GetRole",
        "iam>UpdateAssumeRolePolicy",
        "iamGetInstanceProfile",
        "iam:TagRole",
        "iam>RemoveRoleFromInstanceProfile",
        "iam>CreateRole",
        "iam>DeleteRole",
        "iam:PutRolePolicy",
        "iam>AddRoleToInstanceProfile",
        "iam>CreateOpenIDConnectProvider",
        "iam>ListOpenIDConnectProviders",
        "iam>DeleteRolePolicy",
        "iam:ListOpenIDConnectProviderCertificates"
    ]
}
```

```

    "iam:UpdateRole",
    "iam>DeleteOpenIDConnectProvider",
    "iam:GetRolePolicy"
],
"Resource": "*"
},
{
"Sid": "Route53",
"Effect": "Allow",
>Action": [
    "route53>ListHostedZonesByVPC",
    "route53>CreateHostedZone",
    "route53>ListHostedZones",
    "route53>ChangeResourceRecordSets",
    "route53>ListResourceRecordSets",
    "route53>DeleteHostedZone",
    "route53 AssociateVPCWithHostedZone",
    "route53>ListHostedZonesByName"
],
"Resource": "*"
},
{
"Sid": "S3",
"Effect": "Allow",
>Action": [
    "s3>ListAllMyBuckets",
    "s3>ListBucket",
    "s3>DeleteObject",
    "s3>DeleteBucket"
],
"Resource": "*"
}
]
}
}

```

5. Attach the **policy.json** file to your role by running the following command:

```
$ aws iam put-role-policy \
--role-name <role_name> \①
--policy-name <policy_name> \②
--policy-document file://policy.json \③
```

- ① Replace **<role_name>** with the name of your role.
- ② Replace **<policy_name>** with your policy name.
- ③ The **policy.json** file contains the permission policies for your role.

6. Retrieve STS credentials in a JSON file named **sts-creds.json** by running the following command:

```
$ aws sts get-session-token --output json > sts-creds.json
```

Example **sts-creds.json** file

```
{
  "Credentials": {
    "AccessKeyId": "<access_key_id>",
    "SecretAccessKey": "<secret_access_key>",
    "SessionToken": "<session_token>",
    "Expiration": "<time_stamp>"
  }
}
```

4.1.5. Enabling AWS PrivateLink for hosted control planes

To provision hosted control planes on the Amazon Web Services (AWS) with PrivateLink, enable AWS PrivateLink for hosted control planes.

Procedure

1. Create an AWS credential secret for the HyperShift Operator and name it **hypershift-operator-private-link-credentials**. The secret must reside in the managed cluster namespace that is the namespace of the managed cluster being used as the management cluster. If you used **local-cluster**, create the secret in the **local-cluster** namespace.
2. See the following table to confirm that the secret contains the required fields:

Table 4.2. Required fields for the AWS secret

Field name	Description	Optional or required
region	Region for use with Private Link	Required
aws-access-key-id	The credential access key id.	Required
aws-secret-access-key	The credential access key secret.	Required

To create an AWS secret, run the following command:

```
$ oc create secret generic <secret_name> \
--from-literal=aws-access-key-id=<aws_access_key_id> \
--from-literal=aws-secret-access-key=<aws_secret_access_key> \
--from-literal=region=<region> -n local-cluster
```



NOTE

Disaster recovery backup for the secret is not automatically enabled. Run the following command to add the label that enables the **hypershift-operator-private-link-credentials** secret to be backed up for disaster recovery:

```
$ oc label secret hypershift-operator-private-link-credentials \
-n local-cluster \
cluster.open-cluster-management.io/backup=""
```

4.1.6. Enabling external DNS for hosted control planes on AWS

The control plane and the data plane are separate in hosted control planes. You can configure DNS in two independent areas:

- Ingress for workloads within the hosted cluster, such as the following domain: ***.apps.service-consumer-domain.com**.
- Ingress for service endpoints within the management cluster, such as API or OAuth endpoints through the service provider domain: ***.service-provider-domain.com**.

The input for **hostedCluster.spec.dns** manages the ingress for workloads within the hosted cluster. The input for **hostedCluster.spec.services.servicePublishingStrategy.route.hostname** manages the ingress for service endpoints within the management cluster.

External DNS creates name records for hosted cluster **Services** that specify a publishing type of **LoadBalancer** or **Route** and provide a hostname for that publishing type. For hosted clusters with **Private** or **PublicAndPrivate** endpoint access types, only the **APIServer** and **OAuth** services support hostnames. For **Private** hosted clusters, the DNS record resolves to a private IP address of a Virtual Private Cloud (VPC) endpoint in your VPC.

A hosted control plane exposes the following services:

- **APIServer**
- **OIDC**

You can expose these services by using the **servicePublishingStrategy** field in the **HostedCluster** specification. By default, for the **LoadBalancer** and **Route** types of **servicePublishingStrategy**, you can publish the service in one of the following ways:

- By using the hostname of the load balancer that is in the status of the **Service** with the **LoadBalancer** type.
- By using the **status.host** field of the **Route** resource.

However, when you deploy hosted control planes in a managed service context, those methods can expose the ingress subdomain of the underlying management cluster and limit options for the management cluster lifecycle and disaster recovery.

When a DNS indirection is layered on the **LoadBalancer** and **Route** publishing types, a managed service operator can publish all public hosted cluster services by using a service-level domain. This architecture allows remapping on the DNS name to a new **LoadBalancer** or **Route** and does not expose the ingress domain of the management cluster. Hosted control planes uses external DNS to achieve that indirection layer.

You can deploy **external-dns** alongside the HyperShift Operator in the **hypershift** namespace of the management cluster. External DNS watches for **Services** or **Routes** that have the **external-dns.alpha.kubernetes.io/hostname** annotation. That annotation is used to create a DNS record that points to the **Service**, such as an A record, or the **Route**, such as a CNAME record.

You can use external DNS on cloud environments only. For the other environments, you need to manually configure DNS and services.

For more information about external DNS, see [external DNS](#).

4.1.6.1. Prerequisites

Before you can set up external DNS for hosted control planes on Amazon Web Services (AWS), you must meet the following prerequisites:

- You created an external public domain.
- You have access to the AWS Route53 Management console.
- You enabled AWS PrivateLink for hosted control planes.

4.1.6.2. Setting up external DNS for hosted control planes

You can provision hosted control planes with external DNS or service-level DNS.

1. Create an Amazon Web Services (AWS) credential secret for the HyperShift Operator and name it **hypershift-operator-external-dns-credentials** in the **local-cluster** namespace.
2. See the following table to verify that the secret has the required fields:

Table 4.3. Required fields for the AWS secret

Field name	Description	Optional or required
provider	The DNS provider that manages the service-level DNS zone.	Required
domain-filter	The service-level domain.	Required
credentials	The credential file that supports all external DNS types.	Optional when you use AWS keys
aws-access-key-id	The credential access key id.	Optional when you use the AWS DNS service
aws-secret-access-key	The credential access key secret.	Optional when you use the AWS DNS service

3. To create an AWS secret, run the following command:

```
$ oc create secret generic <secret_name> \
--from-literal=provider=aws \
--from-literal=domain-filter=<domain_name> \
--from-file=credentials=<path_to_aws_credentials_file> -n local-cluster
```

**NOTE**

Disaster recovery backup for the secret is not automatically enabled. To back up the secret for disaster recovery, add the **hypershift-operator-external-dns-credentials** by entering the following command:

```
$ oc label secret hypershift-operator-external-dns-credentials \
-n local-cluster \
cluster.open-cluster-management.io/backup=""
```

4.1.6.3. Creating the public DNS hosted zone

The External DNS Operator uses the public DNS hosted zone to create your public hosted cluster.

You can create the public DNS hosted zone to use as the external DNS domain-filter. Complete the following steps in the AWS Route 53 management console.

Procedure

1. In the Route 53 management console, click **Create hosted zone**.
2. On the **Hosted zone configuration** page, type a domain name, verify that **Public hosted zone** is selected as the type, and click **Create hosted zone**.
3. After the zone is created, on the **Records** tab, note the values in the **Value/Route traffic to** column.
4. In the main domain, create an NS record to redirect the DNS requests to the delegated zone. In the **Value** field, enter the values that you noted in the previous step.
5. Click **Create records**.
6. Verify that the DNS hosted zone is working by creating a test entry in the new subzone and testing it with a **dig** command, such as in the following example:

```
$ dig +short test.user-dest-public.aws.kerberos.com
```

Example output

```
192.168.1.1
```

7. To create a hosted cluster that sets the hostname for the **LoadBalancer** and **Route** services, enter the following command:

```
$ hcp create cluster aws --name=<hosted_cluster_name> \
--endpoint-access=PublicAndPrivate \
--external-dns-domain=<public_hosted_zone> ... 1
```

- 1** Replace **<public_hosted_zone>** with the public hosted zone that you created.

Example services block for the hosted cluster

```
platform:
```

```

aws:
  endpointAccess: PublicAndPrivate
...
services:
- service: APIServer
  servicePublishingStrategy:
    route:
      hostname: api-example.service-provider-domain.com
      type: Route
- service: OAuthServer
  servicePublishingStrategy:
    route:
      hostname: oauth-example.service-provider-domain.com
      type: Route
- service: Konnectivity
  servicePublishingStrategy:
    type: Route
- service: Ignition
  servicePublishingStrategy:
    type: Route

```

The Control Plane Operator creates the **Services** and **Routes** resources and annotates them with the **external-dns.alpha.kubernetes.io/hostname** annotation. For **Services** and **Routes**, the Control Plane Operator uses a value of the **hostname** parameter in the **servicePublishingStrategy** field for the service endpoints. To create the DNS records, you can use a mechanism, such as the **external-dns** deployment.

You can configure service-level DNS indirection for public services only. You cannot set **hostname** for private services because they use the **hypershift.local** private zone.

The following table shows when it is valid to set **hostname** for a service and endpoint combinations:

Table 4.4. Service and endpoint combinations to set `hostname`

Service	Public	PublicAndPrivate	Private
APIServer	Y	Y	N
OAuthServer	Y	Y	N
Konnectivity	Y	N	N
Ignition	Y	N	N

4.1.6.4. Creating a hosted cluster by using the external DNS on AWS

To create a hosted cluster by using the **PublicAndPrivate** or **Public** publishing strategy on Amazon Web Services (AWS), you must have the following artifacts configured in your management cluster:

- The public DNS hosted zone
- The External DNS Operator

- The HyperShift Operator

You can deploy a hosted cluster, by using the **hcp** command-line interface (CLI).

Procedure

- 1 To access your management cluster, enter the following command:

```
$ export KUBECONFIG=<path_to_management_cluster_kubeconfig>
```

- 2 Verify that the External DNS Operator is running by entering the following command:

```
$ oc get pod -n hypershift -lapp=external-dns
```

Example output

NAME	READY	STATUS	RESTARTS	AGE
external-dns-7c89788c69-rn8gp	1/1	Running	0	40s

- 3 To create a hosted cluster by using external DNS, enter the following command:

```
$ hcp create cluster aws \
--role-arn <arn_role> \ ①
--instance-type <instance_type> \ ②
--region <region> \ ③
--auto-repair \
--generate-ssh \
--name <hosted_cluster_name> \ ④
--namespace clusters \
--base-domain <service_consumer_domain> \ ⑤
--node-pool-replicas <node_replica_count> \ ⑥
--pull-secret <path_to_your_pull_secret> \ ⑦
--release-image quay.io/openshift-release-dev/ocp-release:<ocp_release_image> \ ⑧
--external-dns-domain=<service_provider_domain> \ ⑨
--endpoint-access=PublicAndPrivate ⑩
--sts-creds <path_to_sts_credential_file> ⑪
```

- ① Specify the Amazon Resource Name (ARN), for example, **arn:aws:iam::820196288204:role/myrole**.
- ② Specify the instance type, for example, **m6i.xlarge**.
- ③ Specify the AWS region, for example, **us-east-1**.
- ④ Specify your hosted cluster name, for example, **my-external-aws**.
- ⑤ Specify the public hosted zone that the service consumer owns, for example, **service-consumer-domain.com**.
- ⑥ Specify the node replica count, for example, **2**.
- ⑦ Specify the path to your pull secret file.

- 8 Specify the supported OpenShift Container Platform version that you want to use, for example, **4.18.0-multi**.
- 9 Specify the public hosted zone that the service provider owns, for example, **service-provider-domain.com**.
- 10 Set as **PublicAndPrivate**. You can use external DNS with **Public** or **PublicAndPrivate** configurations only.
- 11 Specify the path to your AWS STS credentials file, for example, **/home/user/sts-creds/sts-creds.json**.

4.1.7. Creating a hosted cluster on AWS

You can create a hosted cluster on Amazon Web Services (AWS) by using the **hcp** command-line interface (CLI).

By default for hosted control planes on Amazon Web Services (AWS), you use an AMD64 hosted cluster. However, you can enable hosted control planes to run on an ARM64 hosted cluster. For more information, see "Running hosted clusters on an ARM64 architecture".

For compatible combinations of node pools and hosted clusters, see the following table:

Table 4.5. Compatible architectures for node pools and hosted clusters

Hosted cluster	Node pools
AMD64	AMD64 or ARM64
ARM64	ARM64 or AMD64

Prerequisites

- You have set up the hosted control plane CLI, **hcp**.
- You have enabled the **local-cluster** managed cluster as the management cluster.
- You created an AWS Identity and Access Management (IAM) role and AWS Security Token Service (STS) credentials.

Procedure

1. To create a hosted cluster on AWS, run the following command:

```
$ hcp create cluster aws \
--name <hosted_cluster_name> \①
--infra-id <infra_id> \②
--base-domain <basedomain> \③
--sts-creds <path_to_sts_credential_file> \④
--pull-secret <path_to_pull_secret> \⑤
--region <region> \⑥
--generate-ssh \
```

```
--node-pool-replicas <node_pool_replica_count> \ 7
--namespace <hosted_cluster_namespace> \ 8
--role-arn <role_name> \ 9
--render-into <file_name>.yaml 10
```

- 1 Specify the name of your hosted cluster, for instance, **example**.
- 2 Specify your infrastructure name. You must provide the same value for **<hosted_cluster_name>** and **<infra_id>**. Otherwise the cluster might not appear correctly in the multicloud engine for Kubernetes Operator console.
- 3 Specify your base domain, for example, **example.com**.
- 4 Specify the path to your AWS STS credentials file, for example, **/home/user/sts-creds/sts-creds.json**.
- 5 Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- 6 Specify the AWS region name, for example, **us-east-1**.
- 7 Specify the node pool replica count, for example, **3**.
- 8 By default, all **HostedCluster** and **NodePool** custom resources are created in the **clusters** namespace. You can use the **--namespace <namespace>** parameter, to create the **HostedCluster** and **NodePool** custom resources in a specific namespace.
- 9 Specify the Amazon Resource Name (ARN), for example, **arn:aws:iam::820196288204:role/myrole**.
- 10 If you want to indicate whether the EC2 instance runs on shared or single tenant hardware, include this field. The **--render-into** flag renders Kubernetes resources into the YAML file that you specify in this field. Then, continue to the next step to edit the YAML file.

2. If you included the **--render-into** flag in the previous command, edit the specified YAML file. Edit the **NodePool** specification in the YAML file to indicate whether the EC2 instance should run on shared or single-tenant hardware, similar to the following example:

Example YAML file

```
apiVersion: hypershift.openshift.io/v1beta1
kind: NodePool
metadata:
  name: <nodepool_name> 1
spec:
  platform:
    aws:
      placement:
        tenancy: "default" 2
```

- 1 Specify the name of the **NodePool** resource.
- 2 Specify a valid value for tenancy: **"default"**, **"dedicated"**, or **"host"**. Use **"default"** when node pool instances run on shared hardware. Use **"dedicated"** when each node pool instance runs on single-tenant hardware. Use **"host"** when node pool instances run on your

pre-allocated dedicated hosts.

Verification

1. Verify the status of your hosted cluster to check that the value of **AVAILABLE** is **True**. Run the following command:

```
$ oc get hostedclusters -n <hosted_cluster_namespace>
```

2. Get a list of your node pools by running the following command:

```
$ oc get nodepools --namespace <hosted_cluster_namespace>
```

Additional resources

- [Running hosted clusters on an ARM64 architecture](#)

4.1.7.1. Accessing a hosted cluster on AWS by using the kubeadmin credentials

After creating a hosted cluster on Amazon Web Services (AWS), you can access a hosted cluster by getting the **kubeconfig** file, access secrets, and the **kubeadmin** credentials.

The hosted cluster namespace contains hosted cluster resources and the access secrets. The hosted control plane runs in the hosted control plane namespace.

The secret name formats are as follows:

- The **kubeconfig** secret: **<hosted_cluster_namespace>-<name>-admin-kubeconfig**. For example, **clusters-hypershift-demo-admin-kubeconfig**.
- The **kubeadmin** password secret: **<hosted_cluster_namespace>-<name>-kubeadmin-password**. For example, **clusters-hypershift-demo-kubeadmin-password**.



NOTE

The **kubeadmin** password secret is Base64-encoded and the **kubeconfig** secret contains a Base64-encoded **kubeconfig** configuration. You must decode the Base64-encoded **kubeconfig** configuration and save it into a **<hosted_cluster_name>.kubeconfig** file.

Procedure

- Use your **<hosted_cluster_name>.kubeconfig** file that contains the decoded **kubeconfig** configuration to access the hosted cluster. Enter the following command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get nodes
```

You must decode the **kubeadmin** password secret to log in to the API server or the console of the hosted cluster.

4.1.7.2. Accessing a hosted cluster on AWS by using the hcp CLI

You can access the hosted cluster by using the **hcp** command-line interface (CLI).

Procedure

1. Generate the **kubeconfig** file by entering the following command:

```
$ hcp create kubeconfig --namespace <hosted_cluster_namespace> \
--name <hosted_cluster_name> > <hosted_cluster_name>.kubeconfig
```

2. After you save the **kubeconfig** file, access the hosted cluster by entering the following command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get nodes
```

4.1.8. Creating a hosted cluster in multiple zones on AWS

You can create a hosted cluster in multiple zones on Amazon Web Services (AWS) by using the **hcp** command-line interface (CLI).

Prerequisites

- You created an AWS Identity and Access Management (IAM) role and AWS Security Token Service (STS) credentials.

Procedure

- Create a hosted cluster in multiple zones on AWS by running the following command:

```
$ hcp create cluster aws \
--name <hosted_cluster_name> \①
--node-pool-replicas=<node_pool_replica_count> \②
--base-domain <basedomain> \③
--pull-secret <path_to_pull_secret> \④
--role-arn <arn_role> \⑤
--region <region> \⑥
--zones <zones> \⑦
--sts-creds <path_to_sts_credential_file> \⑧
```

- 1 Specify the name of your hosted cluster, for instance, **example**.
- 2 Specify the node pool replica count, for example, **2**.
- 3 Specify your base domain, for example, **example.com**.
- 4 Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- 5 Specify the Amazon Resource Name (ARN), for example, **arn:aws:iam::820196288204:role/myrole**.
- 6 Specify the AWS region name, for example, **us-east-1**.
- 7 Specify availability zones within your AWS region, for example, **us-east-1a**, and **us-east-1b**.
- 8 Specify the path to your AWS STS credentials file, for example, **/home/user/sts-creds/sts-creds.json**.

For each specified zone, the following infrastructure is created:

- Public subnet
- Private subnet
- NAT gateway
- Private route table

A public route table is shared across public subnets.

One **NodePool** resource is created for each zone. The node pool name is suffixed by the zone name. The private subnet for zone is set in **spec.platform.aws.subnet.id**.

4.1.8.1. Creating a hosted cluster by providing AWS STS credentials

When you create a hosted cluster by using the **hcp create cluster aws** command, you must provide an Amazon Web Services (AWS) account credentials that have permissions to create infrastructure resources for your hosted cluster.

Infrastructure resources include the following examples:

- Virtual Private Cloud (VPC)
- Subnets
- Network address translation (NAT) gateways

You can provide the AWS credentials by using the either of the following ways:

- The AWS Security Token Service (STS) credentials
- The AWS cloud provider secret from multicluster engine Operator

Procedure

- To create a hosted cluster on AWS by providing AWS STS credentials, enter the following command:

```
$ hcp create cluster aws \
--name <hosted_cluster_name> \①
--node-pool-replicas <node_pool_replica_count> \②
--base-domain <basedomain> \③
--pull-secret <path_to_pull_secret> \④
--sts-creds <path_to_sts_credential_file> \⑤
--region <region> \⑥
--role-arn <arn_role> \⑦
```

- ① Specify the name of your hosted cluster, for instance, **example**.
- ② Specify the node pool replica count, for example, **2**.
- ③ Specify your base domain, for example, **example.com**.

- 4 Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- 5 Specify the path to your AWS STS credentials file, for example, **/home/user/sts-creds/sts-creds.json**.
- 6 Specify the AWS region name, for example, **us-east-1**.
- 7 Specify the Amazon Resource Name (ARN), for example, **arn:aws:iam::820196288204:role/myrole**.

4.1.9. Running hosted clusters on an ARM64 architecture

By default for hosted control planes on Amazon Web Services (AWS), you use an AMD64 hosted cluster. However, you can enable hosted control planes to run on an ARM64 hosted cluster.

For compatible combinations of node pools and hosted clusters, see the following table:

Table 4.6. Compatible architectures for node pools and hosted clusters

Hosted cluster	Node pools
AMD64	AMD64 or ARM64
ARM64	ARM64 or AMD64

4.1.9.1. Creating a hosted cluster on an ARM64 OpenShift Container Platform cluster

You can run a hosted cluster on an ARM64 OpenShift Container Platform cluster for Amazon Web Services (AWS) by overriding the default release image with a multi-architecture release image.

If you do not use a multi-architecture release image, the compute nodes in the node pool are not created and reconciliation of the node pool stops until you either use a multi-architecture release image in the hosted cluster or update the **NodePool** custom resource based on the release image.

Prerequisites

- You must have an OpenShift Container Platform cluster with a 64-bit ARM infrastructure that is installed on AWS. For more information, see [Create an OpenShift Container Platform Cluster: AWS \(ARM\)](#).
- You must create an AWS Identity and Access Management (IAM) role and AWS Security Token Service (STS) credentials. For more information, see "Creating an AWS IAM role and STS credentials".

Procedure

- Create a hosted cluster on an ARM64 OpenShift Container Platform cluster by entering the following command:

```
$ hcp create cluster aws \
--name <hosted_cluster_name> \①
--node-pool-replicas <node_pool_replica_count> \②
```

```
--base-domain <basedomain> \③
--pull-secret <path_to_pull_secret> \④
--sts-creds <path_to_sts_credential_file> \⑤
--region <region> \⑥
--release-image quay.io/openshift-release-dev/ocp-release:<ocp_release_image> \⑦
--role-arn <role_name> ⑧
```

- ① Specify the name of your hosted cluster, for instance, **example**.
- ② Specify the node pool replica count, for example, **3**.
- ③ Specify your base domain, for example, **example.com**.
- ④ Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- ⑤ Specify the path to your AWS STS credentials file, for example, **/home/user/sts-creds/sts-creds.json**.
- ⑥ Specify the AWS region name, for example, **us-east-1**.
- ⑦ Specify the supported OpenShift Container Platform version that you want to use, for example, **4.18.0-multi**. If you are using a disconnected environment, replace **<ocp_release_image>** with the digest image. To extract the OpenShift Container Platform release image digest, see "Extracting the OpenShift Container Platform release image digest".
- ⑧ Specify the Amazon Resource Name (ARN), for example, **arn:aws:iam::820196288204:role/myrole**.

4.1.9.2. Creating an ARM or AMD NodePool object on AWS hosted clusters

You can schedule application workloads that is the **NodePool** objects on 64-bit ARM and AMD from the same hosted control plane. You can define the **arch** field in the **NodePool** specification to set the required processor architecture for the **NodePool** object. The valid values for the **arch** field are as follows:

- **arm64**
- **amd64**

Prerequisites

- You must have a multi-architecture image for the **HostedCluster** custom resource to use. You can access [multi-architecture nightly images](#).

Procedure

- Add an ARM or AMD **NodePool** object to the hosted cluster on AWS by running the following command:

```
$ hcp create nodepool aws \
--cluster-name <hosted_cluster_name> \①
--name <node_pool_name> \②
```

```
--node-count <node_pool_replica_count> \③
--arch <architecture> ④
```

- ① Specify the name of your hosted cluster, for instance, **example**.
- ② Specify the node pool name.
- ③ Specify the node pool replica count, for example, **3**.
- ④ Specify the architecture type, such as **arm64** or **amd64**. If you do not specify a value for the **--arch** flag, the **amd64** value is used by default.

Additional resources

- [Extracting the OpenShift Container Platform release image digest](#)

4.1.10. Creating a private hosted cluster on AWS

After you enable the **local-cluster** as the hosting cluster, you can deploy a hosted cluster or a private hosted cluster on Amazon Web Services (AWS).

By default, hosted clusters are publicly accessible through public DNS and the default router for the management cluster.

For private clusters on AWS, all communication with the hosted cluster occurs over AWS PrivateLink.

Prerequisites

- You enabled AWS PrivateLink. For more information, see "Enabling AWS PrivateLink".
- You created an AWS Identity and Access Management (IAM) role and AWS Security Token Service (STS) credentials. For more information, see "Creating an AWS IAM role and STS credentials" and "Identity and Access Management (IAM) permissions".
- You configured a [bastion instance on AWS](#).

Procedure

- Create a private hosted cluster on AWS by entering the following command:

```
$ hcp create cluster aws \
--name <hosted_cluster_name> \①
--node-pool-replicas=<node_pool_replica_count> \②
--base-domain <basedomain> \③
--pull-secret <path_to_pull_secret> \④
--sts-creds <path_to_sts_credential_file> \⑤
--region <region> \⑥
--endpoint-access Private \⑦
--role-arn <role_name> \⑧
```

- ① Specify the name of your hosted cluster, for instance, **example**.
- ② Specify the node pool replica count, for example, **3**.

- 3** Specify your base domain, for example, **example.com**.
- 4** Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- 5** Specify the path to your AWS STS credentials file, for example, **/home/user/sts-creds/sts-creds.json**.
- 6** Specify the AWS region name, for example, **us-east-1**.
- 7** Defines whether a cluster is public or private.
- 8** Specify the Amazon Resource Name (ARN), for example, **arn:aws:iam::820196288204:role/myrole**. For more information about ARN roles, see "Identity and Access Management (IAM) permissions".

The following API endpoints for the hosted cluster are accessible through a private DNS zone:

- **api.<hosted_cluster_name>.hypershift.local**
- ***.apps.<hosted_cluster_name>.hypershift.local**

4.1.10.1. Accessing a private management cluster on AWS

Additional resources

You can access your private management cluster by using the command-line interface (CLI).

Procedure

1. Find the private IPs of nodes by entering the following command:

```
$ aws ec2 describe-instances \
--filter="Name=tag:kubernetes.io/cluster/<infra_id>,Values=owned" \
| jq '.Reservations[] | .Instances[] | select(.PublicDnsName== "")' \
|.PrivateIpAddress'
```

2. Create a **kubeconfig** file for the hosted cluster that you can copy to a node by entering the following command:

```
$ hcp create kubeconfig > <hosted_cluster_kubeconfig>
```

3. To SSH into one of the nodes through the bastion, enter the following command:

```
$ ssh -o ProxyCommand="ssh ec2-user@<bastion_ip> \
-W %h:%p" core@<node_ip>
```

4. From the SSH shell, copy the **kubeconfig** file contents to a file on the node by entering the following command:

```
$ mv <path_to_kubeconfig_file> <new_file_name>
```

5. Export the **kubeconfig** file by entering the following command:

```
$ export KUBECONFIG=<path_to_kubeconfig_file>
```

- Observe the hosted cluster status by entering the following command:

```
$ oc get clusteroperators clusterversion
```

4.2. DEPLOYING HOSTED CONTROL PLANES ON BARE METAL

You can deploy hosted control planes by configuring a cluster to function as a management cluster. The management cluster is the OpenShift Container Platform cluster where the control planes are hosted. In some contexts, the management cluster is also known as the *hosting* cluster.



NOTE

The management cluster is not the same thing as the *managed* cluster. A managed cluster is a cluster that the hub cluster manages.

The hosted control planes feature is enabled by default.

The multicluster engine Operator supports only the default **local-cluster**, which is a hub cluster that is managed, and the hub cluster as the management cluster. If you have Red Hat Advanced Cluster Management installed, you can use the managed hub cluster, also known as the **local-cluster**, as the management cluster.

A *hosted cluster* is an OpenShift Container Platform cluster with its API endpoint and control plane that are hosted on the management cluster. The hosted cluster includes the control plane and its corresponding data plane. You can use the multicluster engine Operator console or the hosted control plane command-line interface, **hcp**, to create a hosted cluster.

The hosted cluster is automatically imported as a managed cluster. If you want to disable this automatic import feature, see *Disabling the automatic import of hosted clusters into multicluster engine Operator*.

4.2.1. Preparing to deploy hosted control planes on bare metal

As you prepare to deploy hosted control planes on bare metal, consider the following information:

- Run the management cluster and workers on the same platform for hosted control planes.
- All bare metal hosts require a manual start with a Discovery Image ISO that the central infrastructure management provides. You can start the hosts manually or through automation by using **Cluster-Baremetal-Operator**. After each host starts, it runs an Agent process to discover the host details and complete the installation. An **Agent** custom resource represents each host.
- When you configure storage for hosted control planes, consider the recommended etcd practices. To ensure that you meet the latency requirements, dedicate a fast storage device to all hosted control plane etcd instances that run on each control-plane node. You can use LVM storage to configure a local storage class for hosted etcd pods. For more information, see "Recommended etcd practices" and "Persistent storage using logical volume manager storage".

4.2.1.1. Prerequisites to configure a management cluster

- You need the multicluster engine for Kubernetes Operator 2.2 and later installed on an OpenShift Container Platform cluster. You can install multicluster engine Operator as an Operator from the OpenShift Container Platform OperatorHub.

- The multicloud engine Operator must have at least one managed OpenShift Container Platform cluster. The **local-cluster** is automatically imported in multicloud engine Operator 2.2 and later. For more information about the **local-cluster**, see *Advanced configuration* in the Red Hat Advanced Cluster Management documentation. You can check the status of your hub cluster by running the following command:

```
$ oc get managedclusters local-cluster
```

- You must add the **topology.kubernetes.io/zone** label to your bare-metal hosts on your management cluster. Ensure that each host has a unique value for **topology.kubernetes.io/zone**. Otherwise, all of the hosted control plane pods are scheduled on a single node, causing a single point of failure.
- To provision hosted control planes on bare metal, you can use the Agent platform. The Agent platform uses the central infrastructure management service to add worker nodes to a hosted cluster. For more information, see *Enabling the central infrastructure management service*.
- You need to install the hosted control plane command-line interface.

Additional resources

- [Advanced configuration](#)
- [Enabling the central infrastructure management service](#)

4.2.1.2. Bare metal firewall, port, and service requirements

You must meet the firewall, port, and service requirements so that ports can communicate between the management cluster, the control plane, and hosted clusters.



NOTE

Services run on their default ports. However, if you use the **NodePort** publishing strategy, services run on the port that is assigned by the **NodePort** service.

Use firewall rules, security groups, or other access controls to restrict access to only required sources. Avoid exposing ports publicly unless necessary. For production deployments, use a load balancer to simplify access through a single IP address.

If your hub cluster has a proxy configuration, ensure that it can reach the hosted cluster API endpoint by adding all hosted cluster API endpoints to the **noProxy** field on the **Proxy** object. For more information, see "Configuring the cluster-wide proxy".

A hosted control plane exposes the following services on bare metal:

- APIServer**
 - The **APIServer** service runs on port 6443 by default and requires ingress access for communication between the control plane components.
 - If you use MetalLB load balancing, allow ingress access to the IP range that is used for load balancer IP addresses.
- OAuthServer**
 - The **OAuthServer** service runs on port 443 by default when you use the route and ingress

to expose the service.

- If you use the **NodePort** publishing strategy, use a firewall rule for the **OAuthServer** service.
- **Konnectivity**
 - The **Konnectivity** service runs on port 443 by default when you use the route and ingress to expose the service.
 - The **Konnectivity** agent establishes a reverse tunnel to allow the control plane to access the network for the hosted cluster. The agent uses egress to connect to the **Konnectivity** server. The server is exposed by using either a route on port 443 or a manually assigned **NodePort**.
 - If the cluster API server address is an internal IP address, allow access from the workload subnets to the IP address on port 6443.
 - If the address is an external IP address, allow egress on port 6443 to that external IP address from the nodes.

● **Ignition**

- The **Ignition** service runs on port 443 by default when you use the route and ingress to expose the service.
- If you use the **NodePort** publishing strategy, use a firewall rule for the **Ignition** service.

You do not need the following services on bare metal:

- **OVNSbDb**
- **OIDC**

Additional resources

- [Configuring the cluster-wide proxy](#)

4.2.1.3. Bare metal infrastructure requirements

The Agent platform does not create any infrastructure, but it does have the following requirements for infrastructure:

- Agents: An *Agent* represents a host that is booted with a discovery image and is ready to be provisioned as an OpenShift Container Platform node.
- DNS: The API and ingress endpoints must be routable.

Additional resources

- [Recommended etcd practices](#)
- [Persistent storage using LVM Storage](#)
- [Disabling the automatic import of hosted clusters into multicluster engine Operator](#)
- [Enabling or disabling the hosted control planes feature](#)

- Configuring Ansible Automation Platform jobs to run on hosted clusters

4.2.2. DNS configurations on bare metal

The API Server for the hosted cluster is exposed as a **NodePort** service. A DNS entry must exist for **api.<hosted_cluster_name>.<base_domain>** that points to destination where the API Server can be reached.

The DNS entry can be as simple as a record that points to one of the nodes in the managed cluster that is running the hosted control plane. The entry can also point to a load balancer that is deployed to redirect incoming traffic to the ingress pods.

Example DNS configuration

```
api.example.krnl.es. IN A 192.168.122.20
api.example.krnl.es. IN A 192.168.122.21
api.example.krnl.es. IN A 192.168.122.22
api-int.example.krnl.es. IN A 192.168.122.20
api-int.example.krnl.es. IN A 192.168.122.21
api-int.example.krnl.es. IN A 192.168.122.22
`*.apps.example.krnl.es. IN A 192.168.122.23
```

If you are configuring DNS for a disconnected environment on an IPv6 network, the configuration looks like the following example.

Example DNS configuration for an IPv6 network

```
api.example.krnl.es. IN A 2620:52:0:1306::5
api.example.krnl.es. IN A 2620:52:0:1306::6
api.example.krnl.es. IN A 2620:52:0:1306::7
api-int.example.krnl.es. IN A 2620:52:0:1306::5
api-int.example.krnl.es. IN A 2620:52:0:1306::6
api-int.example.krnl.es. IN A 2620:52:0:1306::7
`*.apps.example.krnl.es. IN A 2620:52:0:1306::10
```

If you are configuring DNS for a disconnected environment on a dual stack network, be sure to include DNS entries for both IPv4 and IPv6.

Example DNS configuration for a dual stack network

```
host-record=api-int.hub-dual.dns.base.domain.name,192.168.126.10
host-record=api.hub-dual.dns.base.domain.name,192.168.126.10
address=/apps.hub-dual.dns.base.domain.name/192.168.126.11
dhcp-host=aa:aa:aa:aa:10:01,ocp-master-0,192.168.126.20
dhcp-host=aa:aa:aa:aa:10:02,ocp-master-1,192.168.126.21
dhcp-host=aa:aa:aa:aa:10:03,ocp-master-2,192.168.126.22
dhcp-host=aa:aa:aa:aa:10:06,ocp-installer,192.168.126.25
dhcp-host=aa:aa:aa:aa:10:07,ocp-bootstrap,192.168.126.26

host-record=api-int.hub-dual.dns.base.domain.name,2620:52:0:1306::2
host-record=api.hub-dual.dns.base.domain.name,2620:52:0:1306::2
address=/apps.hub-dual.dns.base.domain.name/2620:52:0:1306::3
dhcp-host=aa:aa:aa:aa:10:01,ocp-master-0,[2620:52:0:1306::5]
dhcp-host=aa:aa:aa:aa:10:02,ocp-master-1,[2620:52:0:1306::6]
```

```
dhcp-host=aa:aa:aa:aa:10:03,ocp-master-2,[2620:52:0:1306::7]
dhcp-host=aa:aa:aa:aa:10:06,ocp-installer,[2620:52:0:1306::8]
dhcp-host=aa:aa:aa:aa:10:07,ocp-bootstrap,[2620:52:0:1306::9]
```

4.2.3. Creating a hosted cluster on bare metal

When you create a hosted cluster with the Agent platform, HyperShift installs the Agent Cluster API provider in the hosted control plane namespace. You can create a hosted cluster on bare metal or import one.

As you create a hosted cluster, keep the following guidelines in mind:

- Each hosted cluster must have a cluster-wide unique name. A hosted cluster name cannot be the same as any existing managed cluster in order for multicluster engine Operator to manage it.
- Do not use **clusters** as a hosted cluster name.
- A hosted cluster cannot be created in the namespace of a multicluster engine Operator managed cluster.
- The most common service publishing strategy is to expose services through a load balancer. That strategy is the preferred method for exposing the Kubernetes API server. If you create a hosted cluster by using the web console or by using Red Hat Advanced Cluster Management, to set a publishing strategy for a service besides the Kubernetes API server, you must manually specify the **servicePublishingStrategy** information in the **HostedCluster** custom resource.

Procedure

1. Create the hosted control plane namespace by entering the following command:

```
$ oc create ns <hosted_cluster_namespace>-<hosted_cluster_name>
```

Replace **<hosted_cluster_namespace>** with your hosted cluster namespace name, for example, **clusters**. Replace **<hosted_cluster_name>** with your hosted cluster name.

2. Verify that you have a default storage class configured for your cluster. Otherwise, you might see pending PVCs. Run the following command:

```
$ hcp create cluster agent \
--name=<hosted_cluster_name> \①
--pull-secret=<path_to_pull_secret> \②
--agent-namespace=<hosted_control_plane_namespace> \③
--base-domain=<basedomain> \④
--api-server-address=api.<hosted_cluster_name>.<basedomain> \⑤
--etcd-storage-class=<etcd_storage_class> \⑥
--ssh-key <path_to_ssh_public_key> \⑦
--namespace <hosted_cluster_namespace> \⑧
--control-plane-availability-policy HighlyAvailable \⑨
--release-image=quay.io/openshift-release-dev/ocp-release:<ocp_release_image> \⑩
--node-pool-replicas <node_pool_replica_count> \⑪
```

① Specify the name of your hosted cluster, for instance, **example**.

- 2** Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- 3** Specify your hosted control plane namespace, for example, **clusters-example**. Ensure that agents are available in this namespace by using the **oc get agent -n <hosted_control_plane_namespace>** command.
- 4** Specify your base domain, for example, **krnl.es**.
- 5** The **--api-server-address** flag defines the IP address that is used for the Kubernetes API communication in the hosted cluster. If you do not set the **--api-server-address** flag, you must log in to connect to the management cluster.
- 6** Specify the etcd storage class name, for example, **lvm-storageclass**.
- 7** Specify the path to your SSH public key. The default file path is **~/.ssh/id_rsa.pub**.
- 8** Specify your hosted cluster namespace.
- 9** Specify the availability policy for the hosted control plane components. Supported options are **SingleReplica** and **HighlyAvailable**. The default value is **HighlyAvailable**.
- 10** Specify the supported OpenShift Container Platform version that you want to use, for example, **4.18.0-multi**. If you are using a disconnected environment, replace **<ocp_release_image>** with the digest image. To extract the OpenShift Container Platform release image digest, see *Extracting the OpenShift Container Platform release image digest*.
- 11** Specify the node pool replica count, for example, **3**. You must specify the replica count as **0** or greater to create the same number of replicas. Otherwise, no node pools are created.

3. After a few moments, verify that your hosted control plane pods are up and running by entering the following command:

```
$ oc -n <hosted_cluster_namespace>-<hosted_cluster_name> get pods
```

Example output

NAME	READY	STATUS	RESTARTS	AGE
capi-provider-7dcf5fc4c4-nr9sq	1/1	Running	0	4m32s
catalog-operator-6cd867cc7-phb2q	2/2	Running	0	2m50s
certified-operators-catalog-884c756c4-zdt64	1/1	Running	0	2m51s
cluster-api-f75d86f8c-56wfz	1/1	Running	0	4m32s

Additional resources

- [Manually importing a hosted cluster](#)

4.2.3.1. Creating a hosted cluster on bare metal by using the console

To create a hosted cluster by using the console, complete the following steps.

Procedure

1. Open the OpenShift Container Platform web console and log in by entering your administrator credentials. For instructions to open the console, see [Accessing the web console](#).
2. In the console header, ensure that **All Clusters** is selected.
3. Click **Infrastructure → Clusters**.
4. Click **Create cluster → Host inventory → Hosted control plane**
The **Create cluster** page is displayed.
5. On the **Create cluster** page, follow the prompts to enter details about the cluster, node pools, networking, and automation.



NOTE

As you enter details about the cluster, you might find the following tips useful:

- If you want to use predefined values to automatically populate fields in the console, you can create a host inventory credential. For more information, see [Creating a credential for an on-premises environment](#).
- On the **Cluster details** page, the pull secret is your OpenShift Container Platform pull secret that you use to access OpenShift Container Platform resources. If you selected a host inventory credential, the pull secret is automatically populated.
- On the **Node pools** page, the namespace contains the hosts for the node pool. If you created a host inventory by using the console, the console creates a dedicated namespace.
- On the **Networking** page, you select an API server publishing strategy. The API server for the hosted cluster can be exposed either by using an existing load balancer or as a service of the **NodePort** type. A DNS entry must exist for the `api.<hosted_cluster_name>.<base_domain>` setting that points to the destination where the API server can be reached. This entry can be a record that points to one of the nodes in the management cluster or a record that points to a load balancer that redirects incoming traffic to the Ingress pods.

6. Review your entries and click **Create**.
The **Hosted cluster** view is displayed.
7. Monitor the deployment of the hosted cluster in the **Hosted cluster** view.
8. If you do not see information about the hosted cluster, ensure that **All Clusters** is selected, then click the cluster name.
9. Wait until the control plane components are ready. This process can take a few minutes.
10. To view the node pool status, scroll to the **NodePool** section. The process to install the nodes takes about 10 minutes. You can also click **Nodes** to confirm whether the nodes joined the hosted cluster.

Next steps

- To access the web console, see [Accessing the web console](#).

4.2.3.2. Creating a hosted cluster on bare metal by using a mirror registry

You can use a mirror registry to create a hosted cluster on bare metal by specifying the **--image-content-sources** flag in the **hcp create cluster** command.

Procedure

1. Create a YAML file to define Image Content Source Policies (ICSP). See the following example:

```
- mirrors:
  - brew.registry.redhat.io
    source: registry.redhat.io
- mirrors:
  - brew.registry.redhat.io
    source: registry.stage.redhat.io
- mirrors:
  - brew.registry.redhat.io
    source: registry-proxy.engineering.redhat.com
```

2. Save the file as **icsp.yaml**. This file contains your mirror registries.
3. To create a hosted cluster by using your mirror registries, run the following command:

```
$ hcp create cluster agent \
--name=<hosted_cluster_name> \①
--pull-secret=<path_to_pull_secret> \②
--agent-namespace=<hosted_control_plane_namespace> \③
--base-domain=<basedomain> \④
--api-server-address=api.<hosted_cluster_name>.<basedomain> \⑤
--image-content-sources icsp.yaml \⑥
--ssh-key <path_to_ssh_key> \⑦
--namespace <hosted_cluster_namespace> \⑧
--release-image=quay.io/openshift-release-dev/ocp-release:<ocp_release_image> \⑨
```

- 1 Specify the name of your hosted cluster, for instance, **example**.
- 2 Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- 3 Specify your hosted control plane namespace, for example, **clusters-example**. Ensure that agents are available in this namespace by using the **oc get agent -n <hosted-control-plane-namespace>** command.
- 4 Specify your base domain, for example, **krnl.es**.
- 5 The **--api-server-address** flag defines the IP address that is used for the Kubernetes API communication in the hosted cluster. If you do not set the **--api-server-address** flag, you must log in to connect to the management cluster.
- 6 Specify the **icsp.yaml** file that defines ICSP and your mirror registries.
- 7 Specify the path to your SSH public key. The default file path is **~/.ssh/id_rsa.pub**.
- 8 Specify your hosted cluster namespace.

- 9** Specify the supported OpenShift Container Platform version that you want to use, for example, **4.18.0-multi**. If you are using a disconnected environment, replace

Next steps

- To create credentials that you can reuse when you create a hosted cluster with the console, see [Creating a credential for an on-premises environment](#).
- To access a hosted cluster, see [Accessing the hosted cluster](#).
- To add hosts to the host inventory by using the Discovery Image, see [Adding hosts to the host inventory by using the Discovery Image](#).
- To extract the OpenShift Container Platform release image digest, see [Extracting the OpenShift Container Platform release image digest](#).

4.2.4. Verifying hosted cluster creation

After the deployment process is complete, you can verify that the hosted cluster was created successfully. Follow these steps a few minutes after you create the hosted cluster.

Procedure

1. Obtain the kubeconfig for your new hosted cluster by entering the extract command:

```
$ oc extract -n <hosted-control-plane-namespace> secret/admin-kubeconfig \
--to=- > kubeconfig-<hosted-cluster-name>
```

2. Use the kubeconfig to view the cluster Operators of the hosted cluster. Enter the following command:

```
$ oc get co --kubeconfig=kubeconfig-<hosted-cluster-name>
```

Example output

NAME SINCE MESSAGE	VERSION	AVAILABLE	PROGRESSING	DEGRADED
console	4.10.26	True	False	False 2m38s
dns	4.10.26	True	False	False 2m52s
image-registry	4.10.26	True	False	False 2m8s
ingress	4.10.26	True	False	False 22m

3. You can also view the running pods on your hosted cluster by entering the following command:

```
$ oc get pods -A --kubeconfig=kubeconfig-<hosted-cluster-name>
```

Example output

NAMESPACE	STATUS	RESTARTS	AGE	NAME	READY
kube-system	Running	0	3m52s	konnectivity-agent-khlqv	0/1

openshift-cluster-node-tuning-operator	tuned-dhw5p	1/1
Running 0 109s		
openshift-cluster-storage-operator	cluster-storage-operator-5f784969f5-vwzgz	
1/1 Running 1 (113s ago) 20m		
openshift-cluster-storage-operator	csi-snapshot-controller-6b7687b7d9-7nrfw	
1/1 Running 0 3m8s		
openshift-console	console-5cbf6c7969-6gk6z	1/1
Running 0 119s		
openshift-console	downloads-7bcd756565-6wj5j	1/1
Running 0 4m3s		
openshift-dns-operator	dns-operator-77d755cd8c-xjfbn	2/2
Running 0 21m		
openshift-dns	dns-default-kfqnh	2/2
Running 0 113s		

4.3. DEPLOYING HOSTED CONTROL PLANES ON OPENSHIFT VIRTUALIZATION

With hosted control planes and OpenShift Virtualization, you can create OpenShift Container Platform clusters with worker nodes that are hosted by KubeVirt virtual machines. Hosted control planes on OpenShift Virtualization provides several benefits:

- Enhances resource usage by packing hosted control planes and hosted clusters in the same underlying bare metal infrastructure
- Separates hosted control planes and hosted clusters to provide strong isolation
- Reduces cluster provision time by eliminating the bare metal node bootstrapping process
- Manages many releases under the same base OpenShift Container Platform cluster

The hosted control planes feature is enabled by default.

You can use the hosted control plane command-line interface, **hcp**, to create an OpenShift Container Platform hosted cluster. The hosted cluster is automatically imported as a managed cluster. If you want to disable this automatic import feature, see "Disabling the automatic import of hosted clusters into multicluster engine Operator".

Additional resources

- [Disabling the automatic import of hosted clusters into multicluster engine Operator](#)
- [Enabling or disabling the hosted control planes feature](#)
- [Configuring Ansible Automation Platform jobs to run on hosted clusters](#)

4.3.1. Requirements to deploy hosted control planes on OpenShift Virtualization

As you prepare to deploy hosted control planes on OpenShift Virtualization, consider the following information:

- Run the management cluster on bare metal.
- Each hosted cluster must have a cluster-wide unique name.

- Do not use **clusters** as a hosted cluster name.
- A hosted cluster cannot be created in the namespace of a multicluster engine Operator managed cluster.
- When you configure storage for hosted control planes, consider the recommended etcd practices. To ensure that you meet the latency requirements, dedicate a fast storage device to all hosted control plane etcd instances that run on each control-plane node. You can use LVM storage to configure a local storage class for hosted etcd pods. For more information, see "Recommended etcd practices" and "Persistent storage using Logical Volume Manager storage".

Additional resources

- [Recommended etcd practices](#)
- [Persistent storage using Logical Volume Manager Storage](#)

4.3.1.1. Prerequisites

You must meet the following prerequisites to create an OpenShift Container Platform cluster on OpenShift Virtualization:

- You have administrator access to an OpenShift Container Platform cluster, version 4.14 or later, specified in the **KUBECONFIG** environment variable.
- The OpenShift Container Platform management cluster has wildcard DNS routes enabled, as shown in the following DNS:

```
$ oc patch ingresscontroller -n openshift-ingress-operator default \
--type=json \
-p '[{"op": "add", "path": "/spec/routeAdmission", "value": {"wildcardPolicy": "WildcardsAllowed"}}]'
```

- The OpenShift Container Platform management cluster has OpenShift Virtualization, version 4.14 or later, installed on it. For more information, see "Installing OpenShift Virtualization using the web console".
- The OpenShift Container Platform management cluster is on-premise bare metal.
- The OpenShift Container Platform management cluster is configured with OVNKubernetes as the default pod network CNI.
- The OpenShift Container Platform management cluster has a default storage class. For more information, see "Postinstallation storage configuration". The following example shows how to set a default storage class:

```
$ oc patch storageclass ocs-storagecluster-ceph-rbd \
-p '{"metadata": {"annotations": {"storageclass.kubernetes.io/is-default-class": "true"}}}'
```

- You have a valid pull secret file for the **quay.io/openshift-release-dev** repository. For more information, see "Install OpenShift on any x86_64 platform with user-provisioned infrastructure".
- You have installed the hosted control plane command-line interface.

- You have configured a load balancer. For more information, see "Configuring MetalLB".
- For optimal network performance, you are using a network maximum transmission unit (MTU) of 9000 or greater on the OpenShift Container Platform cluster that hosts the KubeVirt virtual machines. If you use a lower MTU setting, network latency and the throughput of the hosted pods are affected. Enable multiqueue on node pools only when the MTU is 9000 or greater.
- The multicloud engine Operator has at least one managed OpenShift Container Platform cluster. The **local-cluster** is automatically imported. For more information about the **local-cluster**, see "Advanced configuration" in the multicloud engine Operator documentation. You can check the status of your hub cluster by running the following command:

```
$ oc get managedclusters local-cluster
```

- On the OpenShift Container Platform cluster that hosts the OpenShift Virtualization virtual machines, you are using a **ReadWriteMany** (RWX) storage class so that live migration can be enabled.

Additional resources

- [Installing OpenShift Virtualization using the web console](#)
- [Postinstallation storage configuration](#)
- [Install OpenShift on any x86_64 platform with user-provisioned infrastructure](#)
- [Configuring MetalLB](#)
- [Advanced configuration](#)

4.3.1.2. Firewall and port requirements

Ensure that you meet the firewall and port requirements so that ports can communicate between the management cluster, the control plane, and hosted clusters:

- The **kube-apiserver** service runs on port 6443 by default and requires ingress access for communication between the control plane components.
 - If you use the **NodePort** publishing strategy, ensure that the node port that is assigned to the **kube-apiserver** service is exposed.
 - If you use MetalLB load balancing, allow ingress access to the IP range that is used for load balancer IP addresses.
- If you use the **NodePort** publishing strategy, use a firewall rule for the **ignition-server** and **Oauth-server** settings.
- The **konnectivity** agent, which establishes a reverse tunnel to allow bi-directional communication on the hosted cluster, requires egress access to the cluster API server address on port 6443. With that egress access, the agent can reach the **kube-apiserver** service.
 - If the cluster API server address is an internal IP address, allow access from the workload subnets to the IP address on port 6443.
 - If the address is an external IP address, allow egress on port 6443 to that external IP address from the nodes.

- If you change the default port of 6443, adjust the rules to reflect that change.
- Ensure that you open any ports that are required by the workloads that run in the clusters.
- Use firewall rules, security groups, or other access controls to restrict access to only required sources. Avoid exposing ports publicly unless necessary.
- For production deployments, use a load balancer to simplify access through a single IP address.

4.3.2. Live migration for compute nodes

While the management cluster for hosted cluster virtual machines (VMs) is undergoing updates or maintenance, the hosted cluster VMs can be automatically live migrated to prevent disrupting hosted cluster workloads. As a result, the management cluster can be updated without affecting the availability and operation of the KubeVirt platform hosted clusters.



IMPORTANT

The live migration of KubeVirt VMs is enabled by default provided that the VMs use **ReadWriteMany** (RWX) storage for both the root volume and the storage classes that are mapped to the **kubevirt-csi** CSI provider.

You can verify that the VMs in a node pool are capable of live migration by checking the **KubeVirtNodesLiveMigratable** condition in the **status** section of a **NodePool** object.

In the following example, the VMs cannot be live migrated because RWX storage is not used.

Example configuration where VMs cannot be live migrated

```
- lastTransitionTime: "2024-10-08T15:38:19Z"
  message: |
    3 of 3 machines are not live migratable
    Machine user-np-ngst4-gw2hz: DisksNotLiveMigratable: user-np-ngst4-gw2hz is not a live
    migratable machine: cannot migrate VMI: PVC user-np-ngst4-gw2hz-rhcos is not shared, live
    migration requires that all PVCs must be shared (using ReadWriteMany access mode)
    Machine user-np-ngst4-npq7x: DisksNotLiveMigratable: user-np-ngst4-npq7x is not a live
    migratable machine: cannot migrate VMI: PVC user-np-ngst4-npq7x-rhcos is not shared, live
    migration requires that all PVCs must be shared (using ReadWriteMany access mode)
    Machine user-np-ngst4-q5nkb: DisksNotLiveMigratable: user-np-ngst4-q5nkb is not a live
    migratable machine: cannot migrate VMI: PVC user-np-ngst4-q5nkb-rhcos is not shared, live
    migration requires that all PVCs must be shared (using ReadWriteMany access mode)
  observedGeneration: 1
  reason: DisksNotLiveMigratable
  status: "False"
  type: KubeVirtNodesLiveMigratable
```

In the next example, the VMs meet the requirements to be live migrated.

Example configuration where VMs can be live migrated

```
- lastTransitionTime: "2024-10-08T15:38:19Z"
  message: "All is well"
  observedGeneration: 1
```

```
reason: AsExpected
status: "True"
type: KubeVirtNodesLiveMigratable
```

While live migration can protect VMs from disruption in normal circumstances, events such as infrastructure node failure can result in a hard restart of any VMs that are hosted on the failed node. For live migration to be successful, the source node that a VM is hosted on must be working correctly.

When the VMs in a node pool cannot be live migrated, workload disruption might occur on the hosted cluster during maintenance on the management cluster. By default, the hosted control planes controllers try to drain the workloads that are hosted on KubeVirt VMs that cannot be live migrated before the VMs are stopped. Draining the hosted cluster nodes before stopping the VMs allows pod disruption budgets to protect workload availability within the hosted cluster.

4.3.3. Creating a hosted cluster with the KubeVirt platform

With OpenShift Container Platform 4.14 and later, you can create a cluster with KubeVirt, to include creating with an external infrastructure.

4.3.3.1. Creating a hosted cluster with the KubeVirt platform by using the CLI

To create a hosted cluster, you can use the hosted control plane command-line interface, **hcp**.

Procedure

- Create a hosted cluster with the KubeVirt platform by entering the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \①
--node-pool-replicas <node_pool_replica_count> \②
--pull-secret <path_to_pull_secret> \③
--memory <value_for_memory> \④
--cores <value_for_cpu> \⑤
--etcd-storage-class=<etcd_storage_class> \⑥
```

- Specify the name of your hosted cluster, for instance, **example**.
- Specify the node pool replica count, for example, **3**. You must specify the replica count as **0** or greater to create the same number of replicas. Otherwise, no node pools are created.
- Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- Specify a value for memory, for example, **6Gi**.
- Specify a value for CPU, for example, **2**.
- Specify the etcd storage class name, for example, **lvm-storageclass**.



NOTE

You can use the **--release-image** flag to set up the hosted cluster with a specific OpenShift Container Platform release.

A default node pool is created for the cluster with two virtual machine worker replicas according to the **--node-pool-replicas** flag.

- After a few moments, verify that the hosted control plane pods are running by entering the following command:

```
$ oc -n clusters-<hosted-cluster-name> get pods
```

Example output

NAME	READY	STATUS	RESTARTS	AGE
capi-provider-5cc7b74f47-n5gkr	1/1	Running	0	3m
catalog-operator-5f799567b7-fd6jw	2/2	Running	0	69s
certified-operators-catalog-784b9899f9-mrp6p	1/1	Running	0	66s
cluster-api-6bbc867966-l4dw1	1/1	Running	0	66s
.				
.				
.				
redhat-operators-catalog-9d5fd4d44-z8qqk	1/1	Running	0	66s

A hosted cluster that has worker nodes that are backed by KubeVirt virtual machines typically takes 10-15 minutes to be fully provisioned.

- To check the status of the hosted cluster, see the corresponding **HostedCluster** resource by entering the following command:

```
$ oc get --namespace clusters hostedclusters
```

See the following example output, which illustrates a fully provisioned **HostedCluster** object:

NAMESPACE	NAME	VERSION	KUBECONFIG	PROGRESS	AVAILABLE
PROGRESSING	MESSAGE				
clusters	example <4.x.0>	example-admin-kubeconfig	Completed	True	False
	The hosted control plane is available				

Replace **<4.x.0>** with the supported OpenShift Container Platform version that you want to use.

4.3.3.2. Creating a hosted cluster with the KubeVirt platform by using external infrastructure

By default, the HyperShift Operator hosts both the control plane pods of the hosted cluster and the KubeVirt worker VMs within the same cluster. With the external infrastructure feature, you can place the worker node VMs on a separate cluster from the control plane pods.

- The *management cluster* is the OpenShift Container Platform cluster that runs the HyperShift Operator and hosts the control plane pods for a hosted cluster.
- The *infrastructure cluster* is the OpenShift Container Platform cluster that runs the KubeVirt worker VMs for a hosted cluster.
- By default, the management cluster also acts as the infrastructure cluster that hosts VMs. However, for external infrastructure, the management and infrastructure clusters are different.

Prerequisites

- You must have a namespace on the external infrastructure cluster for the KubeVirt nodes to be hosted in.
- You must have a **kubeconfig** file for the external infrastructure cluster.

Procedure

You can create a hosted cluster by using the **hcp** command-line interface.

- To place the KubeVirt worker VMs on the infrastructure cluster, use the **--infra-kubeconfig-file** and **--infra-namespace** arguments, as shown in the following example:

```
$ hcp create cluster kubevirt \
--name <hosted-cluster-name> \ ①
--node-pool-replicas <worker-count> \ ②
--pull-secret <path-to-pull-secret> \ ③
--memory <value-for-memory> \ ④
--cores <value-for-cpu> \ ⑤
--infra-namespace=<hosted-cluster-namespace>-<hosted-cluster-name> \ ⑥
--infra-kubeconfig-file=<path-to-external-infra-kubeconfig> ⑦
```

- ① Specify the name of your hosted cluster, for instance, **example**.
- ② Specify the worker count, for example, **2**.
- ③ Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- ④ Specify a value for memory, for example, **6Gi**.
- ⑤ Specify a value for CPU, for example, **2**.
- ⑥ Specify the infrastructure namespace, for example, **clusters-example**.
- ⑦ Specify the path to your **kubeconfig** file for the infrastructure cluster, for example, **/user/name/external-infra-kubeconfig**.

After you enter that command, the control plane pods are hosted on the management cluster that the HyperShift Operator runs on, and the KubeVirt VMs are hosted on a separate infrastructure cluster.

4.3.3.3. Creating a hosted cluster by using the console

To create a hosted cluster with the KubeVirt platform by using the console, complete the following steps.

Procedure

1. Open the OpenShift Container Platform web console and log in by entering your administrator credentials.
2. In the console header, ensure that **All Clusters** is selected.
3. Click **Infrastructure > Clusters**.
4. Click **Create cluster > Red Hat OpenShift Virtualization > Hosted**

- On the **Create cluster** page, follow the prompts to enter details about the cluster and node pools.



NOTE

- If you want to use predefined values to automatically populate fields in the console, you can create a OpenShift Virtualization credential. For more information, see "Creating a credential for an on-premises environment".
- On the **Cluster details** page, the pull secret is your OpenShift Container Platform pull secret that you use to access OpenShift Container Platform resources. If you selected a OpenShift Virtualization credential, the pull secret is automatically populated.

- On the **Node pools** page, expand the **Networking options** section and configure the networking options for your node pool:

- In the **Additional networks** field, enter a network name in the format of **<namespace>/<name>**; for example, **my-namespace/network1**. The namespace and the name must be valid DNS labels. Multiple networks are supported.
- By default, the **Attach default pod network** checkbox is selected. You can clear this checkbox only if additional networks exist.

- Review your entries and click **Create**.

The **Hosted cluster** view is displayed.

- Monitor the deployment of the hosted cluster in the **Hosted cluster** view. If you do not see information about the hosted cluster, ensure that **All Clusters** is selected, and click the cluster name.
- Wait until the control plane components are ready. This process can take a few minutes.
- To view the node pool status, scroll to the **NodePool** section. The process to install the nodes takes about 10 minutes. You can also click **Nodes** to confirm whether the nodes joined the hosted cluster.

Additional resources

- To create credentials that you can reuse when you create a hosted cluster with the console, see [Creating a credential for an on-premises environment](#).
- To access the hosted cluster, see [Accessing the hosted cluster](#).

4.3.4. Configuring the default ingress and DNS for hosted control planes on OpenShift Virtualization

Every OpenShift Container Platform cluster includes a default application Ingress Controller, which must have an wildcard DNS record associated with it. By default, hosted clusters that are created by using the HyperShift KubeVirt provider automatically become a subdomain of the OpenShift Container Platform cluster that the KubeVirt virtual machines run on.

For example, your OpenShift Container Platform cluster might have the following default ingress DNS entry:

`*.apps.mgmt-cluster.example.com`

As a result, a KubeVirt hosted cluster that is named **guest** and that runs on that underlying OpenShift Container Platform cluster has the following default ingress:

`*.apps.guest.apps.mgmt-cluster.example.com`

Procedure

For the default ingress DNS to work properly, the cluster that hosts the KubeVirt virtual machines must allow wildcard DNS routes.

- You can configure this behavior by entering the following command:

```
$ oc patch ingresscontroller -n openshift-ingress-operator default \
--type=json \
-p '[{"op": "add", "path": "/spec/routeAdmission", "value": {"wildcardPolicy": "WildcardsAllowed"}}]'
```



NOTE

When you use the default hosted cluster ingress, connectivity is limited to HTTPS traffic over port 443. Plain HTTP traffic over port 80 is rejected. This limitation applies to only the default ingress behavior.

4.3.5. Customizing ingress and DNS behavior

If you do not want to use the default ingress and DNS behavior, you can configure a KubeVirt hosted cluster with a unique base domain at creation time. This option requires manual configuration steps during creation and involves three main steps: cluster creation, load balancer creation, and wildcard DNS configuration.

4.3.5.1. Deploying a hosted cluster that specifies the base domain

To create a hosted cluster that specifies a base domain, complete the following steps.

Procedure

- 1 Enter the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \ ①
--node-pool-replicas <worker_count> \ ②
--pull-secret <path_to_pull_secret> \ ③
--memory <value_for_memory> \ ④
--cores <value_for_cpu> \ ⑤
--base-domain <basedomain> ⑥
```

- 1 Specify the name of your hosted cluster.
- 2 Specify the worker count, for example, 2.
- 3 Specify the path to your pull secret, for example, /user/name/pullsecret.

- 4 Specify a value for memory, for example, **6Gi**.
- 5 Specify a value for CPU, for example, **2**.
- 6 Specify the base domain, for example, **hypershift.lab**.

As a result, the hosted cluster has an ingress wildcard that is configured for the cluster name and the base domain, for example, **.apps.example.hypershift.lab**. The hosted cluster remains in **Partial** status because after you create a hosted cluster with unique base domain, you must configure the required DNS records and load balancer.

- 2 View the status of your hosted cluster by entering the following command:

```
$ oc get --namespace clusters hostedclusters
```

Example output

NAME	VERSION	KUBECONFIG	PROGRESS	AVAILABLE
example		example-admin-kubeconfig	Partial	True
		hosted control plane is available		False
			The	

- 3 Access the cluster by entering the following commands:

```
$ hcp create kubeconfig --name <hosted_cluster_name> \
> <hosted_cluster_name>-kubeconfig
```

```
$ oc --kubeconfig <hosted_cluster_name>-kubeconfig get co
```

Example output

NAME	VERSION	AVAILABLE	PROGRESSING	DEGRADED
console	<4.x.0>	False	False	False
	SINCE	30m		
	MESSAGE			
RouteHealthAvailable: failed to GET route (https://console-openshift-console.apps.example.hypershift.lab): Get "https://console-openshift-console.apps.example.hypershift.lab": dial tcp: lookup console-openshift-console.apps.example.hypershift.lab on 172.31.0.10:53: no such host				
ingress	<4.x.0>	True	False	True
	ingress controller reports Degraded=True: DegradedConditions: One or more other status conditions indicate a degraded state: CanaryChecksSucceeding=False (CanaryChecksRepetitiveFailures: Canary route checks for the default ingress controller are failing)	28m	The "default"	

Replace **<4.x.0>** with the supported OpenShift Container Platform version that you want to use.

Next steps

To fix the errors in the output, complete the steps in "Setting up the load balancer" and "Setting up a wildcard DNS".



NOTE

If your hosted cluster is on bare metal, you might need MetallLB to set up load balancer services. For more information, see "Configuring MetallLB".

4.3.5.2. Setting up the load balancer

Set up the load balancer service that routes ingress traffic to the KubeVirt VMs and assigns a wildcard DNS entry to the load balancer IP address.

Procedure

1. A **NodePort** service that exposes the hosted cluster ingress already exists. You can export the node ports and create the load balancer service that targets those ports.

- a. Get the HTTP node port by entering the following command:

```
$ oc --kubeconfig <hosted_cluster_name>-kubeconfig get services \
-n openshift-ingress router-nodeport-default \
-o jsonpath='{.spec.ports[?(@.name=="http")].nodePort}'
```

Note the HTTP node port value to use in the next step.

- b. Get the HTTPS node port by entering the following command:

```
$ oc --kubeconfig <hosted_cluster_name>-kubeconfig get services \
-n openshift-ingress router-nodeport-default \
-o jsonpath='{.spec.ports[?(@.name=="https")].nodePort}'
```

Note the HTTPS node port value to use in the next step.

2. Create the load balancer service by entering the following command:

```
oc apply -f -
apiVersion: v1
kind: Service
metadata:
labels:
  app: <hosted_cluster_name>
  name: <hosted_cluster_name>-apps
  namespace: clusters-<hosted_cluster_name>
spec:
  ports:
    - name: https-443
      port: 443
      protocol: TCP
      targetPort: <https_node_port> ①
    - name: http-80
      port: 80
      protocol: TCP
      targetPort: <http-node-port> ②
  selector:
    kubevirt.io: virt-launcher
  type: LoadBalancer
```

- 1** Specify the HTTPS node port value that you noted in the previous step.
- 2** Specify the HTTP node port value that you noted in the previous step.

4.3.5.3. Setting up a wildcard DNS

Set up a wildcard DNS record or CNAME that references the external IP of the load balancer service.

Procedure

1. Get the external IP address by entering the following command:

```
$ oc -n clusters-<hosted_cluster_name> get service <hosted-cluster-name>-apps \
-o jsonpath='{.status.loadBalancer.ingress[0].ip}'
```

Example output

```
192.168.20.30
```

2. Configure a wildcard DNS entry that references the external IP address. View the following example DNS entry:

```
*.apps.<hosted_cluster_name>.<base_domain>.
```

The DNS entry must be able to route inside and outside of the cluster.

DNS resolutions example

```
dig +short test.apps.example.hypershift.lab
```

```
192.168.20.30
```

3. Check that hosted cluster status has moved from **Partial** to **Completed** by entering the following command:

```
$ oc get --namespace clusters hostedclusters
```

Example output

NAME	VERSION	KUBECONFIG	PROGRESS	AVAILABLE
example	<4.x.0>	example-admin-kubeconfig	Completed	True
				False
The hosted control plane is available				

Replace **<4.x.0>** with the supported OpenShift Container Platform version that you want to use.

4.3.6. Configuring MetalLB

You must install the MetalLB Operator before you configure MetalLB.

Procedure

Complete the following steps to configure MetalLB on your hosted cluster:

1. Create a **MetalLB** resource by saving the following sample YAML content in the **configure-metallb.yaml** file:

```
apiVersion: metallb.io/v1beta1
kind: MetalLB
metadata:
  name: metallb
  namespace: metallb-system
```

2. Apply the YAML content by entering the following command:

```
$ oc apply -f configure-metallb.yaml
```

Example output

```
metallb.metallb.io/metallb created
```

3. Create a **IPAddressPool** resource by saving the following sample YAML content in the **create-ip-address-pool.yaml** file:

```
apiVersion: metallb.io/v1beta1
kind: IPAddressPool
metadata:
  name: metallb
  namespace: metallb-system
spec:
  addresses:
    - 192.168.216.32-192.168.216.122 ①
```

- 1 Create an address pool with an available range of IP addresses within the node network. Replace the IP address range with an unused pool of available IP addresses in your network.

4. Apply the YAML content by entering the following command:

```
$ oc apply -f create-ip-address-pool.yaml
```

Example output

```
ipaddresspool.metallb.io/metallb created
```

5. Create a **L2Advertisement** resource by saving the following sample YAML content in the **l2advertisement.yaml** file:

```
apiVersion: metallb.io/v1beta1
kind: L2Advertisement
metadata:
  name: l2advertisement
  namespace: metallb-system
```

```
spec:
  ipAddressPools:
    - metallb
```

6. Apply the YAML content by entering the following command:

```
$ oc apply -f l2advertisement.yaml
```

Example output

```
l2advertisement.metallb.io/metallb created
```

Additional resources

- For more information about MetalLB, see [Installing the MetalLB Operator](#).

4.3.7. Configuring additional networks, guaranteed CPUs, and VM scheduling for node pools

If you need to configure additional networks for node pools, request a guaranteed CPU access for Virtual Machines (VMs), or manage scheduling of KubeVirt VMs, see the following procedures.

4.3.7.1. Adding multiple networks to a node pool

By default, nodes generated by a node pool are attached to the pod network. You can attach additional networks to the nodes by using Multus and NetworkAttachmentDefinitions.

Procedure

To add multiple networks to nodes, use the **--additional-network** argument by running the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \ ①
--node-pool-replicas <worker_node_count> \ ②
--pull-secret <path_to_pull_secret> \ ③
--memory <memory> \ ④
--cores <cpu> \ ⑤
--additional-network name:<namespace/name> \ ⑥
--additional-network name:<namespace/name>
```

- ① Specify the name of your hosted cluster, for instance, **example**.
- ② Specify your worker node count, for example, **2**.
- ③ Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- ④ Specify the memory value, for example, **8Gi**.
- ⑤ Specify the CPU value, for example, **2**.
- ⑥ Set the value of the **--additional-network** argument to **name:<namespace/name>**. Replace **<namespace/name>** with a namespace and name of your NetworkAttachmentDefinitions.

4.3.7.1.1. Using an additional network as default

You can add your additional network as a default network for the nodes by disabling the default pod network.

Procedure

- To add an additional network as default to your nodes, run the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \ ①
--node-pool-replicas <worker_node_count> \ ②
--pull-secret <path_to_pull_secret> \ ③
--memory <memory> \ ④
--cores <cpu> \ ⑤
--attach-default-network false \ ⑥
--additional-network name:<namespace>/<network_name> ⑦
```

- Specify the name of your hosted cluster, for instance, **example**.
- Specify your worker node count, for example, **2**.
- Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- Specify the memory value, for example, **8Gi**.
- Specify the CPU value, for example, **2**.
- The **--attach-default-network false** argument disables the default pod network.
- Specify the additional network that you want to add to your nodes, for example, **name:my-namespace/my-network**.

4.3.7.2. Requesting guaranteed CPU resources

By default, KubeVirt VMs might share its CPUs with other workloads on a node. This might impact performance of a VM. To avoid the performance impact, you can request a guaranteed CPU access for VMs.

Procedure

- To request guaranteed CPU resources, set the **--qos-class** argument to **Guaranteed** by running the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \ ①
--node-pool-replicas <worker_node_count> \ ②
--pull-secret <path_to_pull_secret> \ ③
--memory <memory> \ ④
--cores <cpu> \ ⑤
--qos-class Guaranteed ⑥
```

- 1 Specify the name of your hosted cluster, for instance, **example**.
- 2 Specify your worker node count, for example, **2**.
- 3 Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- 4 Specify the memory value, for example, **8Gi**.
- 5 Specify the CPU value, for example, **2**.
- 6 The **--qos-class Guaranteed** argument guarantees that the specified number of CPU resources are assigned to VMs.

4.3.7.3. Scheduling KubeVirt VMs on a set of nodes

By default, KubeVirt VMs created by a node pool are scheduled to any available nodes. You can schedule KubeVirt VMs on a specific set of nodes that has enough capacity to run the VM.

Procedure

- To schedule KubeVirt VMs within a node pool on a specific set of nodes, use the **--vm-node-selector** argument by running the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \ ①
--node-pool-replicas <worker_node_count> \ ②
--pull-secret <path_to_pull_secret> \ ③
--memory <memory> \ ④
--cores <cpu> \ ⑤
--vm-node-selector <label_key>=<label_value>,<label_key>=<label_value> \ ⑥
```

- 1 Specify the name of your hosted cluster, for instance, **example**.
- 2 Specify your worker node count, for example, **2**.
- 3 Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- 4 Specify the memory value, for example, **8Gi**.
- 5 Specify the CPU value, for example, **2**.
- 6 The **--vm-node-selector** flag defines a specific set of nodes that contains the key-value pairs. Replace **<label_key>** and **<label_value>** with the key and value of your labels respectively.

4.3.8. Scaling a node pool

You can manually scale a node pool by using the **oc scale** command.

Procedure

1. Run the following command:

```
NODEPOOL_NAME=${CLUSTER_NAME}-work
NODEPOOL_REPLICAS=5

$ oc scale nodepool/$NODEPOOL_NAME --namespace clusters \
--replicas=$NODEPOOL_REPLICAS
```

- After a few moments, enter the following command to see the status of the node pool:

```
$ oc --kubeconfig $CLUSTER_NAME-kubeconfig get nodes
```

Example output

NAME	STATUS	ROLES	AGE	VERSION
example-9jvnf	Ready	worker	97s	v1.27.4+18eadca
example-n6prw	Ready	worker	116m	v1.27.4+18eadca
example-nc6g4	Ready	worker	117m	v1.27.4+18eadca
example-thp29	Ready	worker	4m17s	v1.27.4+18eadca
example-twixns	Ready	worker	88s	v1.27.4+18eadca

4.3.8.1. Adding node pools

You can create node pools for a hosted cluster by specifying a name, number of replicas, and any additional information, such as memory and CPU requirements.

Procedure

- To create a node pool, enter the following information. In this example, the node pool has more CPUs assigned to the VMs:

```
export NODEPOOL_NAME=${CLUSTER_NAME}-extra-cpu
export WORKER_COUNT="2"
export MEM="6Gi"
export CPU="4"
export DISK="16"

$ hcp create nodepool kubevirt \
--cluster-name $CLUSTER_NAME \
--name $NODEPOOL_NAME \
--node-count $WORKER_COUNT \
--memory $MEM \
--cores $CPU \
--root-volume-size $DISK
```

- Check the status of the node pool by listing **nodepool** resources in the **clusters** namespace:

```
$ oc get nodepools --namespace clusters
```

Example output

NAME	CLUSTER	DESIRED NODES	CURRENT NODES	
AUTOSCALING	AUTOREPAIR	VERSION	UPDATINGVERSION	UPDATINGCONFIG
MESSAGE				

example	example	5	5	False	False	<4.x.0>
example-extra-cpu	example	2		False	False	True
True	Minimum availability requires 2 replicas, current 0 available					

Replace <4.x.0> with the supported OpenShift Container Platform version that you want to use.

- After some time, you can check the status of the node pool by entering the following command:

```
$ oc --kubeconfig $CLUSTER_NAME-kubeconfig get nodes
```

Example output

NAME	STATUS	ROLES	AGE	VERSION
example-9jvnf	Ready	worker	97s	v1.27.4+18eadca
example-n6prw	Ready	worker	116m	v1.27.4+18eadca
example-nc6g4	Ready	worker	117m	v1.27.4+18eadca
example-thp29	Ready	worker	4m17s	v1.27.4+18eadca
example-twxns	Ready	worker	88s	v1.27.4+18eadca
example-extra-cpu-zh9l5	Ready	worker	2m6s	v1.27.4+18eadca
example-extra-cpu-zr8mj	Ready	worker	102s	v1.27.4+18eadca

- Verify that the node pool is in the status that you expect by entering this command:

```
$ oc get nodepools --namespace clusters
```

Example output

NAME	CLUSTER	DESIRED NODES	CURRENT NODES	AUTOSCALING	AUTOREPAIR	MESSAGE	VERSION	UPDATINGVERSION	UPDATINGCONFIG
example	example	5	5	False	False	<4.x.0>			
example-extra-cpu	example	2	2	False	False	<4.x.0>			

Replace <4.x.0> with the supported OpenShift Container Platform version that you want to use.

Additional resources

- To scale down the data plane to zero, see [Scaling down the data plane to zero](#).

4.3.9. Verifying hosted cluster creation on OpenShift Virtualization

To verify that your hosted cluster was successfully created, complete the following steps.

Procedure

- Verify that the **HostedCluster** resource transitioned to the **completed** state by entering the following command:

```
$ oc get --namespace clusters hostedclusters <hosted_cluster_name>
```

Example output

NAMESPACE	NAME	VERSION	KUBECONFIG	PROGRESS	AVAILABLE
PROGRESSING	MESSAGE				
clusters	example	4.12.2	example-admin-kubeconfig	Completed	True
The hosted control plane is available					

2. Verify that all the cluster operators in the hosted cluster are online by entering the following commands:

```
$ hcp create kubeconfig --name <hosted_cluster_name> \
> <hosted_cluster_name>-kubeconfig
```

```
$ oc get co --kubeconfig=<hosted_cluster_name>-kubeconfig
```

Example output

NAME	VERSION	AVAILABLE	PROGRESSING	DEGRADED
SINCE	MESSAGE			
console	4.12.2	True	False	False 2m38s
csi-snapshot-controller	4.12.2	True	False	False 4m3s
dns	4.12.2	True	False	False 2m52s
image-registry	4.12.2	True	False	False 2m8s
ingress	4.12.2	True	False	False 22m
kube-apiserver	4.12.2	True	False	False 23m
kube-controller-manager	4.12.2	True	False	False 23m
kube-scheduler	4.12.2	True	False	False 23m
kube-storage-version-migrator	4.12.2	True	False	False 4m52s
monitoring	4.12.2	True	False	False 69s
network	4.12.2	True	False	False 4m3s
node-tuning	4.12.2	True	False	False 2m22s
openshift-apiserver	4.12.2	True	False	False 23m
openshift-controller-manager	4.12.2	True	False	False 23m
openshift-samples	4.12.2	True	False	False 2m15s
operator-lifecycle-manager	4.12.2	True	False	False 22m
operator-lifecycle-manager-catalog	4.12.2	True	False	False 23m
operator-lifecycle-manager-packageserver	4.12.2	True	False	False 23m
service-ca	4.12.2	True	False	False 4m41s
storage	4.12.2	True	False	False 4m43s

4.4. DEPLOYING HOSTED CONTROL PLANES ON NON-BARE-METAL AGENT MACHINES

You can deploy hosted control planes by configuring a cluster to function as a hosting cluster. The hosting cluster is an OpenShift Container Platform cluster where the control planes are hosted. The hosting cluster is also known as the management cluster.



IMPORTANT

Hosted control planes on non-bare-metal agent machines is a Technology Preview feature only. Technology Preview features are not supported with Red Hat production service level agreements (SLAs) and might not be functionally complete. Red Hat does not recommend using them in production. These features provide early access to upcoming product features, enabling customers to test functionality and provide feedback during the development process.

For more information about the support scope of Red Hat Technology Preview features, see [Technology Preview Features Support Scope](#).



NOTE

The management cluster is not the same thing as the *managed* cluster. A managed cluster is a cluster that the hub cluster manages.

The hosted control planes feature is enabled by default.

The multicluster engine Operator supports only the default **local-cluster** managed hub cluster. On Red Hat Advanced Cluster Management (RHACM) 2.10, you can use the **local-cluster** managed hub cluster as the hosting cluster.

A *hosted cluster* is an OpenShift Container Platform cluster with its API endpoint and control plane that are hosted on the hosting cluster. The hosted cluster includes the control plane and its corresponding data plane. You can use the multicluster engine Operator console or the **hcp** command-line interface (CLI) to create a hosted cluster.

The hosted cluster is automatically imported as a managed cluster. If you want to disable this automatic import feature, see "Disabling the automatic import of hosted clusters into multicluster engine Operator".

4.4.1. Preparing to deploy hosted control planes on non-bare-metal agent machines

As you prepare to deploy hosted control planes on bare metal, consider the following information:

- You can add agent machines as a worker node to a hosted cluster by using the Agent platform. Agent machine represents a host booted with a Discovery Image and ready to be provisioned as an OpenShift Container Platform node. The Agent platform is part of the central infrastructure management service. For more information, see [Enabling the central infrastructure management service](#).
- All hosts that are not bare metal require a manual boot with a Discovery Image ISO that the central infrastructure management provides.
- When you scale up the node pool, a machine is created for every replica. For every machine, the Cluster API provider finds and installs an Agent that is approved, is passing validations, is not currently in use, and meets the requirements that are specified in the node pool specification. You can monitor the installation of an Agent by checking its status and conditions.
- When you scale down a node pool, Agents are unbound from the corresponding cluster. Before you can reuse the Agents, you must restart them by using the Discovery image.
- When you configure storage for hosted control planes, consider the recommended etcd practices. To ensure that you meet the latency requirements, dedicate a fast storage device to all hosted control planes etcd instances that run on each control-plane node. You can use LVM

storage to configure a local storage class for hosted etcd pods. For more information, see "Recommended etcd practices" and "Persistent storage using logical volume manager storage" in the OpenShift Container Platform documentation.

4.4.1.1. Prerequisites for deploying hosted control planes on non-bare-metal agent machines

Before you deploy hosted control planes on non-bare-metal agent machines, ensure you meet the following prerequisites:

- You must have multicluster engine for Kubernetes Operator 2.5 or later installed on an OpenShift Container Platform cluster. You can install the multicluster engine Operator as an Operator from the OpenShift Container Platform OperatorHub.
- You must have at least one managed OpenShift Container Platform cluster for the multicluster engine Operator. The **local-cluster** management cluster is automatically imported. For more information about the **local-cluster**, see [Advanced configuration](#) in the Red Hat Advanced Cluster Management documentation. You can check the status of your management cluster by running the following command:

```
$ oc get managedclusters local-cluster
```

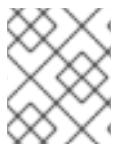
- You have enabled central infrastructure management. For more information, see [Enabling the central infrastructure management service](#) in the Red Hat Advanced Cluster Management documentation.
- You have installed the **hcp** command-line interface.
- Your hosted cluster has a cluster-wide unique name.
- You are running the management cluster and workers on the same infrastructure.

Additional resources

- [Advanced configuration](#)
- [Enabling the central infrastructure management service](#)

4.4.1.2. Firewall, port, and service requirements for non-bare-metal agent machines

You must meet the firewall and port requirements so that ports can communicate between the management cluster, the control plane, and hosted clusters.



NOTE

Services run on their default ports. However, if you use the **NodePort** publishing strategy, services run on the port that is assigned by the **NodePort** service.

Use firewall rules, security groups, or other access controls to restrict access to only required sources. Avoid exposing ports publicly unless necessary. For production deployments, use a load balancer to simplify access through a single IP address.

A hosted control plane exposes the following services on non-bare-metal agent machines:

- **APIServer**

- The **APIServer** service runs on port 6443 by default and requires ingress access for communication between the control plane components.
- If you use MetalLB load balancing, allow ingress access to the IP range that is used for load balancer IP addresses.

- **OAuthServer**

- The **OAuthServer** service runs on port 443 by default when you use the route and ingress to expose the service.
- If you use the **NodePort** publishing strategy, use a firewall rule for the **OAuthServer** service.

- **Konnectivity**

- The **Konnectivity** service runs on port 443 by default when you use the route and ingress to expose the service.
- The **Konnectivity** agent establishes a reverse tunnel to allow the control plane to access the network for the hosted cluster. The agent uses egress to connect to the **Konnectivity** server. The server is exposed by using either a route on port 443 or a manually assigned **NodePort**.
- If the cluster API server address is an internal IP address, allow access from the workload subnets to the IP address on port 6443.
- If the address is an external IP address, allow egress on port 6443 to that external IP address from the nodes.

- **Ignition**

- The **Ignition** service runs on port 443 by default when you use the route and ingress to expose the service.
- If you use the **NodePort** publishing strategy, use a firewall rule for the **Ignition** service.

You do not need the following services on non-bare-metal agent machines:

- **OVNSbDb**
- **OIDC**

4.4.1.3. Infrastructure requirements for non-bare-metal agent machines

The Agent platform does not create any infrastructure, but it has the following infrastructure requirements:

- Agents: An *Agent* represents a host that is booted with a discovery image and is ready to be provisioned as an OpenShift Container Platform node.
- DNS: The API and ingress endpoints must be routable.

Additional resources

- [Recommended etcd practices](#)
- [Persistent storage using logical volume manager storage](#)

- Disabling the automatic import of hosted clusters into multicluster engine Operator
- Manually enabling the hosted control planes feature
- Disabling the hosted control planes feature
- Configuring Ansible Automation Platform jobs to run on hosted clusters

4.4.2. Configuring DNS on non-bare-metal agent machines

The API Server for the hosted cluster is exposed as a **NodePort** service. A DNS entry must exist for **api.<hosted_cluster_name>.<basedomain>** that points to destination where the API Server can be reached.

The DNS entry can be as simple as a record that points to one of the nodes in the managed cluster that is running the hosted control plane. The entry can also point to a load balancer that is deployed to redirect incoming traffic to the ingress pods.

- If you are configuring DNS for a connected environment on an IPv4 network, see the following example DNS configuration:

```
api.example.krnl.es.    IN A 192.168.122.20
api.example.krnl.es.    IN A 192.168.122.21
api.example.krnl.es.    IN A 192.168.122.22
api-int.example.krnl.es. IN A 192.168.122.20
api-int.example.krnl.es. IN A 192.168.122.21
api-int.example.krnl.es. IN A 192.168.122.22
`*.apps.example.krnl.es. IN A 192.168.122.23
```

- If you are configuring DNS for a disconnected environment on an IPv6 network, see the following example DNS configuration:

```
api.example.krnl.es.    IN A 2620:52:0:1306::5
api.example.krnl.es.    IN A 2620:52:0:1306::6
api.example.krnl.es.    IN A 2620:52:0:1306::7
api-int.example.krnl.es. IN A 2620:52:0:1306::5
api-int.example.krnl.es. IN A 2620:52:0:1306::6
api-int.example.krnl.es. IN A 2620:52:0:1306::7
`*.apps.example.krnl.es. IN A 2620:52:0:1306::10
```

- If you are configuring DNS for a disconnected environment on a dual stack network, be sure to include DNS entries for both IPv4 and IPv6. See the following example DNS configuration:

```
host-record=api-int.hub-dual.dns.base.domain.name,192.168.126.10
host-record=api.hub-dual.dns.base.domain.name,192.168.126.10
address=/apps.hub-dual.dns.base.domain.name/192.168.126.11
dhcp-host=aa:aa:aa:aa:10:01,ocp-master-0,192.168.126.20
dhcp-host=aa:aa:aa:aa:10:02,ocp-master-1,192.168.126.21
dhcp-host=aa:aa:aa:aa:10:03,ocp-master-2,192.168.126.22
dhcp-host=aa:aa:aa:aa:10:06,ocp-installer,192.168.126.25
dhcp-host=aa:aa:aa:aa:10:07,ocp-bootstrap,192.168.126.26

host-record=api-int.hub-dual.dns.base.domain.name,2620:52:0:1306::2
host-record=api.hub-dual.dns.base.domain.name,2620:52:0:1306::2
address=/apps.hub-dual.dns.base.domain.name/2620:52:0:1306::3
```

```
dhcp-host=aa:aa:aa:aa:10:01,ocp-master-0,[2620:52:0:1306::5]
dhcp-host=aa:aa:aa:aa:10:02,ocp-master-1,[2620:52:0:1306::6]
dhcp-host=aa:aa:aa:aa:10:03,ocp-master-2,[2620:52:0:1306::7]
dhcp-host=aa:aa:aa:aa:10:06,ocp-installer,[2620:52:0:1306::8]
dhcp-host=aa:aa:aa:aa:10:07,ocp-bootstrap,[2620:52:0:1306::9]
```

4.4.3. Creating a hosted cluster on non-bare-metal agent machines by using the CLI

When you create a hosted cluster with the Agent platform, the HyperShift Operator installs the Agent Cluster API provider in the hosted control plane namespace. You can create a hosted cluster on bare metal or import one.

As you create a hosted cluster, review the following guidelines:

- Each hosted cluster must have a cluster-wide unique name. A hosted cluster name cannot be the same as any existing managed cluster in order for multicluster engine Operator to manage it.
- Do not use **clusters** as a hosted cluster name.
- A hosted cluster cannot be created in the namespace of a multicluster engine Operator managed cluster.

Procedure

1. Create the hosted control plane namespace by entering the following command:

```
$ oc create ns <hosted_cluster_namespace>-<hosted_cluster_name> ①
```

- ① Replace **<hosted_cluster_namespace>** with your hosted cluster namespace name, for example, **clusters**. Replace **<hosted_cluster_name>** with your hosted cluster name.

2. Create a hosted cluster by entering the following command:

```
$ hcp create cluster agent \
--name=<hosted_cluster_name> ①
--pull-secret=<path_to_pull_secret> ②
--agent-namespace=<hosted_control_plane_namespace> ③
--base-domain=<basedomain> ④
--api-server-address=api.<hosted_cluster_name>.<basedomain> ⑤
--etcd-storage-class=<etcd_storage_class> ⑥
--ssh-key <path_to_ssh_key> ⑦
--namespace <hosted_cluster_namespace> ⑧
--control-plane-availability-policy HighlyAvailable ⑨
--release-image=quay.io/openshift-release-dev/ocp-release:<ocp_release> ⑩
--node-pool-replicas <node_pool_replica_count> ⑪
```

- ① Specify the name of your hosted cluster, for instance, **example**.
- ② Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- ③ Specify your hosted control plane namespace, for example, **clusters-example**. Ensure that agents are available in this namespace by using the **oc get agent -n <hosted-control-**

plane-namespace> command.

- 4 Specify your base domain, for example, **krnl.es**.
- 5 The **--api-server-address** flag defines the IP address that is used for the Kubernetes API communication in the hosted cluster. If you do not set the **--api-server-address** flag, you must log in to connect to the management cluster.
- 6 Verify that you have a default storage class configured for your cluster. Otherwise, you might end up with pending PVCs. Specify the etcd storage class name, for example, **lvm-storageclass**.
- 7 Specify the path to your SSH public key. The default file path is **~/.ssh/id_rsa.pub**.
- 8 Specify your hosted cluster namespace.
- 9 Specify the availability policy for the hosted control plane components. Supported options are **SingleReplica** and **HighlyAvailable**. The default value is **HighlyAvailable**.
- 10 Specify the supported OpenShift Container Platform version that you want to use, for example, **4.18.0-multi**.
- 11 Specify the node pool replica count, for example, **3**. You must specify the replica count as **0** or greater to create the same number of replicas. Otherwise, no node pools are created.

Verification

- After a few moments, verify that your hosted control plane pods are up and running by entering the following command:

```
$ oc -n <hosted_cluster_namespace>-<hosted_cluster_name> get pods
```

Example output

NAME	READY	STATUS	RESTARTS	AGE
catalog-operator-6cd867cc7-phb2q	2/2	Running	0	2m50s
control-plane-operator-f6b4c8465-4k5dh	1/1	Running	0	4m32s

Additional resources

- [Manually importing a hosted cluster](#)

4.4.3.1. Creating a hosted cluster on non-bare-metal agent machines by using the web console

You can create a hosted cluster on non-bare-metal agent machines by using the OpenShift Container Platform web console.

Prerequisites

- You have access to the cluster with **cluster-admin** privileges.
- You have access to the OpenShift Container Platform web console.

Procedure

1. Open the OpenShift Container Platform web console and log in by entering your administrator credentials.
2. In the console header, select **All Clusters**.
3. Click **Infrastructure → Clusters**.
4. Click **Create cluster Host inventory → Hosted control plane**
The **Create cluster** page is displayed.
5. On the **Create cluster** page, follow the prompts to enter details about the cluster, node pools, networking, and automation.

As you enter details about the cluster, you might find the following tips useful:

- If you want to use predefined values to automatically populate fields in the console, you can create a host inventory credential. For more information, see *Creating a credential for an on-premises environment*.
- On the **Cluster details** page, the pull secret is your OpenShift Container Platform pull secret that you use to access OpenShift Container Platform resources. If you selected a host inventory credential, the pull secret is automatically populated.
- On the **Node pools** page, the namespace contains the hosts for the node pool. If you created a host inventory by using the console, the console creates a dedicated namespace.
- On the **Networking** page, you select an API server publishing strategy. The API server for the hosted cluster can be exposed either by using an existing load balancer or as a service of the **NodePort** type. A DNS entry must exist for the **api.<hosted_cluster_name>.<basedomain>** setting that points to the destination where the API server can be reached. This entry can be a record that points to one of the nodes in the management cluster or a record that points to a load balancer that redirects incoming traffic to the Ingress pods.

1. Review your entries and click **Create**.

The **Hosted cluster** view is displayed.

1. Monitor the deployment of the hosted cluster in the **Hosted cluster** view. If you do not see information about the hosted cluster, ensure that **All Clusters** is selected, and click the cluster name. Wait until the control plane components are ready. This process can take a few minutes.
2. To view the node pool status, scroll to the **NodePool** section. The process to install the nodes takes about 10 minutes. You can also click **Nodes** to confirm whether the nodes joined the hosted cluster.

Next steps

- To access the web console, see [Accessing the web console](#).

4.4.3.2. Creating a hosted cluster on bare metal by using a mirror registry

You can use a mirror registry to create a hosted cluster on bare metal by specifying the **--image-content-sources** flag in the **hcp create cluster** command.

Procedure

1. Create a YAML file to define Image Content Source Policies (ICSP). See the following example:

```
- mirrors:
  - brew.registry.redhat.io
    source: registry.redhat.io
- mirrors:
  - brew.registry.redhat.io
    source: registry.stage.redhat.io
- mirrors:
  - brew.registry.redhat.io
    source: registry-proxy.engineering.redhat.com
```

2. Save the file as **icsp.yaml**. This file contains your mirror registries.
3. To create a hosted cluster by using your mirror registries, run the following command:

```
$ hcp create cluster agent \
--name=<hosted_cluster_name> \ ①
--pull-secret=<path_to_pull_secret> \ ②
--agent-namespace=<hosted_control_plane_namespace> \ ③
--base-domain=<basedomain> \ ④
--api-server-address=api.<hosted_cluster_name>.<basedomain> \ ⑤
--image-content-sources icsp.yaml \ ⑥
--ssh-key <path_to_ssh_key> \ ⑦
--namespace <hosted_cluster_namespace> \ ⑧
--release-image=quay.io/openshift-release-dev/ocp-release:<ocp_release_image> \ ⑨
```

- 1 Specify the name of your hosted cluster, for instance, **example**.
- 2 Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- 3 Specify your hosted control plane namespace, for example, **clusters-example**. Ensure that agents are available in this namespace by using the **oc get agent -n <hosted-control-plane-namespace>** command.
- 4 Specify your base domain, for example, **krnl.es**.
- 5 The **--api-server-address** flag defines the IP address that is used for the Kubernetes API communication in the hosted cluster. If you do not set the **--api-server-address** flag, you must log in to connect to the management cluster.
- 6 Specify the **icsp.yaml** file that defines ICSP and your mirror registries.
- 7 Specify the path to your SSH public key. The default file path is **~/.ssh/id_rsa.pub**.
- 8 Specify your hosted cluster namespace.
- 9 Specify the supported OpenShift Container Platform version that you want to use, for example, **4.18.0-multi**. If you are using a disconnected environment, replace **<ocp_release_image>** with the digest image. To extract the OpenShift Container Platform release image digest, see *Extracting the OpenShift Container Platform release image digest*.

Next steps

- To create credentials that you can reuse when you create a hosted cluster with the console, see [Creating a credential for an on-premises environment](#).
- To access a hosted cluster, see [Accessing the hosted cluster](#).
- To add hosts to the host inventory by using the Discovery Image, see [Adding hosts to the host inventory by using the Discovery Image](#).
- To extract the OpenShift Container Platform release image digest, see [Extracting the OpenShift Container Platform release image digest](#).

4.4.4. Verifying hosted cluster creation on non-bare-metal agent machines

After the deployment process is complete, you can verify that the hosted cluster was created successfully. Follow these steps a few minutes after you create the hosted cluster.

Procedure

1. Obtain the **kubeconfig** file for your new hosted cluster by entering the following command:

```
$ oc extract -n <hosted_cluster_namespace> \
secret/<hosted_cluster_name>-admin-kubeconfig --to=- \
> kubeconfig-<hosted_cluster_name>
```

2. Use the **kubeconfig** file to view the cluster Operators of the hosted cluster. Enter the following command:

```
$ oc get co --kubeconfig=kubeconfig-<hosted_cluster_name>
```

Example output

NAME SINCE MESSAGE	VERSION	AVAILABLE	PROGRESSING	DEGRADED
console	4.10.26	True	False	False
csi-snapshot-controller	4.10.26	True	False	False
dns	4.10.26	True	False	False

3. View the running pods on your hosted cluster by entering the following command:

```
$ oc get pods -A --kubeconfig=kubeconfig-<hosted_cluster_name>
```

Example output

NAMESPACE	STATUS	RESTARTS	AGE	NAME	READY
kube-system	Running	0	3m52s	konnectivity-agent-khlqv	0/1
openshift-cluster-samples-operator	2/2	Running	0	cluster-samples-operator-6b5bcb9dff-kpnbc	
openshift-monitoring				alertmanager-main-0	6/6

Running	0	100s	
openshift-monitoring			openshift-state-metrics-677b9fb74f-qqp6g
3/3 Running	0	104s	

4.5. DEPLOYING HOSTED CONTROL PLANES ON IBM Z

You can deploy hosted control planes by configuring a cluster to function as a management cluster. The management cluster is the OpenShift Container Platform cluster where the control planes are hosted. The management cluster is also known as the *hosting* cluster.



NOTE

The *management* cluster is not the *managed* cluster. A managed cluster is a cluster that the hub cluster manages.

You can convert a managed cluster to a management cluster by using the **hypershift** add-on to deploy the HyperShift Operator on that cluster. Then, you can start to create the hosted cluster.

The multicluster engine Operator supports only the default **local-cluster**, which is a hub cluster that is managed, and the hub cluster as the management cluster.

To provision hosted control planes on bare metal, you can use the Agent platform. The Agent platform uses the central infrastructure management service to add worker nodes to a hosted cluster. For more information, see "Enabling the central infrastructure management service".

Each IBM Z system host must be started with the PXE images provided by the central infrastructure management. After each host starts, it runs an Agent process to discover the details of the host and completes the installation. An Agent custom resource represents each host.

When you create a hosted cluster with the Agent platform, HyperShift Operator installs the Agent Cluster API provider in the hosted control plane namespace.

4.5.1. Prerequisites to configure hosted control planes on IBM Z

- The multicluster engine for Kubernetes Operator version 2.5 or later must be installed on an OpenShift Container Platform cluster. You can install multicluster engine Operator as an Operator from the OpenShift Container Platform OperatorHub.
- The multicluster engine Operator must have at least one managed OpenShift Container Platform cluster. The **local-cluster** is automatically imported in multicluster engine Operator 2.5 and later. For more information about the **local-cluster**, see *Advanced configuration* in the Red Hat Advanced Cluster Management documentation. You can check the status of your hub cluster by running the following command:

```
$ oc get managedclusters local-cluster
```

- You need a hosting cluster with at least three worker nodes to run the HyperShift Operator.
- You need to enable the central infrastructure management service. For more information, see *Enabling the central infrastructure management service*.
- You need to install the hosted control plane command-line interface. For more information, see *Installing the hosted control plane command-line interface*.

Additional resources

- [Advanced configuration](#)
- [Enabling the central infrastructure management service](#)
- [Installing the hosted control planes command-line interface](#)
- [Enabling or disabling the hosted control planes feature](#)

4.5.2. IBM Z infrastructure requirements

The Agent platform does not create any infrastructure, but requires the following resources for infrastructure:

- Agents: An *Agent* represents a host that is booted with a discovery image, or PXE image and is ready to be provisioned as an OpenShift Container Platform node.
- DNS: The API and Ingress endpoints must be routable.

The hosted control planes feature is enabled by default. If you disabled the feature and want to manually enable it, or if you need to disable the feature, see *Enabling or disabling the hosted control planes feature*.

Additional resources

- [Enabling or disabling the hosted control planes feature](#)

4.5.3. DNS configuration for hosted control planes on IBM Z

The API server for the hosted cluster is exposed as a **NodePort** service. A DNS entry must exist for the **api.<hosted_cluster_name>.<base_domain>** that points to the destination where the API server is reachable.

The DNS entry can be as simple as a record that points to one of the nodes in the managed cluster that is running the hosted control plane.

The entry can also point to a load balancer deployed to redirect incoming traffic to the Ingress pods.

See the following example of a DNS configuration:

```
$ cat /var/named/<example.krnl.es.zone>
```

Example output

```
$ TTL 900
@ IN SOA bastion.example.krnl.es.com. hostmaster.example.krnl.es.com. (
    2019062002
    1D 1H 1W 3H )
IN NS bastion.example.krnl.es.com.
;
;
api          IN A 1xx.2x.2xx.1xx ①
api-int      IN A 1xx.2x.2xx.1xx
;
```

```

;
*.apps    IN A 1xx.2x.2xx.1xx
;
:EOF

```

- 1** The record refers to the IP address of the API load balancer that handles ingress and egress traffic for hosted control planes.

For IBM z/VM, add IP addresses that correspond to the IP address of the agent.

```

compute-0      IN A 1xx.2x.2xx.1yy
compute-1      IN A 1xx.2x.2xx.1yy

```

4.5.4. Creating a hosted cluster on bare metal

When you create a hosted cluster with the Agent platform, HyperShift installs the Agent Cluster API provider in the hosted control plane namespace. You can create a hosted cluster on bare metal or import one.

As you create a hosted cluster, keep the following guidelines in mind:

- Each hosted cluster must have a cluster-wide unique name. A hosted cluster name cannot be the same as any existing managed cluster in order for multicluster engine Operator to manage it.
- Do not use **clusters** as a hosted cluster name.
- A hosted cluster cannot be created in the namespace of a multicluster engine Operator managed cluster.
- The most common service publishing strategy is to expose services through a load balancer. That strategy is the preferred method for exposing the Kubernetes API server. If you create a hosted cluster by using the web console or by using Red Hat Advanced Cluster Management, to set a publishing strategy for a service besides the Kubernetes API server, you must manually specify the **servicePublishingStrategy** information in the **HostedCluster** custom resource.

Procedure

1. Create the hosted control plane namespace by entering the following command:

```
$ oc create ns <hosted_cluster_namespace>-<hosted_cluster_name>
```

Replace **<hosted_cluster_namespace>** with your hosted cluster namespace name, for example, **clusters**. Replace **<hosted_cluster_name>** with your hosted cluster name.

2. Verify that you have a default storage class configured for your cluster. Otherwise, you might see pending PVCs. Run the following command:

```
$ hcp create cluster agent \
--name=<hosted_cluster_name> \b1
--pull-secret=<path_to_pull_secret> \b2
--agent-namespace=<hosted_control_plane_namespace> \b3
--base-domain=<basedomain> \b4
```

```
--api-server-address=api.<hosted_cluster_name>.<basedomain> \ ⑤
--etcd-storage-class=<etcd_storage_class> \ ⑥
--ssh-key <path_to_ssh_public_key> \ ⑦
--namespace <hosted_cluster_namespace> \ ⑧
--control-plane-availability-policy HighlyAvailable \ ⑨
--release-image=quay.io/openshift-release-dev/ocp-release:<ocp_release_image> \ ⑩
--node-pool-replicas <node_pool_replica_count> ⑪
```

- ① Specify the name of your hosted cluster, for instance, **example**.
 - ② Specify the path to your pull secret, for example, **/user/name/pullsecret**.
 - ③ Specify your hosted control plane namespace, for example, **clusters-example**. Ensure that agents are available in this namespace by using the **oc get agent -n <hosted_control_plane_namespace>** command.
 - ④ Specify your base domain, for example, **krnl.es**.
 - ⑤ The **--api-server-address** flag defines the IP address that is used for the Kubernetes API communication in the hosted cluster. If you do not set the **--api-server-address** flag, you must log in to connect to the management cluster.
 - ⑥ Specify the etcd storage class name, for example, **lvm-storageclass**.
 - ⑦ Specify the path to your SSH public key. The default file path is **~/.ssh/id_rsa.pub**.
 - ⑧ Specify your hosted cluster namespace.
 - ⑨ Specify the availability policy for the hosted control plane components. Supported options are **SingleReplica** and **HighlyAvailable**. The default value is **HighlyAvailable**.
 - ⑩ Specify the supported OpenShift Container Platform version that you want to use, for example, **4.18.0-multi**. If you are using a disconnected environment, replace **<ocp_release_image>** with the digest image. To extract the OpenShift Container Platform release image digest, see *Extracting the OpenShift Container Platform release image digest*.
 - ⑪ Specify the node pool replica count, for example, ③. You must specify the replica count as **0** or greater to create the same number of replicas. Otherwise, no node pools are created.
3. After a few moments, verify that your hosted control plane pods are up and running by entering the following command:

```
$ oc -n <hosted_cluster_namespace>-<hosted_cluster_name> get pods
```

Example output

NAME	READY	STATUS	RESTARTS	AGE
capi-provider-7dcf5fc4c4-nr9sq	1/1	Running	0	4m32s
catalog-operator-6cd867cc7-phb2q	2/2	Running	0	2m50s
certified-operators-catalog-884c756c4-zdt64	1/1	Running	0	2m51s
cluster-api-f75d86f8c-56wfz	1/1	Running	0	4m32s

Additional resources

- [Creating a hosted cluster on bare metal by using the console](#)

4.5.5. Creating an InfraEnv resource for hosted control planes on IBM Z

An **InfraEnv** is an environment where hosts that are booted with PXE images can join as agents. In this case, the agents are created in the same namespace as your hosted control plane.

Procedure

1. Create a YAML file to contain the configuration. See the following example:

```
apiVersion: agent-install.openshift.io/v1beta1
kind: InfraEnv
metadata:
  name: <hosted_cluster_name>
  namespace: <hosted_control_plane_namespace>
spec:
  cpuArchitecture: s390x
  pullSecretRef:
    name: pull-secret
  sshAuthorizedKey: <ssh_public_key>
```

2. Save the file as **infraenv-config.yaml**.

3. Apply the configuration by entering the following command:

```
$ oc apply -f infraenv-config.yaml
```

4. To fetch the URL to download the PXE images, such as, **initrd.img**, **kernel.img**, or **rootfs.img**, which allows IBM Z machines to join as agents, enter the following command:

```
$ oc -n <hosted_control_plane_namespace> get InfraEnv <hosted_cluster_name> -o json
```

4.5.6. Adding IBM Z agents to the InfraEnv resource

To attach compute nodes to a hosted control plane, create agents that help you to scale the node pool. Adding agents in an IBM Z environment requires additional steps, which are described in detail in this section.

Unless stated otherwise, these procedures apply to both z/VM and RHEL KVM installations on IBM Z and IBM LinuxONE.

4.5.6.1. Adding IBM Z KVM as agents

For IBM Z with KVM, run the following command to start your IBM Z environment with the downloaded PXE images from the **InfraEnv** resource. After the Agents are created, the host communicates with the Assisted Service and registers in the same namespace as the **InfraEnv** resource on the management cluster.

Procedure

1. Run the following command:

```

virt-install \
    --name "<vm_name>" \ ①
    --autostart \
    --ram=16384 \
    --cpu host \
    --vcpus=4 \
    --location "<path_to_kernel_initrd_image>,kernel=kernel.img,initrd=initrd.img" \ ②
    --disk <qcow_image_path> \ ③
    --network network:macvtap-net,mac=<mac_address> \ ④
    --graphics none \
    --noautoconsole \
    --wait=-1
    --extra-args "rd.neednet=1 nameserver=<nameserver>
coreos.live.rootfs_url=http://<http_server>/rootfs.img random.trust_cpu=on
rd.luks.options=discard ignition.firstboot ignition.platform.id=metal console=tty1
console=ttyS1,115200n8 coreos.inst.persistent-kargs=console=tty1
console=ttyS1,115200n8" ⑤

```

- ① Specify the name of the virtual machine.
- ② Specify the location of the **kernel_initrd_image** file.
- ③ Specify the disk image path.
- ④ Specify the Mac address.
- ⑤ Specify the server name of the agents.

2. For ISO boot, download ISO from the **InfraEnv** resource and boot the nodes by running the following command:

```

virt-install \
    --name "<vm_name>" \ ①
    --autostart \
    --memory=16384 \
    --cpu host \
    --vcpus=4 \
    --network network:macvtap-net,mac=<mac_address> \ ②
    --cdrom "<path_to_image.iso>" \ ③
    --disk <qcow_image_path> \
    --graphics none \
    --noautoconsole \
    --os-variant <os_version> \ ④
    --wait=-1

```

- ① Specify the name of the virtual machine.
- ② Specify the Mac address.
- ③ Specify the location of the **image.iso** file.
- ④ Specify the operating system version that you are using.

4.5.6.2. Adding IBM Z LPAR as agents

You can add the Logical Partition (LPAR) on IBM Z or IBM LinuxONE as a compute node to a hosted control plane.

Procedure

1. Create a boot parameter file for the agents:

Example parameter file

```
rd.neednet=1 cio_ignore=all,!condev \
console=ttySclp0 \
ignition.firstboot ignition.platform.id=metal
coreos.live.rootfs_url=http://<http_server>/rhcos-<version>-live-rootfs.<architecture>.img \ ①
coreos.inst.persistent-kargs=console=ttySclp0
ip=<ip>:<gateway>:<netmask>:<hostname>::none nameserver=<dns> \ ②
rd.znet=qeth,<network_adaptor_range>,layer2=1
rd.<disk_type>=<adapter> \ ③
zfcp.allow_lun_scan=0
ai.ip_cfg_override=1 \ ④
random.trust_cpu=on rd.luks.options=discard
```

- ① For the **coreos.live.rootfs_url** artifact, specify the matching **rootfs** artifact for the **kernel** and **initramfs** that you are starting. Only HTTP and HTTPS protocols are supported.
- ② For the **ip** parameter, manually assign the IP address, as described in *Installing a cluster with z/VM on IBM Z and IBM LinuxONE*.
- ③ For installations on DASD-type disks, use **rd.dasd** to specify the DASD where Red Hat Enterprise Linux CoreOS (RHCOS) is to be installed. For installations on FCP-type disks, use **rd.zfcp=<adapter>,<wwpn>,<lun>** to specify the FCP disk where RHCOS is to be installed.
- ④ Specify this parameter when you use an Open Systems Adapter (OSA) or HiperSockets.

2. Download the **.ins** and **initrd.img.addrsize** files from the **InfraEnv** resource.

By default, the URL for the **.ins** and **initrd.img.addrsize** files is not available in the **InfraEnv** resource. You must edit the URL to fetch those artifacts.

- a. Update the kernel URL endpoint to include **ins-file** by running the followign command:

```
$ curl -k -L -o generic.ins "< url for ins-file >"
```

Example URL

```
https://.../boot-artifacts/ins-file?arch=s390x&version=4.17.0
```

- b. Update the **initrd** URL endpoint to include **s390x-initrd-addrsize**:

Example URL

```
https://.../s390x-initrd-addrsize?api_key=<api-key>&arch=s390x&version=4.17.0
```

3. Transfer the **initrd**, **kernel**, **generic.ins**, and **initrd.img.addrsiz**e parameter files to the file server. For more information about how to transfer the files with FTP and boot, see "Installing in an LPAR".
4. Start the machine.
5. Repeat the procedure for all other machines in the cluster.

Additional resources

- [Installing in an LPAR](#)

4.5.6.3. Adding IBM z/VM as agents

If you want to use a static IP for z/VM guest, you must configure the **NMStateConfig** attribute for the z/VM agent so that the IP parameter persists in the second start.

Complete the following steps to start your IBM Z environment with the downloaded PXE images from the **InfraEnv** resource. After the Agents are created, the host communicates with the Assisted Service and registers in the same namespace as the **InfraEnv** resource on the management cluster.

Procedure

1. Update the parameter file to add the **rootfs_url**, **network_adaptor** and **disk_type** values.

Example parameter file

```
rd.neednet=1 cio_ignore=all,!condev \
console=ttySCLP0 \
ignition.firstboot ignition.platform.id=metal \
coreos.live.rootfs_url=http://<http_server>/rhcos-<version>-live-rootfs.<architecture>.img \①
coreos.inst.persistent-kargs=console=ttySCLP0
ip=<ip>:<gateway>:<netmask>:<hostname>::none nameserver=<dns> \②
rd.znet=qeth,<network_adaptor_range>,layer2=1
rd.<disk_type>=<adapter> \③
zfcp.allow_lun_scan=0
ai.ip_cfg_override=1 \④
```

- 1 For the **coreos.live.rootfs_url** artifact, specify the matching **rootfs** artifact for the **kernel** and **initramfs** that you are starting. Only HTTP and HTTPS protocols are supported.
- 2 For the **ip** parameter, manually assign the IP address, as described in *Installing a cluster with z/VM on IBM Z and IBM LinuxONE*.
- 3 For installations on DASD-type disks, use **rd.dasd** to specify the DASD where Red Hat Enterprise Linux CoreOS (RHCOS) is to be installed. For installations on FCP-type disks, use **rd.zfcp=<adapter>,<wwpn>,<lun>** to specify the FCP disk where RHCOS is to be installed.
- 4 Specify this parameter when you use an Open Systems Adapter (OSA) or HiperSockets.

2. Move **initrd**, kernel images, and the parameter file to the guest VM by running the following commands:

```
vmur pun -r -u -N kernel.img $INSTALLERKERNELLOCATION/<image name>
vmur pun -r -u -N generic.parm $PARMFILELOCATION/paramfilename
vmur pun -r -u -N initrd.img $INSTALLERINITRAMFSLOCATION/<image name>
```

- Run the following command from the guest VM console:

```
cp ipl c
```

- To list the agents and their properties, enter the following command:

```
$ oc -n <hosted_control_plane_namespace> get agents
```

Example output

NAME	CLUSTER	APPROVED	ROLE	STAGE
50c23cda-cedc-9bbd-bcf1-9b3a5c75804d			auto-assign	
5e498cd3-542c-e54f-0c58-ed43e28b568a			auto-assign	

- Run the following command to approve the agent.

```
$ oc -n <hosted_control_plane_namespace> patch agent \
50c23cda-cedc-9bbd-bcf1-9b3a5c75804d -p \
'{"spec":{"installation_disk_id":"/dev/sda","approved":true,"hostname":"worker-zvm-0.hostedn.example.com"}}' \①
--type merge
```

- Optionally, you can set the agent ID **<installation_disk_id>** and **<hostname>** in the specification.

- Run the following command to verify that the agents are approved:

```
$ oc -n <hosted_control_plane_namespace> get agents
```

Example output

NAME	CLUSTER	APPROVED	ROLE	STAGE
50c23cda-cedc-9bbd-bcf1-9b3a5c75804d		true	auto-assign	
5e498cd3-542c-e54f-0c58-ed43e28b568a		true	auto-assign	

4.5.7. Scaling the NodePool object for a hosted cluster on IBM Z

The **NodePool** object is created when you create a hosted cluster. By scaling the **NodePool** object, you can add more compute nodes to the hosted control plane.

When you scale up a node pool, a machine is created. The Cluster API provider finds an Agent that is approved, is passing validations, is not currently in use, and meets the requirements that are specified in the node pool specification. You can monitor the installation of an Agent by checking its status and conditions.

Procedure

- Run the following command to scale the **NodePool** object to two nodes:

```
$ oc -n <clusters_namespace> scale nodepool <nodepool_name> --replicas 2
```

The Cluster API agent provider randomly picks two agents that are then assigned to the hosted cluster. Those agents go through different states and finally join the hosted cluster as OpenShift Container Platform nodes. The agents pass through the transition phases in the following order:

- **binding**
- **discovering**
- **insufficient**
- **installing**
- **installing-in-progress**
- **added-to-existing-cluster**

- Run the following command to see the status of a specific scaled agent:

```
$ oc -n <hosted_control_plane_namespace> get agent -o \
jsonpath='{range .items[*]}BMH: {@.metadata.labels.agent-install\\.openshift\\.io/bmh} \
Agent: {@.metadata.name} State: {@.status.debugInfo.state}{"\n"}{end}'
```

Example output

```
BMH: Agent: 50c23cda-cedc-9bbd-bcf1-9b3a5c75804d State: known-unbound
BMH: Agent: 5e498cd3-542c-e54f-0c58-ed43e28b568a State: insufficient
```

- Run the following command to see the transition phases:

```
$ oc -n <hosted_control_plane_namespace> get agent
```

Example output

NAME	CLUSTER	APPROVED	ROLE	STAGE
50c23cda-cedc-9bbd-bcf1-9b3a5c75804d	hosted-forwarder	true		auto-assign
5e498cd3-542c-e54f-0c58-ed43e28b568a		true		auto-assign
da503cf1-a347-44f2-875c-4960ddb04091	hosted-forwarder	true		auto-assign

- Run the following command to generate the **kubeconfig** file to access the hosted cluster:

```
$ hcp create kubeconfig \
--namespace <clusters_namespace> \
--name <hosted_cluster_namespace> > <hosted_cluster_name>.kubeconfig
```

- After the agents reach the **added-to-existing-cluster** state, verify that you can see the OpenShift Container Platform nodes by entering the following command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get nodes
```

Example output

NAME	STATUS	ROLES	AGE	VERSION
worker-zvm-0.hostedn.example.com	Ready	worker	5m41s	v1.24.0+3882f8f
worker-zvm-1.hostedn.example.com	Ready	worker	6m3s	v1.24.0+3882f8f

Cluster Operators start to reconcile by adding workloads to the nodes.

6. Enter the following command to verify that two machines were created when you scaled up the **NodePool** object:

```
$ oc -n <hosted_control_plane_namespace> get machine.cluster.x-k8s.io
```

Example output

NAME	CLUSTER	NODENAME	PROVIDERID	PHASE	AGE
hosted-forwarder-79558597ff-5tbqp	hosted-forwarder-crqq5	worker-zvm-0.hostedn.example.com	agent://50c23cda-cedc-9bbd-bcf1-9b3a5c75804d	Running	41h 4.15.0
hosted-forwarder-79558597ff-lfjfk	hosted-forwarder-crqq5	worker-zvm-1.hostedn.example.com	agent://5e498cd3-542c-e54f-0c58-ed43e28b568a	Running	41h 4.15.0

7. Run the following command to check the cluster version:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get clusterversion,co
```

Example output

NAME	VERSION	AVAILABLE	PROGRESSING	SINCE
clusterversion.config.openshift.io/version	4.15.0-ec.2	True	False	40h Cluster version is 4.15.0-ec.2

8. Run the following command to check the cluster operator status:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get clusteroperators
```

For each component of your cluster, the output shows the following cluster operator statuses: **NAME**, **VERSION**, **AVAILABLE**, **PROGRESSING**, **DEGRADED**, **SINCE**, and **MESSAGE**.

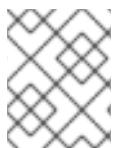
For an output example, see *Initial Operator configuration*.

Additional resources

- [Initial Operator configuration](#)

4.6. DEPLOYING HOSTED CONTROL PLANES ON IBM POWER

You can deploy hosted control planes by configuring a cluster to function as a hosting cluster. The hosting cluster is an OpenShift Container Platform cluster where the control planes are hosted. The hosting cluster is also known as the *management* cluster.



NOTE

The *management* cluster is not the *managed* cluster. A managed cluster is a cluster that the hub cluster manages.

The multicluster engine Operator supports only the default **local-cluster**, which is a hub cluster that is managed, and the hub cluster as the hosting cluster.

To provision hosted control planes on bare metal, you can use the Agent platform. The Agent platform uses the central infrastructure management service to add worker nodes to a hosted cluster. For more information, see "Enabling the central infrastructure management service".

Each IBM Power host must be started with a Discovery Image that the central infrastructure management provides. After each host starts, it runs an Agent process to discover the details of the host and completes the installation. An Agent custom resource represents each host.

When you create a hosted cluster with the Agent platform, HyperShift installs the Agent Cluster API provider in the hosted control plane namespace.

4.6.1. Prerequisites to configure hosted control planes on IBM Power

- The multicluster engine for Kubernetes Operator version 2.7 and later installed on an OpenShift Container Platform cluster. The multicluster engine Operator is automatically installed when you install Red Hat Advanced Cluster Management (RHACM). You can also install the multicluster engine Operator without RHACM as an Operator from the OpenShift Container Platform OperatorHub.
- The multicluster engine Operator must have at least one managed OpenShift Container Platform cluster. The **local-cluster** managed hub cluster is automatically imported in the multicluster engine Operator version 2.7 and later. For more information about **local-cluster**, see *Advanced configuration* in the RHACM documentation. You can check the status of your hub cluster by running the following command:

```
$ oc get managedclusters local-cluster
```

- You need a hosting cluster with at least 3 worker nodes to run the HyperShift Operator.
- You need to enable the central infrastructure management service. For more information, see "Enabling the central infrastructure management service".
- You need to install the hosted control plane command-line interface. For more information, see "Installing the hosted control plane command-line interface".

The hosted control planes feature is enabled by default. If you disabled the feature and want to manually enable it, see "Manually enabling the hosted control planes feature". If you need to disable the feature, see "Disabling the hosted control planes feature".

Additional resources

- [Advanced configuration](#)

- Enabling the central infrastructure management service
- Installing the hosted control planes command-line interface
- Manually enabling the hosted control planes feature
- Disabling the hosted control planes feature

4.6.2. IBM Power infrastructure requirements

The Agent platform does not create any infrastructure, but requires the following resources for infrastructure:

- Agents: An *Agent* represents a host that is booted with a discovery image and is ready to be provisioned as an OpenShift Container Platform node.
- DNS: The API and Ingress endpoints must be routable.

4.6.3. DNS configuration for hosted control planes on IBM Power

The API server for the hosted cluster is exposed. A DNS entry must exist for the **api**.

<hosted_cluster_name>.<basedomain> entry that points to the destination where the API server is reachable.

The DNS entry can be as simple as a record that points to one of the nodes in the managed cluster that is running the hosted control plane.

The entry can also point to a load balancer that is deployed to redirect incoming traffic to the ingress pods.

See the following example of a DNS configuration:

```
$ cat /var/named/<example.krnl.es.zone>
```

Example output

```
$ TTL 900
@ IN SOA bastion.example.krnl.es.com. hostmaster.example.krnl.es.com. (
    2019062002
    1D 1H 1W 3H )
IN NS bastion.example.krnl.es.com.
;
;
api      IN A 1xx.2x.2xx.1xx ①
api-int   IN A 1xx.2x.2xx.1xx
;
;
*.apps.<hosted-cluster-name>.<basedomain>    IN A 1xx.2x.2xx.1xx
;
;EOF
```

① The record refers to the IP address of the API load balancer that handles ingress and egress traffic for hosted control planes.

For IBM Power, add IP addresses that correspond to the IP address of the agent.

Example configuration

```
compute-0      IN A 1xx.2x.2xx.1yy
compute-1      IN A 1xx.2x.2xx.1yy
```

4.6.4. Creating a hosted cluster on bare metal

When you create a hosted cluster with the Agent platform, HyperShift installs the Agent Cluster API provider in the hosted control plane namespace. You can create a hosted cluster on bare metal or import one.

As you create a hosted cluster, keep the following guidelines in mind:

- Each hosted cluster must have a cluster-wide unique name. A hosted cluster name cannot be the same as any existing managed cluster in order for multicluster engine Operator to manage it.
- Do not use **clusters** as a hosted cluster name.
- A hosted cluster cannot be created in the namespace of a multicluster engine Operator managed cluster.
- The most common service publishing strategy is to expose services through a load balancer. That strategy is the preferred method for exposing the Kubernetes API server. If you create a hosted cluster by using the web console or by using Red Hat Advanced Cluster Management, to set a publishing strategy for a service besides the Kubernetes API server, you must manually specify the **servicePublishingStrategy** information in the **HostedCluster** custom resource.

Procedure

1. Create the hosted control plane namespace by entering the following command:

```
$ oc create ns <hosted_cluster_namespace>-<hosted_cluster_name>
```

Replace **<hosted_cluster_namespace>** with your hosted cluster namespace name, for example, **clusters**. Replace **<hosted_cluster_name>** with your hosted cluster name.

2. Verify that you have a default storage class configured for your cluster. Otherwise, you might see pending PVCs. Run the following command:

```
$ hcp create cluster agent \
--name=<hosted_cluster_name> \①
--pull-secret=<path_to_pull_secret> \②
--agent-namespace=<hosted_control_plane_namespace> \③
--base-domain=<basedomain> \④
--api-server-address=api.<hosted_cluster_name>.<basedomain> \⑤
--etcd-storage-class=<etcd_storage_class> \⑥
--ssh-key <path_to_ssh_public_key> \⑦
--namespace <hosted_cluster_namespace> \⑧
```

```
--control-plane-availability-policy HighlyAvailable \ ⑨
--release-image=quay.io/openshift-release-dev/ocp-release:<ocp_release_image> \ ⑩
--node-pool-replicas <node_pool_replica_count> ⑪
```

- ① Specify the name of your hosted cluster, for instance, **example**.
- ② Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- ③ Specify your hosted control plane namespace, for example, **clusters-example**. Ensure that agents are available in this namespace by using the **oc get agent -n <hosted_control_plane_namespace>** command.
- ④ Specify your base domain, for example, **krnl.es**.
- ⑤ The **--api-server-address** flag defines the IP address that is used for the Kubernetes API communication in the hosted cluster. If you do not set the **--api-server-address** flag, you must log in to connect to the management cluster.
- ⑥ Specify the etcd storage class name, for example, **lvm-storageclass**.
- ⑦ Specify the path to your SSH public key. The default file path is **~/.ssh/id_rsa.pub**.
- ⑧ Specify your hosted cluster namespace.
- ⑨ Specify the availability policy for the hosted control plane components. Supported options are **SingleReplica** and **HighlyAvailable**. The default value is **HighlyAvailable**.
- ⑩ Specify the supported OpenShift Container Platform version that you want to use, for example, **4.18.0-multi**. If you are using a disconnected environment, replace **<ocp_release_image>** with the digest image. To extract the OpenShift Container Platform release image digest, see *Extracting the OpenShift Container Platform release image digest*.
- ⑪ Specify the node pool replica count, for example, **3**. You must specify the replica count as **0** or greater to create the same number of replicas. Otherwise, no node pools are created.

3. After a few moments, verify that your hosted control plane pods are up and running by entering the following command:

```
$ oc -n <hosted_cluster_namespace>-<hosted_cluster_name> get pods
```

Example output

NAME	READY	STATUS	RESTARTS	AGE
capi-provider-7dcf5fc4c4-nr9sq	1/1	Running	0	4m32s
catalog-operator-6cd867cc7-phb2q	2/2	Running	0	2m50s
certified-operators-catalog-884c756c4-zdt64	1/1	Running	0	2m51s
cluster-api-f75d86f8c-56wfz	1/1	Running	0	4m32s

CHAPTER 5. MANAGING HOSTED CONTROL PLANES

5.1. MANAGING HOSTED CONTROL PLANES ON AWS

When you use hosted control planes for OpenShift Container Platform on Amazon Web Services (AWS), the infrastructure requirements vary based on your setup.

5.1.1. Prerequisites to manage AWS infrastructure and IAM permissions

To configure hosted control planes for OpenShift Container Platform on Amazon Web Services (AWS), you must meet the following the infrastructure requirements:

- You configured hosted control planes before you can create hosted clusters.
- You created an AWS Identity and Access Management (IAM) role and AWS Security Token Service (STS) credentials.

5.1.1.1. Infrastructure requirements for AWS

When you use hosted control planes on Amazon Web Services (AWS), the infrastructure requirements fit in the following categories:

- Prerequisite and unmanaged infrastructure for the HyperShift Operator in an arbitrary AWS account
- Prerequisite and unmanaged infrastructure in a hosted cluster AWS account
- Hosted control planes-managed infrastructure in a management AWS account
- Hosted control planes-managed infrastructure in a hosted cluster AWS account
- Kubernetes-managed infrastructure in a hosted cluster AWS account

Prerequisite means that hosted control planes requires AWS infrastructure to properly work.

Unmanaged means that no Operator or controller creates the infrastructure for you.

5.1.1.2. Unmanaged infrastructure for the HyperShift Operator in an AWS account

An arbitrary Amazon Web Services (AWS) account depends on the provider of the hosted control planes service.

In self-managed hosted control planes, the cluster service provider controls the AWS account. The cluster service provider is the administrator who hosts cluster control planes and is responsible for uptime. In managed hosted control planes, the AWS account belongs to Red Hat.

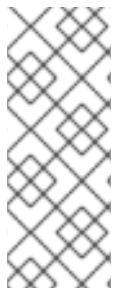
In a prerequisite and unmanaged infrastructure for the HyperShift Operator, the following infrastructure requirements apply for a management cluster AWS account:

- One S3 Bucket
 - OpenID Connect (OIDC)
- Route 53 hosted zones
 - A domain to host private and public entries for hosted clusters

5.1.1.3. Unmanaged infrastructure requirements for a management AWS account

When your infrastructure is prerequired and unmanaged in a hosted cluster Amazon Web Services (AWS) account, the infrastructure requirements for all access modes are as follows:

- One VPC
- One DHCP Option
- Two subnets
 - A private subnet that is an internal data plane subnet
 - A public subnet that enables access to the internet from the data plane
- One internet gateway
- One elastic IP
- One NAT gateway
- One security group (worker nodes)
- Two route tables (one private and one public)
- Two Route 53 hosted zones
- Enough quota for the following items:
 - One Ingress service load balancer for public hosted clusters
 - One private link endpoint for private hosted clusters



NOTE

For private link networking to work, the endpoint zone in the hosted cluster AWS account must match the zone of the instance that is resolved by the service endpoint in the management cluster AWS account. In AWS, the zone names are aliases, such as us-east-2b, which do not necessarily map to the same zone in different accounts. As a result, for private link to work, the management cluster must have subnets or workers in all zones of its region.

5.1.1.4. Infrastructure requirements for a management AWS account

When your infrastructure is managed by hosted control planes in a management AWS account, the infrastructure requirements differ depending on whether your clusters are public, private, or a combination.

For accounts with public clusters, the infrastructure requirements are as follows:

- Network load balancer: a load balancer Kube API server
 - Kubernetes creates a security group
- Volumes
 - For etcd (one or three depending on high availability)

- For OVN-Kube

For accounts with private clusters, the infrastructure requirements are as follows:

- Network load balancer: a load balancer private router
- Endpoint service (private link)

For accounts with public and private clusters, the infrastructure requirements are as follows:

- Network load balancer: a load balancer public router
- Network load balancer: a load balancer private router
- Endpoint service (private link)
- Volumes
 - For etcd (one or three depending on high availability)
 - For OVN-Kube

5.1.1.5. Infrastructure requirements for an AWS account in a hosted cluster

When your infrastructure is managed by hosted control planes in a hosted cluster Amazon Web Services (AWS) account, the infrastructure requirements differ depending on whether your clusters are public, private, or a combination.

For accounts with public clusters, the infrastructure requirements are as follows:

- Node pools must have EC2 instances that have **Role** and **RolePolicy** defined.

For accounts with private clusters, the infrastructure requirements are as follows:

- One private link endpoint for each availability zone
- EC2 instances for node pools

For accounts with public and private clusters, the infrastructure requirements are as follows:

- One private link endpoint for each availability zone
- EC2 instances for node pools

5.1.1.6. Kubernetes-managed infrastructure in a hosted cluster AWS account

When Kubernetes manages your infrastructure in a hosted cluster Amazon Web Services (AWS) account, the infrastructure requirements are as follows:

- A network load balancer for default Ingress
- An S3 bucket for registry

5.1.2. Identity and Access Management (IAM) permissions

In the context of hosted control planes, the consumer is responsible to create the Amazon Resource Name (ARN) roles. The *consumer* is an automated process to generate the permissions files. The

consumer might be the CLI or OpenShift Cluster Manager. Hosted control planes can enable granularity to honor the principle of least-privilege components, which means that every component uses its own role to operate or create Amazon Web Services (AWS) objects, and the roles are limited to what is required for the product to function normally.

The hosted cluster receives the ARN roles as input and the consumer creates an AWS permission configuration for each component. As a result, the component can authenticate through STS and preconfigured OIDC IDP.

The following roles are consumed by some of the components from hosted control planes that run on the control plane and operate on the data plane:

- **controlPlaneOperatorARN**
- **imageRegistryARN**
- **ingressARN**
- **kubeCloudControllerARN**
- **nodePoolManagementARN**
- **storageARN**
- **networkARN**

The following example shows a reference to the IAM roles from the hosted cluster:

```
...
endpointAccess: Public
region: us-east-2
resourceTags:
- key: kubernetes.io/cluster/example-cluster-bz4j5
  value: owned
rolesRef:
  controlPlaneOperatorARN: arn:aws:iam::820196288204:role/example-cluster-bz4j5-control-plane-operator
  imageRegistryARN: arn:aws:iam::820196288204:role/example-cluster-bz4j5-openshift-image-registry
  ingressARN: arn:aws:iam::820196288204:role/example-cluster-bz4j5-openshift-ingress
  kubeCloudControllerARN: arn:aws:iam::820196288204:role/example-cluster-bz4j5-cloud-controller
  networkARN: arn:aws:iam::820196288204:role/example-cluster-bz4j5-cloud-network-config-controller
  nodePoolManagementARN: arn:aws:iam::820196288204:role/example-cluster-bz4j5-node-pool
  storageARN: arn:aws:iam::820196288204:role/example-cluster-bz4j5-aws-ebs-csi-driver-controller
type: AWS
...
```

The roles that hosted control planes uses are shown in the following examples:

- **ingressARN**

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      ...
    }
  ]
}
```

```

    "Effect": "Allow",
    "Action": [
        "elasticloadbalancing:DescribeLoadBalancers",
        "tag:GetResources",
        "route53>ListHostedZones"
    ],
    "Resource": "*"
},
{
    "Effect": "Allow",
    "Action": [
        "route53:ChangeResourceRecordSets"
    ],
    "Resource": [
        "arn:aws:route53:::PUBLIC_ZONE_ID",
        "arn:aws:route53:::PRIVATE_ZONE_ID"
    ]
}
]
}

```

- **imageRegistryARN**

```

{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "s3>CreateBucket",
                "s3>DeleteBucket",
                "s3>PutBucketTagging",
                "s3>GetBucketTagging",
                "s3>PutBucketPublicAccessBlock",
                "s3>GetBucketPublicAccessBlock",
                "s3>PutEncryptionConfiguration",
                "s3>GetEncryptionConfiguration",
                "s3>PutLifecycleConfiguration",
                "s3>GetLifecycleConfiguration",
                "s3>GetBucketLocation",
                "s3>ListBucket",
                "s3>GetObject",
                "s3>PutObject",
                "s3>DeleteObject",
                "s3>ListBucketMultipartUploads",
                "s3>AbortMultipartUpload",
                "s3>ListMultipartUploadParts"
            ],
            "Resource": "*"
        }
    ]
}

```

- **storageARN**

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "ec2:AttachVolume",
        "ec2>CreateSnapshot",
        "ec2>CreateTags",
        "ec2>CreateVolume",
        "ec2>DeleteSnapshot",
        "ec2>DeleteTags",
        "ec2>DeleteVolume",
        "ec2>DescribeInstances",
        "ec2>DescribeSnapshots",
        "ec2>DescribeTags",
        "ec2>DescribeVolumes",
        "ec2>DescribeVolumesModifications",
        "ec2>DetachVolume",
        "ec2>ModifyVolume"
      ],
      "Resource": "*"
    }
  ]
}
```

- **networkARN**

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "ec2>DescribeInstances",
        "ec2>DescribeInstanceStatus",
        "ec2>DescribeInstanceTypes",
        "ec2>UnassignPrivateIpAddresses",
        "ec2>AssignPrivateIpAddresses",
        "ec2>UnassignIpv6Addresses",
        "ec2>AssignIpv6Addresses",
        "ec2>DescribeSubnets",
        "ec2>DescribeNetworkInterfaces"
      ],
      "Resource": "*"
    }
  ]
}
```

- **kubeCloudControllerARN**

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "
```

```

>Action": [
    "ec2:DescribeInstances",
    "ec2:DescribeImages",
    "ec2:DescribeRegions",
    "ec2:DescribeRouteTables",
    "ec2:DescribeSecurityGroups",
    "ec2:DescribeSubnets",
    "ec2:DescribeVolumes",
    "ec2>CreateSecurityGroup",
    "ec2>CreateTags",
    "ec2>CreateVolume",
    "ec2:ModifyInstanceAttribute",
    "ec2:ModifyVolume",
    "ec2:AttachVolume",
    "ec2:AuthorizeSecurityGroupIngress",
    "ec2>CreateRoute",
    "ec2>DeleteRoute",
    "ec2>DeleteSecurityGroup",
    "ec2>DeleteVolume",
    "ec2:DetachVolume",
    "ec2:RevokeSecurityGroupIngress",
    "ec2:DescribeVpcs",
    "elasticloadbalancing:AddTags",
    "elasticloadbalancing:AttachLoadBalancerToSubnets",
    "elasticloadbalancing:ApplySecurityGroupsToLoadBalancer",
    "elasticloadbalancing>CreateLoadBalancer",
    "elasticloadbalancing>CreateLoadBalancerPolicy",
    "elasticloadbalancing>CreateLoadBalancerListeners",
    "elasticloadbalancing:ConfigureHealthCheck",
    "elasticloadbalancing>DeleteLoadBalancer",
    "elasticloadbalancing>DeleteLoadBalancerListeners",
    "elasticloadbalancing:DescribeLoadBalancers",
    "elasticloadbalancing:DescribeLoadBalancerAttributes",
    "elasticloadbalancing:DetachLoadBalancerFromSubnets",
    "elasticloadbalancing:DeregisterInstancesFromLoadBalancer",
    "elasticloadbalancing:ModifyLoadBalancerAttributes",
    "elasticloadbalancing:RegisterInstancesWithLoadBalancer",
    "elasticloadbalancing:SetLoadBalancerPoliciesForBackendServer",
    "elasticloadbalancing:AddTags",
    "elasticloadbalancing>CreateListener",
    "elasticloadbalancing>CreateTargetGroup",
    "elasticloadbalancing>DeleteListener",
    "elasticloadbalancing>DeleteTargetGroup",
    "elasticloadbalancing:DescribeListeners",
    "elasticloadbalancing:DescribeLoadBalancerPolicies",
    "elasticloadbalancing:DescribeTargetGroups",
    "elasticloadbalancing:DescribeTargetHealth",
    "elasticloadbalancing:ModifyListener",
    "elasticloadbalancing:ModifyTargetGroup",
    "elasticloadbalancing:RegisterTargets",
    "elasticloadbalancing:SetLoadBalancerPoliciesOfListener",
    "iam>CreateServiceLinkedRole",
    "kms:DescribeKey"
],
"Resource": [
    "\*"
]

```

```
        ],
        "Effect": "Allow"
    }
]
```

- **nodePoolManagementARN**

```
{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Action": [
                "ec2:AllocateAddress",
                "ec2:AssociateRouteTable",
                "ec2:AttachInternetGateway",
                "ec2:AuthorizeSecurityGroupIngress",
                "ec2>CreateInternetGateway",
                "ec2>CreateNatGateway",
                "ec2>CreateRoute",
                "ec2>CreateRouteTable",
                "ec2>CreateSecurityGroup",
                "ec2>CreateSubnet",
                "ec2>CreateTags",
                "ec2>DeleteInternetGateway",
                "ec2>DeleteNatGateway",
                "ec2>DeleteRouteTable",
                "ec2>DeleteSecurityGroup",
                "ec2>DeleteSubnet",
                "ec2>DeleteTags",
                "ec2:DescribeAccountAttributes",
                "ec2:DescribeAddresses",
                "ec2:DescribeAvailabilityZones",
                "ec2:DescribeImages",
                "ec2:DescribeInstances",
                "ec2:DescribeInternetGateways",
                "ec2:DescribeNatGateways",
                "ec2:DescribeNetworkInterfaces",
                "ec2:DescribeNetworkInterfaceAttribute",
                "ec2:DescribeRouteTables",
                "ec2:DescribeSecurityGroups",
                "ec2:DescribeSubnets",
                "ec2:DescribeVpcs",
                "ec2:DescribeVpcAttribute",
                "ec2:DescribeVolumes",
                "ec2:DetachInternetGateway",
                "ec2:DisassociateRouteTable",
                "ec2:DisassociateAddress",
                "ec2:ModifyInstanceAttribute",
                "ec2:ModifyNetworkInterfaceAttribute",
                "ec2:ModifySubnetAttribute",
                "ec2:ReleaseAddress",
                "ec2:RevokeSecurityGroupIngress",
                "ec2:RunInstances",
                "ec2:TerminateInstances",
                "tag:GetResources",
            ]
        }
    ]
}
```

```

    "ec2:CreateLaunchTemplate",
    "ec2:CreateLaunchTemplateVersion",
    "ec2:DescribeLaunchTemplates",
    "ec2:DescribeLaunchTemplateVersions",
    "ec2>DeleteLaunchTemplate",
    "ec2>DeleteLaunchTemplateVersions"
],
"Resource": [
    "\*"
],
"Effect": "Allow"
},
{
"Condition": {
    "StringLike": {
        "iam:AWSServiceName": "elasticloadbalancing.amazonaws.com"
    }
},
"Action": [
    "iam>CreateServiceLinkedRole"
],
"Resource": [
    "arn:*:iam::role/aws-service-
role/elasticloadbalancing.amazonaws.com/AWSServiceRoleForElasticLoadBalancing"
],
"Effect": "Allow"
},
{
"Action": [
    "iam:PassRole"
],
"Resource": [
    "arn:*:iam::role/*-worker-role"
],
"Effect": "Allow"
}
]
}
}

```

- **controlPlaneOperatorARN**

```

{
    "Version": "2012-10-17",
    "Statement": [
        {
            "Effect": "Allow",
            "Action": [
                "ec2>CreateVpcEndpoint",
                "ec2>DescribeVpcEndpoints",
                "ec2>ModifyVpcEndpoint",
                "ec2>DeleteVpcEndpoints",
                "ec2>CreateTags",
                "route53>ListHostedZones"
            ],
            "Resource": "\*"
        },

```

```
{
    "Effect": "Allow",
    "Action": [
        "route53:ChangeResourceRecordSets",
        "route53>ListResourceRecordSets"
    ],
    "Resource": "arn:aws:route53:::%s"
}
]
```

5.1.3. Creating AWS infrastructure and IAM resources separate

By default, the **hcp create cluster aws** command creates cloud infrastructure with the hosted cluster and applies it. You can create the cloud infrastructure portion separately so that you can use the **hcp create cluster aws** command only to create the cluster, or render it to modify it before you apply it.

To create the cloud infrastructure portion separately, you need to create the Amazon Web Services (AWS) infrastructure, create the AWS Identity and Access (IAM) resources, and create the cluster.

5.1.3.1. Creating the AWS infrastructure separately

To create the Amazon Web Services (AWS) infrastructure, you need to create a Virtual Private Cloud (VPC) and other resources for your cluster. You can use the AWS console or an infrastructure automation and provisioning tool. For instructions to use the AWS console, see [Create a VPC plus other VPC resources](#) in the AWS Documentation.

The VPC must include private and public subnets and resources for external access, such as a network address translation (NAT) gateway and an internet gateway. In addition to the VPC, you need a private hosted zone for the ingress of your cluster. If you are creating clusters that use PrivateLink (**Private** or **PublicAndPrivate** access modes), you need an additional hosted zone for PrivateLink.

Create the AWS infrastructure for your hosted cluster by using the following example configuration:

```
---
apiVersion: v1
kind: Namespace
metadata:
  creationTimestamp: null
  name: clusters
spec: {}
status: {}
---
apiVersion: v1
data:
  .dockerconfigjson: xxxxxxxxxxxx
kind: Secret
metadata:
  creationTimestamp: null
  labels:
    hypershift.openshift.io/safe-to-delete-with-cluster: "true"
  name: <pull_secret_name> ①
  namespace: clusters
---
apiVersion: v1
```

```

data:
  key: xxxxxxxxxxxxxxxxxx
kind: Secret
metadata:
  creationTimestamp: null
  labels:
    hypershift.openshift.io/safe-to-delete-with-cluster: "true"
  name: <etcd_encryption_key_name> ②
  namespace: clusters
type: Opaque
---
apiVersion: v1
data:
  id_rsa: xxxxxxxxx
  id_rsa.pub: xxxxxxxxx
kind: Secret
metadata:
  creationTimestamp: null
  labels:
    hypershift.openshift.io/safe-to-delete-with-cluster: "true"
  name: <ssh-key-name> ③
  namespace: clusters
---
apiVersion: hypershift.openshift.io/v1beta1
kind: HostedCluster
metadata:
  creationTimestamp: null
  name: <hosted_cluster_name> ④
  namespace: clusters
spec:
  autoscaling: {}
  configuration: {}
  controllerAvailabilityPolicy: SingleReplica
  dns:
    baseDomain: <dns_domain> ⑤
    privateZoneID: xxxxxxxx
    publicZoneID: xxxxxxxx
  etcd:
    managed:
      storage:
        persistentVolume:
          size: 8Gi
          storageClassName: gp3-csi
          type: PersistentVolume
        managementType: Managed
    fips: false
  infraID: <infra_id> ⑥
  issuerURL: <issuer_url> ⑦
  networking:
    clusterNetwork:
      - cidr: 10.132.0.0/14
    machineNetwork:
      - cidr: 10.0.0.0/16
  networkType: OVNKubernetes
  serviceNetwork:

```

```

- cidr: 172.31.0.0/16
olmCatalogPlacement: management
platform:
aws:
  cloudProviderConfig:
    subnet:
      id: <subnet_xxx> 8
    vpc: <vpc_xxx> 9
    zone: us-west-1b
  endpointAccess: Public
  multiArch: false
  region: us-west-1
  rolesRef:
    controlPlaneOperatorARN: arn:aws:iam::820196288204:role/<infra_id>-control-plane-operator
    imageRegistryARN: arn:aws:iam::820196288204:role/<infra_id>-openshift-image-registry
    ingressARN: arn:aws:iam::820196288204:role/<infra_id>-openshift-ingress
    kubeCloudControllerARN: arn:aws:iam::820196288204:role/<infra_id>-cloud-controller
    networkARN: arn:aws:iam::820196288204:role/<infra_id>-cloud-network-config-controller
    nodePoolManagementARN: arn:aws:iam::820196288204:role/<infra_id>-node-pool
    storageARN: arn:aws:iam::820196288204:role/<infra_id>-aws-ebs-csi-driver-controller
  type: AWS
pullSecret:
  name: <pull_secret_name>
release:
  image: quay.io/openshift-release-dev/ocp-release:4.16-x86_64
secretEncryption:
  aescbc:
    activeKey:
      name: <etcd_encryption_key_name>
  type: aescbc
services:
- service: APIServer
  servicePublishingStrategy:
    type: LoadBalancer
- service: OAuthServer
  servicePublishingStrategy:
    type: Route
- service: Konnectivity
  servicePublishingStrategy:
    type: Route
- service: Ignition
  servicePublishingStrategy:
    type: Route
- service: OVNSbDb
  servicePublishingStrategy:
    type: Route
sshKey:
  name: <ssh_key_name>
status:
  controlPlaneEndpoint:
    host: ""
    port: 0
---
apiVersion: hypershift.openshift.io/v1beta1
kind: NodePool
metadata:

```

```

creationTimestamp: null
name: <node_pool_name> 10
namespace: clusters
spec:
  arch: amd64
  clusterName: <hosted_cluster_name>
  management:
    autoRepair: true
    upgradeType: Replace
  nodeDrainTimeout: 0s
  platform:
    aws:
      instanceProfile: <instance_profile_name> 11
      instanceType: m6i.xlarge
      rootVolume:
        size: 120
        type: gp3
      subnet:
        id: <subnet_xxx>
      type: AWS
  release:
    image: quay.io/openshift-release-dev/ocp-release:4.16-x86_64
    replicas: 2
  status:
    replicas: 0

```

- 1** Replace **<pull_secret_name>** with the name of your pull secret.
- 2** Replace **<etcd_encryption_key_name>** with the name of your etcd encryption key.
- 3** Replace **<ssh_key_name>** with the name of your SSH key.
- 4** Replace **<hosted_cluster_name>** with the name of your hosted cluster.
- 5** Replace **<dns_domain>** with your base DNS domain, such as **example.com**.
- 6** Replace **<infra_id>** with the value that identifies the IAM resources that are associated with the hosted cluster.
- 7** Replace **<issuer_url>** with your issuer URL, which ends with your **infra_id** value. For example, <https://example-hosted-us-west-1.s3.us-west-1.amazonaws.com/example-hosted-infra-id>.
- 8** Replace **<subnet_xxx>** with your subnet ID. Both private and public subnets need to be tagged. For public subnets, use **kubernetes.io/role/elb=1**. For private subnets, use **kubernetes.io/role/internal-elb=1**.
- 9** Replace **<vpc_xxx>** with your VPC ID.
- 10** Replace **<node_pool_name>** with the name of your **NodePool** resource.
- 11** Replace **<instance_profile_name>** with the name of your AWS instance.

5.1.3.2. Creating the AWS IAM resources

In Amazon Web Services (AWS), you must create the following IAM resources:

- An [OpenID Connect \(OIDC\) identity provider in IAM](#), which is required to enable STS authentication.
- [Seven roles](#), which are separate for every component that interacts with the provider, such as the Kubernetes controller manager, cluster API provider, and registry
- The [instance profile](#), which is the profile that is assigned to all worker instances of the cluster

5.1.3.3. Creating a hosted cluster separately

You can create a hosted cluster separately on Amazon Web Services (AWS).

To create a hosted cluster separately, enter the following command:

```
$ hcp create cluster aws \
--infra-id <infra_id> \①
--name <hosted_cluster_name> \②
--sts-creds <path_to_sts_credential_file> \③
--pull-secret <path_to_pull_secret> \④
--generate-ssh \⑤
--node-pool-replicas 3
--role-arn <role_name> \⑥
```

- ① Replace **<infra_id>** with the same ID that you specified in the **create infra aws** command. This value identifies the IAM resources that are associated with the hosted cluster.
- ② Replace **<hosted_cluster_name>** with the name of your hosted cluster.
- ③ Replace **<path_to_sts_credential_file>** with the same name that you specified in the **create infra aws** command.
- ④ Replace **<path_to_pull_secret>** with the name of the file that contains a valid OpenShift Container Platform pull secret.
- ⑤ The **--generate-ssh** flag is optional, but is good to include in case you need to SSH to your workers. An SSH key is generated for you and is stored as a secret in the same namespace as the hosted cluster.
- ⑥ Replace **<role_name>** with the Amazon Resource Name (ARN), for example, **arn:aws:iam::820196288204:role/myrole**. Specify the Amazon Resource Name (ARN), for example, **arn:aws:iam::820196288204:role/myrole**. For more information about ARN roles, see "Identity and Access Management (IAM) permissions".

You can also add the **--render** flag to the command and redirect output to a file where you can edit the resources before you apply them to the cluster.

After you run the command, the following resources are applied to your cluster:

- A namespace
- A secret with your pull secret
- A **HostedCluster**
- A **NodePool**

- Three AWS STS secrets for control plane components
- One SSH key secret if you specified the **--generate-ssh** flag.

5.1.4. Transitioning a hosted cluster from single-architecture to multi-architecture

You can transition your single-architecture 64-bit AMD hosted cluster to a multi-architecture hosted cluster on Amazon Web Services (AWS), to reduce the cost of running workloads on your cluster. For example, you can run existing workloads on 64-bit AMD while transitioning to 64-bit ARM and you can manage these workloads from a central Kubernetes cluster.

A single-architecture hosted cluster can manage node pools of only one particular CPU architecture. However, a multi-architecture hosted cluster can manage node pools with different CPU architectures. On AWS, a multi-architecture hosted cluster can manage both 64-bit AMD and 64-bit ARM node pools.

Prerequisites

- You have installed an OpenShift Container Platform management cluster for AWS on Red Hat Advanced Cluster Management (RHACM) with the multicluster engine for Kubernetes Operator.
- You have an existing single-architecture hosted cluster that uses 64-bit AMD variant of the OpenShift Container Platform release payload.
- An existing node pool that uses the same 64-bit AMD variant of the OpenShift Container Platform release payload and is managed by an existing hosted cluster.
- Ensure that you installed the following command-line tools:
 - **oc**
 - **kubectl**
 - **hcp**
 - **skopeo**

Procedure

- 1 Review an existing OpenShift Container Platform release image of the single-architecture hosted cluster by running the following command:

```
$ oc get hostedcluster/<hosted_cluster_name> ①
-o jsonpath='{.spec.release.image}'
```

- 1 Replace **<hosted_cluster_name>** with your hosted cluster name.

Example output

```
quay.io/openshift-release-dev/ocp-release:<4.y.z>-x86_64 ①
```

- 1 Replace **<4.y.z>** with the supported OpenShift Container Platform version that you use.

2. In your OpenShift Container Platform release image, if you use the digest instead of a tag, find the multi-architecture tag version of your release image:

- Set the **OCP_VERSION** environment variable for the OpenShift Container Platform version by running the following command:

```
$ OCP_VERSION=$(oc image info quay.io/openshift-release-dev/ocp-release@sha256:ac78ebf77f95ab8ff52847ecd22592b545415e1ff6c7ff7f66bf81f158ae4f5e \
-o jsonpath='{.config.config.Labels["io.openshift.release"]}'")
```

- Set the **MULTI_ARCH_TAG** environment variable for the multi-architecture tag version of your release image by running the following command:

```
$ MULTI_ARCH_TAG=$(skopeo inspect docker://quay.io/openshift-release-dev/ocp-release@sha256:ac78ebf77f95ab8ff52847ecd22592b545415e1ff6c7ff7f66bf81f158ae4f5e \
| jq -r '.RepoTags' | sed 's://"//g' | sed 's/,//g' \
| grep -w "$OCP_VERSION-multi$" | xargs)
```

- Set the **IMAGE** environment variable for the multi-architecture release image name by running the following command:

```
$ IMAGE=quay.io/openshift-release-dev/ocp-release:$MULTI_ARCH_TAG
```

- To see the list of multi-architecture image digests, run the following command:

```
$ oc image info $IMAGE
```

Example output

```
OS          DIGEST
linux/amd64
sha256:b4c7a91802c09a5a748fe19ddd99a8ffab52d8a31db3a081a956a87f22a22ff8
linux/ppc64le
sha256:66fd42ff6bd7704f1ba72be8bfe3e399c323de92262f594f8e482d110ec37388
linux/s390x
sha256:b1c1072dc639aaa2b50ec99b530012e3ceac19ddc28adcbcdc9643f2dfd14f34
linux/arm64
sha256:7b046404572ac96202d82b6cb029b421dddd40e88c73bbf35f602ffc13017f21
```

3. Transition the hosted cluster from single-architecture to multi-architecture:

- Set the multi-architecture OpenShift Container Platform release image for the hosted cluster by ensuring that you use the same OpenShift Container Platform version as the hosted cluster. Run the following command:

```
$ oc patch -n clusters hostedclusters/<hosted_cluster_name> -p \
  '{"spec":{"release":{"image":"quay.io/openshift-release-dev/ocp-release:<4.x.y>-multi"}}}' \
  -t=merge
```

- 1** Replace **<4.y.z>** with the supported OpenShift Container Platform version that you use.

- b. Confirm that the multi-architecture image is set in your hosted cluster by running the following command:

```
$ oc get hostedcluster/<hosted_cluster_name> \
-o jsonpath='{.spec.release.image}'
```

4. Check that the status of the **HostedControlPlane** resource is **Progressing** by running the following command:

```
$ oc get hostedcontrolplane -n <hosted_control_plane_namespace> -oyaml
```

Example output

```
#...
- lastTransitionTime: "2024-07-28T13:07:18Z"
  message: HostedCluster is deploying, upgrading, or reconfiguring
  observedGeneration: 5
  reason: Progressing
  status: "True"
  type: Progressing
#...
```

5. Check that the status of the **HostedCluster** resource is **Progressing** by running the following command:

```
$ oc get hostedcluster <hosted_cluster_name> \
-n <hosted_cluster_namespace> -oyaml
```

Verification

- Verify that a node pool is using the multi-architecture release image in your **HostedControlPlane** resource by running the following command:

```
$ oc get hostedcontrolplane -n clusters-example -oyaml
```

Example output

```
#...
version:
  availableUpdates: null
  desired:
    image: quay.io/openshift-release-dev/ocp-release:<4.x.y>-multi ①
    url: https://access.redhat.com/errata/RHBA-2024:4855
    version: 4.16.5
  history:
    - completionTime: "2024-07-28T13:10:58Z"
      image: quay.io/openshift-release-dev/ocp-release:<4.x.y>-multi
      startedTime: "2024-07-28T13:10:27Z"
      state: Completed
      verified: false
      version: <4.x.y>
```

- 1 Replace <4.y.z> with the supported OpenShift Container Platform version that you use.



NOTE

The multi-architecture OpenShift Container Platform release image is updated in your **HostedCluster**, **HostedControlPlane** resources, and hosted control plane pods. However, your existing node pools do not transition with the multi-architecture image automatically, because the release image transition is decoupled between the hosted cluster and node pools. You must create new node pools on your new multi-architecture hosted cluster.

Next steps

- Creating node pools on the multi-architecture hosted cluster

5.1.5. Creating node pools on the multi-architecture hosted cluster

After transitioning your hosted cluster from single-architecture to multi-architecture, create node pools on compute machines based on 64-bit AMD and 64-bit ARM architectures.

Procedure

- 1 Create node pools based on 64-bit ARM architecture by entering the following command:

```
$ hcp create nodepool aws \
--cluster-name <hosted_cluster_name> \①
--name <nodepool_name> \②
--node-count=<node_count> \③
--arch arm64
```

- 1 Replace <hosted_cluster_name> with your hosted cluster name.
- 2 Replace <nodepool_name> with your node pool name.
- 3 Replace <node_count> with integer for your node count, for example, 2.

- 2 Create node pools based on 64-bit AMD architecture by entering the following command:

```
$ hcp create nodepool aws \
--cluster-name <hosted_cluster_name> \①
--name <nodepool_name> \②
--node-count=<node_count> \③
--arch amd64
```

- 1 Replace <hosted_cluster_name> with your hosted cluster name.
- 2 Replace <nodepool_name> with your node pool name.
- 3 Replace <node_count> with integer for your node count, for example, 2.

Verification

- Verify that a node pool is using the multi-architecture release image by entering the following command:

```
$ oc get nodepool/<nodepool_name> -oyaml
```

Example output for 64-bit AMD node pools

```
#...
spec:
  arch: amd64
#...
release:
  image: quay.io/openshift-release-dev/ocp-release:<4.x.y>-multi ①
```

- Replace <4.y.z> with the supported OpenShift Container Platform version that you use.

Example output for 64-bit ARM node pools

```
#...
spec:
  arch: arm64
#...
release:
  image: quay.io/openshift-release-dev/ocp-release:<4.x.y>-multi
```

5.2. MANAGING HOSTED CONTROL PLANES ON BARE METAL

After you deploy hosted control planes on bare metal, you can manage a hosted cluster by completing the following tasks.

5.2.1. Accessing the hosted cluster

You can access the hosted cluster by either getting the **kubeconfig** file and **kubeadm** credential directly from resources, or by using the **hcp** command-line interface to generate a **kubeconfig** file.

Prerequisites

To access the hosted cluster by getting the **kubeconfig** file and credentials directly from resources, you must be familiar with the access secrets for hosted clusters. The *hosted cluster (hosting)* namespace contains hosted cluster resources and the access secrets. The *hosted control plane* namespace is where the hosted control plane runs.

The secret name formats are as follows:

- kubeconfig** secret: <hosted_cluster_namespace>-<name>-admin-kubeconfig. For example, **clusters-hypershift-demo-admin-kubeconfig**.
- kubeadm** password secret: <hosted_cluster_namespace>-<name>-kubeadm-password. For example, **clusters-hypershift-demo-kubeadm-password**.

The **kubeconfig** secret contains a Base64-encoded **kubeconfig** field, which you can decode and save into a file to use with the following command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get nodes
```

The **kubeadmin** password secret is also Base64-encoded. You can decode it and use the password to log in to the API server or console of the hosted cluster.

Procedure

- To access the hosted cluster by using the **hcp** CLI to generate the **kubeconfig** file, take the following steps:

1. Generate the **kubeconfig** file by entering the following command:

```
$ hcp create kubeconfig --namespace <hosted_cluster_namespace> \  
--name <hosted_cluster_name> > <hosted_cluster_name>.kubeconfig
```

2. After you save the **kubeconfig** file, you can access the hosted cluster by entering the following example command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get nodes
```

5.2.2. Scaling the NodePool object for a hosted cluster

You can scale up the **NodePool** object by adding nodes to your hosted cluster. When you scale a node pool, consider the following information:

- When you scale a replica by the node pool, a machine is created. For every machine, the Cluster API provider finds and installs an Agent that meets the requirements that are specified in the node pool specification. You can monitor the installation of an Agent by checking its status and conditions.
- When you scale down a node pool, Agents are unbound from the corresponding cluster. Before you can reuse the Agents, you must restart them by using the Discovery image.

Procedure

1. Scale the **NodePool** object to two nodes:

```
$ oc -n <hosted_cluster_namespace> scale nodepool <nodepool_name> --replicas 2
```

The Cluster API agent provider randomly picks two agents that are then assigned to the hosted cluster. Those agents go through different states and finally join the hosted cluster as OpenShift Container Platform nodes. The agents pass through states in the following order:

- **binding**
- **discovering**
- **insufficient**
- **installing**
- **installing-in-progress**
- **added-to-existing-cluster**

- Enter the following command:

```
$ oc -n <hosted_control_plane_namespace> get agent
```

Example output

NAME	CLUSTER	APPROVED	ROLE	STAGE
4dac1ab2-7dd5-4894-a220-6a3473b67ee6	hypercluster1	true	auto-assign	
d9198891-39f4-4930-a679-65fb142b108b		true	auto-assign	
da503cf1-a347-44f2-875c-4960ddb04091	hypercluster1	true	auto-assign	

- Enter the following command:

```
$ oc -n <hosted_control_plane_namespace> get agent \
-o jsonpath='{range .items[*]}BMH: {@.metadata.labels.agent-install\\.openshift\\.io/bmh} \
Agent: {@.metadata.name} State: {@.status.debugInfo.state}{"\n"}{end}'
```

Example output

```
BMH: ocp-worker-2 Agent: 4dac1ab2-7dd5-4894-a220-6a3473b67ee6 State: binding
BMH: ocp-worker-0 Agent: d9198891-39f4-4930-a679-65fb142b108b State: known-unbound
BMH: ocp-worker-1 Agent: da503cf1-a347-44f2-875c-4960ddb04091 State: insufficient
```

- Obtain the kubeconfig for your new hosted cluster by entering the extract command:

```
$ oc extract -n <hosted_cluster_namespace> \
secret/<hosted_cluster_name>-admin-kubeconfig --to=- \
> kubeconfig-<hosted_cluster_name>
```

- After the agents reach the **added-to-existing-cluster** state, verify that you can see the OpenShift Container Platform nodes in the hosted cluster by entering the following command:

```
$ oc --kubeconfig kubeconfig-<hosted_cluster_name> get nodes
```

Example output

NAME	STATUS	ROLES	AGE	VERSION
ocp-worker-1	Ready	worker	5m41s	v1.24.0+3882f8f
ocp-worker-2	Ready	worker	6m3s	v1.24.0+3882f8f

Cluster Operators start to reconcile by adding workloads to the nodes.

- Enter the following command to verify that two machines were created when you scaled up the **NodePool** object:

```
$ oc -n <hosted_control_plane_namespace> get machines
```

Example output

NAME	CLUSTER	NODENAME	PROVIDERID
PHASE	AGE	VERSION	
hypercluster1-c96b6f675-m5vch	hypercluster1-b2qhl	ocp-worker-1	agent://da503cf1-

```
a347-44f2-875c-4960ddb04091 Running 15m 4.x.z
hypercluster1-c96b6f675-tl42p hypercluster1-b2qhl ocp-worker-2 agent://4dac1ab2-
7dd5-4894-a220-6a3473b67ee6 Running 15m 4.x.z
```

The **clusterversion** reconcile process eventually reaches a point where only Ingress and Console cluster operators are missing.

- Enter the following command:

```
$ oc --kubeconfig kubeconfig-<hosted_cluster_name> get clusterversion,co
```

Example output

NAME	VERSION	AVAILABLE	PROGRESSING	SINCE
NAME	VERSION	AVAILABLE	PROGRESSING	DEGRADED SINCE MESSAGE
clusterversion.config.openshift.io/version	False	True	40m	Unable to apply 4.x.z: the cluster operator console has not yet successfully rolled out
clusteroperator.config.openshift.io/console	4.12z	False	False	False 11m RouteHealthAvailable: failed to GET route (https://console-openshift-console.apps.hypercluster1.domain.com): Get "https://console-openshift-console.apps.hypercluster1.domain.com": dial tcp 10.19.3.29:443: connect: connection refused
clusteroperator.config.openshift.io/csi-snapshot-controller	4.12z	True	False	False 10m
clusteroperator.config.openshift.io/dns	4.12z	True	False	False 9m16s

5.2.2.1. Adding node pools

You can create node pools for a hosted cluster by specifying a name, number of replicas, and any additional information, such as an agent label selector.

Procedure

- To create a node pool, enter the following information:

```
$ hcp create nodepool agent \
--cluster-name <hosted_cluster_name> \①
--name <nodepool_name> \②
--node-count <worker_node_count> \③
--agentLabelSelector size=medium \④
```

- Replace **<hosted_cluster_name>** with your hosted cluster name.
- Replace **<nodepool_name>** with the name of your node pool, for example, **<hosted_cluster_name>-extra-cpu**.
- Replace **<worker_node_count>** with the worker node count, for example, **2**.
- The **--agentLabelSelector** flag is optional. The node pool uses agents with the **size=medium** label.

- Check the status of the node pool by listing **nodepool** resources in the **clusters** namespace:

```
$ oc get nodepools --namespace clusters
```

- Extract the **admin-kubeconfig** secret by entering the following command:

```
$ oc extract -n <hosted_control_plane_namespace> secret/admin-kubeconfig --to=../hostedcluster-secrets --confirm
```

Example output

```
hostedcluster-secrets/kubeconfig
```

- After some time, you can check the status of the node pool by entering the following command:

```
$ oc --kubeconfig ./hostedcluster-secrets get nodes
```

Verification

- Verify that the number of available node pools match the number of expected node pools by entering this command:

```
$ oc get nodepools --namespace clusters
```

5.2.2.2. Enabling node auto-scaling for the hosted cluster

When you need more capacity in your hosted cluster and spare agents are available, you can enable auto-scaling to install new worker nodes.

Procedure

- To enable auto-scaling, enter the following command:

```
$ oc -n <hosted_cluster_namespace> patch nodepool <hosted_cluster_name> \
--type=json \
-p '[{"op": "remove", "path": "/spec/replicas"}, {"op": "add", "path": "/spec/autoScaling", \
"value": { "max": 5, "min": 2 }}]'
```



NOTE

In the example, the minimum number of nodes is 2, and the maximum is 5. The maximum number of nodes that you can add might be bound by your platform. For example, if you use the Agent platform, the maximum number of nodes is bound by the number of available agents.

- Create a workload that requires a new node.

- Create a YAML file that contains the workload configuration, by using the following example:

```
apiVersion: apps/v1
kind: Deployment
```

```

metadata:
  creationTimestamp: null
  labels:
    app: reversewords
    name: reversewords
    namespace: default
spec:
  replicas: 40
  selector:
    matchLabels:
      app: reversewords
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: reversewords
    spec:
      containers:
        - image: quay.io/mavazque/reversewords:latest
          name: reversewords
      resources:
        requests:
          memory: 2Gi
  status: {}

```

- b. Save the file as **workload-config.yaml**.
- c. Apply the YAML by entering the following command:

```
$ oc apply -f workload-config.yaml
```

3. Extract the **admin-kubeconfig** secret by entering the following command:

```
$ oc extract -n <hosted_cluster_namespace> \
secret/<hosted_cluster_name>-admin-kubeconfig \
--to=../hostedcluster-secrets --confirm
```

Example output

```
hostedcluster-secrets/kubeconfig
```

4. You can check if new nodes are in the **Ready** status by entering the following command:

```
$ oc --kubeconfig ./hostedcluster-secrets get nodes
```

5. To remove the node, delete the workload by entering the following command:

```
$ oc --kubeconfig ./hostedcluster-secrets -n <namespace> \
delete deployment <deployment_name>
```

6. Wait for several minutes to pass without requiring the additional capacity. On the Agent platform, the agent is decommissioned and can be reused. You can confirm that the node was removed by entering the following command:

```
$ oc --kubeconfig ./hostedcluster-secrets get nodes
```



NOTE

For IBM Z® agents, if you are using an OSA network device in Processor Resource/Systems Manager (PR/SM) mode, auto scaling is not supported. You must delete the old agent manually and scale up the node pool because the new agent joins during the scale down process.

5.2.2.3. Disabling node auto-scaling for the hosted cluster

To disable node auto-scaling, complete the following procedure.

Procedure

- Enter the following command to disable node auto-scaling for the hosted cluster:

```
$ oc -n <hosted_cluster_namespace> patch nodepool <hosted_cluster_name> \
--type=json \
-p '[{"op": "remove", "path": "/spec/autoScaling"}, {"op": "add", "path": "/spec/replicas",
"value": <specify_value_to_scale_replicas>}]
```

The command removes "**spec.autoScaling**" from the YAML file, adds "**spec.replicas**", and sets "**spec.replicas**" to the integer value that you specify.

Additional resources

- [Scaling down the data plane to zero](#)

5.2.3. Handling ingress in a hosted cluster on bare metal

Every OpenShift Container Platform cluster has a default application Ingress Controller that typically has an external DNS record associated with it. For example, if you create a hosted cluster named **example** with the base domain **krnl.es**, you can expect the wildcard domain ***.apps.example.krnl.es** to be routable.

Procedure

To set up a load balancer and wildcard DNS record for the ***.apps** domain, perform the following actions on your guest cluster:

- Deploy MetalLB by creating a YAML file that contains the configuration for the MetalLB Operator:

```
apiVersion: v1
kind: Namespace
metadata:
  name: metallb
  labels:
    openshift.io/cluster-monitoring: "true"
  annotations:
    workload.openshift.io/allowed: management
---
apiVersion: operators.coreos.com/v1
```

```

kind: OperatorGroup
metadata:
  name: metallb-operator-operatorgroup
  namespace: metallb
---
apiVersion: operators.coreos.com/v1alpha1
kind: Subscription
metadata:
  name: metallb-operator
  namespace: metallb
spec:
  channel: "stable"
  name: metallb-operator
  source: redhat-operators
  sourceNamespace: openshift-marketplace

```

2. Save the file as **metallb-operator-config.yaml**.
3. Enter the following command to apply the configuration:

```
$ oc apply -f metallb-operator-config.yaml
```

4. After the Operator is running, create the MetalLB instance:
 - a. Create a YAML file that contains the configuration for the MetalLB instance:

```

apiVersion: metallb.io/v1beta1
kind: MetalLB
metadata:
  name: metallb
  namespace: metallb

```

- b. Save the file as **metallb-instance-config.yaml**.
- c. Create the MetalLB instance by entering this command:

```
$ oc apply -f metallb-instance-config.yaml
```

5. Create an **IPAddressPool** resource with a single IP address. This IP address must be on the same subnet as the network that the cluster nodes use.
 - a. Create a file, such as **ipaddresspool.yaml**, with content like the following example:

```

apiVersion: metallb.io/v1beta1
kind: IPAddressPool
metadata:
  namespace: metallb
  name: <ip_address_pool_name> ①
spec:
  addresses:
    - <ingress_ip>-<ingress_ip> ②
  autoAssign: false

```

- ① Specify the **IPAddressPool** resource name.

2 Specify the IP address for your environment. For example, **192.168.122.23**.

b. Apply the configuration for the IP address pool by entering the following command:

```
$ oc apply -f ipaddresspool.yaml
```

6. Create a L2 advertisement.

a. Create a file, such as **l2advertisement.yaml**, with content like the following example:

```
apiVersion: metallb.io/v1beta1
kind: L2Advertisement
metadata:
  name: <l2_advertisement_name> 1
  namespace: metallb
spec:
  ipAddressPools:
    - <ip_address_pool_name> 2
```

1 Specify the **L2Advertisement** resource name.

2 Specify the **IPAddressPool** resource name.

b. Apply the configuration by entering the following command:

```
$ oc apply -f l2advertisement.yaml
```

7. After creating a service of the **LoadBalancer** type, MetalLB adds an external IP address for the service.

a. Configure a new load balancer service that routes ingress traffic to the ingress deployment by creating a YAML file named **metallb-loadbalancer-service.yaml**:

```
kind: Service
apiVersion: v1
metadata:
  annotations:
    metallb.io/address-pool: ingress-public-ip
  name: metallb-ingress
  namespace: openshift-ingress
spec:
  ports:
    - name: http
      protocol: TCP
      port: 80
      targetPort: 80
    - name: https
      protocol: TCP
      port: 443
      targetPort: 443
  selector:
    ingresscontroller.operator.openshift.io/deployment-ingresscontroller: default
  type: LoadBalancer
```

- b. Save the **metallb-loadbalancer-service.yaml** file.
- c. Enter the following command to apply the YAML configuration:

```
$ oc apply -f metallb-loadbalancer-service.yaml
```

- d. Enter the following command to reach the OpenShift Container Platform console:

```
$ curl -kI https://console-openshift-console.apps.example.krnl.es
```

Example output

```
HTTP/1.1 200 OK
```

- e. Check the **clusterversion** and **clusteroperator** values to verify that everything is running. Enter the following command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get clusterversion,co
```

Example output

NAME	VERSION	AVAILABLE	PROGRESSING	SINCE		
STATUS						
clusterversion.config.openshift.io/version	4.x.y	True	False	3m32s	Cluster version is 4.x.y	
NAME	VERSION	AVAILABLE	PROGRESSING	DEGRADED	SINCE	MESSAGE
clusteroperator.config.openshift.io/console	4.x.y	True	False		3m50s	
clusteroperator.config.openshift.io/ingress	4.x.y	True	False		53m	

Replace **<4.x.y>** with the supported OpenShift Container Platform version that you want to use, for example, **4.18.0-multi**.

Additional resources

- [About MetalLB and the MetalLB Operator](#)

5.2.4. Enabling machine health checks on bare metal

You can enable machine health checks on bare metal to repair and replace unhealthy managed cluster nodes automatically. You must have additional agent machines that are ready to install in the managed cluster.

Consider the following limitations before enabling machine health checks:

- You cannot modify the **MachineHealthCheck** object.
- Machine health checks replace nodes only when at least two nodes stay in the **False** or **Unknown** status for more than 8 minutes.

After you enable machine health checks for the managed cluster nodes, the **MachineHealthCheck** object is created in your hosted cluster.

Procedure

To enable machine health checks in your hosted cluster, modify the **NodePool** resource. Complete the following steps:

- Verify that the **spec.nodeDrainTimeout** value in your **NodePool** resource is greater than **0s**. Replace **<hosted_cluster_namespace>** with the name of your hosted cluster namespace and **<nodepool_name>** with the node pool name. Run the following command:

```
$ oc get nodepool -n <hosted_cluster_namespace> <nodepool_name> -o yaml | grep
nodeDrainTimeout
```

Example output

```
nodeDrainTimeout: 30s
```

- If the **spec.nodeDrainTimeout** value is not greater than **0s**, modify the value by running the following command:

```
$ oc patch nodepool -n <hosted_cluster_namespace> <nodepool_name> -p '{"spec":'
{"nodeDrainTimeout": "30m"}' --type=merge
```

- Enable machine health checks by setting the **spec.management.autoRepair** field to **true** in the **NodePool** resource. Run the following command:

```
$ oc patch nodepool -n <hosted_cluster_namespace> <nodepool_name> -p '{"spec":'
{"management": {"autoRepair":true}}}' --type=merge
```

- Verify that the **NodePool** resource is updated with the **autoRepair: true** value by running the following command:

```
$ oc get nodepool -n <hosted_cluster_namespace> <nodepool_name> -o yaml | grep
autoRepair
```

5.2.5. Disabling machine health checks on bare metal

To disable machine health checks for the managed cluster nodes, modify the **NodePool** resource.

Procedure

- Disable machine health checks by setting the **spec.management.autoRepair** field to **false** in the **NodePool** resource. Run the following command:

```
$ oc patch nodepool -n <hosted_cluster_namespace> <nodepool_name> -p '{"spec":'
{"management": {"autoRepair":false}}}' --type=merge
```

- Verify that the **NodePool** resource is updated with the **autoRepair: false** value by running the following command:

```
$ oc get nodepool -n <hosted_cluster_namespace> <nodepool_name> -o yaml | grep autoRepair
```

Additional resources

- [Deploying machine health checks](#)

5.3. MANAGING HOSTED CONTROL PLANES ON OPENSHIFT VIRTUALIZATION

After you deploy a hosted cluster on OpenShift Virtualization, you can manage the cluster by completing the following procedures.

5.3.1. Accessing the hosted cluster

You can access the hosted cluster by either getting the **kubeconfig** file and **kubeadmin** credential directly from resources, or by using the **hcp** command-line interface to generate a **kubeconfig** file.

Prerequisites

To access the hosted cluster by getting the **kubeconfig** file and credentials directly from resources, you must be familiar with the access secrets for hosted clusters. The *hosted cluster (hosting)* namespace contains hosted cluster resources and the access secrets. The *hosted control plane* namespace is where the hosted control plane runs.

The secret name formats are as follows:

- **kubeconfig** secret: **<hosted_cluster_namespace>-<name>-admin-kubeconfig** (clusters-hypershift-demo-admin-kubeconfig)
- **kubeadmin** password secret: **<hosted_cluster_namespace>-<name>-kubeadmin-password** (clusters-hypershift-demo-kubeadmin-password)

The **kubeconfig** secret contains a Base64-encoded **kubeconfig** field, which you can decode and save into a file to use with the following command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get nodes
```

The **kubeadmin** password secret is also Base64-encoded. You can decode it and use the password to log in to the API server or console of the hosted cluster.

Procedure

- To access the hosted cluster by using the **hcp** CLI to generate the **kubeconfig** file, take the following steps:
 1. Generate the **kubeconfig** file by entering the following command:

```
$ hcp create kubeconfig --namespace <hosted_cluster_namespace> \
--name <hosted_cluster_name> > <hosted_cluster_name>.kubeconfig
```
 2. After you save the **kubeconfig** file, you can access the hosted cluster by entering the following example command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get nodes
```

5.3.2. Enabling node auto-scaling for the hosted cluster

When you need more capacity in your hosted cluster and spare agents are available, you can enable auto-scaling to install new worker nodes.

Procedure

1. To enable auto-scaling, enter the following command:

```
$ oc -n <hosted_cluster_namespace> patch nodepool <hosted_cluster_name> \
--type=json \
-p '[{"op": "remove", "path": "/spec/replicas"}, {"op": "add", "path": "/spec/autoScaling",
"value": { "max": 5, "min": 2 }}]'
```



NOTE

In the example, the minimum number of nodes is 2, and the maximum is 5. The maximum number of nodes that you can add might be bound by your platform. For example, if you use the Agent platform, the maximum number of nodes is bound by the number of available agents.

2. Create a workload that requires a new node.

- a. Create a YAML file that contains the workload configuration, by using the following example:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: reversewords
    name: reversewords
    namespace: default
  spec:
    replicas: 40
    selector:
      matchLabels:
        app: reversewords
    strategy: {}
    template:
      metadata:
        creationTimestamp: null
        labels:
          app: reversewords
      spec:
        containers:
        - image: quay.io/mavazque/reversewords:latest
          name: reversewords
        resources:
```

```

  requests:
    memory: 2Gi
  status: {}

```

- b. Save the file as **workload-config.yaml**.
- c. Apply the YAML by entering the following command:

```
$ oc apply -f workload-config.yaml
```

3. Extract the **admin-kubeconfig** secret by entering the following command:

```

$ oc extract -n <hosted_cluster_namespace> \
secret/<hosted_cluster_name>-admin-kubeconfig \
--to=./hostedcluster-secrets --confirm

```

Example output

```
hostedcluster-secrets/kubeconfig
```

4. You can check if new nodes are in the **Ready** status by entering the following command:

```
$ oc --kubeconfig ./hostedcluster-secrets get nodes
```

5. To remove the node, delete the workload by entering the following command:

```

$ oc --kubeconfig ./hostedcluster-secrets -n <namespace> \
delete deployment <deployment_name>

```

6. Wait for several minutes to pass without requiring the additional capacity. On the Agent platform, the agent is decommissioned and can be reused. You can confirm that the node was removed by entering the following command:

```
$ oc --kubeconfig ./hostedcluster-secrets get nodes
```



NOTE

For IBM Z® agents, if you are using an OSA network device in Processor Resource/Systems Manager (PR/SM) mode, auto scaling is not supported. You must delete the old agent manually and scale up the node pool because the new agent joins during the scale down process.

5.3.3. Configuring storage for hosted control planes on OpenShift Virtualization

If you do not provide any advanced storage configuration, the default storage class is used for the KubeVirt virtual machine (VM) images, the KubeVirt Container Storage Interface (CSI) mapping, and the etcd volumes.

The following table lists the capabilities that the infrastructure must provide to support persistent storage in a hosted cluster:

Table 5.1. Persistent storage modes in a hosted cluster

Infrastructure CSI provider	Hosted cluster CSI provider	Hosted cluster capabilities
Any RWX Block CSI provider	kubevirt-csi	<ul style="list-style-type: none"> Basic RWO Block and File Basic RWX Block and Snapshot CSI volume cloning
Any RWX Block CSI provider	Red Hat OpenShift Data Foundation	Red Hat OpenShift Data Foundation feature set. External mode has a smaller footprint and uses a standalone Red Hat Ceph Storage. Internal mode has a larger footprint, but is self-contained and suitable for use cases that require expanded capabilities such as RWX File .



NOTE

OpenShift Virtualization handles storage on hosted clusters, which especially helps customers whose requirements are limited to block storage.

5.3.3.1. Mapping KubeVirt CSI storage classes

KubeVirt CSI supports mapping a infrastructure storage class that is capable of **ReadWriteMany** (RWX) access. You can map the infrastructure storage class to hosted storage class during cluster creation.

Procedure

- To map the infrastructure storage class to the hosted storage class, use the **--infra-storage-class-mapping** argument by running the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \ ①
--node-pool-replicas <worker_node_count> \ ②
--pull-secret <path_to_pull_secret> \ ③
--memory <memory> \ ④
--cores <cpu> \ ⑤
--infra-storage-class-mapping=<infrastructure_storage_class>/<hosted_storage_class> \ ⑥
```

- Specify the name of your hosted cluster, for instance, **example**.
- Specify the worker count, for example, **2**.
- Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- Specify a value for memory, for example, **8Gi**.

- 5** Specify a value for CPU, for example, **2**.
- 6** Replace **<infrastructure_storage_class>** with the infrastructure storage class name and **<hosted_storage_class>** with the hosted cluster storage class name. You can use the **--infra-storage-class-mapping** argument multiple times within the **hcp create cluster** command.

After you create the hosted cluster, the infrastructure storage class is visible within the hosted cluster. When you create a Persistent Volume Claim (PVC) within the hosted cluster that uses one of those storage classes, KubeVirt CSI provisions that volume by using the infrastructure storage class mapping that you configured during cluster creation.



NOTE

KubeVirt CSI supports mapping only an infrastructure storage class that is capable of RWX access.

The following table shows how volume and access mode capabilities map to KubeVirt CSI storage classes:

Table 5.2. Mapping KubeVirt CSI storage classes to access and volume modes

Infrastructure CSI capability	Hosted cluster CSI capability	VM live migration support	Notes
RWX: Block or Filesystem	ReadWriteOnce (RWO) Block or Filesystem RWX Block only	Supported	Use Block mode because Filesystem volume mode results in degraded hosted Block mode performance. RWX Block volume mode is supported only when the hosted cluster is OpenShift Container Platform 4.16 or later.
RWO Block storage	RWO Block storage or Filesystem	Not supported	Lack of live migration support affects the ability to update the underlying infrastructure cluster that hosts the KubeVirt VMs.

Infrastructure CSI capability	Hosted cluster CSI capability	VM live migration support	Notes
RWO FileSystem	RWO Block or Filesystem	Not supported	Lack of live migration support affects the ability to update the underlying infrastructure cluster that hosts the KubeVirt VMs. Use of the infrastructure Filesystem volume mode results in degraded hosted Block mode performance.

5.3.3.2. Mapping a single KubeVirt CSI volume snapshot class

You can expose your infrastructure volume snapshot class to the hosted cluster by using KubeVirt CSI.

Procedure

- To map your volume snapshot class to the hosted cluster, use the **--infra-volumesnapshot-class-mapping** argument when creating a hosted cluster. Run the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \ ①
--node-pool-replicas <worker_node_count> \ ②
--pull-secret <path_to_pull_secret> \ ③
--memory <memory> \ ④
--cores <cpu> \ ⑤
--infra-storage-class-mapping=<infrastructure_storage_class>/<hosted_storage_class> \
⑥
--infra-volumesnapshot-class-mapping=
<infrastructure_volume_snapshot_class>/<hosted_volume_snapshot_class> ⑦
```

- Specify the name of your hosted cluster, for instance, **example**.
- Specify the worker count, for example, **2**.
- Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- Specify a value for memory, for example, **8Gi**.
- Specify a value for CPU, for example, **2**.
- Replace **<infrastructure_storage_class>** with the storage class present in the infrastructure cluster. Replace **<hosted_storage_class>** with the storage class present in the hosted cluster.
- Replace **<infrastructure_volume_snapshot_class>** with the volume snapshot class

**NOTE**

If you do not use the **--infra-storage-class-mapping** and **--infra-volumesnapshot-class-mapping** arguments, a hosted cluster is created with the default storage class and the volume snapshot class. Therefore, you must set the default storage class and the volume snapshot class in the infrastructure cluster.

5.3.3.3. Mapping multiple KubeVirt CSI volume snapshot classes

You can map multiple volume snapshot classes to the hosted cluster by assigning them to a specific group. The infrastructure storage class and the volume snapshot class are compatible with each other only if they belong to a same group.

Procedure

- To map multiple volume snapshot classes to the hosted cluster, use the **group** option when creating a hosted cluster. Run the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \ ①
--node-pool-replicas <worker_node_count> \ ②
--pull-secret <path_to_pull_secret> \ ③
--memory <memory> \ ④
--cores <cpu> \ ⑤
--infra-storage-class-mapping=
<infrastructure_storage_class>/<hosted_storage_class>,group=<group_name> \ ⑥
--infra-storage-class-mapping=
<infrastructure_storage_class>/<hosted_storage_class>,group=<group_name> \
--infra-storage-class-mapping=
<infrastructure_storage_class>/<hosted_storage_class>,group=<group_name> \
--infra-volumesnapshot-class-mapping=
<infrastructure_volume_snapshot_class>/<hosted_volume_snapshot_class>,group=
<group_name> \ ⑦
--infra-volumesnapshot-class-mapping=
<infrastructure_volume_snapshot_class>/<hosted_volume_snapshot_class>,group=
<group_name>
```

- Specify the name of your hosted cluster, for instance, **example**.
- Specify the worker count, for example, **2**.
- Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- Specify a value for memory, for example, **8Gi**.
- Specify a value for CPU, for example, **2**.
- Replace **<infrastructure_storage_class>** with the storage class present in the infrastructure cluster. Replace **<hosted_storage_class>** with the storage class present in the hosted cluster. Replace **<group_name>** with the group name. For example, **infra-storage-class-mygroup/hosted-storage-class-mygroup,group=mygroup** and **infra-storage-class-mymap/hosted-storage-class-mymap,group=mymap**.
-

Replace `<infrastructure_volume_snapshot_class>` with the volume snapshot class present in the infrastructure cluster. Replace `<hosted_volume_snapshot_class>` with

5.3.3.4. Configuring KubeVirt VM root volume

At cluster creation time, you can configure the storage class that is used to host the KubeVirt VM root volumes by using the `--root-volume-storage-class` argument.

Procedure

- To set a custom storage class and volume size for KubeVirt VMs, run the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \ ①
--node-pool-replicas <worker_node_count> \ ②
--pull-secret <path_to_pull_secret> \ ③
--memory <memory> \ ④
--cores <cpu> \ ⑤
--root-volume-storage-class <root_volume_storage_class> \ ⑥
--root-volume-size <volume_size> ⑦
```

- Specify the name of your hosted cluster, for instance, **example**.
- Specify the worker count, for example, **2**.
- Specify the path to your pull secret, for example, `/user/name/pullsecret`.
- Specify a value for memory, for example, **8Gi**.
- Specify a value for CPU, for example, **2**.
- Specify a name of the storage class to host the KubeVirt VM root volumes, for example, **ocs-storagecluster-ceph-rbd**.
- Specify the volume size, for example, **64**.

As a result, you get a hosted cluster created with VMs hosted on PVCs.

5.3.3.5. Enabling KubeVirt VM image caching

You can use KubeVirt VM image caching to optimize both cluster startup time and storage usage. KubeVirt VM image caching supports the use of a storage class that is capable of smart cloning and the **ReadWriteMany** access mode. For more information about smart cloning, see [Cloning a data volume using smart-cloning](#).

Image caching works as follows:

- The VM image is imported to a PVC that is associated with the hosted cluster.
- A unique clone of that PVC is created for every KubeVirt VM that is added as a worker node to the cluster.

Image caching reduces VM startup time by requiring only a single image import. It can further reduce overall cluster storage usage when the storage class supports copy-on-write cloning.

Procedure

- To enable image caching, during cluster creation, use the **--root-volume-cache-strategy=PVC** argument by running the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \ ①
--node-pool-replicas <worker_node_count> \ ②
--pull-secret <path_to_pull_secret> \ ③
--memory <memory> \ ④
--cores <cpu> \ ⑤
--root-volume-cache-strategy=PVC ⑥
```

- Specify the name of your hosted cluster, for instance, **example**.
- Specify the worker count, for example, **2**.
- Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- Specify a value for memory, for example, **8Gi**.
- Specify a value for CPU, for example, **2**.
- Specify a strategy for image caching, for example, **PVC**.

5.3.3.6. KubeVirt CSI storage security and isolation

KubeVirt Container Storage Interface (CSI) extends the storage capabilities of the underlying infrastructure cluster to hosted clusters. The CSI driver ensures secure and isolated access to the infrastructure storage classes and hosted clusters by using the following security constraints:

- The storage of a hosted cluster is isolated from the other hosted clusters.
- Worker nodes in a hosted cluster do not have a direct API access to the infrastructure cluster. The hosted cluster can provision storage on the infrastructure cluster only through the controlled KubeVirt CSI interface.
- The hosted cluster does not have access to the KubeVirt CSI cluster controller. As a result, the hosted cluster cannot access arbitrary storage volumes on the infrastructure cluster that are not associated with the hosted cluster. The KubeVirt CSI cluster controller runs in a pod in the hosted control plane namespace.
- Role-based access control (RBAC) of the KubeVirt CSI cluster controller limits the persistent volume claim (PVC) access to only the hosted control plane namespace. Therefore, KubeVirt CSI components cannot access storage from the other namespaces.

Additional resources

- [Cloning a data volume using smart-cloning](#)

5.3.3.7. Configuring etcd storage

At cluster creation time, you can configure the storage class that is used to host etcd data by using the **--etcd-storage-class** argument.

Procedure

- To configure a storage class for etcd, run the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \ ①
--node-pool-replicas <worker_node_count> \ ②
--pull-secret <path_to_pull_secret> \ ③
--memory <memory> \ ④
--cores <cpu> \ ⑤
--etcd-storage-class=<etcd_storage_class_name> ⑥
```

- Specify the name of your hosted cluster, for instance, **example**.
- Specify the worker count, for example, **2**.
- Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- Specify a value for memory, for example, **8Gi**.
- Specify a value for CPU, for example, **2**.
- Specify the etcd storage class name, for example, **lvm-storageclass**. If you do not provide an **--etcd-storage-class** argument, the default storage class is used.

5.3.4. Attaching NVIDIA GPU devices by using the hcp CLI

You can attach one or more NVIDIA graphics processing unit (GPU) devices to node pools by using the **hcp** command-line interface (CLI) in a hosted cluster on OpenShift Virtualization.



IMPORTANT

Attaching NVIDIA GPU devices to node pools is a Technology Preview feature only. Technology Preview features are not supported with Red Hat production service level agreements (SLAs) and might not be functionally complete. Red Hat does not recommend using them in production. These features provide early access to upcoming product features, enabling customers to test functionality and provide feedback during the development process.

For more information about the support scope of Red Hat Technology Preview features, see [Technology Preview Features Support Scope](#).

Prerequisites

- You have exposed the NVIDIA GPU device as a resource on the node where the GPU device resides. For more information, see [NVIDIA GPU Operator with OpenShift Virtualization](#).
- You have exposed the NVIDIA GPU device as an [extended resource](#) on the node to assign it to node pools.

Procedure

- You can attach the GPU device to node pools during cluster creation by running the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \①
--node-pool-replicas <worker_node_count> \②
--pull-secret <path_to_pull_secret> \③
--memory <memory> \④
--cores <cpu> \⑤
--host-device-name=",count:<value>" \⑥
```

- ① Specify the name of your hosted cluster, for instance, **example**.
- ② Specify the worker count, for example, **3**.
- ③ Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- ④ Specify a value for memory, for example, **16Gi**.
- ⑤ Specify a value for CPU, for example, **2**.
- ⑥ Specify the GPU device name and the count, for example, **--host-device-name="nvidia-a100,count:2"**. The **--host-device-name** argument takes the name of the GPU device from the infrastructure node and an optional count that represents the number of GPU devices you want to attach to each virtual machine (VM) in node pools. The default count is **1**. For example, if you attach 2 GPU devices to 3 node pool replicas, all 3 VMs in the node pool are attached to the 2 GPU devices.

TIP

You can use the **--host-device-name** argument multiple times to attach multiple devices of different types.

5.3.5. Attaching NVIDIA GPU devices by using the NodePool resource

You can attach one or more NVIDIA graphics processing unit (GPU) devices to node pools by configuring the **nodepool.spec.platform.kubevirt.hostDevices** field in the **NodePool** resource.



IMPORTANT

Attaching NVIDIA GPU devices to node pools is a Technology Preview feature only. Technology Preview features are not supported with Red Hat production service level agreements (SLAs) and might not be functionally complete. Red Hat does not recommend using them in production. These features provide early access to upcoming product features, enabling customers to test functionality and provide feedback during the development process.

For more information about the support scope of Red Hat Technology Preview features, see [Technology Preview Features Support Scope](#).

Procedure

- Attach one or more GPU devices to node pools:

- To attach a single GPU device, configure the **NodePool** resource by using the following example configuration:

```
apiVersion: hypershift.openshift.io/v1beta1
kind: NodePool
metadata:
  name: <hosted_cluster_name> ①
  namespace: <hosted_cluster_namespace> ②
spec:
  arch: amd64
  clusterName: <hosted_cluster_name>
  management:
    autoRepair: false
    upgradeType: Replace
  nodeDrainTimeout: 0s
  nodeVolumeDetachTimeout: 0s
  platform:
    kubevirt:
      attachDefaultNetwork: true
      compute:
        cores: <cpu> ③
        memory: <memory> ④
      hostDevices: ⑤
        - count: <count> ⑥
          deviceName: <gpu_device_name> ⑦
      networkInterfaceMultiqueue: Enable
      rootVolume:
        persistent:
          size: 32Gi
        type: Persistent
      type: KubeVirt
  replicas: <worker_node_count> ⑧
```

- ① Specify the name of your hosted cluster, for instance, **example**.
- ② Specify the name of the hosted cluster namespace, for example, **clusters**.
- ③ Specify a value for CPU, for example, **2**.
- ④ Specify a value for memory, for example, **16Gi**.
- ⑤ The **hostDevices** field defines a list of different types of GPU devices that you can attach to node pools.
- ⑥ Specify the number of GPU devices you want to attach to each virtual machine (VM) in node pools. For example, if you attach 2 GPU devices to 3 node pool replicas, all 3 VMs in the node pool are attached to the 2 GPU devices. The default count is **1**.
- ⑦ Specify the GPU device name, for example, **nvidia-a100**.
- ⑧ Specify the worker count, for example, **3**.

- To attach multiple GPU devices, configure the **NodePool** resource by using the following example configuration:

```
apiVersion: hypershift.openshift.io/v1beta1
kind: NodePool
metadata:
  name: <hosted_cluster_name>
  namespace: <hosted_cluster_namespace>
spec:
  arch: amd64
  clusterName: <hosted_cluster_name>
  management:
    autoRepair: false
    upgradeType: Replace
    nodeDrainTimeout: 0s
    nodeVolumeDetachTimeout: 0s
  platform:
    kubevirt:
      attachDefaultNetwork: true
    compute:
      cores: <cpu>
      memory: <memory>
    hostDevices:
      - count: <count>
        deviceName: <gpu_device_name>
      - count: <count>
        deviceName: <gpu_device_name>
      - count: <count>
        deviceName: <gpu_device_name>
      - count: <count>
        deviceName: <gpu_device_name>
    networkInterfaceMultiqueue: Enable
  rootVolume:
    persistent:
      size: 32Gi
    type: Persistent
  type: KubeVirt
  replicas: <worker_node_count>
```

5.3.6. Spreading node pool VMs by using **topologySpreadConstraint**

By default, KubeVirt virtual machines (VMs) created by a node pool are scheduled on any available nodes that have the capacity to run the VMs. By default, the **topologySpreadConstraint** constraint is set to schedule VMs on multiple nodes.

In some scenarios, node pool VMs might run on the same node, which can cause availability issues. To avoid distribution of VMs on a single node, use the descheduler to continuously honor the **topologySpreadConstraint** constraint to spread VMs on multiple nodes.

Prerequisites

- You installed the Kube Descheduler Operator. For more information, see "Installing the descheduler".

Procedure

- Open the **KubeDescheduler** custom resource (CR) by entering the following command, and then modify the **KubeDescheduler** CR to use the **SoftTopologyAndDuplicates** and **DevKubeVirtRelieveAndMigrate** profiles so that you maintain the **topologySpreadConstraint** constraint settings.
The **KubeDescheduler** CR named **cluster** runs in the **openshift-kube-descheduler-operator** namespace.

```
$ oc edit kubedescheduler cluster -n openshift-kube-descheduler-operator
```

Example KubeDescheduler configuration

```
apiVersion: operator.openshift.io/v1
kind: KubeDescheduler
metadata:
  name: cluster
  namespace: openshift-kube-descheduler-operator
spec:
  mode: Automatic
  managementState: Managed
  deschedulingIntervalSeconds: 30 1
  profiles:
    - SoftTopologyAndDuplicates 2
    - DevKubeVirtRelieveAndMigrate 3
  profileCustomizations:
    devDeviationThresholds: AsymmetricLow
    devActualUtilizationProfile: PrometheusCPUCombined
# ...
```

- 1** Sets the number of seconds between the descheduler running cycles.
- 2** This profile evicts pods that follow the soft topology constraint: **whenUnsatisfiable: ScheduleAnyway**.
- 3** This profile balances resource usage between nodes and enables the strategies, such as **RemovePodsHavingTooManyRestarts** and **LowNodeUtilization**.

Additional resources

- [Installing the descheduler](#)

5.4. MANAGING HOSTED CONTROL PLANES ON NON-BARE-METAL AGENT MACHINES

After you deploy hosted control planes on non-bare-metal agent machines, you can manage a hosted cluster by completing the following tasks.

5.4.1. Accessing the hosted cluster

You can access the hosted cluster by either getting the **kubeconfig** file and **kubeadmin** credential directly from resources, or by using the **hcp** command-line interface to generate a **kubeconfig** file.

Prerequisites

To access the hosted cluster by getting the **kubeconfig** file and credentials directly from resources, you must be familiar with the access secrets for hosted clusters. The *hosted cluster (hosting)* namespace contains hosted cluster resources and the access secrets. The *hosted control plane* namespace is where the hosted control plane runs.

The secret name formats are as follows:

- **kubeconfig** secret: <hosted_cluster_namespace>-<name>-admin-kubeconfig. For example, **clusters-hypershift-demo-admin-kubeconfig**.
- **kubeadmin** password secret: <hosted_cluster_namespace>-<name>-kubeadmin-password. For example, **clusters-hypershift-demo-kubeadmin-password**.

The **kubeconfig** secret contains a Base64-encoded **kubeconfig** field, which you can decode and save into a file to use with the following command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get nodes
```

The **kubeadmin** password secret is also Base64-encoded. You can decode it and use the password to log in to the API server or console of the hosted cluster.

Procedure

- To access the hosted cluster by using the **hcp** CLI to generate the **kubeconfig** file, take the following steps:

1. Generate the **kubeconfig** file by entering the following command:

```
$ hcp create kubeconfig --namespace <hosted_cluster_namespace> \
--name <hosted_cluster_name> > <hosted_cluster_name>.kubeconfig
```

2. After you save the **kubeconfig** file, you can access the hosted cluster by entering the following example command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get nodes
```

5.4.2. Scaling the NodePool object for a hosted cluster

You can scale up the **NodePool** object by adding nodes to your hosted cluster. When you scale a node pool, consider the following information:

- When you scale a replica by the node pool, a machine is created. For every machine, the Cluster API provider finds and installs an Agent that meets the requirements that are specified in the node pool specification. You can monitor the installation of an Agent by checking its status and conditions.
- When you scale down a node pool, Agents are unbound from the corresponding cluster. Before you can reuse the Agents, you must restart them by using the Discovery image.

Procedure

1. Scale the **NodePool** object to two nodes:

```
$ oc -n <hosted_cluster_namespace> scale nodepool <nodepool_name> --replicas 2
```

The Cluster API agent provider randomly picks two agents that are then assigned to the hosted cluster. Those agents go through different states and finally join the hosted cluster as OpenShift Container Platform nodes. The agents pass through states in the following order:

- **binding**
- **discovering**
- **insufficient**
- **installing**
- **installing-in-progress**
- **added-to-existing-cluster**

2. Enter the following command:

```
$ oc -n <hosted_control_plane_namespace> get agent
```

Example output

NAME	CLUSTER	APPROVED	ROLE	STAGE
4dac1ab2-7dd5-4894-a220-6a3473b67ee6	hypercluster1	true		auto-assign
d9198891-39f4-4930-a679-65fb142b108b		true		auto-assign
da503cf1-a347-44f2-875c-4960ddb04091	hypercluster1	true		auto-assign

3. Enter the following command:

```
$ oc -n <hosted_control_plane_namespace> get agent \
-o jsonpath='{range .items[*]}BMH: {@.metadata.labels.agent-install\\.openshift\\.io/bmh} \
Agent: {@.metadata.name} State: {@.status.debugInfo.state}{\"\\n\"}{end}'
```

Example output

```
BMH: ocp-worker-2 Agent: 4dac1ab2-7dd5-4894-a220-6a3473b67ee6 State: binding
BMH: ocp-worker-0 Agent: d9198891-39f4-4930-a679-65fb142b108b State: known-unbound
BMH: ocp-worker-1 Agent: da503cf1-a347-44f2-875c-4960ddb04091 State: insufficient
```

4. Obtain the kubeconfig for your new hosted cluster by entering the extract command:

```
$ oc extract -n <hosted_cluster_namespace> \
secret/<hosted_cluster_name>-admin-kubeconfig --to=- \
> kubeconfig-<hosted_cluster_name>
```

5. After the agents reach the **added-to-existing-cluster** state, verify that you can see the OpenShift Container Platform nodes in the hosted cluster by entering the following command:

```
$ oc --kubeconfig kubeconfig-<hosted_cluster_name> get nodes
```

Example output

NAME	STATUS	ROLES	AGE	VERSION
ocp-worker-1	Ready	worker	5m41s	v1.24.0+3882f8f
ocp-worker-2	Ready	worker	6m3s	v1.24.0+3882f8f

Cluster Operators start to reconcile by adding workloads to the nodes.

- Enter the following command to verify that two machines were created when you scaled up the **NodePool** object:

```
$ oc -n <hosted_control_plane_namespace> get machines
```

Example output

NAME	CLUSTER	NODENAME	PROVIDERID
PHASE	AGE	VERSION	
hypercluster1-c96b6f675-m5vch	hypercluster1-b2qhl	ocp-worker-1	agent://da503cf1-a347-44f2-875c-4960ddb04091
hypercluster1-c96b6f675-tl42p	hypercluster1-b2qhl	ocp-worker-2	agent://4dac1ab2-7dd5-4894-a220-6a3473b67ee6
	Running	15m	4.x.z

The **clusterversion** reconcile process eventually reaches a point where only Ingress and Console cluster operators are missing.

- Enter the following command:

```
$ oc --kubeconfig kubeconfig-<hosted_cluster_name> get clusterversion,co
```

Example output

NAME	VERSION	AVAILABLE	PROGRESSING	SINCE
STATUS				
clusterversion.config.openshift.io/version	False	True	40m	Unable to apply 4.x.z: the cluster operator console has not yet successfully rolled out
NAME	VERSION	AVAILABLE		
PROGRESSING	DEGRADED	SINCE	MESSAGE	
clusteroperator.config.openshift.io/console	4.12z	False	False	
False	11m	RouteHealthAvailable: failed to GET route (https://console-openshift-console.apps.hypercluster1.domain.com): Get "https://console-openshift-console.apps.hypercluster1.domain.com": dial tcp 10.19.3.29:443: connect: connection refused		
clusteroperator.config.openshift.io/csi-snapshot-controller	4.12z	True		
False	False	10m		
clusteroperator.config.openshift.io/dns	4.12z	True	False	
False	9m16s			

5.4.2.1. Adding node pools

You can create node pools for a hosted cluster by specifying a name, number of replicas, and any additional information, such as an agent label selector.

Procedure

- To create a node pool, enter the following information:

```
$ hcp create nodepool agent \
--cluster-name <hosted_cluster_name> \ ①
--name <nodepool_name> \ ②
--node-count <worker_node_count> \ ③
--agentLabelSelector size=medium ④
```

- Replace **<hosted_cluster_name>** with your hosted cluster name.
- Replace **<nodepool_name>** with the name of your node pool, for example, **<hosted_cluster_name>-extra-cpu**.
- Replace **<worker_node_count>** with the worker node count, for example, **2**.
- The **--agentLabelSelector** flag is optional. The node pool uses agents with the **size=medium** label.

- Check the status of the node pool by listing **nodepool** resources in the **clusters** namespace:

```
$ oc get nodepools --namespace clusters
```

- Extract the **admin-kubeconfig** secret by entering the following command:

```
$ oc extract -n <hosted_control_plane_namespace> secret/admin-kubeconfig --
to=./hostedcluster-secrets --confirm
```

Example output

```
hostedcluster-secrets/kubeconfig
```

- After some time, you can check the status of the node pool by entering the following command:

```
$ oc --kubeconfig ./hostedcluster-secrets get nodes
```

Verification

- Verify that the number of available node pools match the number of expected node pools by entering this command:

```
$ oc get nodepools --namespace clusters
```

5.4.2.2. Enabling node auto-scaling for the hosted cluster

When you need more capacity in your hosted cluster and spare agents are available, you can enable auto-scaling to install new worker nodes.

Procedure

- To enable auto-scaling, enter the following command:

```
$ oc -n <hosted_cluster_namespace> patch nodepool <hosted_cluster_name> \
--type=json \
-p '[{"op": "remove", "path": "/spec/replicas"}, {"op": "add", "path": "/spec/autoScaling", \
"value": { "max": 5, "min": 2 }}]'
```



NOTE

In the example, the minimum number of nodes is 2, and the maximum is 5. The maximum number of nodes that you can add might be bound by your platform. For example, if you use the Agent platform, the maximum number of nodes is bound by the number of available agents.

2. Create a workload that requires a new node.
 - a. Create a YAML file that contains the workload configuration, by using the following example:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: reversewords
    name: reversewords
    namespace: default
spec:
  replicas: 40
  selector:
    matchLabels:
      app: reversewords
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: reversewords
    spec:
      containers:
        - image: quay.io/mavazque/reversewords:latest
          name: reversewords
      resources:
        requests:
          memory: 2Gi
status: {}
```

- b. Save the file as **workload-config.yaml**.
- c. Apply the YAML by entering the following command:

```
$ oc apply -f workload-config.yaml
```

3. Extract the **admin-kubeconfig** secret by entering the following command:

```
$ oc extract -n <hosted_cluster_namespace> \
secret/<hosted_cluster_name>-admin-kubeconfig \
--to=../hostedcluster-secrets --confirm
```

Example output

```
hostedcluster-secrets/kubeconfig
```

4. You can check if new nodes are in the **Ready** status by entering the following command:

```
$ oc --kubeconfig ./hostedcluster-secrets get nodes
```

5. To remove the node, delete the workload by entering the following command:

```
$ oc --kubeconfig ./hostedcluster-secrets -n <namespace> \
delete deployment <deployment_name>
```

6. Wait for several minutes to pass without requiring the additional capacity. On the Agent platform, the agent is decommissioned and can be reused. You can confirm that the node was removed by entering the following command:

```
$ oc --kubeconfig ./hostedcluster-secrets get nodes
```



NOTE

For IBM Z® agents, if you are using an OSA network device in Processor Resource/Systems Manager (PR/SM) mode, auto scaling is not supported. You must delete the old agent manually and scale up the node pool because the new agent joins during the scale down process.

5.4.2.3. Disabling node auto-scaling for the hosted cluster

To disable node auto-scaling, complete the following procedure.

Procedure

- Enter the following command to disable node auto-scaling for the hosted cluster:

```
$ oc -n <hosted_cluster_namespace> patch nodepool <hosted_cluster_name> \
--type=json \
-p '[{"op": "remove", "path": "/spec/autoScaling"}, {"op": "add", "path": "/spec/replicas", \
"value": <specify_value_to_scale_replicas>}]
```

The command removes "**spec.autoScaling**" from the YAML file, adds "**spec.replicas**", and sets "**spec.replicas**" to the integer value that you specify.

Additional resources

- [Scaling down the data plane to zero](#)

5.4.3. Handling ingress in a hosted cluster on non-bare-metal agent machines

Every OpenShift Container Platform cluster has a default application Ingress Controller that typically has an external DNS record associated with it. For example, if you create a hosted cluster named **example** with the base domain **krnl.es**, you can expect the wildcard domain ***.apps.example.krnl.es** to be routable.

Procedure

To set up a load balancer and wildcard DNS record for the ***.apps** domain, perform the following actions on your guest cluster:

1. Deploy MetalLB by creating a YAML file that contains the configuration for the MetalLB Operator:

```
apiVersion: v1
kind: Namespace
metadata:
  name: metallb
  labels:
    openshift.io/cluster-monitoring: "true"
  annotations:
    workload.openshift.io/allowed: management
---
apiVersion: operators.coreos.com/v1
kind: OperatorGroup
metadata:
  name: metallb-operator-operatorgroup
  namespace: metallb
---
apiVersion: operators.coreos.com/v1alpha1
kind: Subscription
metadata:
  name: metallb-operator
  namespace: metallb
spec:
  channel: "stable"
  name: metallb-operator
  source: redhat-operators
  sourceNamespace: openshift-marketplace
```

2. Save the file as **metallb-operator-config.yaml**.

3. Enter the following command to apply the configuration:

```
$ oc apply -f metallb-operator-config.yaml
```

4. After the Operator is running, create the MetalLB instance:

- a. Create a YAML file that contains the configuration for the MetalLB instance:

```
apiVersion: metallb.io/v1beta1
kind: MetalLB
metadata:
  name: metallb
  namespace: metallb
```

- b. Save the file as **metallb-instance-config.yaml**.

- c. Create the MetalLB instance by entering this command:

```
$ oc apply -f metallb-instance-config.yaml
```

5. Create an **IPAddressPool** resource with a single IP address. This IP address must be on the same subnet as the network that the cluster nodes use.

- a. Create a file, such as **ipaddresspool.yaml**, with content like the following example:

```
apiVersion: metallb.io/v1beta1
kind: IPAddressPool
metadata:
  namespace: metallb
  name: <ip_address_pool_name> ①
spec:
  addresses:
    - <ingress_ip>-<ingress_ip> ②
  autoAssign: false
```

- ① Specify the **IPAddressPool** resource name.
- ② Specify the IP address for your environment. For example, **192.168.122.23**.

- b. Apply the configuration for the IP address pool by entering the following command:

```
$ oc apply -f ipaddresspool.yaml
```

6. Create a L2 advertisement.

- a. Create a file, such as **l2advertisement.yaml**, with content like the following example:

```
apiVersion: metallb.io/v1beta1
kind: L2Advertisement
metadata:
  name: <l2_advertisement_name> ①
  namespace: metallb
spec:
  ipAddressPools:
    - <ip_address_pool_name> ②
```

- ① Specify the **L2Advertisement** resource name.
- ② Specify the **IPAddressPool** resource name.

- b. Apply the configuration by entering the following command:

```
$ oc apply -f l2advertisement.yaml
```

7. After creating a service of the **LoadBalancer** type, MetalLB adds an external IP address for the service.

- a. Configure a new load balancer service that routes ingress traffic to the ingress deployment by creating a YAML file named **metallb-loadbalancer-service.yaml**:

```

kind: Service
apiVersion: v1
metadata:
  annotations:
    metallb.io/address-pool: ingress-public-ip
  name: metallb-ingress
  namespace: openshift-ingress
spec:
  ports:
    - name: http
      protocol: TCP
      port: 80
      targetPort: 80
    - name: https
      protocol: TCP
      port: 443
      targetPort: 443
  selector:
    ingresscontroller.operator.openshift.io/deployment-ingresscontroller: default
  type: LoadBalancer

```

- b. Save the **metallb-loadbalancer-service.yaml** file.
- c. Enter the following command to apply the YAML configuration:

```
$ oc apply -f metallb-loadbalancer-service.yaml
```

- d. Enter the following command to reach the OpenShift Container Platform console:

```
$ curl -kI https://console-openshift-console.apps.example.krnl.es
```

Example output

```
HTTP/1.1 200 OK
```

- e. Check the **clusterversion** and **clusteroperator** values to verify that everything is running. Enter the following command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get clusterversion,co
```

Example output

NAME	VERSION	AVAILABLE	PROGRESSING	SINCE
STATUS				
clusterversion.config.openshift.io/version	4.x.y	True	False	3m32s Cluster version is 4.x.y
NAME	VERSION	AVAILABLE		
PROGRESSING DEGRADED SINCE MESSAGE				
clusteroperator.config.openshift.io/console	4.x.y	True	False	
False 3m50s				
clusteroperator.config.openshift.io/ingress	4.x.y	True	False	
False 53m				

Replace <4.x.y> with the supported OpenShift Container Platform version that you want to use, for example, **4.18.0-multi**.

Additional resources

- [About MetalLB and the MetalLB Operator](#)

5.4.4. Enabling machine health checks on non-bare-metal agent machines

You can enable machine health checks on bare metal to repair and replace unhealthy managed cluster nodes automatically. You must have additional agent machines that are ready to install in the managed cluster.

Consider the following limitations before enabling machine health checks:

- You cannot modify the **MachineHealthCheck** object.
- Machine health checks replace nodes only when at least two nodes stay in the **False** or **Unknown** status for more than 8 minutes.

After you enable machine health checks for the managed cluster nodes, the **MachineHealthCheck** object is created in your hosted cluster.

Procedure

To enable machine health checks in your hosted cluster, modify the **NodePool** resource. Complete the following steps:

1. Verify that the **spec.nodeDrainTimeout** value in your **NodePool** resource is greater than **0s**. Replace <hosted_cluster_namespace> with the name of your hosted cluster namespace and <nodepool_name> with the node pool name. Run the following command:

```
$ oc get nodepool -n <hosted_cluster_namespace> <nodepool_name> -o yaml | grep nodeDrainTimeout
```

Example output

```
nodeDrainTimeout: 30s
```

2. If the **spec.nodeDrainTimeout** value is not greater than **0s**, modify the value by running the following command:

```
$ oc patch nodepool -n <hosted_cluster_namespace> <nodepool_name> -p '{"spec": {"nodeDrainTimeout": "30m"}}' --type=merge
```

3. Enable machine health checks by setting the **spec.management.autoRepair** field to **true** in the **NodePool** resource. Run the following command:

```
$ oc patch nodepool -n <hosted_cluster_namespace> <nodepool_name> -p '{"spec": {"management": {"autoRepair":true}}}' --type=merge
```

4. Verify that the **NodePool** resource is updated with the **autoRepair: true** value by running the following command:

```
$ oc get nodepool -n <hosted_cluster_namespace> <nodepool_name> -o yaml | grep autoRepair
```

5.4.5. Disabling machine health checks on non-bare-metal agent machines

To disable machine health checks for the managed cluster nodes, modify the **NodePool** resource.

Procedure

1. Disable machine health checks by setting the **spec.management.autoRepair** field to **false** in the **NodePool** resource. Run the following command:

```
$ oc patch nodepool -n <hosted_cluster_namespace> <nodepool_name> -p '{"spec": {"management": {"autoRepair": false}}}' --type=merge
```

2. Verify that the **NodePool** resource is updated with the **autoRepair: false** value by running the following command:

```
$ oc get nodepool -n <hosted_cluster_namespace> <nodepool_name> -o yaml | grep autoRepair
```

Additional resources

- [Deploying machine health checks](#)

5.5. MANAGING HOSTED CONTROL PLANES ON IBM POWER

After you deploy hosted control planes on IBM Power, you can manage a hosted cluster by completing the following tasks.

5.5.1. Creating an InfraEnv resource for hosted control planes on IBM Power

An **InfraEnv** is a environment where hosts that are starting the live ISO can join as agents. In this case, the agents are created in the same namespace as your hosted control plane.

You can create an **InfraEnv** resource for hosted control planes on 64-bit x86 bare metal for IBM Power compute nodes.

Procedure

1. Create a YAML file to configure an **InfraEnv** resource. See the following example:

```
apiVersion: agent-install.openshift.io/v1beta1
kind: InfraEnv
metadata:
  name: <hosted_cluster_name> \①
  namespace: <hosted_control_plane_namespace> \②
spec:
  cpuArchitecture: ppc64le
  pullSecretRef:
    name: pull-secret
  sshAuthorizedKey: <path_to_ssh_public_key> \③
```

- 1 Replace <hosted_cluster_name> with the name of your hosted cluster.
- 2 Replace <hosted_control_plane_namespace> with the name of the hosted control plane namespace, for example, **clusters-hosted**.
- 3 Replace <path_to_ssh_public_key> with the path to your SSH public key. The default file path is `~/.ssh/id_rsa.pub`.

2. Save the file as **infraenv-config.yaml**.

3. Apply the configuration by entering the following command:

```
$ oc apply -f infraenv-config.yaml
```

4. To fetch the URL to download the live ISO, which allows IBM Power machines to join as agents, enter the following command:

```
$ oc -n <hosted_control_plane_namespace> get InfraEnv <hosted_cluster_name> \
-o json
```

5.5.2. Adding IBM Power agents to the InfraEnv resource

You can add agents by manually configuring the machine to start with the live ISO.

Procedure

1. Download the live ISO and use it to start a bare metal or a virtual machine (VM) host. You can find the URL for the live ISO in the **status.isoDownloadURL** field, in the **InfraEnv** resource. At startup, the host communicates with the Assisted Service and registers as an agent in the same namespace as the **InfraEnv** resource.
2. To list the agents and some of their properties, enter the following command:

```
$ oc -n <hosted_control_plane_namespace> get agents
```

Example output

NAME	CLUSTER	APPROVED	ROLE	STAGE
86f7ac75-4fc4-4b36-8130-40fa12602218			auto-assign	
e57a637f-745b-496e-971d-1abbf03341ba			auto-assign	

3. After each agent is created, you can optionally set the **installation_disk_id** and **hostname** for an agent:
 - a. To set the **installation_disk_id** field for an agent, enter the following command:

```
$ oc -n <hosted_control_plane_namespace> patch agent <agent_name> -p '{"spec": \
{"installation_disk_id":"<installation_disk_id>","approved":true}}' --type merge
```

- b. To set the **hostname** field for an agent, enter the following command:

```
$ oc -n <hosted_control_plane_namespace> patch agent <agent_name> -p '{"spec": \
{"hostname":"<hostname>","approved":true}}' --type merge
```

Verification

- To verify that the agents are approved for use, enter the following command:

```
$ oc -n <hosted_control_plane_namespace> get agents
```

Example output

NAME	CLUSTER	APPROVED	ROLE	STAGE
86f7ac75-4fc4-4b36-8130-40fa12602218		true	auto-assign	
e57a637f-745b-496e-971d-1abbf03341ba		true	auto-assign	

5.5.3. Scaling the NodePool object for a hosted cluster on IBM Power

The **NodePool** object is created when you create a hosted cluster. By scaling the **NodePool** object, you can add more compute nodes to hosted control planes.

Procedure

- Run the following command to scale the **NodePool** object to two nodes:

```
$ oc -n <hosted_cluster_namespace> scale nodepool <nodepool_name> --replicas 2
```

The Cluster API agent provider randomly picks two agents that are then assigned to the hosted cluster. Those agents go through different states and finally join the hosted cluster as OpenShift Container Platform nodes. The agents pass through the transition phases in the following order:

- binding**
- discovering**
- insufficient**
- installing**
- installing-in-progress**
- added-to-existing-cluster**

- Run the following command to see the status of a specific scaled agent:

```
$ oc -n <hosted_control_plane_namespace> get agent \
-o jsonpath='{range .items[*]}BMH: {@.metadata.labels.agent-install\\.openshift\\.io/bmh} \
Agent: {@.metadata.name} State: {@.status.debugInfo.state}{"\n"}{end}'
```

Example output

```
BMH: Agent: 50c23cda-cedc-9bbd-bcf1-9b3a5c75804d State: known-unbound
BMH: Agent: 5e498cd3-542c-e54f-0c58-ed43e28b568a State: insufficient
```

- Run the following command to see the transition phases:

```
$ oc -n <hosted_control_plane_namespace> get agent
```

Example output

NAME	CLUSTER	APPROVED	ROLE	STAGE
50c23cda-cedc-9bbd-bcf1-9b3a5c75804d	hosted-forwarder	true		auto-assign
5e498cd3-542c-e54f-0c58-ed43e28b568a		true		auto-assign
da503cf1-a347-44f2-875c-4960ddb04091	hosted-forwarder	true		auto-assign

4. Run the following command to generate the **kubeconfig** file to access the hosted cluster:

```
$ hcp create kubeconfig --namespace <hosted_cluster_namespace> \
--name <hosted_cluster_name> > <hosted_cluster_name>.kubeconfig
```

5. After the agents reach the **added-to-existing-cluster** state, verify that you can see the OpenShift Container Platform nodes by entering the following command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get nodes
```

Example output

NAME	STATUS	ROLES	AGE	VERSION
worker-zvm-0.hostedn.example.com	Ready	worker	5m41s	v1.24.0+3882f8f
worker-zvm-1.hostedn.example.com	Ready	worker	6m3s	v1.24.0+3882f8f

6. Enter the following command to verify that two machines were created when you scaled up the **NodePool** object:

```
$ oc -n <hosted_control_plane_namespace> get machine.cluster.x-k8s.io
```

Example output

NAME	CLUSTER	NODENAME	PROVIDERID
PHASE	AGE	VERSION	
hosted-forwarder-79558597ff-5tbqp	hosted-forwarder-crqq5	worker-zvm-0.hostedn.example.com	agent://50c23cda-cedc-9bbd-bcf1-9b3a5c75804d
	4.15.0		Running 41h
hosted-forwarder-79558597ff-lfjfk	hosted-forwarder-crqq5	worker-zvm-1.hostedn.example.com	agent://5e498cd3-542c-e54f-0c58-ed43e28b568a
	4.15.0		Running 41h

7. Run the following command to check the cluster version:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get clusterversion
```

Example output

NAME	VERSION	AVAILABLE	PROGRESSING	SINCE
STATUS				
clusterversion.config.openshift.io/version	4.15.0	True	False	40h Cluster
version is 4.15.0				

8. Run the following command to check the Cluster Operator status:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get clusteroperators
```

For each component of your cluster, the output shows the following Cluster Operator statuses:

- **NAME**
- **VERSION**
- **AVAILABLE**
- **PROGRESSING**
- **DEGRADED**
- **SINCE**
- **MESSAGE**

Additional resources

- [Initial Operator configuration](#)
- [Scaling down the data plane to zero](#)

CHAPTER 6. DEPLOYING HOSTED CONTROL PLANES IN A DISCONNECTED ENVIRONMENT

6.1. INTRODUCTION TO HOSTED CONTROL PLANES IN A DISCONNECTED ENVIRONMENT

In the context of hosted control planes, a disconnected environment is an OpenShift Container Platform deployment that is not connected to the internet and that uses hosted control planes as a base. You can deploy hosted control planes in a disconnected environment on bare metal or OpenShift Virtualization.

Hosted control planes in disconnected environments function differently than in standalone OpenShift Container Platform:

- The control plane is in the management cluster. The control plane is where the pods of the hosted control plane are run and managed by the Control Plane Operator.
- The data plane is in the workers of the hosted cluster. The data plane is where the workloads and other pods run, all managed by the HostedClusterConfig Operator.

Depending on where the pods are running, they are affected by the **ImageDigestMirrorSet** (IDMS) or **ImageContentSourcePolicy** (ICSP) that is created in the management cluster or by the **ImageContentSource** that is set in the **spec** field of the manifest for the hosted cluster. The **spec** field is translated into an IDMS object on the hosted cluster.

You can deploy hosted control planes in a disconnected environment on IPv4, IPv6, and dual-stack networks. IPv4 is one of the simplest network configurations to deploy hosted control planes in a disconnected environment. IPv4 ranges require fewer external components than IPv6 or dual-stack setups. For hosted control planes on OpenShift Virtualization in a disconnected environment, use either an IPv4 or a dual-stack network.



IMPORTANT

Hosted control planes in a disconnected environment on a dual-stack network is a Technology Preview feature only. Technology Preview features are not supported with Red Hat production service level agreements (SLAs) and might not be functionally complete. Red Hat does not recommend using them in production. These features provide early access to upcoming product features, enabling customers to test functionality and provide feedback during the development process.

For more information about the support scope of Red Hat Technology Preview features, see [Technology Preview Features Support Scope](#).

6.2. DEPLOYING HOSTED CONTROL PLANES ON OPENSHIFT VIRTUALIZATION IN A DISCONNECTED ENVIRONMENT

When you deploy hosted control planes in a disconnected environment, some of the steps differ depending on the platform you use. The following procedures are specific to deployments on OpenShift Virtualization.

6.2.1. Prerequisites

- You have a disconnected OpenShift Container Platform environment serving as your management cluster.

- You have an internal registry to mirror images on. For more information, see [About disconnected installation mirroring](#).

6.2.2. Configuring image mirroring for hosted control planes in a disconnected environment

Image mirroring is the process of fetching images from external registries, such as [registry.redhat.com](#) or [quay.io](#), and storing them in your private registry.

In the following procedures, the **oc-mirror** tool is used, which is a binary that uses the **ImageSetConfiguration** object. In the file, you can specify the following information:

- The OpenShift Container Platform versions to mirror. The versions are in [quay.io](#).
- The additional Operators to mirror. Select packages individually.
- The extra images that you want to add to the repository.

Prerequisites

- Ensure that the registry server is running before you start the mirroring process.

Procedure

To configure image mirroring, complete the following steps:

1. Ensure that your **`\${HOME}/.docker/config.json`** file is updated with the registries that you are going to mirror from and with the private registry that you plan to push the images to.
2. By using the following example, create an **ImageSetConfiguration** object to use for mirroring. Replace values as needed to match your environment:

```
apiVersion: mirror.openshift.io/v2alpha1
kind: ImageSetConfiguration
mirror:
  platform:
    channels:
      - name: candidate-4.18
        minVersion: <4.x.y-build> 1
        maxVersion: <4.x.y-build> 2
      type: ocp
    kubeVirtContainer: true 3
    graph: true
  operators:
    - catalog: registry.redhat.io/redhat/redhat-operator-index:v4.18
      packages:
        - name: lvms-operator
        - name: local-storage-operator
        - name: odf-csi-addons-operator
        - name: odf-operator
        - name: mcg-operator
        - name: ocs-operator
        - name: metallb-operator
        - name: kubevirt-hyperconverged 4
```

- 1** **2** Replace <4.x.y-build> with the supported OpenShift Container Platform version you want to use.
- 3** Set this optional flag to **true** if you want to also mirror the container disk image for the Red Hat Enterprise Linux CoreOS (RHCOS) boot image for the KubeVirt provider. This flag is available with oc-mirror v2 only.
- 4** For deployments that use the KubeVirt provider, include this line.

3. Start the mirroring process by entering the following command:

```
$ oc-mirror --v2 --config imagesetconfig.yaml \
--workspace file://mirror-file docker://<registry>
```

After the mirroring process is finished, you have a new folder named **mirror-file**, which contains the **ImageDigestMirrorSet** (IDMS), **ImageTagMirrorSet** (ITMS), and the catalog sources to apply on the hosted cluster.

4. Mirror the nightly or CI versions of OpenShift Container Platform by configuring the **imagesetconfig.yaml** file as follows:

```
apiVersion: mirror.openshift.io/v2alpha1
kind: ImageSetConfiguration
mirror:
platform:
graph: true
release: registry.ci.openshift.org/ocp/release:<4.x.y-build> 1
kubeVirtContainer: true 2
# ...
```

- 1** Replace <4.x.y-build> with the supported OpenShift Container Platform version you want to use.
- 2** Set this optional flag to **true** if you want to also mirror the container disk image for the Red Hat Enterprise Linux CoreOS (RHCOS) boot image for the KubeVirt provider. This flag is available with oc-mirror v2 only.

5. If you have a partially disconnected environment, mirror the images from the image set configuration to a registry by entering the following command:

```
$ oc mirror -c imagesetconfig.yaml \
--workspace file://<file_path> docker://<mirror_registry_url> --v2
```

For more information, see "Mirroring an image set in a partially disconnected environment".

6. If you have a fully disconnected environment, perform the following steps:

- a. Mirror the images from the specified image set configuration to the disk by entering the following command:

```
$ oc mirror -c imagesetconfig.yaml file://<file_path> --v2
```

For more information, see "Mirroring an image set in a fully disconnected environment".

- b. Process the image set file on the disk and mirror the contents to a target mirror registry by entering the following command:

```
$ oc mirror -c imagesetconfig.yaml \
--from file://<file_path> docker://<mirror_registry_url> --v2
```

7. Mirror the latest multicluster engine Operator images by following the steps in [Install on disconnected networks](#).

Additional resources

- [Mirroring an image set in a partially disconnected environment](#)
- [Mirroring an image set in a fully disconnected environment](#)

6.2.3. Applying objects in the management cluster

After the mirroring process is complete, you need to apply two objects in the management cluster:

- **ImageContentSourcePolicy** (ICSP) or **ImageDigestMirrorSet** (IDMS)
- Catalog sources

When you use the **oc-mirror** tool, the output artifacts are in a folder named **oc-mirror-workspace/results-XXXXXX/**.

The ICSP or IDMS initiates a **MachineConfig** change that does not restart your nodes but restarts the kubelet on each of them. After the nodes are marked as **READY**, you need to apply the newly generated catalog sources.

The catalog sources initiate actions in the **openshift-marketplace** Operator, such as downloading the catalog image and processing it to retrieve all the **PackageManifests** that are included in that image.

Procedure

1. To check the new sources, run the following command by using the new **CatalogSource** as a source:

```
$ oc get packagemanifest
```

2. To apply the artifacts, complete the following steps:

- a. Create the ICSP or IDMS artifacts by entering the following command:

```
$ oc apply -f oc-mirror-workspace/results-XXXXXX/imageContentSourcePolicy.yaml
```

- b. Wait for the nodes to become ready, and then enter the following command:

```
$ oc apply -f catalogSource-XXXXXXXXX-index.yaml
```

3. Mirror the OLM catalogs and configure the hosted cluster to point to the mirror.

When you use the **management** (default) OLMCatalogPlacement mode, the image stream that is used for OLM catalogs is not automatically amended with override information from the ICSP on the management cluster.

- a. If the OLM catalogs are properly mirrored to an internal registry by using the original name and tag, add the **hypershift.openshift.io/olm-catalogs-is-registry-overrides** annotation to the **HostedCluster** resource. The format is "**sr1=dr1,sr2=dr2**", where the source registry string is a key and the destination registry is a value.
- b. To bypass the OLM catalog image stream mechanism, use the following four annotations on the **HostedCluster** resource to directly specify the addresses of the four images to use for OLM Operator catalogs:
 - **hypershift.openshift.io/certified-operators-catalog-image**
 - **hypershift.openshift.io/community-operators-catalog-image**
 - **hypershift.openshift.io/redhat-marketplace-catalog-image**
 - **hypershift.openshift.io/redhat-operators-catalog-image**

In this case, the image stream is not created, and you must update the value of the annotations when the internal mirror is refreshed to pull in Operator updates.

Next steps

Deploy the multicloud engine Operator by completing the steps in *Deploying multicloud engine Operator for a disconnected installation of hosted control planes*.

Additional resources

- [Mirroring images for a disconnected installation by using the oc-mirror plugin v2](#)

6.2.4. Deploying multicloud engine Operator for a disconnected installation of hosted control planes

The multicloud engine for Kubernetes Operator plays a crucial role in deploying clusters across providers. If you do not have multicloud engine Operator installed, review the following documentation to understand the prerequisites and steps to install it:

- [About cluster lifecycle with multicloud engine operator](#)
- [Installing and upgrading multicloud engine operator](#)

6.2.5. Configuring TLS certificates for a disconnected installation of hosted control planes

To ensure proper function in a disconnected deployment, you need to configure the registry CA certificates in the management cluster and the worker nodes for the hosted cluster.

6.2.5.1. Adding the registry CA to the management cluster

To add the registry CA to the management cluster, complete the following steps.

Procedure

1. Create a config map that resembles the following example:

```
apiVersion: v1
```

```

kind: ConfigMap
metadata:
  name: <config_map_name> ①
  namespace: <config_map_namespace> ②
data: ③
  <registry_name>..<port>: | ④
    -----BEGIN CERTIFICATE-----
    -----END CERTIFICATE-----
  <registry_name>..<port>: |
    -----BEGIN CERTIFICATE-----
    -----END CERTIFICATE-----
  <registry_name>..<port>: |
    -----BEGIN CERTIFICATE-----
    -----END CERTIFICATE-----

```

- ① Specify the name of the config map.
- ② Specify the namespace for the config map.
- ③ In the **data** field, specify the registry names and the registry certificate content. Replace **<port>** with the port where the registry server is running; for example, **5000**.
- ④ Ensure that the data in the config map is defined by using | only instead of other methods, such as | -. If you use other methods, issues can occur when the pod reads the certificates.

2. Patch the cluster-wide object, **image.config.openshift.io** to include the following specification:

```

spec:
  additionalTrustedCA:
  - name: registry-config

```

As a result of this patch, the control plane nodes can retrieve images from the private registry and the HyperShift Operator can extract the OpenShift Container Platform payload for hosted cluster deployments.

The process to patch the object might take several minutes to be completed.

6.2.5.2. Adding the registry CA to the worker nodes for the hosted cluster

In order for the data plane workers in the hosted cluster to be able to retrieve images from the private registry, you need to add the registry CA to the worker nodes.

Procedure

1. In the **hc.spec.additionalTrustBundle** file, add the following specification:

```

spec:
  additionalTrustBundle:
  - name: user-ca-bundle ①

```

- ① The **user-ca-bundle** entry is a config map that you create in the next step.

2. In the same namespace where the **HostedCluster** object is created, create the **user-ca-bundle** config map. The config map resembles the following example:

```

apiVersion: v1
data:
  ca-bundle.crt: |
    // Registry1 CA
    -----BEGIN CERTIFICATE-----
    -----END CERTIFICATE-----

    // Registry2 CA
    -----BEGIN CERTIFICATE-----
    -----END CERTIFICATE-----

    // Registry3 CA
    -----BEGIN CERTIFICATE-----
    -----END CERTIFICATE-----

kind: ConfigMap
metadata:
  name: user-ca-bundle
  namespace: <hosted_cluster_namespace> ①

```

- 1 Specify the namespace where the **HostedCluster** object is created.

6.2.6. Creating a hosted cluster on OpenShift Virtualization

A hosted cluster is an OpenShift Container Platform cluster with its control plane and API endpoint hosted on a management cluster. The hosted cluster includes the control plane and its corresponding data plane.

6.2.6.1. Requirements to deploy hosted control planes on OpenShift Virtualization

As you prepare to deploy hosted control planes on OpenShift Virtualization, consider the following information:

- Run the management cluster on bare metal.
- Each hosted cluster must have a cluster-wide unique name.
- Do not use **clusters** as a hosted cluster name.
- A hosted cluster cannot be created in the namespace of a multicluster engine Operator managed cluster.
- When you configure storage for hosted control planes, consider the recommended etcd practices. To ensure that you meet the latency requirements, dedicate a fast storage device to all hosted control plane etcd instances that run on each control-plane node. You can use LVM storage to configure a local storage class for hosted etcd pods. For more information, see "Recommended etcd practices" and "Persistent storage using Logical Volume Manager storage".

6.2.6.2. Creating a hosted cluster with the KubeVirt platform by using the CLI

To create a hosted cluster, you can use the hosted control plane command-line interface, **hcp**.

Procedure

- 1 Create a hosted cluster with the KubeVirt platform by entering the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \①
--node-pool-replicas <node_pool_replica_count> \②
--pull-secret <path_to_pull_secret> \③
--memory <value_for_memory> \④
--cores <value_for_cpu> \⑤
--etcd-storage-class=<etcd_storage_class> \⑥
```

- 1 Specify the name of your hosted cluster, for instance, **example**.
- 2 Specify the node pool replica count, for example, 3. You must specify the replica count as 0 or greater to create the same number of replicas. Otherwise, no node pools are created.
- 3 Specify the path to your pull secret, for example, /**user/name/pullsecret**.
- 4 Specify a value for memory, for example, **6Gi**.
- 5 Specify a value for CPU, for example, **2**.
- 6 Specify the etcd storage class name, for example, **lvm-storageclass**.



NOTE

You can use the **--release-image** flag to set up the hosted cluster with a specific OpenShift Container Platform release.

A default node pool is created for the cluster with two virtual machine worker replicas according to the **--node-pool-replicas** flag.

- 2 After a few moments, verify that the hosted control plane pods are running by entering the following command:

```
$ oc -n clusters-<hosted-cluster-name> get pods
```

Example output

NAME	READY	STATUS	RESTARTS	AGE
capi-provider-5cc7b74f47-n5gkr	1/1	Running	0	3m
catalog-operator-5f799567b7-fd6jw	2/2	Running	0	69s
certified-operators-catalog-784b9899f9-mrp6p	1/1	Running	0	66s
cluster-api-6bbc867966-l4dwl	1/1	Running	0	66s
.				
.				
.				
redhat-operators-catalog-9d5fd4d44-z8qqk	1/1	Running	0	66s

A hosted cluster that has worker nodes that are backed by KubeVirt virtual machines typically takes 10-15 minutes to be fully provisioned.

- To check the status of the hosted cluster, see the corresponding **HostedCluster** resource by entering the following command:

```
$ oc get --namespace clusters hostedclusters
```

See the following example output, which illustrates a fully provisioned **HostedCluster** object:

NAMESPACE	NAME	VERSION	KUBECONFIG	PROGRESS	AVAILABLE
PROGRESSING	MESSAGE				
clusters	example	<4.x.0>	example-admin-kubeconfig	Completed	True
					False

The hosted control plane is available

Replace **<4.x.0>** with the supported OpenShift Container Platform version that you want to use.

6.2.6.3. Configuring the default ingress and DNS for hosted control planes on OpenShift Virtualization

Every OpenShift Container Platform cluster includes a default application Ingress Controller, which must have an wildcard DNS record associated with it. By default, hosted clusters that are created by using the HyperShift KubeVirt provider automatically become a subdomain of the OpenShift Container Platform cluster that the KubeVirt virtual machines run on.

For example, your OpenShift Container Platform cluster might have the following default ingress DNS entry:

```
*.apps.mgmt-cluster.example.com
```

As a result, a KubeVirt hosted cluster that is named **guest** and that runs on that underlying OpenShift Container Platform cluster has the following default ingress:

```
*.apps.guest.apps.mgmt-cluster.example.com
```

Procedure

For the default ingress DNS to work properly, the cluster that hosts the KubeVirt virtual machines must allow wildcard DNS routes.

- You can configure this behavior by entering the following command:

```
$ oc patch ingresscontroller -n openshift-ingress-operator default \
--type=json \
-p '[{"op": "add", "path": "/spec/routeAdmission", "value": {"wildcardPolicy": "WildcardsAllowed"}}]'
```



NOTE

When you use the default hosted cluster ingress, connectivity is limited to HTTPS traffic over port 443. Plain HTTP traffic over port 80 is rejected. This limitation applies to only the default ingress behavior.

6.2.6.4. Customizing ingress and DNS behavior

If you do not want to use the default ingress and DNS behavior, you can configure a KubeVirt hosted cluster with a unique base domain at creation time. This option requires manual configuration steps during creation and involves three main steps: cluster creation, load balancer creation, and wildcard DNS configuration.

6.2.6.4.1. Deploying a hosted cluster that specifies the base domain

To create a hosted cluster that specifies a base domain, complete the following steps.

Procedure

- 1 Enter the following command:

```
$ hcp create cluster kubevirt \
--name <hosted_cluster_name> \ ①
--node-pool-replicas <worker_count> \ ②
--pull-secret <path_to_pull_secret> \ ③
--memory <value_for_memory> \ ④
--cores <value_for_cpu> \ ⑤
--base-domain <basedomain> ⑥
```

- 1 Specify the name of your hosted cluster.
- 2 Specify the worker count, for example, **2**.
- 3 Specify the path to your pull secret, for example, **/user/name/pullsecret**.
- 4 Specify a value for memory, for example, **6Gi**.
- 5 Specify a value for CPU, for example, **2**.
- 6 Specify the base domain, for example, **hypershift.lab**.

As a result, the hosted cluster has an ingress wildcard that is configured for the cluster name and the base domain, for example, **.apps.example.hypershift.lab**. The hosted cluster remains in **Partial** status because after you create a hosted cluster with unique base domain, you must configure the required DNS records and load balancer.

- 2 View the status of your hosted cluster by entering the following command:

```
$ oc get --namespace clusters hostedclusters
```

Example output

NAME	VERSION	KUBECONFIG	PROGRESS	AVAILABLE
PROGRESSING	MESSAGE			
example		example-admin-kubeconfig	Partial	True
		hosted control plane is available		False
			The	

- 3 Access the cluster by entering the following commands:

```
$ hcp create kubeconfig --name <hosted_cluster_name> \
> <hosted_cluster_name>-kubeconfig

$ oc --kubeconfig <hosted_cluster_name>-kubeconfig get co
```

Example output

	NAME	VERSION	AVAILABLE	PROGRESSING	DEGRADED
	SINCE	MESSAGE			
console	<4.x.0>	False False False	30m		
	RouteHealthAvailable: failed to GET route (https://console-openshift-console.apps.example.hypershift.lab): Get "https://console-openshift-console.apps.example.hypershift.lab": dial tcp: lookup console-openshift-console.apps.example.hypershift.lab on 172.31.0.10:53: no such host				
ingress	<4.x.0>	True False True	28m	The "default"	
	ingress controller reports Degraded=True: DegradedConditions: One or more other status conditions indicate a degraded state: CanaryChecksSucceeding=False (CanaryChecksRepetitiveFailures: Canary route checks for the default ingress controller are failing)				

Replace **<4.x.0>** with the supported OpenShift Container Platform version that you want to use.

Next steps

To fix the errors in the output, complete the steps in "Setting up the load balancer" and "Setting up a wildcard DNS".



NOTE

If your hosted cluster is on bare metal, you might need MetalLB to set up load balancer services. For more information, see "Configuring MetalLB".

6.2.6.4.2. Setting up the load balancer

Set up the load balancer service that routes ingress traffic to the KubeVirt VMs and assigns a wildcard DNS entry to the load balancer IP address.

Procedure

1. A **NodePort** service that exposes the hosted cluster ingress already exists. You can export the node ports and create the load balancer service that targets those ports.

- a. Get the HTTP node port by entering the following command:

```
$ oc --kubeconfig <hosted_cluster_name>-kubeconfig get services \
-n openshift-ingress router-nodeport-default \
-o jsonpath='{.spec.ports[?(@.name=="http")].nodePort}'
```

Note the HTTP node port value to use in the next step.

- b. Get the HTTPS node port by entering the following command:

```
$ oc --kubeconfig <hosted_cluster_name>-kubeconfig get services \
-n openshift-ingress router-nodeport-default \
-o jsonpath='{.spec.ports[?(@.name=="https")].nodePort}'
```

Note the HTTPS node port value to use in the next step.

2. Create the load balancer service by entering the following command:

```
oc apply -f -
apiVersion: v1
kind: Service
metadata:
labels:
app: <hosted_cluster_name>
name: <hosted_cluster_name>-apps
namespace: clusters-<hosted_cluster_name>
spec:
ports:
- name: https-443
port: 443
protocol: TCP
targetPort: <https_node_port> ①
- name: http-80
port: 80
protocol: TCP
targetPort: <http-node-port> ②
selector:
kubevirt.io: virt-launcher
type: LoadBalancer
```

- 1 Specify the HTTPS node port value that you noted in the previous step.
- 2 Specify the HTTP node port value that you noted in the previous step.

6.2.6.4.3. Setting up a wildcard DNS

Set up a wildcard DNS record or CNAME that references the external IP of the load balancer service.

Procedure

1. Get the external IP address by entering the following command:

```
$ oc -n clusters-<hosted_cluster_name> get service <hosted-cluster-name>-apps \
-o jsonpath='{.status.loadBalancer.ingress[0].ip}'
```

Example output

```
192.168.20.30
```

2. Configure a wildcard DNS entry that references the external IP address. View the following example DNS entry:

```
*.apps.<hosted_cluster_name>.<base_domain>.
```

The DNS entry must be able to route inside and outside of the cluster.

DNS resolutions example

```
dig +short test.apps.example.hypershift.lab
192.168.20.30
```

- Check that hosted cluster status has moved from **Partial** to **Completed** by entering the following command:

```
$ oc get --namespace clusters hostedclusters
```

Example output

NAME	VERSION	KUBECONFIG	PROGRESS	AVAILABLE
PROGRESSING	MESSAGE			
example	<4.x.0>	example-admin-kubeconfig	Completed	True
The hosted control plane is available				

Replace **<4.x.0>** with the supported OpenShift Container Platform version that you want to use.

6.2.7. Finishing the deployment

You can monitor the deployment of a hosted cluster from two perspectives: the control plane and the data plane.

6.2.7.1. Monitoring the control plane

While the deployment proceeds, you can monitor the control plane by gathering information about the following artifacts:

- The HyperShift Operator
- The **HostedControlPlane** pod
- The bare metal hosts
- The agents
- The **InfraEnv** resource
- The **HostedCluster** and **NodePool** resources

Procedure

- Enter the following commands to monitor the control plane:

```
$ export KUBECONFIG=/root/.kcli/clusters/hub-ipv4/auth/kubeconfig
$ watch "oc get pod -n hypershift;echo;echo; \
  oc get pod -n clusters-hosted-ipv4;echo;echo; \
  oc get bmh -A;echo;echo;\"
```

```
oc get agent -A;echo;echo;\  
oc get infraenv -A;echo;echo;\  
oc get hostedcluster -A;echo;echo;\  
oc get nodepool -A;echo;echo;"
```

6.2.7.2. Monitoring the data plane

While the deployment proceeds, you can monitor the data plane by gathering information about the following artifacts:

- The cluster version
- The nodes, specifically, about whether the nodes joined the cluster
- The cluster Operators

Procedure

- Enter the following commands:

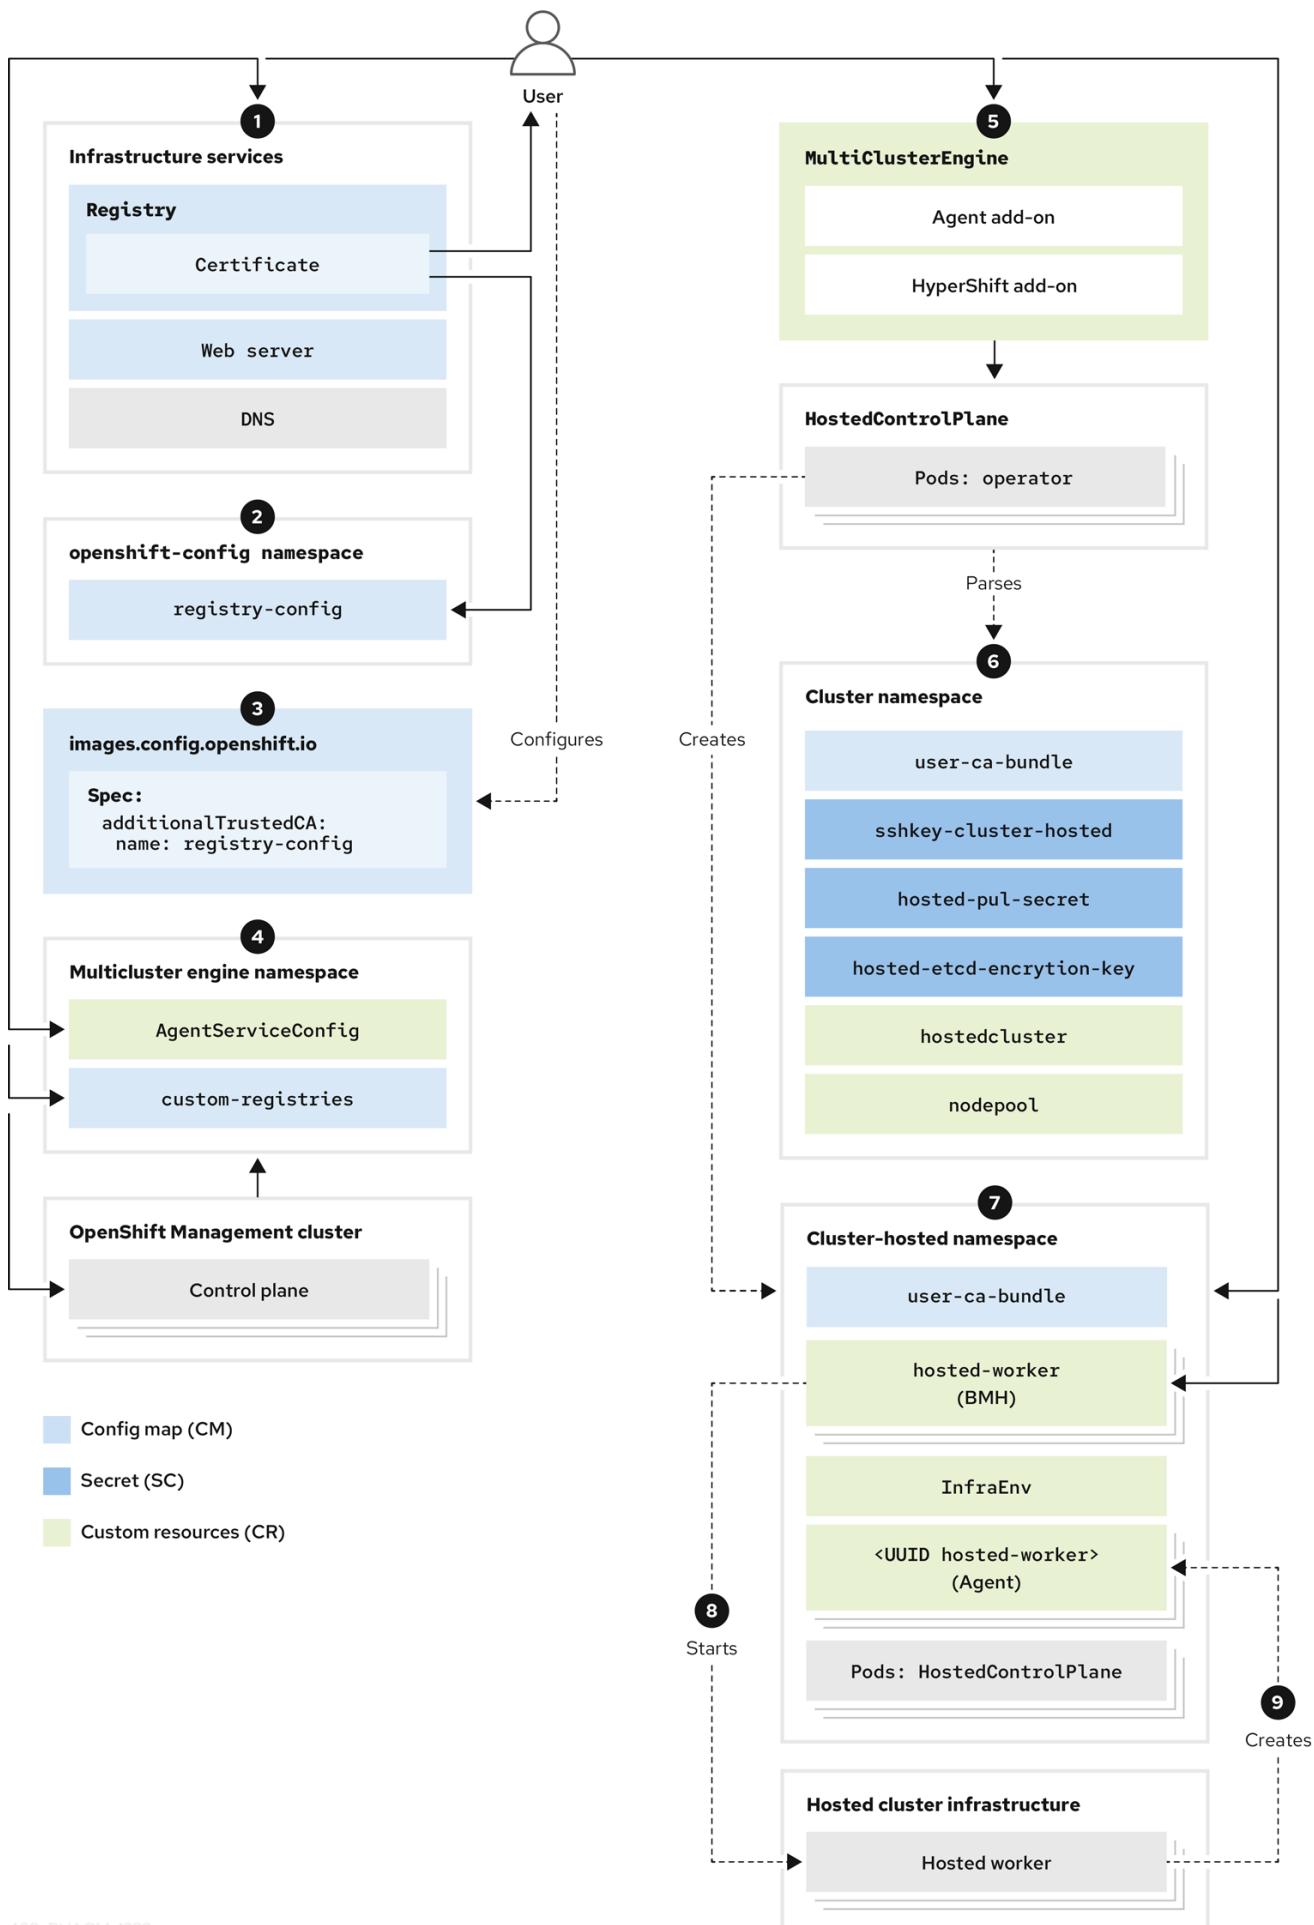
```
$ oc get secret -n clusters-hosted-ipv4 admin-kubeconfig \  
-o jsonpath='{.data.kubeconfig}' | base64 -d > /root/hc_admin_kubeconfig.yaml  
  
$ export KUBECONFIG=/root/hc_admin_kubeconfig.yaml  
  
$ watch "oc get clusterversion,nodes,co"
```

6.3. DEPLOYING HOSTED CONTROL PLANES ON BARE METAL IN A DISCONNECTED ENVIRONMENT

When you provision hosted control planes on bare metal, you use the Agent platform. The Agent platform and multicluster engine for Kubernetes Operator work together to enable disconnected deployments. The Agent platform uses the central infrastructure management service to add worker nodes to a hosted cluster. For an introduction to the central infrastructure management service, see [Enabling the central infrastructure management service](#).

6.3.1. Disconnected environment architecture for bare metal

The following diagram illustrates an example architecture of a disconnected environment:



1. Configure infrastructure services, including the registry certificate deployment with TLS support, web server, and DNS, to ensure that the disconnected deployment works.
2. Create a config map in the **openshift-config** namespace. In this example, the config map is named **registry-config**. The content of the config map is the Registry CA certificate. The data field of the config map must contain the following key/value:
 - Key: **<registry_dns_domain_name>..<port>**, for example, **registry.hypershiftdomain.lab..5000**: Ensure that you place .. after the registry DNS domain name when you specify a port.
 - Value: The certificate content
For more information about creating a config map, see *Configuring TLS certificates for a disconnected installation of hosted control planes*.
3. Modify the **images.config.openshift.io** custom resource (CR) specification and adds a new field named **additionalTrustedCA** with a value of **name: registry-config**.
4. Create a config map that contains two data fields. One field contains the **registries.conf** file in **RAW** format, and the other field contains the Registry CA and is named **ca-bundle.crt**. The config map belongs to the **multicluster-engine** namespace, and the config map name is referenced in other objects. For an example of a config map, see the following sample configuration:

```

apiVersion: v1
kind: ConfigMap
metadata:
  name: custom-registries
  namespace: multicluster-engine
  labels:
    app: assisted-service
data:
  ca-bundle.crt: |
    -----BEGIN CERTIFICATE-----
    # ...
    -----END CERTIFICATE-----
  registries.conf: |
    unqualified-search-registries = ["registry.access.redhat.com", "docker.io"]

    [[registry]]
    prefix = ""
    location = "registry.redhat.io/openshift4"
    mirror-by-digest-only = true

    [[registry.mirror]]
    location = "registry.ocp-edge-cluster-0.qe.lab.redhat.com:5000/openshift4"

    [[registry]]
    prefix = ""
    location = "registry.redhat.io/rhacm2"
    mirror-by-digest-only = true
    # ...
    # ...
  
```

5. In the multicluster engine Operator namespace, you create the **multiclusterengine** CR, which enables both the Agent and **hypershift-addon** add-ons. The multicluster engine Operator

namespace must contain the config maps to modify behavior in a disconnected deployment. The namespace also contains the **multicluster-engine**, **assisted-service**, and **hypershift-addon-manager** pods.

6. Create the following objects that are necessary to deploy the hosted cluster:
 - Secrets: Secrets contain the pull secret, SSH key, and etcd encryption key.
 - Config map: The config map contains the CA certificate of the private registry.
 - **HostedCluster**: The **HostedCluster** resource defines the configuration of the cluster that the user intends to create.
 - **NodePool**: The **NodePool** resource identifies the node pool that references the machines to use for the data plane.
7. After you create the hosted cluster objects, the HyperShift Operator establishes the **HostedControlPlane** namespace to accommodate control plane pods. The namespace also hosts components such as Agents, bare metal hosts (BMHs), and the **InfraEnv** resource. Later, you create the **InfraEnv** resource, and after ISO creation, you create the BMHs and their secrets that contain baseboard management controller (BMC) credentials.
8. The Metal3 Operator in the **openshift-machine-api** namespace inspects the new BMHs. Then, the Metal3 Operator tries to connect to the BMCs to start them by using the configured **LiveISO** and **RootFS** values that are specified through the **AgentServiceConfig** CR in the multicluster engine Operator namespace.
9. After the worker nodes of the **HostedCluster** resource are started, an Agent container is started. This agent establishes contact with the Assisted Service, which orchestrates the actions to complete the deployment. Initially, you need to scale the **NodePool** resource to the number of worker nodes for the **HostedCluster** resource. The Assisted Service manages the remaining tasks.
10. At this point, you wait for the deployment process to be completed.

6.3.2. Requirements to deploy hosted control planes on bare metal in a disconnected environment

To configure hosted control planes in a disconnected environment, you must meet the following prerequisites:

- CPU: The number of CPUs provided determines how many hosted clusters can run concurrently. In general, use 16 CPUs for each node for 3 nodes. For minimal development, you can use 12 CPUs for each node for 3 nodes.
- Memory: The amount of RAM affects how many hosted clusters can be hosted. Use 48 GB of RAM for each node. For minimal development, 18 GB of RAM might be sufficient.
- Storage: Use SSD storage for multicluster engine Operator.
 - Management cluster: 250 GB.
 - Registry: The storage needed depends on the number of releases, operators, and images that are hosted. An acceptable number might be 500 GB, preferably separated from the disk that hosts the hosted cluster.

- Web server: The storage needed depends on the number of ISOs and images that are hosted. An acceptable number might be 500 GB.
- Production: For a production environment, separate the management cluster, the registry, and the web server on different disks. This example illustrates a possible configuration for production:
 - Registry: 2 TB
 - Management cluster: 500 GB
 - Web server: 2 TB

6.3.3. Extracting the release image digest

You can extract the OpenShift Container Platform release image digest by using the tagged image.

Procedure

- Obtain the image digest by running the following command:

```
$ oc adm release info <tagged_openshift_release_image> | grep "Pull From"
```

Replace **<tagged_openshift_release_image>** with the tagged image for the supported OpenShift Container Platform version, for example, **quay.io/openshift-release-dev/ocp-release:4.14.0-x8_64**.

Example output

```
Pull From: quay.io/openshift-release-dev/ocp-release@sha256:69d1292f64a2b67227c5592c1a7d499c7d00376e498634ff8e1946bc9ccdddf
```

6.3.4. DNS configurations on bare metal

The API Server for the hosted cluster is exposed as a **NodePort** service. A DNS entry must exist for **api.<hosted_cluster_name>.<base_domain>** that points to destination where the API Server can be reached.

The DNS entry can be as simple as a record that points to one of the nodes in the managed cluster that is running the hosted control plane. The entry can also point to a load balancer that is deployed to redirect incoming traffic to the ingress pods.

Example DNS configuration

```
api.example.krnl.es. IN A 192.168.122.20
api.example.krnl.es. IN A 192.168.122.21
api.example.krnl.es. IN A 192.168.122.22
api-int.example.krnl.es. IN A 192.168.122.20
api-int.example.krnl.es. IN A 192.168.122.21
api-int.example.krnl.es. IN A 192.168.122.22
`*.apps.example.krnl.es. IN A 192.168.122.23
```

If you are configuring DNS for a disconnected environment on an IPv6 network, the configuration looks like the following example.

Example DNS configuration for an IPv6 network

```
api.example.krnl.es. IN A 2620:52:0:1306::5
api.example.krnl.es. IN A 2620:52:0:1306::6
api.example.krnl.es. IN A 2620:52:0:1306::7
api-int.example.krnl.es. IN A 2620:52:0:1306::5
api-int.example.krnl.es. IN A 2620:52:0:1306::6
api-int.example.krnl.es. IN A 2620:52:0:1306::7
`*.apps.example.krnl.es. IN A 2620:52:0:1306::10
```

If you are configuring DNS for a disconnected environment on a dual stack network, be sure to include DNS entries for both IPv4 and IPv6.

Example DNS configuration for a dual stack network

```
host-record=api-int.hub-dual.dns.base.domain.name,192.168.126.10
host-record=api.hub-dual.dns.base.domain.name,192.168.126.10
address=/apps.hub-dual.dns.base.domain.name/192.168.126.11
dhcp-host=aa:aa:aa:aa:10:01,ocp-master-0,192.168.126.20
dhcp-host=aa:aa:aa:aa:10:02,ocp-master-1,192.168.126.21
dhcp-host=aa:aa:aa:aa:10:03,ocp-master-2,192.168.126.22
dhcp-host=aa:aa:aa:aa:10:06,ocp-installer,192.168.126.25
dhcp-host=aa:aa:aa:aa:10:07,ocp-bootstrap,192.168.126.26

host-record=api-int.hub-dual.dns.base.domain.name,2620:52:0:1306::2
host-record=api.hub-dual.dns.base.domain.name,2620:52:0:1306::2
address=/apps.hub-dual.dns.base.domain.name/2620:52:0:1306::3
dhcp-host=aa:aa:aa:aa:10:01,ocp-master-0,[2620:52:0:1306::5]
dhcp-host=aa:aa:aa:aa:10:02,ocp-master-1,[2620:52:0:1306::6]
dhcp-host=aa:aa:aa:aa:10:03,ocp-master-2,[2620:52:0:1306::7]
dhcp-host=aa:aa:aa:aa:10:06,ocp-installer,[2620:52:0:1306::8]
dhcp-host=aa:aa:aa:aa:10:07,ocp-bootstrap,[2620:52:0:1306::9]
```

6.3.5. Deploying a registry for hosted control planes in a disconnected environment

For development environments, deploy a small, self-hosted registry by using a Podman container. For production environments, deploy an enterprise-hosted registry, such as Red Hat Quay, Nexus, or Artifactory.

Procedure

To deploy a small registry by using Podman, complete the following steps:

- As a privileged user, access the `${HOME}` directory and create the following script:

```
#!/usr/bin/env bash

set -euo pipefail

PRIMARY_NIC=$(ls -1 /sys/class/net | grep -v podman | head -1)
export PATH=/root/bin:$PATH
export PULL_SECRET="/root/baremetal/hub/openshift_pull.json" ①

if [[ ! -f $PULL_SECRET ]];then
    echo "Pull Secret not found, exiting..."

```

```

    exit 1
fi

dnf -y install podman httpd httpd-tools jq skopeo libseccomp-devel
export IP=$(ip -o addr show $PRIMARY_NIC | head -1 | awk '{print $4}' | cut -d'/' -f1)
REGISTRY_NAME=registry.$(hostname --long)
REGISTRY_USER=dummy
REGISTRY_PASSWORD=dummy
KEY=$(echo -n $REGISTRY_USER:$REGISTRY_PASSWORD | base64)
echo "{\"auths\": {\"$REGISTRY_NAME:5000\": {\"auth\": \"$KEY\", \"email\": \"sample-email@domain.ltd\"}}}" > /root/disconnected_pull.json
mv ${PULL_SECRET} /root/openshift_pull.json.old
jq ".auths += {\"$REGISTRY_NAME:5000\": {\"auth\": \"$KEY\", \"email\": \"sample-email@domain.ltd\"}}" < /root/openshift_pull.json.old > $PULL_SECRET
mkdir -p /opt/registry/{auth,certs,data,conf}
cat <<EOF > /opt/registry/conf/config.yml
version: 0.1
log:
  fields:
    service: registry
storage:
  cache:
    blobdescriptor: inmemory
  filesystem:
    rootdirectory: /var/lib/registry
delete:
  enabled: true
http:
  addr: :5000
  headers:
    X-Content-Type-Options: [nosniff]
health:
  storagedriver:
    enabled: true
    interval: 10s
    threshold: 3
compatibility:
  schema1:
    enabled: true
EOF
openssl req -newkey rsa:4096 -nodes -sha256 -keyout /opt/registry/certs/domain.key -x509 -days 3650 -out /opt/registry/certs/domain.crt -subj "/C=US/ST=Madrid/L=San Bernardo/O=Karmalabs/OU=Guitar/CN=$REGISTRY_NAME" -addext "subjectAltName=DNS:$REGISTRY_NAME"
cp /opt/registry/certs/domain.crt /etc/pki/ca-trust/source/anchors/
update-ca-trust extract
htpasswd -bBc /opt/registry/auth/htpasswd $REGISTRY_USER $REGISTRY_PASSWORD
podman create --name registry --net host --security-opt label=disable --replace -v /opt/registry/data:/var/lib/registry:z -v /opt/registry/auth:/auth:z -v /opt/registry/conf/config.yml:/etc/docker/registry/config.yml -e "REGISTRY_AUTH=htpasswd" -e "REGISTRY_AUTH_HTPASSWD_REALM=Registry" -e "REGISTRY_HTTP_SECRET=ALongRandomSecretForRegistry" -e REGISTRY_AUTH_HTPASSWD_PATH=/auth/htpasswd -v /opt/registry/certs:/certs:z -e REGISTRY_HTTP_CERTIFICATE=/certs/domain.crt -e

```

```
REGISTRY_HTTP_TLS_KEY=/certs/domain.key docker.io/library/registry:latest
[ "$?" == "0" ] || !!
systemctl enable --now registry
```

- 1** Replace the location of the **PULL_SECRET** with the appropriate location for your setup.
2. Name the script file **registry.sh** and save it. When you run the script, it pulls in the following information:
 - The registry name, based on the hypervisor hostname
 - The necessary credentials and user access details
3. Adjust permissions by adding the execution flag as follows:

```
$ chmod u+x ${HOME}/registry.sh
```

4. To run the script without any parameters, enter the following command:

```
$ ${HOME}/registry.sh
```

The script starts the server. The script uses a **systemd** service for management purposes.

5. If you need to manage the script, you can use the following commands:

```
$ systemctl status
```

```
$ systemctl start
```

```
$ systemctl stop
```

The root folder for the registry is in the **/opt/registry** directory and contains the following subdirectories:

- **certs** contains the TLS certificates.
- **auth** contains the credentials.
- **data** contains the registry images.
- **conf** contains the registry configuration.

6.3.6. Setting up a management cluster for hosted control planes in a disconnected environment

To set up an OpenShift Container Platform management cluster, you need to ensure that the multicluster engine for Kubernetes Operator is installed. The multicluster engine Operator plays a crucial role in deploying clusters across providers.

Prerequisites

- There must be bidirectional connectivity between the management cluster and the Baseboard Management Controller (BMC) of the target Bare Metal Host (BMH). As an alternative, you follow a Boot It Yourself approach through the Agent provider.

- The hosted cluster must be able to resolve and reach the API hostname of the management cluster hostname and ***.apps** hostname. Here is an example of the API hostname of the management cluster and ***.apps** hostname:
 - **api.management-cluster.internal.domain.com**
 - **console-openshift-console.apps.management-cluster.internal.domain.com**
- The management cluster must be able to resolve and reach the API and ***.apps** hostname of the hosted cluster. Here is an example of the API hostname of the hosted cluster and ***.apps** hostname:
 - **api.sno-hosted-cluster-1.internal.domain.com**
 - **console-openshift-console.apps.sno-hosted-cluster-1.internal.domain.com**

Procedure

1. Install multicluster engine Operator 2.4 or later on an OpenShift Container Platform cluster. You can install multicluster engine Operator as an Operator from the OpenShift Container Platform OperatorHub. The HyperShift Operator is included with multicluster engine Operator. For more information about installing multicluster engine Operator, see "Installing and upgrading multicluster engine operator" in the Red Hat Advanced Cluster Management documentation.
2. Ensure that the HyperShift Operator is installed. The HyperShift Operator is automatically included with multicluster engine Operator, but if you need to manually install it, follow the steps in "Manually enabling the hypershift-addon managed cluster add-on for local-cluster".

Next steps

Next, configure the web server.

Additional resources

- [Installing and upgrading multicluster engine operator](#)
- [Manually enabling the hypershift-addon managed cluster add-on for local-cluster](#)
- [About cluster lifecycle with multicluster engine operator](#)

6.3.7. Configuring the web server for hosted control planes in a disconnected environment

You need to configure an additional web server to host the Red Hat Enterprise Linux CoreOS (RHCOS) images that are associated with the OpenShift Container Platform release that you are deploying as a hosted cluster.

Procedure

To configure the web server, complete the following steps:

1. Extract the **openshift-install** binary from the OpenShift Container Platform release that you want to use by entering the following command:

```
$ oc adm -a ${LOCAL_SECRET_JSON} release extract --command=openshift-install \
"${LOCAL_REGISTRY}/${LOCAL_REPOSITORY}:${OCP_RELEASE}-"
${ARCHITECTURE}"
```

- Run the following script. The script creates a folder in the `/opt/srv` directory. The folder contains the RHCOS images to provision the worker nodes.

```
#!/bin/bash

WEBSRV_FOLDER=/opt/srv
ROOTFS_IMG_URL="$(./openshift-install coreos print-stream-json | jq -r
'.architectures.x86_64.artifacts.metal.formats.pxe.rootfs.location')"  
①
LIVE_ISO_URL="$(./openshift-install coreos print-stream-json | jq -r
'.architectures.x86_64.artifacts.metal.formats.iso.disk.location')"  
②

mkdir -p ${WEBSRV_FOLDER}/images
curl -Lk ${ROOTFS_IMG_URL} -o
${WEBSRV_FOLDER}/images/${ROOTFS_IMG_URL##*/}
curl -Lk ${LIVE_ISO_URL} -o ${WEBSRV_FOLDER}/images/${LIVE_ISO_URL##*/}
chmod -R 755 ${WEBSRV_FOLDER}/*

## Run Webserver
podman ps --noheading | grep -q websrv-ai
if [[ $? == 0 ]];then
    echo "Launching Registry pod..."
    /usr/bin/podman run --name websrv-ai --net host -v /opt/srv:/usr/local/apache2/htdocs:z
    quay.io/atosadag/httpd:p8080
fi
```

- ① You can find the `ROOTFS_IMG_URL` value on the OpenShift CI Release page.
- ② You can find the `LIVE_ISO_URL` value on the OpenShift CI Release page.

After the download is completed, a container runs to host the images on a web server. The container uses a variation of the official HTTPd image, which also enables it to work with IPv6 networks.

6.3.8. Configuring image mirroring for hosted control planes in a disconnected environment

Image mirroring is the process of fetching images from external registries, such as registry.redhat.com or quay.io, and storing them in your private registry.

In the following procedures, the `oc-mirror` tool is used, which is a binary that uses the `ImageSetConfiguration` object. In the file, you can specify the following information:

- The OpenShift Container Platform versions to mirror. The versions are in quay.io.
- The additional Operators to mirror. Select packages individually.
- The extra images that you want to add to the repository.

Prerequisites

- Ensure that the registry server is running before you start the mirroring process.

Procedure

To configure image mirroring, complete the following steps:

1. Ensure that your `${HOME}/.docker/config.json` file is updated with the registries that you are going to mirror from and with the private registry that you plan to push the images to.
2. By using the following example, create an **ImageSetConfiguration** object to use for mirroring. Replace values as needed to match your environment:

```
apiVersion: mirror.openshift.io/v2alpha1
kind: ImageSetConfiguration
mirror:
  platform:
    channels:
      - name: candidate-4.18
        minVersion: <4.x.y-build> ①
        maxVersion: <4.x.y-build> ②
    type: ocp
    kubeVirtContainer: true ③
    graph: true
  operators:
    - catalog: registry.redhat.io/redhat/redhat-operator-index:v4.18
    packages:
      - name: lvms-operator
      - name: local-storage-operator
      - name: odf-csi-addons-operator
      - name: odf-operator
      - name: mcg-operator
      - name: ocs-operator
      - name: metallb-operator
      - name: kubevirt-hyperconverged ④
```

① ② Replace `<4.x.y-build>` with the supported OpenShift Container Platform version you want to use.

- ③ Set this optional flag to `true` if you want to also mirror the container disk image for the Red Hat Enterprise Linux CoreOS (RHCOS) boot image for the KubeVirt provider. This flag is available with `oc-mirror v2` only.
- ④ For deployments that use the KubeVirt provider, include this line.

3. Start the mirroring process by entering the following command:

```
$ oc-mirror --v2 --config imagesetconfig.yaml \
--workspace file://mirror-file docker://<registry>
```

After the mirroring process is finished, you have a new folder named **mirror-file**, which contains the **ImageDigestMirrorSet** (IDMS), **ImageTagMirrorSet** (ITMS), and the catalog sources to apply on the hosted cluster.

4. Mirror the nightly or CI versions of OpenShift Container Platform by configuring the **imagesetconfig.yaml** file as follows:

```
apiVersion: mirror.openshift.io/v2alpha1
kind: ImageSetConfiguration
mirror:
  platform:
```

```
graph: true
release: registry.ci.openshift.org/ocp/release:<4.x.y-build> ①
kubeVirtContainer: true ②
# ...
```

- ① Replace <4.x.y-build> with the supported OpenShift Container Platform version you want to use.
- ② Set this optional flag to **true** if you want to also mirror the container disk image for the Red Hat Enterprise Linux CoreOS (RHCOS) boot image for the KubeVirt provider. This flag is available with oc-mirror v2 only.

5. If you have a partially disconnected environment, mirror the images from the image set configuration to a registry by entering the following command:

```
$ oc mirror -c imagesetconfig.yaml \
--workspace file://<file_path> docker://<mirror_registry_url> --v2
```

For more information, see "Mirroring an image set in a partially disconnected environment".

6. If you have a fully disconnected environment, perform the following steps:

- a. Mirror the images from the specified image set configuration to the disk by entering the following command:

```
$ oc mirror -c imagesetconfig.yaml file://<file_path> --v2
```

For more information, see "Mirroring an image set in a fully disconnected environment".

- b. Process the image set file on the disk and mirror the contents to a target mirror registry by entering the following command:

```
$ oc mirror -c imagesetconfig.yaml \
--from file://<file_path> docker://<mirror_registry_url> --v2
```

7. Mirror the latest multicluster engine Operator images by following the steps in [Install on disconnected networks](#).

Additional resources

- [Mirroring an image set in a partially disconnected environment](#)
- [Mirroring an image set in a fully disconnected environment](#)

6.3.9. Applying objects in the management cluster

After the mirroring process is complete, you need to apply two objects in the management cluster:

- **ImageContentSourcePolicy** (ICSP) or **ImageDigestMirrorSet** (IDMS)
- Catalog sources

When you use the **oc-mirror** tool, the output artifacts are in a folder named **oc-mirror-workspace/results-XXXXXX/**.

The ICSP or IDMS initiates a **MachineConfig** change that does not restart your nodes but restarts the kubelet on each of them. After the nodes are marked as **READY**, you need to apply the newly generated catalog sources.

The catalog sources initiate actions in the **openshift-marketplace** Operator, such as downloading the catalog image and processing it to retrieve all the **PackageManifests** that are included in that image.

Procedure

1. To check the new sources, run the following command by using the new **CatalogSource** as a source:

```
$ oc get packagemanifest
```

2. To apply the artifacts, complete the following steps:

- a. Create the ICSP or IDMS artifacts by entering the following command:

```
$ oc apply -f oc-mirror-workspace/results-XXXXXX/imageContentSourcePolicy.yaml
```

- b. Wait for the nodes to become ready, and then enter the following command:

```
$ oc apply -f catalogSource-XXXXXXXXX-index.yaml
```

3. Mirror the OLM catalogs and configure the hosted cluster to point to the mirror.

When you use the **management** (default) OLMCatalogPlacement mode, the image stream that is used for OLM catalogs is not automatically amended with override information from the ICSP on the management cluster.

- a. If the OLM catalogs are properly mirrored to an internal registry by using the original name and tag, add the **hypershift.openshift.io/olm-catalogs-is-registry-overrides** annotation to the **HostedCluster** resource. The format is "**sr1=dr1,sr2=dr2**", where the source registry string is a key and the destination registry is a value.

- b. To bypass the OLM catalog image stream mechanism, use the following four annotations on the **HostedCluster** resource to directly specify the addresses of the four images to use for OLM Operator catalogs:

- **hypershift.openshift.io/certified-operators-catalog-image**
- **hypershift.openshift.io/community-operators-catalog-image**
- **hypershift.openshift.io/redhat-marketplace-catalog-image**
- **hypershift.openshift.io/redhat-operators-catalog-image**

In this case, the image stream is not created, and you must update the value of the annotations when the internal mirror is refreshed to pull in Operator updates.

Next steps

Deploy the multicluster engine Operator by completing the steps in *Deploying multicluster engine Operator for a disconnected installation of hosted control planes*.

Additional resources

- Mirroring images for a disconnected installation by using the oc-mirror plugin v2

6.3.10. Deploying AgentServiceConfig resources

The **AgentServiceConfig** custom resource is an essential component of the Assisted Service add-on that is part of multicluster engine Operator. It is responsible for bare metal cluster deployment. When the add-on is enabled, you deploy the **AgentServiceConfig** resource to configure the add-on.

In addition to configuring the **AgentServiceConfig** resource, you need to include additional config maps to ensure that multicluster engine Operator functions properly in a disconnected environment.

Procedure

- 1 Configure the custom registries by adding the following config map, which contains the disconnected details to customize the deployment:

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: custom-registries
  namespace: multicluster-engine
  labels:
    app: assisted-service
data:
  ca-bundle.crt: |
    -----BEGIN CERTIFICATE-----
    -----END CERTIFICATE-----
  registries.conf: |
    unqualified-search-registries = ["registry.access.redhat.com", "docker.io"]

    [[registry]]
    prefix = ""
    location = "registry.redhat.io/openshift4"
    mirror-by-digest-only = true

    [[registry.mirror]]
    location = "registry.dns.base.domain.name:5000/openshift4" ①

    [[registry]]
    prefix = ""
    location = "registry.redhat.io/rhacm2"
    mirror-by-digest-only = true
    # ...
    # ...
```

- 1 Replace **dns.base.domain.name** with the DNS base domain name.

The object contains two fields:

- Custom CAs: This field contains the Certificate Authorities (CAs) that are loaded into the various processes of the deployment.
- Registries: The **Registries.conf** field contains information about images and namespaces that need to be consumed from a mirror registry rather than the original source registry.

- Configure the Assisted Service by adding the **AssistedServiceConfig** object, as shown in the following example:

```

apiVersion: agent-install.openshift.io/v1beta1
kind: AgentServiceConfig
metadata:
  annotations:
    unsupported.agent-install.openshift.io/assisted-service-configmap: assisted-service-config
  1
  name: agent
  namespace: multicluster-engine
spec:
  mirrorRegistryRef:
    name: custom-registries 2
  databaseStorage:
    storageClassName: lvms-vg1
    accessModes:
      - ReadWriteOnce
  resources:
    requests:
      storage: 10Gi
  filesystemStorage:
    storageClassName: lvms-vg1
    accessModes:
      - ReadWriteOnce
  resources:
    requests:
      storage: 20Gi
  oslImages: 3
    - cpuArchitecture: x86_64 4
      openshiftVersion: "4.14"
      rootFSUrl: http://registry.dns.base.domain.name:8080/images/rhcos-414.92.202308281054-0-live-rootfs.x86_64.img 5
      url: http://registry.dns.base.domain.name:8080/images/rhcos-414.92.202308281054-0-live.x86_64.iso
      version: 414.92.202308281054-0
    - cpuArchitecture: x86_64
      openshiftVersion: "4.15"
      rootFSUrl: http://registry.dns.base.domain.name:8080/images/rhcos-415.92.202403270524-0-live-rootfs.x86_64.img
      url: http://registry.dns.base.domain.name:8080/images/rhcos-415.92.202403270524-0-live.x86_64.iso
      version: 415.92.202403270524-0

```

- The **metadata.annotations["unsupported.agent-install.openshift.io/assisted-service-configmap"]** annotation references the config map name that the Operator consumes to customize behavior.
- The **spec.mirrorRegistryRef.name** annotation points to the config map that contains disconnected registry information that the Assisted Service Operator consumes. This config map adds those resources during the deployment process.
- The **spec.oslImages** field contains different versions available for deployment by this Operator. This field is mandatory. This example assumes that you already downloaded the **RootFS** and **LiveISO** files.

- 4 Add a **cpuArchitecture** subsection for every OpenShift Container Platform release that you want to deploy. In this example, **cpuArchitecture** subsections are included for 4.14 and
- 5 In the **rootFSUrl** and **url** fields, replace **dns.base.domain.name** with the DNS base domain name.

- 3 Deploy all of the objects by concatenating them into a single file and applying them to the management cluster. To do so, enter the following command:

```
$ oc apply -f agentServiceConfig.yaml
```

The command triggers two pods.

Example output

assisted-image-service-0	1/1	Running	2	11d	1
assisted-service-668b49548-9m7xw	2/2	Running	5	11d	2

- 1 The **assisted-image-service** pod is responsible for creating the Red Hat Enterprise Linux CoreOS (RHCOS) boot image template, which is customized for each cluster that you deploy.
- 2 The **assisted-service** refers to the Operator.

Next steps

Configure TLS certificates.

6.3.11. Configuring TLS certificates for a disconnected installation of hosted control planes

To ensure proper function in a disconnected deployment, you need to configure the registry CA certificates in the management cluster and the worker nodes for the hosted cluster.

6.3.11.1. Adding the registry CA to the management cluster

To add the registry CA to the management cluster, complete the following steps.

Procedure

- 1 Create a config map that resembles the following example:

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: <config_map_name> 1
  namespace: <config_map_namespace> 2
data: 3
  <registry_name>..<port>: | 4
    -----BEGIN CERTIFICATE-----
    -----END CERTIFICATE-----
  <registry_name>..<port>: |
```

```
-----BEGIN CERTIFICATE-----
-----END CERTIFICATE-----
<registry_name>..<port>: |
-----BEGIN CERTIFICATE-----
-----END CERTIFICATE-----
```

- 1 Specify the name of the config map.
- 2 Specify the namespace for the config map.
- 3 In the **data** field, specify the registry names and the registry certificate content. Replace **<port>** with the port where the registry server is running; for example, **5000**.
- 4 Ensure that the data in the config map is defined by using | only instead of other methods, such as | -. If you use other methods, issues can occur when the pod reads the certificates.

2. Patch the cluster-wide object, **image.config.openshift.io** to include the following specification:

```
spec:
  additionalTrustedCA:
    - name: registry-config
```

As a result of this patch, the control plane nodes can retrieve images from the private registry and the HyperShift Operator can extract the OpenShift Container Platform payload for hosted cluster deployments.

The process to patch the object might take several minutes to be completed.

6.3.11.2. Adding the registry CA to the worker nodes for the hosted cluster

In order for the data plane workers in the hosted cluster to be able to retrieve images from the private registry, you need to add the registry CA to the worker nodes.

Procedure

1. In the **hc.spec.additionalTrustBundle** file, add the following specification:

```
spec:
  additionalTrustBundle:
    - name: user-ca-bundle ①
```

- 1 The **user-ca-bundle** entry is a config map that you create in the next step.

2. In the same namespace where the **HostedCluster** object is created, create the **user-ca-bundle** config map. The config map resembles the following example:

```
apiVersion: v1
data:
  ca-bundle.crt: |
    // Registry1 CA
    -----BEGIN CERTIFICATE-----
    -----END CERTIFICATE-----
```

```
// Registry2 CA
-----BEGIN CERTIFICATE-----
-----END CERTIFICATE-----

// Registry3 CA
-----BEGIN CERTIFICATE-----
-----END CERTIFICATE-----

kind: ConfigMap
metadata:
  name: user-ca-bundle
  namespace: <hosted_cluster_namespace> ①
```

- ① Specify the namespace where the **HostedCluster** object is created.

6.3.12. Creating a hosted cluster on bare metal

A hosted cluster is an OpenShift Container Platform cluster with its control plane and API endpoint hosted on a management cluster. The hosted cluster includes the control plane and its corresponding data plane.

6.3.12.1. Deploying hosted cluster objects

Typically, the HyperShift Operator creates the **HostedControlPlane** namespace. However, in this case, you want to include all the objects before the HyperShift Operator begins to reconcile the **HostedCluster** object. Then, when the Operator starts the reconciliation process, it can find all of the objects in place.

Procedure

- Create a YAML file with the following information about the namespaces:

```
---
apiVersion: v1
kind: Namespace
metadata:
  creationTimestamp: null
  name: <hosted_cluster_namespace>-<hosted_cluster_name> ①
spec: {}
status: {}

---
apiVersion: v1
kind: Namespace
metadata:
  creationTimestamp: null
  name: <hosted_cluster_namespace> ②
spec: {}
status: {}
```

- ① Replace **<hosted_cluster_name>** with your hosted cluster.
- ② Replace **<hosted_cluster_namespace>** with the name of your hosted cluster namespace.

2. Create a YAML file with the following information about the config maps and secrets to include in the **HostedCluster** deployment:

```

---
apiVersion: v1
data:
  ca-bundle.crt: |
    -----BEGIN CERTIFICATE-----
    -----END CERTIFICATE-----
kind: ConfigMap
metadata:
  name: user-ca-bundle
  namespace: <hosted_cluster_namespace> ①
---
apiVersion: v1
data:
  .dockerconfigjson: xxxxxxxxx
kind: Secret
metadata:
  creationTimestamp: null
  name: <hosted_cluster_name>-pull-secret ②
  namespace: <hosted_cluster_namespace> ③
---
apiVersion: v1
kind: Secret
metadata:
  name: sshkey-cluster-<hosted_cluster_name> ④
  namespace: <hosted_cluster_namespace> ⑤
stringData:
  id_rsa.pub: ssh-rsa xxxxxxxxx
---
apiVersion: v1
data:
  key: nTPtVBEt03owkrKhIdmSW8jrWRxU57KO/fnZa8oaG0Y=
kind: Secret
metadata:
  creationTimestamp: null
  name: <hosted_cluster_name>-etcd-encryption-key ⑥
  namespace: <hosted_cluster_namespace> ⑦
type: Opaque

```

① ③ ⑤ ⑦ Replace **<hosted_cluster_namespace>** with the name of your hosted cluster namespace.

② ④ ⑥ Replace **<hosted_cluster_name>** with your hosted cluster.

3. Create a YAML file that contains the RBAC roles so that Assisted Service agents can be in the same **HostedControlPlane** namespace as the hosted control plane and still be managed by the cluster API:

```

apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  creationTimestamp: null

```

```

name: capi-provider-role
namespace: <hosted_cluster_namespace>-<hosted_cluster_name> ① ②
rules:
- apiGroups:
  - agent-install.openshift.io
resources:
- agents
verbs:
- '*'

```

- ① Replace **<hosted_cluster_namespace>** with the name of your hosted cluster namespace.
- ② Replace **<hosted_cluster_name>** with your hosted cluster.

4. Create a YAML file with information about the **HostedCluster** object, replacing values as necessary:

```

apiVersion: hypershift.openshift.io/v1beta1
kind: HostedCluster
metadata:
  name: <hosted_cluster_name> ①
  namespace: <hosted_cluster_namespace> ②
spec:
  additionalTrustBundle:
    name: "user-ca-bundle"
  olmCatalogPlacement: guest
  imageContentSources: ③
    - source: quay.io/openshift-release-dev/ocp-v4.0-art-dev
      mirrors:
        - registry.<dns.base.domain.name>:5000/openshift/release ④
    - source: quay.io/openshift-release-dev/ocp-release
      mirrors:
        - registry.<dns.base.domain.name>:5000/openshift/release-images ⑤
    - mirrors:
      ...
      ...
  autoscaling: {}
  controllerAvailabilityPolicy: SingleReplica
  dns:
    baseDomain: <dns.base.domain.name> ⑥
  etcd:
    managed:
      storage:
        persistentVolume:
          size: 8Gi
        restoreSnapshotURL: null
        type: PersistentVolume
      managementType: Managed
    fips: false
  networking:
    clusterNetwork:
      - cidr: 10.132.0.0/14
      - cidr: fd01::/48
    networkType: OVNKubernetes

```

```

serviceNetwork:
- cidr: 172.31.0.0/16
- cidr: fd02::/112
platform:
agent:
  agentNamespace: <hosted_cluster_namespace>-<hosted_cluster_name> 7 8
  type: Agent
pullSecret:
  name: <hosted_cluster_name>-pull-secret 9
release:
  image: registry.<dns.base.domain.name>:5000/openshift/release-images:<4.x.y>-x86_64
10 11
secretEncryption:
aescbc:
  activeKey:
    name: <hosted_cluster_name>-etcd-encryption-key 12
type: aescbc
services:
- service: APIServer
  servicePublishingStrategy:
    type: LoadBalancer
- service: OAuthServer
  servicePublishingStrategy:
    type: Route
- service: OIDC
  servicePublishingStrategy:
    type: Route
- service: Konnectivity
  servicePublishingStrategy:
    type: Route
- service: Ignition
  servicePublishingStrategy:
    type: Route
sshKey:
  name: sshkey-cluster-<hosted_cluster_name> 13
status:
controlPlaneEndpoint:
  host: ""
  port: 0

```

1 7 9 12 13 Replace <hosted_cluster_name> with your hosted cluster.

2 8 Replace <hosted_cluster_namespace> with the name of your hosted cluster namespace.

3 The **imageContentSources** section contains mirror references for user workloads within the hosted cluster.

4 5 6 10 Replace <dns.base.domain.name> with the DNS base domain name.

11 Replace <4.x.y> with the supported OpenShift Container Platform version you want to use.

5. Add an annotation in the **HostedCluster** object that points to the HyperShift Operator release in the OpenShift Container Platform release:

- Obtain the image payload by entering the following command:

```
$ oc adm release info \
  registry.<dns.base.domain.name>:5000/openshift-release-dev/ocp-release:<4.x.y>-
x86_64 \
| grep hypershift
```

where **<dns.base.domain.name>** is the DNS base domain name and **<4.x.y>** is the supported OpenShift Container Platform version you want to use.

Example output

```
hypershift
sha256:31149e3e5f8c5e5b5b100ff2d89975cf5f7a73801b2c06c639bf6648766117f8
```

- By using the OpenShift Container Platform Images namespace, check the digest by entering the following command:

```
podman pull registry.<dns.base.domain.name>:5000/openshift-release-dev/ocp-v4.0-art-
dev@sha256:31149e3e5f8c5e5b5b100ff2d89975cf5f7a73801b2c06c639bf6648766117f8
```

where **<dns.base.domain.name>** is the DNS base domain name.

Example output

```
podman pull
registry.dns.base.domain.name:5000/openshift/release@sha256:31149e3e5f8c5e5b5b100
ff2d89975cf5f7a73801b2c06c639bf6648766117f8
Trying to pull
registry.dns.base.domain.name:5000/openshift/release@sha256:31149e3e5f8c5e5b5b100
ff2d89975cf5f7a73801b2c06c639bf6648766117f8...
Getting image source signatures
Copying blob d8190195889e skipped: already exists
Copying blob c71d2589fba7 skipped: already exists
Copying blob d4dc6e74b6ce skipped: already exists
Copying blob 97da74cc6d8f skipped: already exists
Copying blob b70007a560c9 done
Copying config 3a62961e6e done
Writing manifest to image destination
Storing signatures
3a62961e6ed6edab46d5ec8429ff1f41d6bb68de51271f037c6cb8941a007fde
```

The release image that is set in the **HostedCluster** object must use the digest rather than the tag; for example, **quay.io/openshift-release-dev/ocp-
release@sha256:e3ba11bd1e5e8ea5a0b36a75791c90f29afb0fdbbe4125be4e48f69c76a5c
47a0**.

- Create all of the objects that you defined in the YAML files by concatenating them into a file and applying them against the management cluster. To do so, enter the following command:

```
$ oc apply -f 01-4.14-hosted_cluster-nodeport.yaml
```

Example output

NAME	READY	STATUS	RESTARTS	AGE
capi-provider-5b57dbd6d5-pxlqc	1/1	Running	0	3m57s
catalog-operator-9694884dd-m7zzv	2/2	Running	0	93s
cluster-api-f98b9467c-9hfqrq	1/1	Running	0	3m57s
cluster-autoscaler-d7f95dd5-d8m5d	1/1	Running	0	93s
cluster-image-registry-operator-5ff5944b4b-648ht	1/2	Running	0	93s
cluster-network-operator-77b896ddc-wpkq8	1/1	Running	0	94s
cluster-node-tuning-operator-84956cd484-4hfqf	1/1	Running	0	94s
cluster-policy-controller-5fd8595d97-rhbwf	1/1	Running	0	95s
cluster-storage-operator-54dcf584b5-xrnts	1/1	Running	0	93s
cluster-version-operator-9c554b999-l22s7	1/1	Running	0	95s
control-plane-operator-6fdc9c569-t7hr4	1/1	Running	0	3m57s
csi-snapshot-controller-785c6dc77c-8ljmr	1/1	Running	0	77s
csi-snapshot-controller-operator-7c6674bc5b-d9dtp	1/1	Running	0	93s
csi-snapshot-webhook-5b8584875f-2492j	1/1	Running	0	77s
dns-operator-6874b577f-9tc6b	1/1	Running	0	94s
etcd-0	3/3	Running	0	3m39s
hosted-cluster-config-operator-f5cf5c464-4nmbh	1/1	Running	0	93s
ignition-server-6b689748fc-zdqzk	1/1	Running	0	95s
ignition-server-proxy-54d4bb9b9b-6zkg7	1/1	Running	0	95s
ingress-operator-6548dc758b-f9gtg	1/2	Running	0	94s
konnektivity-agent-7767cdc6f5-tw782	1/1	Running	0	95s
kube-apiserver-7b5799b6c8-9f5bp	4/4	Running	0	3m7s
kube-controller-manager-5465bc4dd6-zpdllk	1/1	Running	0	44s
kube-scheduler-5dd5f78b94-bbbck	1/1	Running	0	2m36s
machine approver-846c69f56-jxvfr	1/1	Running	0	92s
oauth-openshift-79c7bf44bf-j975g	2/2	Running	0	62s
olm-operator-767f9584c-4lcl2	2/2	Running	0	93s
openshift-apiserver-5d469778c6-pl8tj	3/3	Running	0	2m36s
openshift-controller-manager-6475fdff58-hl4f7	1/1	Running	0	95s
openshift-oauth-apiserver-db5cc5f-98574	2/2	Running	0	95s
openshift-route-controller-manager-5f6997b48f-s9vdc	1/1	Running	0	95s
packageserver-67c87d4d4f-kl7qh	2/2	Running	0	93s

When the hosted cluster is available, the output looks like the following example.

Example output

NAMESPACE	NAME	VERSION	KUBECONFIG	PROGRESS	AVAILABLE
PROGRESSING	MESSAGE				
clusters	hosted-dual	hosted-admin-kubeconfig	Partial	True	False
					The hosted control plane is available

6.3.12.2. Creating a NodePool object for the hosted cluster

A **NodePool** is a scalable set of worker nodes that is associated with a hosted cluster. **NodePool** machine architectures remain consistent within a specific pool and are independent of the machine architecture of the control plane.

Procedure

- Create a YAML file with the following information about the **NodePool** object, replacing values as necessary:

```

apiVersion: hypershift.openshift.io/v1beta1
kind: NodePool
metadata:
  creationTimestamp: null
  name: <hosted_cluster_name> \①
  namespace: <hosted_cluster_namespace> \②
spec:
  arch: amd64
  clusterName: <hosted_cluster_name>
  management:
    autoRepair: false \③
    upgradeType: InPlace \④
    nodeDrainTimeout: 0s
  platform:
    type: Agent
  release:
    image: registry.<dns.base.domain.name>:5000/openshift/release-images:4.x.y-x86_64 \
      \⑤
  replicas: 2 \⑥
status:
  replicas: 2

```

- ① Replace **<hosted_cluster_name>** with your hosted cluster.
- ② Replace **<hosted_cluster_namespace>** with the name of your hosted cluster namespace.
- ③ The **autoRepair** field is set to **false** because the node will not be re-created if it is removed.
- ④ The **upgradeType** is set to **InPlace**, which indicates that the same bare metal node is reused during an upgrade.
- ⑤ All of the nodes included in this **NodePool** are based on the following OpenShift Container Platform version: **4.x.y-x86_64**. Replace the **<dns.base.domain.name>** value with your DNS base domain name and the **4.x.y** value with the supported OpenShift Container Platform version you want to use.
- ⑥ You can set the **replicas** value to **2** to create two node pool replicas in your hosted cluster.

2. Create the **NodePool** object by entering the following command:

```
$ oc apply -f 02-nodepool.yaml
```

Example output

NAMESPACE	NAME	CLUSTER	DESIRED NODES	CURRENT NODES	AUTOSCALING	AUTOREPAIR	VERSION	UPDATINGVERSION
clusters	hosted-dual	hosted	0			False	False	4.x.y-x86_64

6.3.12.3. Creating an InfraEnv resource for the hosted cluster

The **InfraEnv** resource is an Assisted Service object that includes essential details, such as the **pullSecretRef** and the **sshAuthorizedKey**. Those details are used to create the Red Hat Enterprise Linux CoreOS (RHCOS) boot image that is customized for the hosted cluster.

You can host more than one **InfraEnv** resource, and each one can adopt certain types of hosts. For example, you might want to divide your server farm between a host that has greater RAM capacity.

Procedure

- 1 Create a YAML file with the following information about the **InfraEnv** resource, replacing values as necessary:

```
apiVersion: agent-install.openshift.io/v1beta1
kind: InfraEnv
metadata:
  name: <hosted_cluster_name>
  namespace: <hosted-cluster-namespace>-<hosted_cluster_name> ① ②
spec:
  pullSecretRef: ③
    name: pull-secret
  sshAuthorizedKey: ssh-rsa
  AAAAAB3NzaC1yc2EAAAQABAAABgQDk7ICaUE+/k4zTpxLk4+xFdHi4ZuDi5qjeF52afsNk
  w0w/gIILHhwpl5gnp5WkRuL8GwJuZ1VqLC9EKrdmgn4MrmUlq7WTsP0VFOZFBfq2XRUXo
  1wrRdor2z0Bbh93ytR+ZsDbbLIGngXaMa0Vbt+z74FqlcajbHTZ6zBmTpBVq5RHtDPgKITdpE1f
  ongp7+ZXQNBIkaavaqv8bnyrP4BWahLP4iO9/xJF9IQYboYwEEDzmnKLMW1VtCE6nJzEgWC
  ufACTbxpNS7GvKtoHT/OVzw8ArEXhZXQUS1UY8zKsX2iXwmhw5Sj6YboA8WICs4z+TrFP8
  9LmxXY0j6536TQFyRz1iB4WWvCbH5n6W+ABV2e8ssJB1AmEy8QYNwpJQJNpSxzoKBjl73X>
  vPYYC/IjPFMySwZqrSZCkJYqQ023ySkaQxWZT7in4KeMu7eS2tC+Kn4deJ7KwwUycx8n6RH
  MeD8Qg9flTHCv3gmab8JKZJqN3hW1D378JuvmlX4V0= ④
```

- ① Replace **<hosted_cluster_name>** with your hosted cluster.
- ② Replace **<hosted_cluster_namespace>** with the name of your hosted cluster namespace.
- ③ The **pullSecretRef** refers to the config map reference in the same namespace as the **InfraEnv**, where the pull secret is used.
- ④ The **sshAuthorizedKey** represents the SSH public key that is placed in the boot image. The SSH key allows access to the worker nodes as the **core** user.

- 2 Create the **InfraEnv** resource by entering the following command:

```
$ oc apply -f 03-infraenv.yaml
```

Example output

NAMESPACE	NAME	ISO CREATED AT
clusters-hosted-dual	hosted	2023-09-11T15:14:10Z

6.3.12.4. Creating bare metal hosts for the hosted cluster

A *bare metal host* is an **openshift-machine-api** object that encompasses physical and logical details so that it can be identified by a Metal3 Operator. Those details are associated with other Assisted Service objects, known as *agents*.

Prerequisites

- Before you create the bare metal host and destination nodes, you must have the destination machines ready.
- You have installed a Red Hat Enterprise Linux CoreOS (RHCOS) compute machine on bare-metal infrastructure for use in the cluster.

Procedure

To create a bare metal host, complete the following steps:

- 1 Create a YAML file with the following information. For more information about what details to enter for the bare metal host, see "Provisioning new hosts in a user-provisioned cluster by using the BMO".

Because you have at least one secret that holds the bare metal host credentials, you need to create at least two objects for each worker node.

```
apiVersion: v1
kind: Secret
metadata:
  name: <hosted_cluster_name>-worker0-bmc-secret ①
  namespace: <hosted_cluster_namespace>-<hosted_cluster_name> ②
data:
  password: YWRtaW4= ③
  username: YWRtaW4= ④
type: Opaque
# ...
apiVersion: metal3.io/v1alpha1
kind: BareMetalHost
metadata:
  name: <hosted_cluster_name>-worker0
  namespace: <hosted_cluster_namespace>-<hosted_cluster_name> ⑤
  labels:
    infraenvs.agent-install.openshift.io: <hosted_cluster_name> ⑥
  annotations:
    inspect.metal3.io: disabled
    bmac.agent-install.openshift.io/hostname: <hosted_cluster_name>-worker0 ⑦
spec:
  automatedCleaningMode: disabled ⑧
  bmc:
    disableCertificateVerification: true ⑨
    address: redfish-
    virtualmedia://[192.168.126.1]:9000/redfish/v1/Systems/local/<hosted_cluster_name>-worker0 ⑩
    credentialsName: <hosted_cluster_name>-worker0-bmc-secret ⑪
    bootMACAddress: aa:aa:aa:aa:02:11 ⑫
    online: true ⑬
```

① Replace **<hosted_cluster_name>** with your hosted cluster.

- 2** **5** Replace `<hosted_cluster_name>` with your hosted cluster. Replace `<hosted_cluster_namespace>` with the name of your hosted cluster namespace.
- 3** Specify the password of the baseboard management controller (BMC) in Base64 format.
- 4** Specify the user name of the BMC in Base64 format.
- 6** Replace `<hosted_cluster_name>` with your hosted cluster. The `infraenvs.agent-install.openshift.io` field serves as the link between the Assisted Installer and the `BareMetalHost` objects.
- 7** Replace `<hosted_cluster_name>` with your hosted cluster. The `bmac.agent-install.openshift.io/hostname` field represents the node name that is adopted during deployment.
- 8** The `automatedCleaningMode` field prevents the node from being erased by the Metal3 Operator.
- 9** The `disableCertificateVerification` field is set to `true` to bypass certificate validation from the client.
- 10** Replace `<hosted_cluster_name>` with your hosted cluster. The `address` field denotes the BMC address of the worker node.
- 11** Replace `<hosted_cluster_name>` with your hosted cluster. The `credentialsName` field points to the secret where the user and password credentials are stored.
- 12** The `bootMACAddress` field indicates the interface MAC address that the node starts from.
- 13** The `online` field defines the state of the node after the `BareMetalHost` object is created.

2. Deploy the `BareMetalHost` object by entering the following command:

```
$ oc apply -f 04-bmh.yaml
```

During the process, you can view the following output:

- This output indicates that the process is trying to reach the nodes:

Example output

NAMESPACE	NAME	STATE	CONSUMER	ONLINE	ERROR	AGE
clusters-hosted	hosted-worker0	registering	true	2s		
clusters-hosted	hosted-worker1	registering	true	2s		
clusters-hosted	hosted-worker2	registering	true	2s		

- This output indicates that the nodes are starting:

Example output

NAMESPACE	NAME	STATE	CONSUMER	ONLINE	ERROR	AGE
clusters-hosted	hosted-worker0	provisioning	true	16s		
clusters-hosted	hosted-worker1	provisioning	true	16s		
clusters-hosted	hosted-worker2	provisioning	true	16s		

- This output indicates that the nodes started successfully:

Example output

NAMESPACE	NAME	STATE	CONSUMER	ONLINE	ERROR	AGE
clusters-hosted	hosted-worker0	provisioned		true		67s
clusters-hosted	hosted-worker1	provisioned		true		67s
clusters-hosted	hosted-worker2	provisioned		true		67s

3. After the nodes start, notice the agents in the namespace, as shown in this example:

Example output

NAMESPACE	NAME	CLUSTER	APPROVED	ROLE
STAGE				
clusters-hosted	aaaaaaaa-aaaa-aaaa-aaaa-aaaaaaaa0411		true	auto-assign
clusters-hosted	aaaaaaaa-aaaa-aaaa-aaaa-aaaaaaaa0412		true	auto-assign
clusters-hosted	aaaaaaaa-aaaa-aaaa-aaaa-aaaaaaaa0413		true	auto-assign

The agents represent nodes that are available for installation. To assign the nodes to a hosted cluster, scale up the node pool.

Additional resources

- Provisioning new hosts in a user-provisioned cluster by using the BMO
- Understanding secrets

6.3.12.5. Scaling up the node pool

After you create the bare metal hosts, their statuses change from **Registering** to **Provisioning** to **Provisioned**. The nodes start with the **LiveISO** of the agent and a default pod that is named **agent**. That agent is responsible for receiving instructions from the Assisted Service Operator to install the OpenShift Container Platform payload.

Procedure

1. To scale up the node pool, enter the following command:

```
$ oc -n <hosted_cluster_namespace> scale nodepool <hosted_cluster_name> \
--replicas 3
```

where:

- **<hosted_cluster_namespace>** is the name of the hosted cluster namespace.
- **<hosted_cluster_name>** is the name of the hosted cluster.

2. After the scaling process is complete, notice that the agents are assigned to a hosted cluster:

Example output

NAMESPACE	NAME	CLUSTER	APPROVED	ROLE
STAGE				

```

clusters-hosted aaaaaaaaaaaaaaaa-aaaa-aaaa-aaaa-aaaaaaaaaa0411 hosted true auto-assign
clusters-hosted aaaaaaaaaaaaaaaa-aaaa-aaaa-aaaa-aaaaaaaaaa0412 hosted true auto-assign
clusters-hosted aaaaaaaaaaaaaaaa-aaaa-aaaa-aaaa-aaaaaaaaaa0413 hosted true auto-assign

```

- Also notice that the node pool replicas are set:

Example output

```

NAMESPACE NAME CLUSTER DESIRED NODES CURRENT NODES
AUTOSCALING AUTOREPAIR VERSION UPDATINGVERSION
UPDATINGCONFIG MESSAGE
clusters hosted hosted 3 False False <4.x.y>-x86_64
Minimum availability requires 3 replicas, current 0 available

```

Replace **<4.x.y>** with the supported OpenShift Container Platform version that you want to use.

- Wait until the nodes join the cluster. During the process, the agents provide updates on their stage and status.

6.4. DEPLOYING HOSTED CONTROL PLANES ON IBM Z IN A DISCONNECTED ENVIRONMENT

Hosted control planes deployments in disconnected environments function differently than in a standalone OpenShift Container Platform.

Hosted control planes involves two distinct environments:

- Control plane: Located in the management cluster, where the hosted control planes pods are run and managed by the Control Plane Operator.
- Data plane: Located in the workers of the hosted cluster, where the workload and a few other pods run, managed by the Hosted Cluster Config Operator.

The **ImageContentSourcePolicy** (ICSP) custom resource for the data plane is managed through the **ImageContentSources** API in the hosted cluster manifest.

For the control plane, ICSP objects are managed in the management cluster. These objects are parsed by the HyperShift Operator and are shared as **registry-overrides** entries with the Control Plane Operator. These entries are injected into any one of the available deployments in the hosted control planes namespace as an argument.

To work with disconnected registries in the hosted control planes, you must first create the appropriate ICSP in the management cluster. Then, to deploy disconnected workloads in the data plane, you need to add the entries that you want into the **ImageContentSources** field in the hosted cluster manifest.

6.4.1. Prerequisites to deploy hosted control planes on IBM Z in a disconnected environment

- A mirror registry. For more information, see "Creating a mirror registry with mirror registry for Red Hat OpenShift".
- A mirrored image for a disconnected installation. For more information, see "Mirroring images for a disconnected installation using the oc-mirror plugin".

Additional resources

- [Creating a mirror registry with mirror registry for Red Hat OpenShift](#)
- [Mirroring images for a disconnected installation by using the oc-mirror plugin v2](#)

6.4.2. Adding credentials and the registry certificate authority to the management cluster

To pull the mirror registry images from the management cluster, you must first add credentials and the certificate authority of the mirror registry to the management cluster. Use the following procedure:

Procedure

1. Create a **ConfigMap** with the certificate of the mirror registry by running the following command:

```
$ oc apply -f registry-config.yaml
```

Example registry-config.yaml file

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: registry-config
  namespace: openshift-config
data:
  <mirror_registry>: |
    -----BEGIN CERTIFICATE-----
    -----END CERTIFICATE-----
```

2. Patch the **image.config.openshift.io** cluster-wide object to include the following entries:

```
spec:
  additionalTrustedCA:
  - name: registry-config
```

3. Update the management cluster pull secret to add the credentials of the mirror registry.

- a. Fetch the pull secret from the cluster in a JSON format by running the following command:

```
$ oc get secret/pull-secret -n openshift-config -o json \
| jq -r '.data.".dockerconfigjson"' \
| base64 -d > authfile
```

- b. Edit the fetched secret JSON file to include a section with the credentials of the certificate authority:

```
"auths": {
  "<mirror_registry>": {
    "auth": "<credentials>", ❷
    "email": "you@example.com"
  }
},
```

- - 1 Provide the name of the mirror registry.
 - 2 Provide the credentials for the mirror registry to allow fetch of images.

c. Update the pull secret on the cluster by running the following command:

```
$ oc set data secret/pull-secret -n openshift-config \
--from-file=.dockerconfigjson=authfile
```

6.4.3. Update the registry certificate authority in the AgentServiceConfig resource with the mirror registry

When you use a mirror registry for images, agents need to trust the registry's certificate to securely pull images. You can add the certificate authority of the mirror registry to the **AgentServiceConfig** custom resource by creating a **ConfigMap**.

Prerequisites

- You must have installed multicluster engine for Kubernetes Operator.

Procedure

- 1 In the same namespace where you installed multicluster engine Operator, create a **ConfigMap** resource with the mirror registry details. This **ConfigMap** resource ensures that you grant the hosted cluster workers the capability to retrieve images from the mirror registry.

Example ConfigMap file

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: mirror-config
  namespace: multicluster-engine
  labels:
    app: assisted-service
data:
  ca-bundle.crt: |
    -----BEGIN CERTIFICATE-----
    -----END CERTIFICATE-----
  registries.conf: |
[[registry]]
location = "registry.stage.redhat.io"
insecure = false
blocked = false
mirror-by-digest-only = true
prefix = ""

[[registry.mirror]]
location = "<mirror_registry>"
insecure = false

[[registry]]
```

```

location = "registry.redhat.io/multicluster-engine"
insecure = false
blocked = false
mirror-by-digest-only = true
prefix = ""

[[registry.mirror]]
location = "<mirror_registry>/multicluster-engine" ①
insecure = false

```

- ① Where: <**mirror_registry**> is the name of the mirror registry.

- Patch the **AgentServiceConfig** resource to include the **ConfigMap** resource that you created. If the **AgentServiceConfig** resource is not present, create the **AgentServiceConfig** resource with the following content embedded into it:

```

spec:
  mirrorRegistryRef:
    name: mirror-config

```

6.4.4. Adding the registry certificate authority to the hosted cluster

When you are deploying hosted control planes on IBM Z in a disconnected environment, include the **additional-trust-bundle** and **image-content-sources** resources. Those resources allow the hosted cluster to inject the certificate authority into the data plane workers so that the images are pulled from the registry.

- Create the **icsp.yaml** file with the **image-content-sources** information.

The **image-content-sources** information is available in the **ImageContentSourcePolicy** YAML file that is generated after you mirror the images by using **oc-mirror**.

Example ImageContentSourcePolicy file

```

# cat icsp.yaml
- mirrors:
  - <mirror_registry>/openshift/release
    source: quay.io/openshift-release-dev/ocp-v4.0-art-dev
  - mirrors:
    - <mirror_registry>/openshift/release-images
      source: quay.io/openshift-release-dev/ocp-release

```

- Create a hosted cluster and provide the **additional-trust-bundle** certificate to update the compute nodes with the certificates as in the following example:

```

$ hcp create cluster agent \
--name=<hosted_cluster_name> \ ①
--pull-secret=<path_to_pull_secret> \ ②
--agent-namespace=<hosted_control_plane_namespace> \ ③
--base-domain=<basedomain> \ ④
--api-server-address=api.<hosted_cluster_name>.<basedomain> \
--etcd-storage-class=<etcd_storage_class> \ ⑤
--ssh-key <path_to_ssh_public_key> \ ⑥

```

```
--namespace <hosted_cluster_namespace> \ 7
--control-plane-availability-policy SingleReplica \
--release-image=quay.io/openshift-release-dev/ocp-release:<ocp_release_image> \ 8
--additional-trust-bundle <path for cert> \ 9
--image-content-sources icsp.yaml
```

- 1 Replace **<hosted_cluster_name>** with the name of your hosted cluster.
- 2 Replace the path to your pull secret, for example, **/user/name/pullsecret**.
- 3 Replace **<hosted_control_plane_namespace>** with the name of the hosted control plane namespace, for example, **clusters-hosted**.
- 4 Replace the name with your base domain, for example, **example.com**.
- 5 Replace the etcd storage class name, for example, **lvm-storageclass**.
- 6 Replace the path to your SSH public key. The default file path is **~/.ssh/id_rsa.pub**.
- 7 8 Replace with the supported OpenShift Container Platform version that you want to use, for example, **4.18.0-multi**.
- 9 Replace the path to Certificate Authority of mirror registry.

6.5. MONITORING USER WORKLOAD IN A DISCONNECTED ENVIRONMENT

The **hypershift-addon** managed cluster add-on enables the **--enable-uwm-telemetry-remote-write** option in the HyperShift Operator. By enabling that option, you ensure that user workload monitoring is enabled and that it can remotely write telemetry metrics from control planes.

6.5.1. Resolving user workload monitoring issues

If you installed multicluster engine Operator on OpenShift Container Platform clusters that are not connected to the internet, when you try to run the user workload monitoring feature of the HyperShift Operator by entering the following command, the feature fails with an error:

```
$ oc get events -n hypershift
```

Example error

LAST SEEN	TYPE	REASON	OBJECT	MESSAGE
4m46s	Warning	ReconcileError	deployment/operator	Failed to ensure UWM telemetry remote write: cannot get telemeter client secret: Secret "telemeter-client" not found

To resolve the error, you must disable the user workload monitoring option by creating a config map in the **local-cluster** namespace. You can create the config map either before or after you enable the add-on. The add-on agent reconfigures the HyperShift Operator.

Procedure

- 1 Create the following config map:

▪

```

kind: ConfigMap
apiVersion: v1
metadata:
  name: hypershift-operator-install-flags
  namespace: local-cluster
data:
  installFlagsToAdd: ""
  installFlagsToRemove: "--enable-uwm-telemetry-remote-write"

```

2. Apply the config map by running the following command:

```
$ oc apply -f <filename>.yaml
```

6.5.2. Verifying the status of the hosted control plane feature

The hosted control plane feature is enabled by default.

Procedure

1. If the feature is disabled and you want to enable it, enter the following command. Replace **<multiclusterengine>** with the name of your multicluster engine Operator instance:

```
$ oc patch mce <multiclusterengine> --type=merge -p \
  '{"spec":{"overrides":{"components":[{"name":"hypershift","enabled":true}]}}}'
```

When you enable the feature, the **hypershift-addon** managed cluster add-on is installed in the **local-cluster** managed cluster, and the add-on agent installs the HyperShift Operator on the multicluster engine Operator hub cluster.

2. Confirm that the **hypershift-addon** managed cluster add-on is installed by entering the following command:

```
$ oc get managedclusteraddons -n local-cluster hypershift-addon
```

Example output

NAME	AVAILABLE	DEGRADED	PROGRESSING
hypershift-addon	True	False	

3. To avoid a timeout during this process, enter the following commands:

```
$ oc wait --for=condition=Degraded=True managedclusteraddons/hypershift-addon \
-n local-cluster --timeout=5m
```

```
$ oc wait --for=condition=Available=True managedclusteraddons/hypershift-addon \
-n local-cluster --timeout=5m
```

When the process is complete, the **hypershift-addon** managed cluster add-on and the HyperShift Operator are installed, and the **local-cluster** managed cluster is available to host and manage hosted clusters.

6.5.3. Configuring the hypershift-addon managed cluster add-on to run on an infrastructure node

By default, no node placement preference is specified for the **hypershift-addon** managed cluster add-on. Consider running the add-ons on the infrastructure nodes, because by doing so, you can prevent incurring billing costs against subscription counts and separate maintenance and management tasks.

Procedure

1. Log in to the hub cluster.
2. Open the **hypershift-addon-deploy-config** add-on deployment configuration specification for editing by entering the following command:

```
$ oc edit addondescription hypershift-addon-deploy-config \
-n multicluster-engine
```

3. Add the **nodePlacement** field to the specification, as shown in the following example:

```
apiVersion: addon.open-cluster-management.io/v1alpha1
kind: AddOnDeploymentConfig
metadata:
  name: hypershift-addon-deploy-config
  namespace: multicluster-engine
spec:
  nodePlacement:
    nodeSelector:
      node-role.kubernetes.io/infra: ""
    tolerations:
      - effect: NoSchedule
        key: node-role.kubernetes.io/infra
        operator: Exists
```

4. Save the changes. The **hypershift-addon** managed cluster add-on is deployed on an infrastructure node for new and existing managed clusters.

CHAPTER 7. UPDATING HOSTED CONTROL PLANES

Updates for hosted control planes involve updating the hosted cluster and the node pools. For a cluster to remain fully operational during an update process, you must meet the requirements of the [Kubernetes version skew policy](#) while completing the control plane and node updates.

7.1. REQUIREMENTS TO UPGRADE HOSTED CONTROL PLANES

The multicluster engine for Kubernetes Operator can manage one or more OpenShift Container Platform clusters. After you create a hosted cluster on OpenShift Container Platform, you must import your hosted cluster in the multicluster engine Operator as a managed cluster. Then, you can use the OpenShift Container Platform cluster as a management cluster.

Consider the following requirements before you start updating hosted control planes:

- You must use the bare metal platform for an OpenShift Container Platform cluster when using OpenShift Virtualization as a provider.
- You must use bare metal or OpenShift Virtualization as the cloud platform for the hosted cluster. You can find the platform type of your hosted cluster in the **spec.Platform.type** specification of the **HostedCluster** custom resource (CR).



IMPORTANT

You must update hosted control planes in the following order:

1. Upgrade an OpenShift Container Platform cluster to the latest version. For more information, see "Updating a cluster using the web console" or "Updating a cluster using the CLI".
2. Upgrade the multicluster engine Operator to the latest version. For more information, see "Updating installed Operators".
3. Upgrade the hosted cluster and node pools from the previous OpenShift Container Platform version to the latest version. For more information, see "Updating a control plane in a hosted cluster" and "Updating node pools in a hosted cluster".

Additional resources

- [Updating a cluster using the web console](#)
- [Updating a cluster using the CLI](#)
- [Updating installed Operators](#)
- [Updating a control plane in a hosted cluster](#)
- [Updating node pools in a hosted cluster](#)

7.2. SETTING CHANNELS IN A HOSTED CLUSTER

You can see available updates in the **HostedCluster.Status** field of the **HostedCluster** custom resource (CR).

The available updates are not fetched from the Cluster Version Operator (CVO) of a hosted cluster. The list of the available updates can be different from the available updates from the following fields of the **HostedCluster** custom resource (CR):

- **status.version.availableUpdates**
- **status.version.conditionalUpdates**

The initial **HostedCluster** CR does not have any information in the **status.version.availableUpdates** and **status.version.conditionalUpdates** fields. After you set the **spec.channel** field to the stable OpenShift Container Platform release version, the HyperShift Operator reconciles the **HostedCluster** CR and updates the **status.version** field with the available and conditional updates.

See the following example of the **HostedCluster** CR that contains the channel configuration:

```
spec:
  autoscaling: {}
  channel: stable-4.y ①
  clusterID: d6d42268-7dff-4d37-92cf-691bd2d42f41
  configuration: {}
  controllerAvailabilityPolicy: SingleReplica
  dns:
    baseDomain: dev11.red-chesterfield.com
    privateZoneID: Z0180092I0DQRKL55LN0
    publicZoneID: Z00206462VG6ZP0H2QLWK
```

- ① Replace <4.y> with the OpenShift Container Platform release version you specified in **spec.release**. For example, if you set the **spec.release** to **ocp-release:4.16.4-multi**, you must set **spec.channel** to **stable-4.16**.

After you configure the channel in the **HostedCluster** CR, to view the output of the **status.version.availableUpdates** and **status.version.conditionalUpdates** fields, run the following command:

```
$ oc get -n <hosted_cluster_namespace> hostedcluster <hosted_cluster_name> -o yaml
```

Example output

```
version:
  availableUpdates:
    - channels:
        - candidate-4.16
        - candidate-4.17
        - eus-4.16
        - fast-4.16
        - stable-4.16
    image: quay.io/openshift-release-dev/ocp-
release@sha256:b7517d13514c6308ae16c5fd8108133754eb922cd37403ed27c846c129e67a9a
    url: https://access.redhat.com/errata/RHBA-2024:6401
  version: 4.16.11
  channels:
    - candidate-4.16
    - candidate-4.17
    - eus-4.16
```

- fast-4.16
- stable-4.16

image: quay.io/openshift-release-dev/ocp-release@sha256:d08e7c8374142c239a07d7b27d1170eae2b0d9f00ccf074c3f13228a1761c162
 url: <https://access.redhat.com/errata/RHSA-2024:6004>
 version: 4.16.10

- channels:
- candidate-4.16
- candidate-4.17
- eus-4.16
- fast-4.16
- stable-4.16

image: quay.io/openshift-release-dev/ocp-release@sha256:6a80ac72a60635a313ae511f0959cc267a21a89c7654f1c15ee16657aafa41a0
 url: <https://access.redhat.com/errata/RHBA-2024:5757>
 version: 4.16.9

- channels:
- candidate-4.16
- candidate-4.17
- eus-4.16
- fast-4.16
- stable-4.16

image: quay.io/openshift-release-dev/ocp-release@sha256:ea624ae7d91d3f15094e9e15037244679678bdc89e5a29834b2ddb7e1d9b57e6
 url: <https://access.redhat.com/errata/RHSA-2024:5422>
 version: 4.16.8

- channels:
- candidate-4.16
- candidate-4.17
- eus-4.16
- fast-4.16
- stable-4.16

image: quay.io/openshift-release-dev/ocp-release@sha256:e4102eb226130117a0775a83769fe8edb029f0a17b6cbca98a682e3f1225d6b7
 url: <https://access.redhat.com/errata/RHSA-2024:4965>
 version: 4.16.6

- channels:
- candidate-4.16
- candidate-4.17
- eus-4.16
- fast-4.16
- stable-4.16

image: quay.io/openshift-release-dev/ocp-release@sha256:f828eda3eaac179e9463ec7b1ed6baeba2cd5bd3f1dd56655796c86260db819b
 url: <https://access.redhat.com/errata/RHBA-2024:4855>
 version: 4.16.5

conditionalUpdates:

- conditions:
- lastTransitionTime: "2024-09-23T22:33:38Z"
- message: |-
 Could not evaluate exposure to update risk SRIOVFailedToConfigureVF (creating PromQL round-tripper: unable to load specified CA cert /etc/tls/service-ca/service-ca.crt: open /etc/tls/service-ca/service-ca.crt: no such file or directory)

SRIOVFailedToConfigureVF description: OCP Versions 4.14.34, 4.15.25, 4.16.7 and ALL subsequent versions include kernel datastructure changes which are not compatible with older versions of the SR-IOV operator. Please update SR-IOV operator to versions dated 20240826 or

newer before updating OCP.

```

SRIOVFailedToConfigureVF URL: https://issues.redhat.com/browse/NHE-1171
reason: EvaluationFailed
status: Unknown
type: Recommended
release:
channels:
- candidate-4.16
- candidate-4.17
- eus-4.16
- fast-4.16
- stable-4.16
image: quay.io/openshift-release-dev/ocp-
release@sha256:fb321a3f50596b43704dbbed2e51fdefd7a7fd488ee99655d03784d0cd02283f
url: https://access.redhat.com/errata/RHSA-2024:5107
version: 4.16.7
risks:
- matchingRules:
  - promql:
    promql: |
      group(csv_succeeded{_id="d6d42268-7dff-4d37-92cf-691bd2d42f41", name=~"sriv-network-operator[.].*"})
      or
      0 * group(csv_count{_id="d6d42268-7dff-4d37-92cf-691bd2d42f41"})
type: PromQL
message: OCP Versions 4.14.34, 4.15.25, 4.16.7 and ALL subsequent versions
include kernel datastructure changes which are not compatible with older
versions of the SR-IOV operator. Please update SR-IOV operator to versions
dated 20240826 or newer before updating OCP.
name: SRIOVFailedToConfigureVF
url: https://issues.redhat.com/browse/NHE-1171

```

7.3. UPDATING THE OPENSHIFT CONTAINER PLATFORM VERSION IN A HOSTED CLUSTER

Hosted control planes enables the decoupling of updates between the control plane and the data plane.

As a cluster service provider or cluster administrator, you can manage the control plane and the data separately.

You can update a control plane by modifying the **HostedCluster** custom resource (CR) and a node by modifying its **NodePool** CR. Both the **HostedCluster** and **NodePool** CRs specify an OpenShift Container Platform release image in a **.release** field.

To keep your hosted cluster fully operational during an update process, the control plane and the node updates must follow the [Kubernetes version skew policy](#).

7.3.1. The multicluster engine Operator hub management cluster

The multicluster engine for Kubernetes Operator requires a specific OpenShift Container Platform version for the management cluster to remain in a supported state. You can install the multicluster engine Operator from OperatorHub in the OpenShift Container Platform web console.

See the following support matrices for the multicluster engine Operator versions:

- [multicluster engine Operator 2.7](#)
- [multicluster engine Operator 2.6](#)
- [multicluster engine Operator 2.5](#)
- [multicluster engine Operator 2.4](#)

The multicluster engine Operator supports the following OpenShift Container Platform versions:

- The latest unreleased version
- The latest released version
- Two versions before the latest released version

You can also get the multicluster engine Operator version as a part of Red Hat Advanced Cluster Management (RHACM).

7.3.2. Supported OpenShift Container Platform versions in a hosted cluster

When deploying a hosted cluster, the OpenShift Container Platform version of the management cluster does not affect the OpenShift Container Platform version of a hosted cluster.

The HyperShift Operator creates the **supported-versions** ConfigMap in the **hypershift** namespace. The **supported-versions** ConfigMap describes the range of supported OpenShift Container Platform versions that you can deploy.

See the following example of the **supported-versions** ConfigMap:

```
apiVersion: v1
data:
  server-version: 2f6cfe21a0861dea3130f3bed0d3ae5553b8c28b
  supported-versions: '{"versions":["4.17","4.16","4.15","4.14"]}'
kind: ConfigMap
metadata:
  creationTimestamp: "2024-06-20T07:12:31Z"
  labels:
    hypershift.openshift.io/supported-versions: "true"
  name: supported-versions
  namespace: hypershift
  resourceVersion: "927029"
  uid: f6336f91-33d3-472d-b747-94abae725f70
```



IMPORTANT

To create a hosted cluster, you must use the OpenShift Container Platform version from the support version range. However, the multicluster engine Operator can manage only between **n+1** and **n-2** OpenShift Container Platform versions, where **n** defines the current minor version. You can check the multicluster engine Operator support matrix to ensure the hosted clusters managed by the multicluster engine Operator are within the supported OpenShift Container Platform range.

To deploy a higher version of a hosted cluster on OpenShift Container Platform, you must update the multicluster engine Operator to a new minor version release to deploy a new version of the Hypershift

Operator. Upgrading the multicluster engine Operator to a new patch, or z-stream, release does not update the HyperShift Operator to the next version.

See the following example output of the **hcp version** command that shows the supported OpenShift Container Platform versions for OpenShift Container Platform 4.16 in the management cluster:

```
Client Version: openshift/hypershift: fe67b47fb60e483fe60e4755a02b3be393256343. Latest
supported OCP: 4.17.0
Server Version: 05864f61f24a8517731664f8091cedfcf5f9b60d
Server Supports OCP Versions: 4.17, 4.16, 4.15, 4.14
```

7.4. UPDATES FOR THE HOSTED CLUSTER

The **spec.release.image** value dictates the version of the control plane. The **HostedCluster** object transmits the intended **spec.release.image** value to the **HostedControlPlane.spec.releaseImage** value and runs the appropriate Control Plane Operator version.

The hosted control plane manages the rollout of the new version of the control plane components along with any OpenShift Container Platform components through the new version of the Cluster Version Operator (CVO).



IMPORTANT

In hosted control planes, the **NodeHealthCheck** resource cannot detect the status of the CVO. A cluster administrator must manually pause the remediation triggered by **NodeHealthCheck**, before performing critical operations, such as updating the cluster, to prevent new remediation actions from interfering with cluster updates.

To pause the remediation, enter the array of strings, for example, **pause-test-cluster**, as a value of the **pauseRequests** field in the **NodeHealthCheck** resource. For more information, see [About the Node Health Check Operator](#).

After the cluster update is complete, you can edit or delete the remediation. Navigate to the **Compute → NodeHealthCheck** page, click your node health check, and then click **Actions**, which shows a drop-down list.

7.5. UPDATES FOR NODE POOLS

With node pools, you can configure the software that is running in the nodes by exposing the **spec.release** and **spec.config** values. You can start a rolling node pool update in the following ways:

- Changing the **spec.release** or **spec.config** values.
- Changing any platform-specific field, such as the AWS instance type. The result is a set of new instances with the new type.
- Changing the cluster configuration, if the change propagates to the node.

Node pools support replace updates and in-place updates. The **nodepool.spec.release** value dictates the version of any particular node pool. A **NodePool** object completes a replace or an in-place rolling update according to the **.spec.management.upgradeType** value.

After you create a node pool, you cannot change the update type. If you want to change the update type, you must create a node pool and delete the other one.

7.5.1. Replace updates for node pools

A *replace* update creates instances in the new version while it removes old instances from the previous version. This update type is effective in cloud environments where this level of immutability is cost effective.

Replace updates do not preserve any manual changes because the node is entirely re-provisioned.

7.5.2. In place updates for node pools

An *in-place* update directly updates the operating systems of the instances. This type is suitable for environments where the infrastructure constraints are higher, such as bare metal.

In-place updates can preserve manual changes, but will report errors if you make manual changes to any file system or operating system configuration that the cluster directly manages, such as kubelet certificates.

7.6. UPDATING NODE POOLS IN A HOSTED CLUSTER

You can update your version of OpenShift Container Platform by updating the node pools in your hosted cluster. The node pool version must not surpass the hosted control plane version.

The **.spec.release** field in the **NodePool** custom resource (CR) shows the version of a node pool.

Procedure

- Change the **spec.release.image** value in the node pool by entering the following command:

```
$ oc patch nodepool <node_pool_name> -n <hosted_cluster_namespace> \ ①
  --type=merge \
  -p '{"spec":{"nodeDrainTimeout":"60s","release":{"image":"<openshift_release_image>"}}}' ②
```

- Replace **<node_pool_name>** and **<hosted_cluster_namespace>** with your node pool name and hosted cluster namespace, respectively.
- The **<openshift_release_image>** variable specifies the new OpenShift Container Platform release image that you want to upgrade to, for example, **quay.io/openshift-release-dev/ocp-release:4.y.z-x86_64**. Replace **<4.y.z>** with the supported OpenShift Container Platform version.

Verification

- To verify that the new version was rolled out, check the **.status.conditions** value in the node pool by running the following command:

```
$ oc get -n <hosted_cluster_namespace> nodepool <node_pool_name> -o yaml
```

Example output

```
status:
conditions:
- lastTransitionTime: "2024-05-20T15:00:40Z"
```

1 message: 'Using release image: quay.io/openshift-release-dev/ocp-release:4.y.z-x86_64'
 reason: AsExpected
 status: "True"
 type: ValidReleaseImage

- 1 Replace <4.y.z> with the supported OpenShift Container Platform version.

7.7. UPDATING A CONTROL PLANE IN A HOSTED CLUSTER

On hosted control planes, you can upgrade your version of OpenShift Container Platform by updating the hosted cluster. The **.spec.release** in the **HostedCluster** custom resource (CR) shows the version of the control plane. The **HostedCluster** updates the **.spec.release** field to the **HostedControlPlane.spec.release** and runs the appropriate Control Plane Operator version.

The **HostedControlPlane** resource orchestrates the rollout of the new version of the control plane components along with the OpenShift Container Platform component in the data plane through the new version of the Cluster Version Operator (CVO). The **HostedControlPlane** includes the following artifacts:

- CVO
- Cluster Network Operator (CNO)
- Cluster Ingress Operator
- Manifests for the Kube API server, scheduler, and manager
- Machine approver
- Autoscaler
- Infrastructure resources to enable ingress for control plane endpoints such as the Kube API server, ignition, and konnectivity

You can set the **.spec.release** field in the **HostedCluster** CR to update the control plane by using the information from the **status.version.availableUpdates** and **status.version.conditionalUpdates** fields.

Procedure

- 1 Add the **hypershift.openshift.io/force-upgrade-to=<openshift_release_image>** annotation to the hosted cluster by entering the following command:

```
$ oc annotate hostedcluster \
-n <hosted_cluster_namespace> <hosted_cluster_name> \ ①
"hypershift.openshift.io/force-upgrade-to=<openshift_release_image>" \ ②
--overwrite
```

- 1 Replace <hosted_cluster_name> and <hosted_cluster_namespace> with your hosted cluster name and hosted cluster namespace, respectively.
- 2 The <openshift_release_image> variable specifies the new OpenShift Container Platform release image that you want to upgrade to, for example, **quay.io/openshift-release-dev/ocp-release:4.y.z-x86_64**. Replace <4.y.z> with the supported OpenShift

Container Platform version.

- Change the **spec.release.image** value in the hosted cluster by entering the following command:

```
$ oc patch hostedcluster <hosted_cluster_name> -n <hosted_cluster_namespace> \
--type=merge \
-p '{"spec":{"release":{"image":"<openshift_release_image>"}}}'
```

Verification

- To verify that the new version was rolled out, check the **.status.conditions** and **.status.version** values in the hosted cluster by running the following command:

```
$ oc get -n <hosted_cluster_namespace> hostedcluster <hosted_cluster_name> \
-o yaml
```

Example output

```
status:
conditions:
- lastTransitionTime: "2024-05-20T15:01:01Z"
  message: Payload loaded version="4.y.z" image="quay.io/openshift-release-dev/ocp-
release:4.y.z-x86_64" ①
  status: "True"
  type: ClusterVersionReleaseAccepted
#...
version:
  availableUpdates: null
  desired:
    image: quay.io/openshift-release-dev/ocp-release:4.y.z-x86_64 ②
    version: 4.y.z
```

① ② Replace <4.y.z> with the supported OpenShift Container Platform version.

7.8. UPDATING A HOSTED CLUSTER BY USING THE MULTICLUSTER ENGINE OPERATOR CONSOLE

You can update your hosted cluster by using the multicluster engine Operator console.



IMPORTANT

Before updating a hosted cluster, you must refer to the available and conditional updates of a hosted cluster. Choosing a wrong release version might break the hosted cluster.

Procedure

- Select All clusters.
- Navigate to Infrastructure → Clusters to view managed hosted clusters.

3. Click the **Upgrade available** link to update the control plane and node pools.

CHAPTER 8. HIGH AVAILABILITY FOR HOSTED CONTROL PLANES

8.1. ABOUT HIGH AVAILABILITY FOR HOSTED CONTROL PLANES

You can maintain high availability (HA) of hosted control planes by implementing the following actions:

- Recover etcd members for a hosted cluster.
- Back up and restore etcd for a hosted cluster.
- Perform a disaster recovery process for a hosted cluster.

8.1.1. Impact of the failed management cluster component

If the management cluster component fails, your workload remains unaffected. In the OpenShift Container Platform management cluster, the control plane is decoupled from the data plane to provide resiliency.

The following table covers the impact of a failed management cluster component on the control plane and the data plane. However, the table does not cover all scenarios for the management cluster component failures.

Table 8.1. Impact of the failed component on hosted control planes

Name of the failed component	Hosted control plane API status	Hosted cluster data plane status
Worker node	Available	Available
Availability zone	Available	Available
Management cluster control plane	Available	Available
Management cluster control plane and worker nodes	Not available	Available

8.2. RECOVERING AN UNHEALTHY ETCD CLUSTER

In a highly available control plane, three etcd pods run as a part of a stateful set in an etcd cluster. To recover an etcd cluster, identify unhealthy etcd pods by checking the etcd cluster health.

8.2.1. Checking the status of an etcd cluster

You can check the status of the etcd cluster health by logging into any etcd pod.

Procedure

1. Log in to an etcd pod by entering the following command:

```
$ oc rsh -n openshift-etcd -c etcd <etcd_pod_name>
```

- Print the health status of an etcd cluster by entering the following command:

```
sh-4.4# etcdctl endpoint status -w table
```

Example output

ENDPOINT	ID	VERSION	DB SIZE	IS LEADER	IS LEARNER	RAFT TERM	RAFT INDEX	RAFT APPLIED INDEX	ERRORS
https://192.168.1xxx.20:2379 8fxxxxxxxxxx 3.5.12 123 MB false false 10 180156 180156									
https://192.168.1xxx.21:2379 a5xxxxxxxxxx 3.5.12 122 MB false false 10 180156 180156									
https://192.168.1xxx.22:2379 7cxxxxxxxxxxx 3.5.12 124 MB true false 10 180156 180156									
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									

8.2.2. Recovering a failing etcd pod

Each etcd pod of a 3-node cluster has its own persistent volume claim (PVC) to store its data. An etcd pod might fail because of corrupted or missing data. You can recover a failing etcd pod and its PVC.

Procedure

- To confirm that the etcd pod is failing, enter the following command:

```
$ oc get pods -l app=etcd -n openshift-etcd
```

Example output

NAME	READY	STATUS	RESTARTS	AGE
etcd-0	2/2	Running	0	64m
etcd-1	2/2	Running	0	45m
etcd-2	1/2	CrashLoopBackOff	1 (5s ago)	64m

The failing etcd pod might have the **CrashLoopBackOff** or **Error** status.

- Delete the failing pod and its PVC by entering the following command:

```
$ oc delete pods etcd-2 -n openshift-etcd
```

Verification

- Verify that a new etcd pod is up and running by entering the following command:

```
$ oc get pods -l app=etcd -n openshift-etcd
```

Example output

NAME	READY	STATUS	RESTARTS	AGE
etcd-0	2/2	Running	0	67m
etcd-1	2/2	Running	0	48m
etcd-2	2/2	Running	0	2m2s

8.3. BACKING UP AND RESTORING ETCD IN AN ON-PREMISE ENVIRONMENT

You can back up and restore etcd on a hosted cluster in an on-premise environment to fix failures.

8.3.1. Backing up and restoring etcd on a hosted cluster in an on-premise environment

By backing up and restoring etcd on a hosted cluster, you can fix failures, such as corrupted or missing data in an etcd member of a three node cluster. If multiple members of the etcd cluster encounter data loss or have a **CrashLoopBackOff** status, this approach helps prevent an etcd quorum loss.



IMPORTANT

This procedure requires API downtime.

Prerequisites

- The **oc** and **jq** binaries have been installed.

Procedure

- First, set up your environment variables:

- Set up environment variables for your hosted cluster by entering the following commands, replacing values as necessary:

```
$ CLUSTER_NAME=my-cluster
```

```
$ HOSTED_CLUSTER_NAMESPACE=clusters
```

```
$ CONTROL_PLANE_NAMESPACE="${HOSTED_CLUSTER_NAMESPACE}-${CLUSTER_NAME}"
```

- Pause reconciliation of the hosted cluster by entering the following command, replacing values as necessary:

```
$ oc patch -n ${HOSTED_CLUSTER_NAMESPACE} hostedclusters/${CLUSTER_NAME} \
-p '{"spec":{"pausedUntil":"true"}}' --type=merge
```

- Next, take a snapshot of etcd by using one of the following methods:

- Use a previously backed-up snapshot of etcd.
- If you have an available etcd pod, take a snapshot from the active etcd pod by completing the following steps:

- i. List etcd pods by entering the following command:

```
$ oc get -n ${CONTROL_PLANE_NAMESPACE} pods -l app=etcd
```

- ii. Take a snapshot of the pod database and save it locally to your machine by entering the following commands:

```
$ ETCD_POD=etcd-0
```

```
$ oc exec -n ${CONTROL_PLANE_NAMESPACE} -c etcd -t ${ETCD_POD} -- \
env ETCDCTL_API=3 /usr/bin/etcdctl \
--cacert /etc/etcd/tls/etcd-ca/ca.crt \
--cert /etc/etcd/tls/client/etcd-client.crt \
--key /etc/etcd/tls/client/etcd-client.key \
--endpoints=https://localhost:2379 \
snapshot save /var/lib/snapshot.db
```

- iii. Verify that the snapshot is successful by entering the following command:

```
$ oc exec -n ${CONTROL_PLANE_NAMESPACE} -c etcd -t ${ETCD_POD} -- \
env ETCDCTL_API=3 /usr/bin/etcdctl -w table snapshot status \
/var/lib/snapshot.db
```

- c. Make a local copy of the snapshot by entering the following command:

```
$ oc cp -c etcd \
${CONTROL_PLANE_NAMESPACE}/${ETCD_POD}:/var/lib/snapshot.db \
/tmp/etcd.snapshot.db
```

- i. Make a copy of the snapshot database from etcd persistent storage:

- A. List etcd pods by entering the following command:

```
$ oc get -n ${CONTROL_PLANE_NAMESPACE} pods -l app=etcd
```

- B. Find a pod that is running and set its name as the value of **ETCD_POD**: **ETCD_POD=etcd-0**, and then copy its snapshot database by entering the following command:

```
$ oc cp -c etcd \
${CONTROL_PLANE_NAMESPACE}/${ETCD_POD}:/var/lib/data/member/snap/
db \
/tmp/etcd.snapshot.db
```

3. Next, scale down the etcd statefulset by entering the following command:

```
$ oc scale -n ${CONTROL_PLANE_NAMESPACE} statefulset/etcd --replicas=0
```

- a. Delete volumes for second and third members by entering the following command:

```
$ oc delete -n ${CONTROL_PLANE_NAMESPACE} pvc/data-etcd-1 pvc/data-etcd-2
```

- b. Create a pod to access the first etcd member's data:

- i. Get the etcd image by entering the following command:

```
$ ETCD_IMAGE=$(oc get -n ${CONTROL_PLANE_NAMESPACE} statefulset/etcd \
-o jsonpath='{ .spec.template.spec.containers[0].image }')
```

- ii. Create a pod that allows access to etcd data:

```
$ cat << EOF | oc apply -n ${CONTROL_PLANE_NAMESPACE} -f -
apiVersion: apps/v1
kind: Deployment
metadata:
  name: etcd-data
spec:
  replicas: 1
  selector:
    matchLabels:
      app: etcd-data
  template:
    metadata:
      labels:
        app: etcd-data
    spec:
      containers:
        - name: access
          image: $ETCD_IMAGE
          volumeMounts:
            - name: data
              mountPath: /var/lib
          command:
            - /usr/bin/bash
          args:
            - -c
            - |
              while true; do
                sleep 1000
              done
          volumes:
            - name: data
          persistentVolumeClaim:
            claimName: data-etcd-0
EOF
```

- iii. Check the status of the **etcd-data** pod and wait for it to be running by entering the following command:

```
$ oc get -n ${CONTROL_PLANE_NAMESPACE} pods -l app=etcd-data
```

- iv. Get the name of the **etcd-data** pod by entering the following command:

```
$ DATA_POD=$(oc get -n ${CONTROL_PLANE_NAMESPACE} pods --no-headers \
\\
-l app=etcd-data -o name | cut -d/ -f2)
```

- c. Copy an etcd snapshot into the pod by entering the following command:

```
$ oc cp /tmp/etcd.snapshot.db \
${CONTROL_PLANE_NAMESPACE}/${DATA_POD}:/var/lib/restored.snap.db
```

- d. Remove old data from the **etcd-data** pod by entering the following commands:

```
$ oc exec -n ${CONTROL_PLANE_NAMESPACE} ${DATA_POD} -- rm -rf /var/lib/data
```

```
$ oc exec -n ${CONTROL_PLANE_NAMESPACE} ${DATA_POD} -- mkdir -p
/var/lib/data
```

- e. Restore the etcd snapshot by entering the following command:

```
$ oc exec -n ${CONTROL_PLANE_NAMESPACE} ${DATA_POD} -- \
etcdutil snapshot restore /var/lib/restored.snap.db \
--data-dir=/var/lib/data --skip-hash-check \
--name etcd-0 \
--initial-cluster-token=etcd-cluster \
--initial-cluster etcd-0=https://etcd-0.etcd-
discovery.${CONTROL_PLANE_NAMESPACE}.svc:2380,etcd-1=https://etcd-1.etcd-
discovery.${CONTROL_PLANE_NAMESPACE}.svc:2380,etcd-2=https://etcd-2.etcd-
discovery.${CONTROL_PLANE_NAMESPACE}.svc:2380 \
--initial-advertise-peer-urls https://etcd-0.etcd-
discovery.${CONTROL_PLANE_NAMESPACE}.svc:2380
```

- f. Remove the temporary etcd snapshot from the pod by entering the following command:

```
$ oc exec -n ${CONTROL_PLANE_NAMESPACE} ${DATA_POD} -- \
rm /var/lib/restored.snap.db
```

- g. Delete data access deployment by entering the following command:

```
$ oc delete -n ${CONTROL_PLANE_NAMESPACE} deployment/etcd-data
```

- h. Scale up the etcd cluster by entering the following command:

```
$ oc scale -n ${CONTROL_PLANE_NAMESPACE} statefulset/etcd --replicas=3
```

- i. Wait for the etcd member pods to return and report as available by entering the following command:

```
$ oc get -n ${CONTROL_PLANE_NAMESPACE} pods -l app=etcd -w
```

4. Restore reconciliation of the hosted cluster by entering the following command:

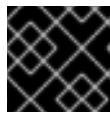
```
$ oc patch -n ${HOSTED_CLUSTER_NAMESPACE} hostedclusters/${CLUSTER_NAME} \
-p '{"spec":{"pausedUntil":"null"}}' --type=merge
```

8.4. BACKING UP AND RESTORING ETCD ON AWS

You can back up and restore etcd on a hosted cluster on Amazon Web Services (AWS) to fix failures.

8.4.1. Taking a snapshot of etcd for a hosted cluster

To back up etcd for a hosted cluster, you must take a snapshot of etcd. Later, you can restore etcd by using the snapshot.



IMPORTANT

This procedure requires API downtime.

Procedure

1. Pause reconciliation of the hosted cluster by entering the following command:

```
$ oc patch -n clusters hostedclusters/<hosted_cluster_name> \
-p '{"spec":{"pausedUntil":"true"}}' --type=merge
```

2. Stop all etcd-writer deployments by entering the following command:

```
$ oc scale deployment -n <hosted_cluster_namespace> --replicas=0 \
kube-apiserver openshift-apiserver openshift-oauth-apiserver
```

3. To take an etcd snapshot, use the **exec** command in each etcd container by entering the following command:

```
$ oc exec -it <etcd_pod_name> -n <hosted_cluster_namespace> -- \
env ETCDCTL_API=3 /usr/bin/etcdctl \
--cacert /etc/etcd/tls/etcd-ca/ca.crt \
--cert /etc/etcd/tls/client/etcd-client.crt \
--key /etc/etcd/tls/client/etcd-client.key \
--endpoints=localhost:2379 \
snapshot save /var/lib/data/snapshot.db
```

4. To check the snapshot status, use the **exec** command in each etcd container by running the following command:

```
$ oc exec -it <etcd_pod_name> -n <hosted_cluster_namespace> -- \
env ETCDCTL_API=3 /usr/bin/etcdctl -w table snapshot status \
/var/lib/data/snapshot.db
```

5. Copy the snapshot data to a location where you can retrieve it later, such as an S3 bucket. See the following example.



NOTE

The following example uses signature version 2. If you are in a region that supports signature version 4, such as the **us-east-2** region, use signature version 4. Otherwise, when copying the snapshot to an S3 bucket, the upload fails.

Example

```
BUCKET_NAME=somebucket
CLUSTER_NAME=cluster_name
FILEPATH="/${BUCKET_NAME}/${CLUSTER_NAME}-snapshot.db"
```

```

CONTENT_TYPE="application/x-compressed-tar"
DATE_VALUE=`date -R`
SIGNATURE_STRING="PUT\n\n${CONTENT_TYPE}\n${DATE_VALUE}\n${FILEPATH}"
ACCESS_KEY=accesskey
SECRET_KEY=secret
SIGNATURE_HASH=`echo -en ${SIGNATURE_STRING} | openssl sha1 -hmac
${SECRET_KEY} -binary | base64`
HOSTED_CLUSTER_NAMESPACE=hosted_cluster_namespace

oc exec -it etcd-0 -n ${HOSTED_CLUSTER_NAMESPACE} -- curl -X PUT -T
"/var/lib/data/snapshot.db" \
-H "Host: ${BUCKET_NAME}.s3.amazonaws.com" \
-H "Date: ${DATE_VALUE}" \
-H "Content-Type: ${CONTENT_TYPE}" \
-H "Authorization: AWS ${ACCESS_KEY}:${SIGNATURE_HASH}" \
https://${BUCKET_NAME}.s3.amazonaws.com/${CLUSTER_NAME}-snapshot.db

```

- To restore the snapshot on a new cluster later, save the encryption secret that the hosted cluster references.

- Get the secret encryption key by entering the following command:

```
$ oc get hostedcluster <hosted_cluster_name> \
-o=jsonpath='{.spec.secretEncryption.aescbc}' \
{"activeKey":{"name":"<hosted_cluster_name>-etcd-encryption-key"}}
```

- Save the secret encryption key by entering the following command:

```
$ oc get secret <hosted_cluster_name>-etcd-encryption-key \
-o=jsonpath='{.data.key}'
```

You can decrypt this key when restoring a snapshot on a new cluster.

- Restart all etcd-writer deployments by entering the following command:

```
$ oc scale deployment -n <control_plane_namespace> --replicas=3 \
kube-apiserver openshift-apiserver openshift-oauth-apiserver
```

- Resume the reconciliation of the hosted cluster by entering the following command:

```
$ oc patch -n <hosted_cluster_namespace> \
-p '[{"op": "remove", "path": "/spec/pausedUntil"}]' --type=json
```

Next steps

Restore the etcd snapshot.

8.4.2. Restoring an etcd snapshot on a hosted cluster

If you have a snapshot of etcd from your hosted cluster, you can restore it. Currently, you can restore an etcd snapshot only during cluster creation.

To restore an etcd snapshot, you modify the output from the **create cluster --render** command and define a **restoreSnapshotURL** value in the etcd section of the **HostedCluster** specification.

**NOTE**

The **--render** flag in the **hcp create** command does not render the secrets. To render the secrets, you must use both the **--render** and the **--render-sensitive** flags in the **hcp create** command.

Prerequisites

You took an etcd snapshot on a hosted cluster.

Procedure

1. On the **aws** command-line interface (CLI), create a pre-signed URL so that you can download your etcd snapshot from S3 without passing credentials to the etcd deployment:

```
ETCD_SNAPSHOT=${ETCD_SNAPSHOT:-"s3://${BUCKET_NAME}/${CLUSTER_NAME}-snapshot.db"}
ETCD_SNAPSHOT_URL=$(aws s3 presign ${ETCD_SNAPSHOT})
```

2. Modify the **HostedCluster** specification to refer to the URL:

```
spec:
etcd:
managed:
storage:
persistentVolume:
size: 4Gi
type: PersistentVolume
restoreSnapshotURL:
- "${ETCD_SNAPSHOT_URL}"
managementType: Managed
```

3. Ensure that the secret that you referenced from the **spec.secretEncryption.aescbc** value contains the same AES key that you saved in the previous steps.

8.5. BACKING UP AND RESTORING A HOSTED CLUSTER ON OPENSHIFT VIRTUALIZATION

You can back up and restore a hosted cluster on OpenShift Virtualization to fix failures.

8.5.1. Backing up a hosted cluster on OpenShift Virtualization

When you back up a hosted cluster on OpenShift Virtualization, the hosted cluster can remain running. The backup contains the hosted control plane components and the etcd for the hosted cluster.

When the hosted cluster is not running compute nodes on external infrastructure, hosted cluster workload data that is stored in persistent volume claims (PVCs) that are provisioned by KubeVirt CSI are also backed up. The backup does not contain any KubeVirt virtual machines (VMs) that are used as compute nodes. Those VMs are automatically re-created after the restore process is completed.

Procedure

1. Create a Velero backup resource by creating a YAML file that is similar to the following example:

```

apiVersion: velero.io/v1
kind: Backup
metadata:
  name: hc-clusters-hosted-backup
  namespace: openshift-adp
  labels:
    velero.io/storage-location: default
spec:
  includedNamespaces: ①
    - clusters
    - clusters-hosted
  includedResources:
    - sa
    - role
    - rolebinding
    - deployment
    - statefulset
    - pv
    - pvc
    - bmh
    - configmap
    - infraenv
    - priorityclasses
    - pdb
    - hostedcluster
    - nodepool
    - secrets
    - hostedcontrolplane
    - cluster
    - datavolume
    - service
    - route
  excludedResources: [ ]
  labelSelector: ②
    matchExpressions:
      - key: 'hypershift.openshift.io/is-kubevirt-rhcos'
        operator: 'DoesNotExist'
  storageLocation: default
  preserveNodePorts: true
  ttl: 4h0m0s
  snapshotMoveData: true ③
  datamover: "velero" ④
  defaultVolumesToFsBackup: false ⑤

```

- ① This field selects the namespaces from the objects to back up. Include namespaces from both the hosted cluster and the hosted control plane. In this example, **clusters** is a namespace from the hosted cluster and **clusters-hosted** is a namespace from the hosted control plane. By default, the **HostedControlPlane** namespace is **clusters-<hosted_cluster_name>**.
- ② The boot image of the VMs that are used as the hosted cluster nodes are stored in large PVCs. To reduce backup time and storage size, you can filter those PVCs out of the backup by adding this label selector.
- ③ This field and the **datamover** field enable automatically uploading the CSI **VolumeSnapshots** to remote cloud storage.

VolumeSnapshots to remote cloud storage.

- ④ This field and the **snapshotMoveData** field enable automatically uploading the CSI **VolumeSnapshots** to remote cloud storage.
- ⑤ This field indicates whether pod volume file system backup is used for all volumes by default. Set this value to **false** to back up the PVCs that you want.

2. Apply the changes to the YAML file by entering the following command:

```
$ oc apply -f <backup_file_name>.yaml
```

Replace **<backup_file_name>** with the name of your file.

3. Monitor the backup process in the backup object status and in the Velero logs.

- To monitor the backup object status, enter the following command:

```
$ watch "oc get backups.velero.io -n openshift-adp <backup_file_name> -o jsonpath='{.status}' | jq"
```

- To monitor the Velero logs, enter the following command:

```
$ oc logs -n openshift-adp -ldeploy=velero -f
```

Verification

- When the **status.phase** field is **Completed**, the backup process is considered complete.

8.5.2. Restoring a hosted cluster on OpenShift Virtualization

After you back up a hosted cluster on OpenShift Virtualization, you can restore the backup.



NOTE

The restore process can be completed only on the same management cluster where you created the backup.

Procedure

1. Ensure that no pods or persistent volume claims (PVCs) are running in the **HostedControlPlane** namespace.
2. Delete the following objects from the management cluster:
 - **HostedCluster**
 - **NodePool**
 - PVCs
3. Create a restoration manifest YAML file that is similar to the following example:

```
apiVersion: velero.io/v1
```

```

kind: Restore
metadata:
  name: hc-clusters-hosted-restore
  namespace: openshift-adp
spec:
  backupName: hc-clusters-hosted-backup
  restorePVs: true 1
  existingResourcePolicy: update 2
  excludedResources:
    - nodes
    - events
    - events.events.k8s.io
    - backups.velero.io
    - restores.velero.io
    - resticrepositories.velero.io

```

- 1 This field starts the recovery of pods with the included persistent volumes.
- 2 Setting **existingResourcePolicy** to **update** ensures that any existing objects are overwritten with backup content. This action can cause issues with objects that contain immutable fields, which is why you deleted the **HostedCluster**, node pools, and PVCs. If you do not set this policy, the Velero engine skips the restoration of objects that already exist.

4. Apply the changes to the YAML file by entering the following command:

```
$ oc apply -f <restore_resource_file_name>.yaml
```

Replace **<restore_resource_file_name>** with the name of your file.

5. Monitor the restore process by checking the restore status field and the Velero logs.

- To check the restore status field, enter the following command:

```
$ watch "oc get restores.velero.io -n openshift-adp <backup_file_name> -o jsonpath='{.status}' | jq"
```

- To check the Velero logs, enter the following command:

```
$ oc logs -n openshift-adp -ldeploy=velero -f
```

Verification

- When the **status.phase** field is **Completed**, the restore process is considered complete.

Next steps

- After some time, the KubeVirt VMs are created and join the hosted cluster as compute nodes. Make sure that the hosted cluster workloads are running again as expected.

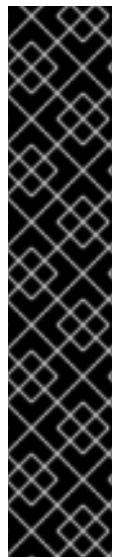
8.6. DISASTER RECOVERY FOR A HOSTED CLUSTER IN AWS

You can recover a hosted cluster to the same region within Amazon Web Services (AWS). For example, you need disaster recovery when the upgrade of a management cluster fails and the hosted cluster is in a read-only state.

The disaster recovery process involves the following steps:

1. Backing up the hosted cluster on the source management cluster
2. Restoring the hosted cluster on a destination management cluster
3. Deleting the hosted cluster from the source management cluster

Your workloads remain running during the process. The Cluster API might be unavailable for a period, but that does not affect the services that are running on the worker nodes.



IMPORTANT

Both the source management cluster and the destination management cluster must have the **--external-dns** flags to maintain the API server URL. That way, the server URL ends with <https://api-sample-hosted.sample-hosted.aws.openshift.com>. See the following example:

Example: External DNS flags

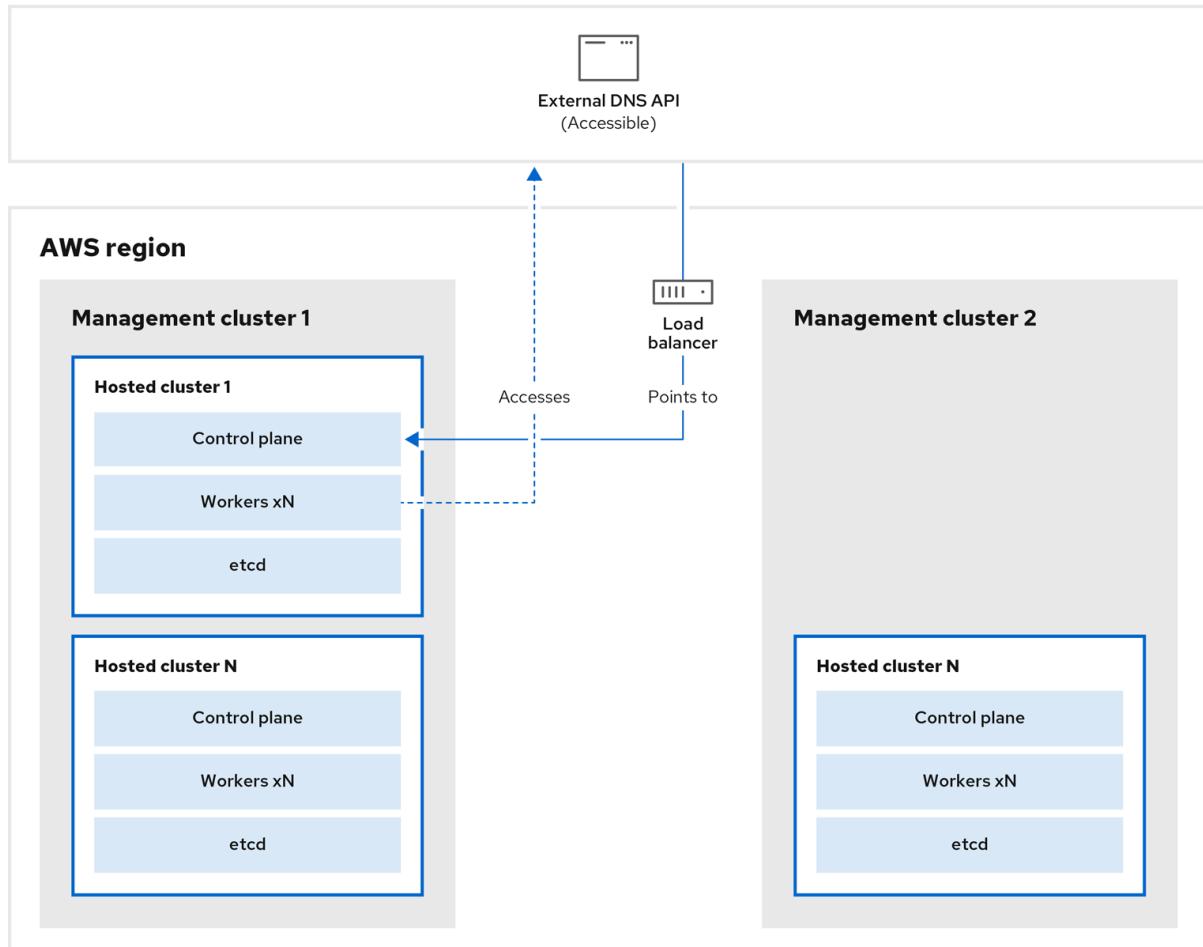
```
--external-dns-provider=aws \
--external-dns-credentials=<path_to_aws_credentials_file> \
--external-dns-domain-filter=<basedomain>
```

If you do not include the **--external-dns** flags to maintain the API server URL, you cannot migrate the hosted cluster.

8.6.1. Overview of the backup and restore process

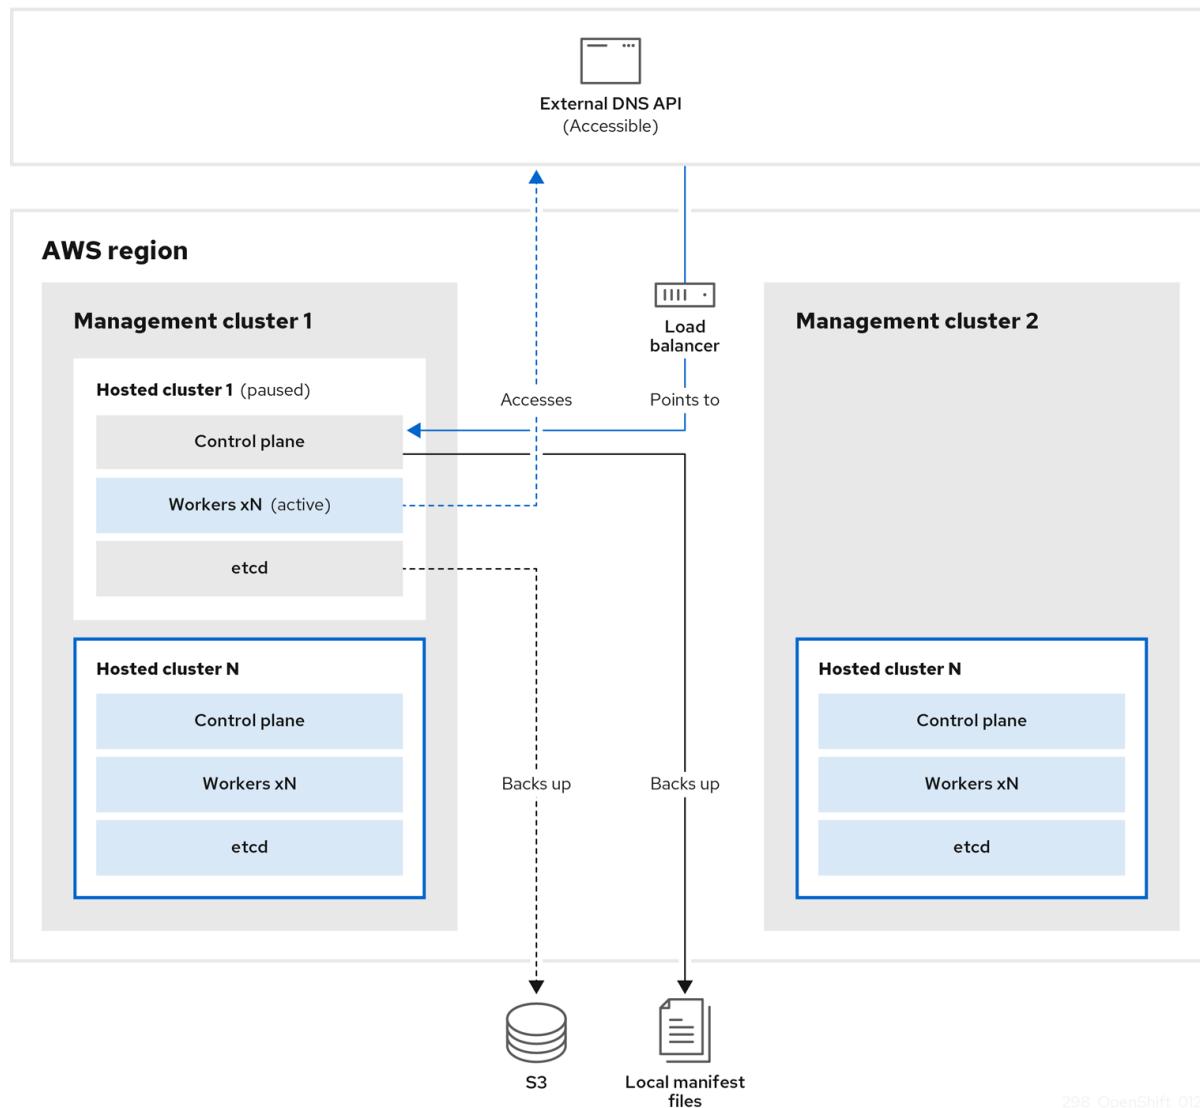
The backup and restore process works as follows:

1. On management cluster 1, which you can think of as the source management cluster, the control plane and workers interact by using the external DNS API. The external DNS API is accessible, and a load balancer sits between the management clusters.



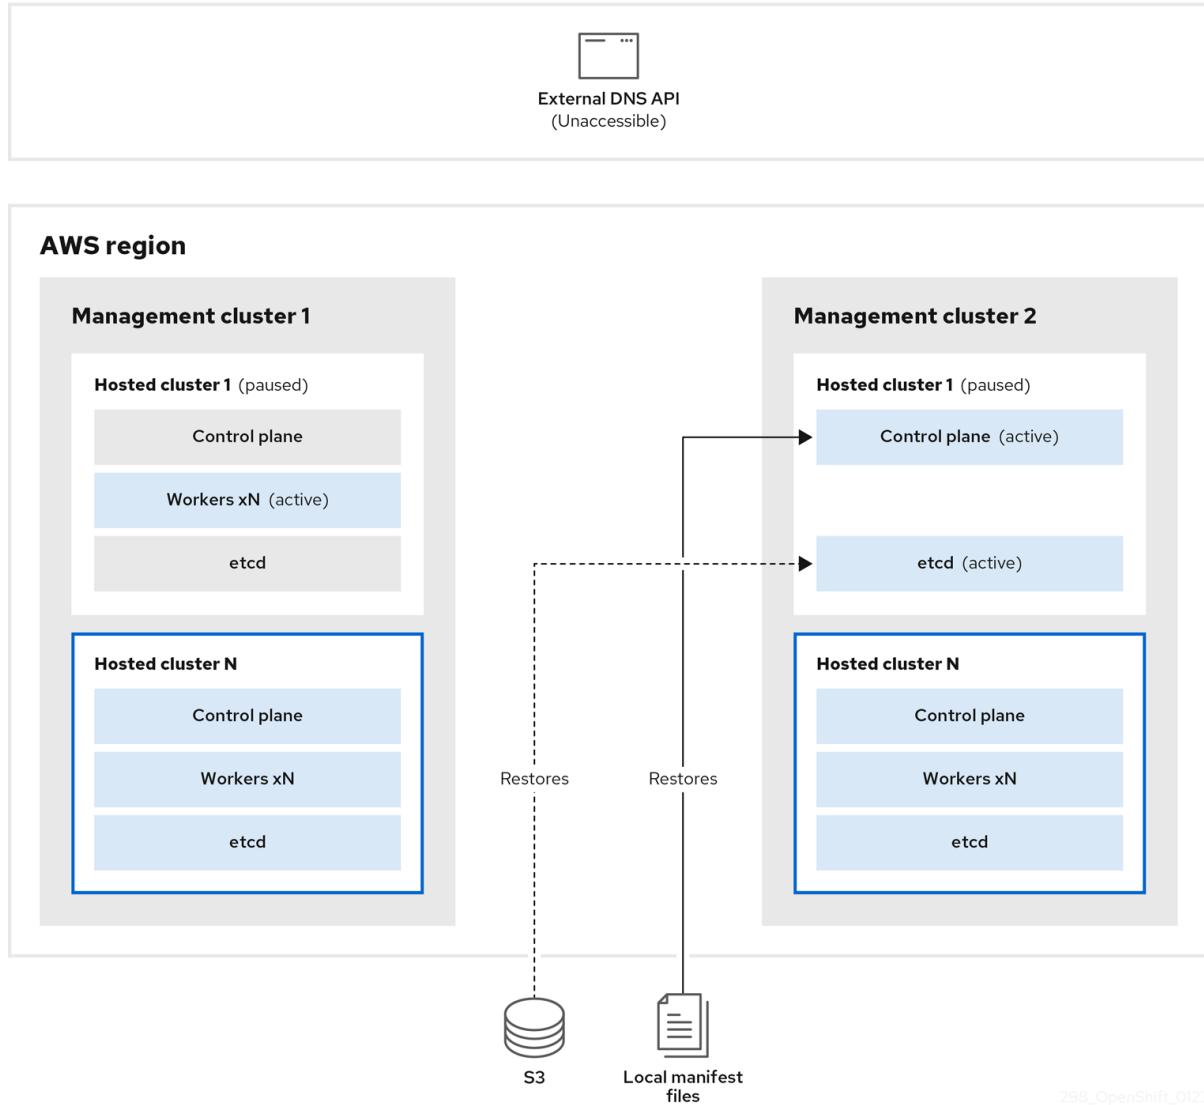
298_OpenShift_0123

2. You take a snapshot of the hosted cluster, which includes etcd, the control plane, and the worker nodes. During this process, the worker nodes continue to try to access the external DNS API even if it is not accessible, the workloads are running, the control plane is saved in a local manifest file, and etcd is backed up to an S3 bucket. The data plane is active and the control plane is paused.

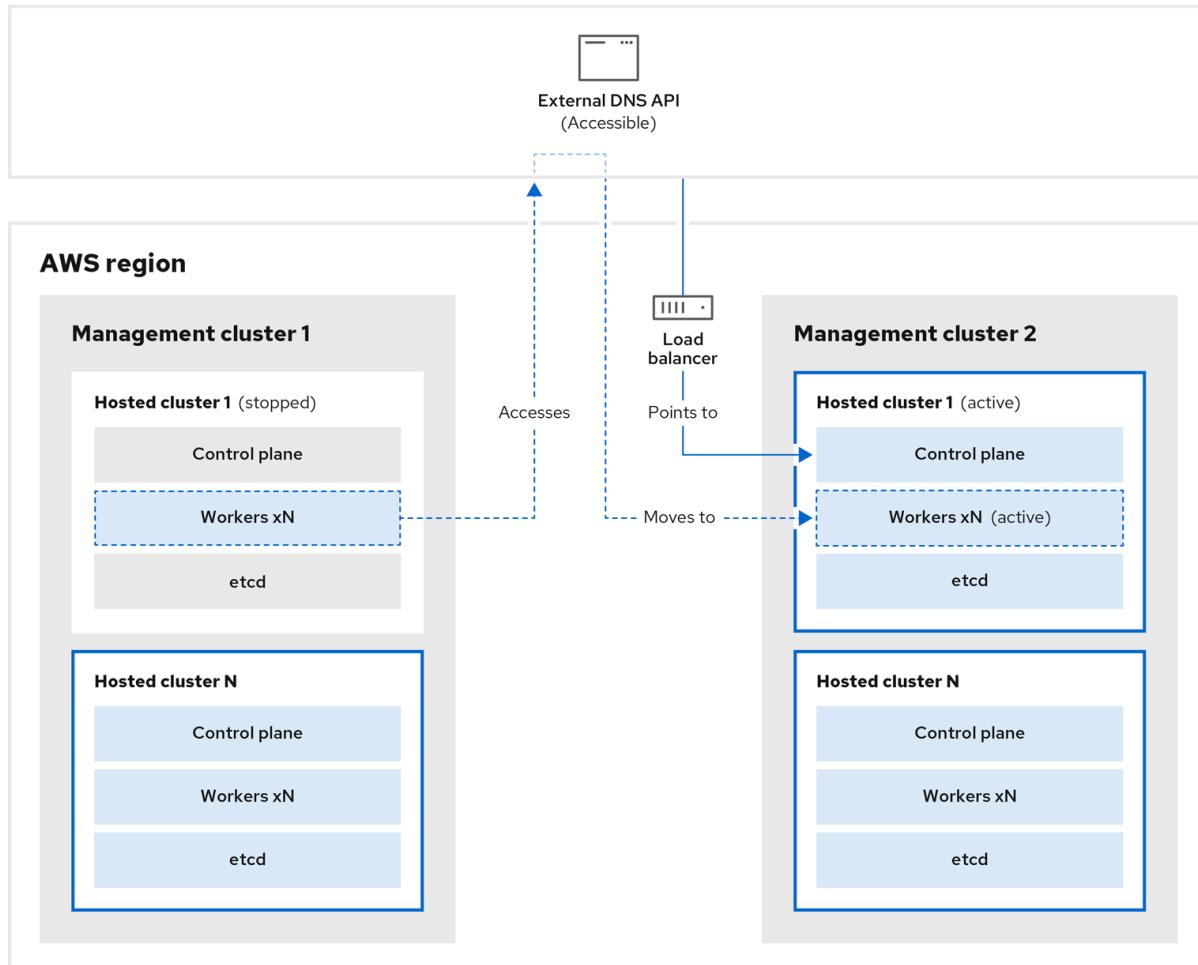


3. On management cluster 2, which you can think of as the destination management cluster, you restore etcd from the S3 bucket and restore the control plane from the local manifest file. During this process, the external DNS API is stopped, the hosted cluster API becomes inaccessible, and any workers that use the API are unable to update their manifest files, but the workloads are still running.

298_OpenShift_0123

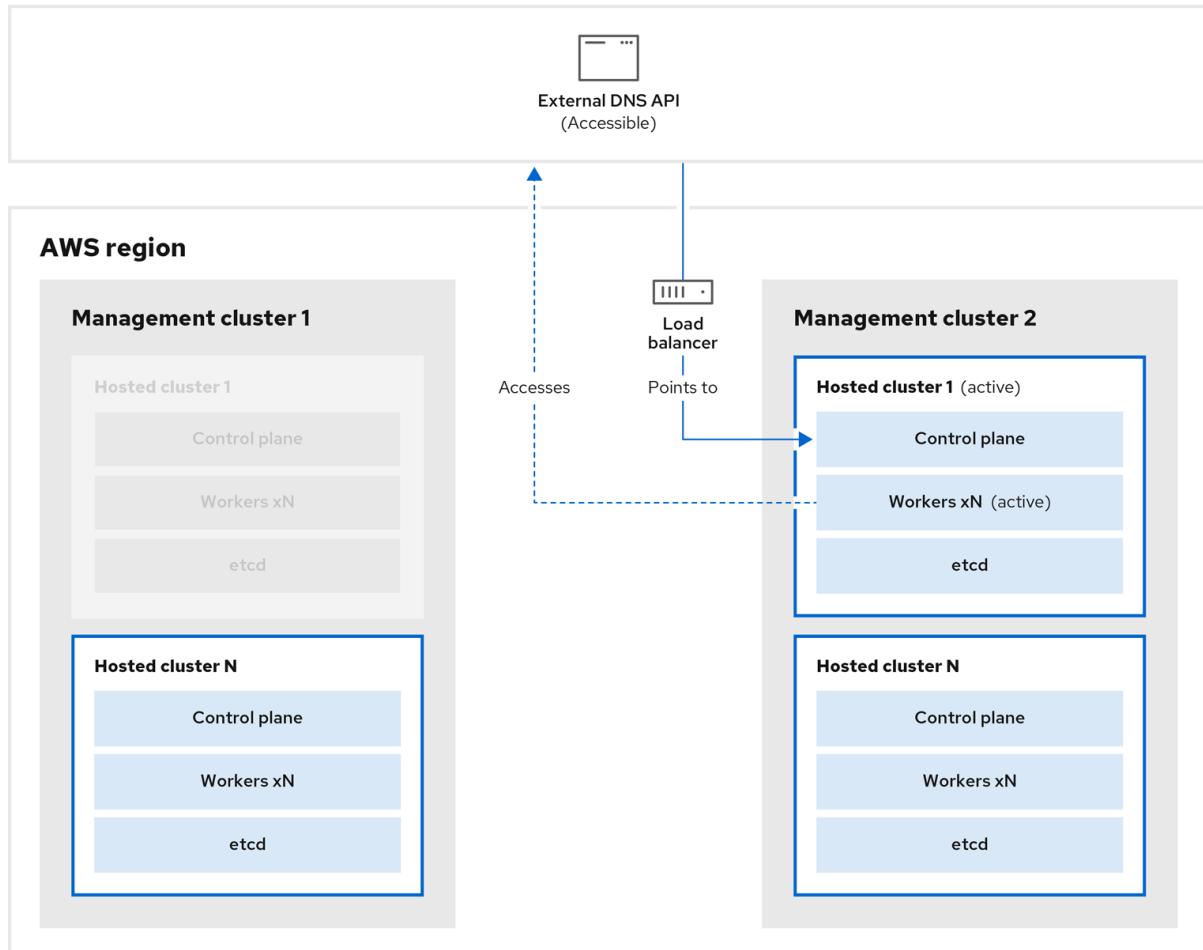


4. The external DNS API is accessible again, and the worker nodes use it to move to management cluster 2. The external DNS API can access the load balancer that points to the control plane.



298_OpenShift_0123

5. On management cluster 2, the control plane and worker nodes interact by using the external DNS API. The resources are deleted from management cluster 1, except for the S3 backup of etcd. If you try to set up the hosted cluster again on management cluster 1, it will not work.



298_OpenShift_0123

8.6.2. Backing up a hosted cluster on AWS

To recover your hosted cluster in your target management cluster, you first need to back up all of the relevant data.

Procedure

1. Create a config map file to declare the source management cluster by entering the following command:

```
$ oc create configmap mgmt-parent-cluster -n default \
--from-literal=from=${MGMT_CLUSTER_NAME}
```

2. Shut down the reconciliation in the hosted cluster and in the node pools by entering the following commands:

```
$ PAUSED_UNTIL="true"
```

```
$ oc patch -n ${HC_CLUSTER_NS} hostedclusters/${HC_CLUSTER_NAME} \
-p '{"spec":{"pausedUntil":"'${PAUSED_UNTIL}'"}}' --type=merge
```

```
$ oc patch -n ${HC_CLUSTER_NS} nodepools/${NODEPOOLS} \
-p '{"spec":{"pausedUntil":"'${PAUSED_UNTIL}'"}}' --type=merge
```

```
$ oc scale deployment -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} --replicas=0 \
  kube-apiserver openshift-apiserver openshift-oauth-apiserver control-plane-operator
```

3. Back up etcd and upload the data to an S3 bucket by running the following bash script:

TIP

Wrap this script in a function and call it from the main function.

```
# ETCD Backup
ETCD_PODS="etcd-0"
if [ "${CONTROL_PLANE_AVAILABILITY_POLICY}" = "HighlyAvailable" ]; then
  ETCD_PODS="etcd-0 etcd-1 etcd-2"
fi

for POD in ${ETCD_PODS}; do
  # Create an etcd snapshot
  oc exec -it ${POD} -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} -- env
  ETCDCTL_API=3 /usr/bin/etcdctl --cacert /etc/etcd/tls/client/etcd-client-ca.crt --cert
  /etc/etcd/tls/client/etcd-client.crt --key /etc/etcd/tls/client/etcd-client.key --
  endpoints=localhost:2379 snapshot save /var/lib/data/snapshot.db
  oc exec -it ${POD} -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} -- env
  ETCDCTL_API=3 /usr/bin/etcdctl -w table snapshot status /var/lib/data/snapshot.db

  FILEPATH="/${BUCKET_NAME}/${HC_CLUSTER_NAME}-${POD}-snapshot.db"
  CONTENT_TYPE="application/x-compressed-tar"
  DATE_VALUE=`date -R`
  SIGNATURE_STRING="PUT\n\n${CONTENT_TYPE}\n${DATE_VALUE}\n${FILEPATH}"

  set +x
  ACCESS_KEY=$(grep aws_access_key_id ${AWS_CREDS} | head -n1 | cut -d= -f2 | sed
"s/ //g")
  SECRET_KEY=$(grep aws_secret_access_key ${AWS_CREDS} | head -n1 | cut -d= -f2 | sed
"s/ //g")
  SIGNATURE_HASH=$(echo -en ${SIGNATURE_STRING} | openssl sha1 -hmac
"${SECRET_KEY}" -binary | base64)
  set -x

  # FIXME: this is pushing to the OIDC bucket
  oc exec -it etcd-0 -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} -- curl -X PUT -T
  "/var/lib/data/snapshot.db" \
  -H "Host: ${BUCKET_NAME}.s3.amazonaws.com" \
  -H "Date: ${DATE_VALUE}" \
  -H "Content-Type: ${CONTENT_TYPE}" \
  -H "Authorization: AWS ${ACCESS_KEY}:${SIGNATURE_HASH}" \
  https://${BUCKET_NAME}.s3.amazonaws.com/${HC_CLUSTER_NAME}-${POD}-
  snapshot.db
done
```

For more information about backing up etcd, see "Backing up and restoring etcd on a hosted cluster".

4. Back up Kubernetes and OpenShift Container Platform objects by entering the following commands. You need to back up the following objects:

- **HostedCluster** and **NodePool** objects from the HostedCluster namespace
- **HostedCluster** secrets from the HostedCluster namespace
- **HostedControlPlane** from the Hosted Control Plane namespace
- **Cluster** from the Hosted Control Plane namespace
- **AWSCluster**, **AWSMachineTemplate**, and **AWSMachine** from the Hosted Control Plane namespace
- **MachineDeployments**, **MachineSets**, and **Machines** from the Hosted Control Plane namespace
- **ControlPlane** secrets from the Hosted Control Plane namespace

a. Enter the following commands:

```
$ mkdir -p ${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS} \
${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-${HC_CLUSTER_NAME}
```

```
$ chmod 700 ${BACKUP_DIR}/namespaces/
```

b. Back up the **HostedCluster** objects from the **HostedCluster** namespace by entering the following commands:

```
$ echo "Backing Up HostedCluster Objects:"
```

```
$ oc get hc ${HC_CLUSTER_NAME} -n ${HC_CLUSTER_NS} -o yaml > \
${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}/hc-
${HC_CLUSTER_NAME}.yaml
```

```
$ echo "--> HostedCluster"
```

```
$ sed -i " -e '/^status:$/,$d' \
${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}/hc-
${HC_CLUSTER_NAME}.yaml
```

c. Back up the **NodePool** objects from the **HostedCluster** namespace by entering the following commands:

```
$ oc get np ${NODEPOOLS} -n ${HC_CLUSTER_NS} -o yaml > \
${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}/np-${NODEPOOLS}.yaml
```

```
$ echo "--> NodePool"
```

```
$ sed -i " -e '/^status:$/,$d' \
${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}/np-${NODEPOOLS}.yaml
```

d. Back up the secrets in the **HostedCluster** namespace by running the following shell script:

```
$ echo "--> HostedCluster Secrets:"
for s in $(oc get secret -n ${HC_CLUSTER_NS} | grep "^${HC_CLUSTER_NAME}" | awk '{print $1}'); do
    oc get secret -n ${HC_CLUSTER_NS} $s -o yaml >
${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}/secret-$s.yaml
done
```

- e. Back up the secrets in the **HostedCluster** control plane namespace by running the following shell script:

```
$ echo "--> HostedCluster ControlPlane Secrets:"
for s in $(oc get secret -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} | egrep -v
"docker|service-account-token|oauth-openshift|NAME|token-
${HC_CLUSTER_NAME}" | awk '{print $1}'); do
    oc get secret -n ${HC_CLUSTER_NS}-$s -o yaml >
${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/secret-$s.yaml
done
```

- f. Back up the hosted control plane by entering the following commands:

```
$ echo "--> HostedControlPlane:"
```



```
$ oc get hcp ${HC_CLUSTER_NAME} -n ${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME} -o yaml > \
${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/hcp-${HC_CLUSTER_NAME}.yaml
```

- g. Back up the cluster by entering the following commands:

```
$ echo "--> Cluster:"
```



```
$ CL_NAME=$(oc get hcp ${HC_CLUSTER_NAME} -n ${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME} \
-o jsonpath={.metadata.labels.\*} | grep ${HC_CLUSTER_NAME})
```



```
$ oc get cluster ${CL_NAME} -n ${HC_CLUSTER_NS}-$s -o yaml > \
${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-$s/${HC_CLUSTER_NAME}/cl-
${HC_CLUSTER_NAME}.yaml
```

- h. Back up the AWS cluster by entering the following commands:

```
$ echo "--> AWS Cluster:"
```



```
$ oc get awscluster ${HC_CLUSTER_NAME} -n ${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME} -o yaml > \
${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/awscl-${HC_CLUSTER_NAME}.yaml
```

- i. Back up the AWS **MachineTemplate** objects by entering the following commands:

```
$ echo "--> AWS Machine Template:"
$ oc get awsmachinetemplate ${NODEPOOLS} -n ${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME} -o yaml > \
${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/awsmt-${HC_CLUSTER_NAME}.yaml
```

- j. Back up the AWS **Machines** objects by running the following shell script:

```
$ echo "--> AWS Machine:"
$ CL_NAME=$(oc get hcp ${HC_CLUSTER_NAME} -n ${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME} -o jsonpath={.metadata.labels.\*} | grep
${HC_CLUSTER_NAME})
for s in $(oc get awsmachines -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} --no-headers | grep ${CL_NAME} | cut -f1 -d\ ); do
    oc get -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} awsmachines $s -o
    yaml > ${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/awsm-$s.yaml
done
```

- k. Back up the **MachineDeployments** objects by running the following shell script:

```
$ echo "--> HostedCluster MachineDeployments:"
for s in $(oc get machinedeployment -n ${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME} -o name); do
    mdp_name=$(echo ${s} | cut -f 2 -d /)
    oc get -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} $s -o yaml >
${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/machinedeployment-${mdp_name}.yaml
done
```

- l. Back up the **MachineSets** objects by running the following shell script:

```
$ echo "--> HostedCluster MachineSets:"
for s in $(oc get machineset -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} -o
name); do
    ms_name=$(echo ${s} | cut -f 2 -d /)
    oc get -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} $s -o yaml >
${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/machineset-${ms_name}.yaml
done
```

- m. Back up the **Machines** objects from the Hosted Control Plane namespace by running the following shell script:

```
$ echo "--> HostedCluster Machine:"
for s in $(oc get machine -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} -o
name); do
    m_name=$(echo ${s} | cut -f 2 -d /)
    oc get -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} $s -o yaml >
```

```

${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/machine-${m_name}.yaml
done

```

- Clean up the **ControlPlane** routes by entering the following command:

```
$ oc delete routes -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} --all
```

By entering that command, you enable the ExternalDNS Operator to delete the Route53 entries.

- Verify that the Route53 entries are clean by running the following script:

```

function clean_routes() {

if [[ -z "${1}" ]];then
    echo "Give me the NS where to clean the routes"
    exit 1
fi

# Constants
if [[ -z "${2}" ]];then
    echo "Give me the Route53 zone ID"
    exit 1
fi

ZONE_ID=${2}
ROUTES=10
timeout=40
count=0

# This allows us to remove the ownership in the AWS for the API route
oc delete route -n ${1} --all

while [ ${ROUTES} -gt 2 ]
do
    echo "Waiting for ExternalDNS Operator to clean the DNS Records in AWS Route53
where the zone id is: ${ZONE_ID}..."
    echo "Try: (${count}/${timeout})"
    sleep 10
    if [[ $count -eq timeout ]];then
        echo "Timeout waiting for cleaning the Route53 DNS records"
        exit 1
    fi
    count=$((count+1))
    ROUTES=$(aws route53 list-resource-record-sets --hosted-zone-id ${ZONE_ID} --max-
items 10000 --output json | grep -c ${EXTERNAL_DNS_DOMAIN})
    done
}

# SAMPLE: clean_routes "<HC ControlPlane Namespace>" "<AWS_ZONE_ID>"
clean_routes "${HC_CLUSTER_NS}-${HC_CLUSTER_NAME}" "${AWS_ZONE_ID}"

```

Verification

Check all of the OpenShift Container Platform objects and the S3 bucket to verify that everything looks as expected.

Next steps

Restore your hosted cluster.

8.6.3. Restoring a hosted cluster

Gather all of the objects that you backed up and restore them in your destination management cluster.

Prerequisites

You backed up the data from your source management cluster.

TIP

Ensure that the **kubeconfig** file of the destination management cluster is placed as it is set in the **KUBECONFIG** variable or, if you use the script, in the **MGMT2_KUBECONFIG** variable. Use **export KUBECONFIG=<Kubeconfig FilePath>** or, if you use the script, use **export KUBECONFIG=\${MGMT2_KUBECONFIG}**.

Procedure

1. Verify that the new management cluster does not contain any namespaces from the cluster that you are restoring by entering these commands:

```
# Just in case
$ export KUBECONFIG=${MGMT2_KUBECONFIG}
$ BACKUP_DIR=${HC_CLUSTER_DIR}/backup

# Namespace deletion in the destination Management cluster
$ oc delete ns ${HC_CLUSTER_NS} || true
$ oc delete ns ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} || true
```

2. Re-create the deleted namespaces by entering these commands:

```
# Namespace creation
$ oc new-project ${HC_CLUSTER_NS}
$ oc new-project ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME}
```

3. Restore the secrets in the HC namespace by entering this command:

```
$ oc apply -f ${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}/secret-*
```

4. Restore the objects in the **HostedCluster** control plane namespace by entering these commands:

```
# Secrets
$ oc apply -f ${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-${HC_CLUSTER_NAME}/secret-*

# Cluster
$ oc apply -f ${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
```

```

${HC_CLUSTER_NAME}/hcp-*
$ oc apply -f ${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/cl-*

```

- If you are recovering the nodes and the node pool to reuse AWS instances, restore the objects in the HC control plane namespace by entering these commands:

```

# AWS
$ oc apply -f ${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/awscl-*
$ oc apply -f ${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/awsmt-*
$ oc apply -f ${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/awsm-*

# Machines
$ oc apply -f ${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/machinedeployment-*
$ oc apply -f ${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/machineset-*
$ oc apply -f ${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME}/machine-*

```

- Restore the etcd data and the hosted cluster by running this bash script:

```

ETCD_PODS="etcd-0"
if [ "${CONTROL_PLANE_AVAILABILITY_POLICY}" = "HighlyAvailable" ]; then
    ETCD_PODS="etcd-0 etcd-1 etcd-2"
fi

HC_RESTORE_FILE=${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}/hc-
${HC_CLUSTER_NAME}-restore.yaml
HC_BACKUP_FILE=${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}/hc-
${HC_CLUSTER_NAME}.yaml
HC_NEW_FILE=${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}/hc-
${HC_CLUSTER_NAME}-new.yaml
cat ${HC_BACKUP_FILE} > ${HC_NEW_FILE}
cat > ${HC_RESTORE_FILE} <<EOF
    restoreSnapshotURL:
EOF

for POD in ${ETCD_PODS}; do
    # Create a pre-signed URL for the etcd snapshot
    ETCD_SNAPSHOT="s3://${BUCKET_NAME}/${HC_CLUSTER_NAME}-${POD}-
snapshot.db"
    ETCD_SNAPSHOT_URL=$(AWS_DEFAULT_REGION=${MGMT2_REGION} aws s3
presign ${ETCD_SNAPSHOT})

    # FIXME no CLI support for restoreSnapshotURL yet
    cat >> ${HC_RESTORE_FILE} <<EOF
        - "${ETCD_SNAPSHOT_URL}"
EOF
done

cat ${HC_RESTORE_FILE}

```

```

if ! grep ${HC_CLUSTER_NAME}-snapshot.db ${HC_NEW_FILE}; then
    sed -i " -e "/type: PersistentVolume/r ${HC_RESTORE_FILE}" ${HC_NEW_FILE}
    sed -i " -e '/pausedUntil:/d' ${HC_NEW_FILE}
fi

HC=$(oc get hc -n ${HC_CLUSTER_NS} ${HC_CLUSTER_NAME} -o name || true)
if [[ ${HC} == "" ]];then
    echo "Deploying HC Cluster: ${HC_CLUSTER_NAME} in ${HC_CLUSTER_NS} namespace"
    oc apply -f ${HC_NEW_FILE}
else
    echo "HC Cluster ${HC_CLUSTER_NAME} already exists, avoiding step"
fi

```

7. If you are recovering the nodes and the node pool to reuse AWS instances, restore the node pool by entering this command:

```
$ oc apply -f ${BACKUP_DIR}/namespaces/${HC_CLUSTER_NS}/np-*
```

Verification

- To verify that the nodes are fully restored, use this function:

```

timeout=40
count=0
NODE_STATUS=$(oc get nodes --kubeconfig=${HC_KUBECONFIG} | grep -v NotReady |
grep -c "worker") || NODE_STATUS=0

while [ ${NODE_POOL_REPLICAS} != ${NODE_STATUS} ]
do
    echo "Waiting for Nodes to be Ready in the destination MGMT Cluster:
${MGMT2_CLUSTER_NAME}"
    echo "Try: (${count}/${timeout})"
    sleep 30
    if [[ $count -eq timeout ]];then
        echo "Timeout waiting for Nodes in the destination MGMT Cluster"
        exit 1
    fi
    count=$((count+1))
    NODE_STATUS=$(oc get nodes --kubeconfig=${HC_KUBECONFIG} | grep -v NotReady |
grep -c "worker") || NODE_STATUS=0
done

```

Next steps

Shut down and delete your cluster.

8.6.4. Deleting a hosted cluster from your source management cluster

After you back up your hosted cluster and restore it to your destination management cluster, you shut down and delete the hosted cluster on your source management cluster.

Prerequisites

You backed up your data and restored it to your source management cluster.

TIP

Ensure that the **kubeconfig** file of the destination management cluster is placed as it is set in the **KUBECONFIG** variable or, if you use the script, in the **MGMT_KUBECONFIG** variable. Use **export KUBECONFIG=<FilePath>** or, if you use the script, use **export KUBECONFIG=\${MGMT_KUBECONFIG}**.

Procedure

1. Scale the **deployment** and **statefulset** objects by entering these commands:



IMPORTANT

Do not scale the stateful set if the value of its **spec.persistentVolumeClaimRetentionPolicy.whenScaled** field is set to **Delete**, because this could lead to a loss of data.

As a workaround, update the value of the **spec.persistentVolumeClaimRetentionPolicy.whenScaled** field to **Retain**. Ensure that no controllers exist that reconcile the stateful set and would return the value back to **Delete**, which could lead to a loss of data.

```
# Just in case
$ export KUBECONFIG=${MGMT_KUBECONFIG}

# Scale down deployments
$ oc scale deployment -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} --replicas=0 --all
$ oc scale statefulset.apps -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} --replicas=0 -all
$ sleep 15
```

2. Delete the **NodePool** objects by entering these commands:

```
NODEPOOLS=$(oc get nodepools -n ${HC_CLUSTER_NS} -o=jsonpath='{.items[?(@.spec.clusterName=="${HC_CLUSTER_NAME}")].metadata.name}')
if [[ ! -z "${NODEPOOLS}" ]];then
    oc patch -n "${HC_CLUSTER_NS}" nodepool ${NODEPOOLS} --type=json --patch='[{"op":"remove", "path": "/metadata/finalizers"}]'
    oc delete np -n ${HC_CLUSTER_NS} ${NODEPOOLS}
fi
```

3. Delete the **machine** and **machineset** objects by entering these commands:

```
# Machines
for m in $(oc get machines -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} -o name); do
    oc patch -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} ${m} --type=json --patch='[{"op":"remove", "path": "/metadata/finalizers"}]' || true
    oc delete -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} ${m} || true
done

$ oc delete machineset -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} --all || true
```

4. Delete the cluster object by entering these commands:

```
# Cluster
$ C_NAME=$(oc get cluster -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} -o name)
$ oc patch -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} ${C_NAME} --type=json --
patch='[ { "op":"remove", "path": "/metadata/finalizers" } ]'
$ oc delete cluster.cluster.x-k8s.io -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} --all
```

5. Delete the AWS machines (Kubernetes objects) by entering these commands. Do not worry about deleting the real AWS machines. The cloud instances will not be affected.

```
# AWS Machines
for m in $(oc get awsmachine.infrastructure.cluster.x-k8s.io -n ${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME} -o name)
do
  oc patch -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} ${m} --type=json --patch='[ {
"op":"remove", "path": "/metadata/finalizers" } ]' || true
  oc delete -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} ${m} || true
done
```

6. Delete the **HostedControlPlane** and **ControlPlane** HC namespace objects by entering these commands:

```
# Delete HCP and ControlPlane HC NS
$ oc patch -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME}
hostedcontrolplane.hypershift.openshift.io ${HC_CLUSTER_NAME} --type=json --patch='[ {
"op":"remove", "path": "/metadata/finalizers" } ]'
$ oc delete hostedcontrolplane.hypershift.openshift.io -n ${HC_CLUSTER_NS}-
${HC_CLUSTER_NAME} --all
$ oc delete ns ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME} || true
```

7. Delete the **HostedCluster** and HC namespace objects by entering these commands:

```
# Delete HC and HC Namespace
$ oc -n ${HC_CLUSTER_NS} patch hostedclusters ${HC_CLUSTER_NAME} -p
'{"metadata":{"finalizers":null}}' --type merge || true
$ oc delete hc -n ${HC_CLUSTER_NS} ${HC_CLUSTER_NAME} || true
$ oc delete ns ${HC_CLUSTER_NS} || true
```

Verification

- To verify that everything works, enter these commands:

```
# Validations
$ export KUBECONFIG=${MGMT2_KUBECONFIG}

$ oc get hc -n ${HC_CLUSTER_NS}
$ oc get np -n ${HC_CLUSTER_NS}
$ oc get pod -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME}
$ oc get machines -n ${HC_CLUSTER_NS}-${HC_CLUSTER_NAME}

# Inside the HostedCluster
```

```
$ export KUBECONFIG=${HC_KUBECONFIG}
$ oc get clusterversion
$ oc get nodes
```

Next steps

Delete the OVN pods in the hosted cluster so that you can connect to the new OVN control plane that runs in the new management cluster:

1. Load the **KUBECONFIG** environment variable with the hosted cluster's kubeconfig path.
2. Enter this command:

```
$ oc delete pod -n openshift-ovn-kubernetes --all
```

8.7. DISASTER RECOVERY FOR A HOSTED CLUSTER BY USING OADP

You can use the OpenShift API for Data Protection (OADP) Operator to perform disaster recovery on Amazon Web Services (AWS) and bare metal.

The disaster recovery process with OpenShift API for Data Protection (OADP) involves the following steps:

1. Preparing your platform, such as Amazon Web Services or bare metal, to use OADP
2. Backing up the data plane workload
3. Backing up the control plane workload
4. Restoring a hosted cluster by using OADP

8.7.1. Prerequisites

You must meet the following prerequisites on the management cluster:

- You [installed the OADP Operator](#).
- You created a storage class.
- You have access to the cluster with **cluster-admin** privileges.
- You have access to the OADP subscription through a catalog source.
- You have access to a cloud storage provider that is compatible with OADP, such as S3, Microsoft Azure, Google Cloud Platform, or MinIO.
- In a disconnected environment, you have access to a self-hosted storage provider, for example [Red Hat OpenShift Data Foundation](#) or [MinIO](#), that is compatible with OADP.
- Your hosted control planes pods are up and running.

8.7.2. Preparing AWS to use OADP

To perform disaster recovery for a hosted cluster, you can use OpenShift API for Data Protection (OADP) on Amazon Web Services (AWS) S3 compatible storage. After creating the

DataProtectionApplication object, new **velero** deployment and **node-agent** pods are created in the **openshift-adp** namespace.

To prepare AWS to use OADP, see "Configuring the OpenShift API for Data Protection with Multicloud Object Gateway".

Additional resources

- [Configuring the OpenShift API for Data Protection with Multicloud Object Gateway](#)

Next steps

- Backing up the data plane workload
- Backing up the control plane workload

8.7.3. Preparing bare metal to use OADP

To perform disaster recovery for a hosted cluster, you can use OpenShift API for Data Protection (OADP) on bare metal. After creating the **DataProtectionApplication** object, new **velero** deployment and **node-agent** pods are created in the **openshift-adp** namespace.

To prepare bare metal to use OADP, see "Configuring the OpenShift API for Data Protection with AWS S3 compatible storage".

Additional resources

- [Configuring the OpenShift API for Data Protection with AWS S3 compatible storage](#)

Next steps

- Backing up the data plane workload
- Backing up the control plane workload

8.7.4. Backing up the data plane workload

If the data plane workload is not important, you can skip this procedure. To back up the data plane workload by using the OADP Operator, see "Backing up applications".

Additional resources

- [Backing up applications](#)

Next steps

- Restoring a hosted cluster by using OADP

8.7.5. Backing up the control plane workload

You can back up the control plane workload by creating the **Backup** custom resource (CR).

To monitor and observe the backup process, see "Observing the backup and restore process".

Procedure

- Pause the reconciliation of the **HostedCluster** resource by running the following command:

```
$ oc --kubeconfig <management_cluster_kubeconfig_file> \
patch hostedcluster -n <hosted_cluster_namespace> <hosted_cluster_name> \
--type json -p '[{"op": "add", "path": "/spec/pausedUntil", "value": "true"}]'
```

- Get the infrastructure ID of your hosted cluster by running the following command:

```
$ oc get hostedcluster -n local-cluster <hosted_cluster_name> -o=jsonpath=".spec.infraID"
```

Note the infrastructure ID to use in the next step.

- Pause the reconciliation of the **cluster.cluster.x-k8s.io** resource by running the following command:

```
$ oc --kubeconfig <management_cluster_kubeconfig_file> \
patch cluster.cluster.x-k8s.io \
-n local-cluster-<hosted_cluster_name> <hosted_cluster_infra_id> \
--type json -p '[{"op": "add", "path": "/spec/paused", "value": true}]'
```

- Pause the reconciliation of the **NodePool** resource by running the following command:

```
$ oc --kubeconfig <management_cluster_kubeconfig_file> \
patch nodepool -n <hosted_cluster_namespace> <node_pool_name> \
--type json -p '[{"op": "add", "path": "/spec/pausedUntil", "value": "true"}]'
```

- Pause the reconciliation of the **AgentCluster** resource by running the following command:

```
$ oc --kubeconfig <management_cluster_kubeconfig_file> \
annotate agentcluster -n <hosted_control_plane_namespace> \
cluster.x-k8s.io/paused=true --all'
```

- Pause the reconciliation of the **AgentMachine** resource by running the following command:

```
$ oc --kubeconfig <management_cluster_kubeconfig_file> \
annotate agentmachine -n <hosted_control_plane_namespace> \
cluster.x-k8s.io/paused=true --all'
```

- Annotate the **HostedCluster** resource to prevent the deletion of the hosted control plane namespace by running the following command:

```
$ oc --kubeconfig <management_cluster_kubeconfig_file> \
annotate hostedcluster -n <hosted_cluster_namespace> <hosted_cluster_name> \
hypershift.openshift.io/skip-delete-hosted-controlplane-namespace=true
```

- Create a YAML file that defines the **Backup** CR:

Example 8.1. Example backup-control-plane.yaml file

```
apiVersion: velero.io/v1
kind: Backup
metadata:
```

```

name: <backup_resource_name> ①
namespace: openshift-adp
labels:
  velero.io/storage-location: default
spec:
  hooks: {}
  includedNamespaces: ②
    - <hosted_cluster_namespace> ③
    - <hosted_control_plane_namespace> ④
  includedResources:
    - sa
    - role
    - rolebinding
    - pod
    - pvc
    - pv
    - bmh
    - configmap
    - infraenv ⑤
    - priorityclasses
    - pdb
    - agents
    - hostedcluster
    - nodepool
    - secrets
    - hostedcontrolplane
    - cluster
    - agentcluster
    - agentmachinetemplate
    - agentmachine
    - machinedeployment
    - machineset
    - machine
  excludedResources: []
  storageLocation: default
  ttl: 2h0m0s
  snapshotMoveData: true ⑥
  datamover: "velero" ⑦
  defaultVolumesToFsBackup: true ⑧

```

- ① Replace **backup_resource_name** with the name of your **Backup** resource.
- ② Selects specific namespaces to back up objects from them. You must include your hosted cluster namespace and the hosted control plane namespace.
- ③ Replace **<hosted_cluster_namespace>** with the name of the hosted cluster namespace, for example, **clusters**.
- ④ Replace **<hosted_control_plane_namespace>** with the name of the hosted control plane namespace, for example, **clusters-hosted**.
- ⑤ You must create the **infraenv** resource in a separate namespace. Do not delete the **infraenv** resource during the backup process.

- 6** **7** Enables the CSI volume snapshots and uploads the control plane workload automatically to the cloud storage.
- 8** Sets the **fs-backup** backing up method for persistent volumes (PVs) as default. This setting is useful when you use a combination of Container Storage Interface (CSI) volume snapshots and the **fs-backup** method.

**NOTE**

If you want to use CSI volume snapshots, you must add the **backup.velero.io/backup-volumes-excludes=<pv_name>** annotation to your PVs.

9. Apply the **Backup** CR by running the following command:

```
$ oc apply -f backup-control-plane.yaml
```

Verification

- Verify if the value of the **status.phase** is **Completed** by running the following command:

```
$ oc get backups.velero.io <backup_resource_name> -n openshift-adp \
-o jsonpath='{.status.phase}'
```

Next steps

- Restoring a hosted cluster by using OADP

8.7.6. Restoring a hosted cluster by using OADP

You can restore the hosted cluster by creating the **Restore** custom resource (CR).

- If you are using an *in-place* update, InfraEnv does not need spare nodes. You need to re-provision the worker nodes from the new management cluster.
- If you are using a *replace* update, you need some spare nodes for InfraEnv to deploy the worker nodes.

**IMPORTANT**

After you back up your hosted cluster, you must destroy it to initiate the restoring process. To initiate node provisioning, you must back up workloads in the data plane before deleting the hosted cluster.

Prerequisites

- You completed the steps in [Removing a cluster by using the console](#) to delete your hosted cluster.
- You completed the steps in [Removing remaining resources after removing a cluster](#).

To monitor and observe the backup process, see "Observing the backup and restore process".

Procedure

- Verify that no pods and persistent volume claims (PVCs) are present in the hosted control plane namespace by running the following command:

```
$ oc get pod pvc -n <hosted_control_plane_namespace>
```

Expected output

```
No resources found
```

- Create a YAML file that defines the **Restore** CR:

Example `restore-hosted-cluster.yaml` file

```
apiVersion: velero.io/v1
kind: Restore
metadata:
  name: <restore_resource_name> ①
  namespace: openshift-adp
spec:
  backupName: <backup_resource_name> ②
  restorePVs: true ③
  existingResourcePolicy: update ④
  excludedResources:
    - nodes
    - events
    - events.events.k8s.io
    - backups.velero.io
    - restores.velero.io
    - resticrepositories.velero.io
```

- ① Replace `<restore_resource_name>` with the name of your **Restore** resource.
- ② Replace `<backup_resource_name>` with the name of your **Backup** resource.
- ③ Initiates the recovery of persistent volumes (PVs) and its pods.
- ④ Ensures that the existing objects are overwritten with the backed up content.



IMPORTANT

You must create the **infraenv** resource in a separate namespace. Do not delete the **infraenv** resource during the restore process. The **infraenv** resource is mandatory for the new nodes to be reprovisioned.

- Apply the **Restore** CR by running the following command:

```
$ oc apply -f restore-hosted-cluster.yaml
```

- Verify if the value of the **status.phase** is **Completed** by running the following command:

```
$ oc get hostedcluster <hosted_cluster_name> -n <hosted_cluster_namespace> \
-o jsonpath='{.status.phase}'
```

5. After the restore process is complete, start the reconciliation of the **HostedCluster** and **NodePool** resources that you paused during backing up of the control plane workload:
 - a. Start the reconciliation of the **HostedCluster** resource by running the following command:

```
$ oc --kubeconfig <management_cluster_kubeconfig_file> \
patch hostedcluster -n <hosted_cluster_namespace> <hosted_cluster_name> \
--type json \
-p '[{"op": "add", "path": "/spec/pausedUntil", "value": "false"}]'
```

- b. Start the reconciliation of the **NodePool** resource by running the following command:

```
$ oc --kubeconfig <management_cluster_kubeconfig_file> \
patch nodepool -n <hosted_cluster_namespace> <node_pool_name> \
--type json \
-p '[{"op": "add", "path": "/spec/pausedUntil", "value": "false"}]'
```

6. Start the reconciliation of the Agent provider resources that you paused during backing up of the control plane workload:

- a. Start the reconciliation of the **AgentCluster** resource by running the following command:

```
$ oc --kubeconfig <management_cluster_kubeconfig_file> \
annotate agentcluster -n <hosted_control_plane_namespace> \
cluster.x-k8s.io/paused- --overwrite=true --all
```

- b. Start the reconciliation of the **AgentMachine** resource by running the following command:

```
$ oc --kubeconfig <management_cluster_kubeconfig_file> \
annotate agentmachine -n <hosted_control_plane_namespace> \
cluster.x-k8s.io/paused- --overwrite=true --all
```

7. Remove the **hypershift.openshift.io/skip-delete-hosted-controlplane-namespace-** annotation in the **HostedCluster** resource to avoid manually deleting the hosted control plane namespace by running the following command:

```
$ oc --kubeconfig <management_cluster_kubeconfig_file> \
annotate hostedcluster -n <hosted_cluster_namespace> <hosted_cluster_name> \
hypershift.openshift.io/skip-delete-hosted-controlplane-namespace- \
--overwrite=true --all
```

8. Scale the **NodePool** resource to the desired number of replicas by running the following command:

```
$ oc --kubeconfig <management_cluster_kubeconfig_file> \
scale nodepool -n <hosted_cluster_namespace> <node_pool_name> \
--replicas <replica_count> 1
```

1 Replace **<replica_count>** by an integer value, for example, **3**.

8.7.7. Observing the backup and restore process

When using OpenShift API for Data Protection (OADP) to backup and restore a hosted cluster, you can monitor and observe the process.

Procedure

1. Observe the backup process by running the following command:

```
$ watch "oc get backups.velero.io -n openshift-adp <backup_resource_name> -o jsonpath='{.status}'"
```

2. Observe the restore process by running the following command:

```
$ watch "oc get restores.velero.io -n openshift-adp <backup_resource_name> -o jsonpath='{.status}'"
```

3. Observe the Velero logs by running the following command:

```
$ oc logs -n openshift-adp -ldeploy=velero -f
```

4. Observe the progress of all of the OADP objects by running the following command:

```
$ watch "echo BackupRepositories:;echo;oc get backuprepositories.velero.io -A;echo; echo BackupStorageLocations: ;echo; oc get backupstoragelocations.velero.io -A;echo;echo DataUploads: ;echo;oc get datauploads.velero.io -A;echo;echo DataDownloads: ;echo;oc get datadownloads.velero.io -n openshift-adp; echo;echo VolumeSnapshotLocations: ;echo;oc get volumesnapshotlocations.velero.io -A;echo;echo Backups:;echo;oc get backup -A; echo;echo Restores:;echo;oc get restore -A"
```

8.7.8. Using the velero CLI to describe the Backup and Restore resources

When using OpenShift API for Data Protection, you can get more details of the **Backup** and **Restore** resources by using the **velero** command-line interface (CLI).

Procedure

1. Create an alias to use the **velero** CLI from a container by running the following command:

```
$ alias velero='oc -n openshift-adp exec deployment/velero -c velero -it -- ./velero'
```

2. Get details of your **Restore** custom resource (CR) by running the following command:

```
$ velero restore describe <restore_resource_name> --details ①
```

① Replace **<restore_resource_name>** with the name of your **Restore** resource.

3. Get details of your **Backup** CR by running the following command:

```
$ velero restore describe <backup_resource_name> --details ①
```

- 1 Replace <backup_resource_name> with the name of your **Backup** resource.

CHAPTER 9. AUTHENTICATION AND AUTHORIZATION FOR HOSTED CONTROL PLANES

The OpenShift Container Platform control plane includes a built-in OAuth server. You can obtain OAuth access tokens to authenticate to the OpenShift Container Platform API. After you create your hosted cluster, you can configure OAuth by specifying an identity provider.

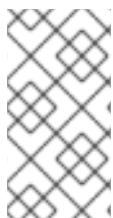
9.1. CONFIGURING THE OAUTH SERVER FOR A HOSTED CLUSTER BY USING THE CLI

You can configure the internal OAuth server for your hosted cluster by using the CLI.

You can configure OAuth for the following supported identity providers:

- **oidc**
- **htpasswd**
- **keystone**
- **ldap**
- **basic-authentication**
- **request-header**
- **github**
- **gitlab**
- **google**

Adding any identity provider in the OAuth configuration removes the default **kubeadmin** user provider.



NOTE

When you configure identity providers, you must configure at least one **NodePool** replica in your hosted cluster in advance. Traffic for DNS resolution is sent through the worker nodes. You do not need to configure the **NodePool** replicas in advance for the **htpasswd** and **request-header** identity providers.

Prerequisites

- You created your hosted cluster.

Procedure

1. Edit the **HostedCluster** custom resource (CR) on the hosting cluster by running the following command:

```
$ oc edit <hosted_cluster_name> -n <hosted_cluster_namespace>
```

2. Add the OAuth configuration in the **HostedCluster** CR by using the following example:

```

apiVersion: hypershift.openshift.io/v1alpha1
kind: HostedCluster
metadata:
  name: <hosted_cluster_name> ①
  namespace: <hosted_cluster_namespace> ②
spec:
  configuration:
    oauth:
      identityProviders:
        - openID: ③
          claims:
            email: ④
            - <email_address>
          name: ⑤
          - <display_name>
          preferredUsername: ⑥
          - <preferred_username>
        clientID: <client_id> ⑦
        clientSecret:
          name: <client_id_secret_name> ⑧
        issuer: https://example.com/identity ⑨
      mappingMethod: lookup ⑩
      name: IAM
      type: OpenID

```

- ① Specifies your hosted cluster name.
- ② Specifies your hosted cluster namespace.
- ③ This provider name is prefixed to the value of the identity claim to form an identity name. The provider name is also used to build the redirect URL.
- ④ Defines a list of attributes to use as the email address.
- ⑤ Defines a list of attributes to use as a display name.
- ⑥ Defines a list of attributes to use as a preferred user name.
- ⑦ Defines the ID of a client registered with the OpenID provider. You must allow the client to redirect to the **https://oauth-openshift.apps.<cluster_name><cluster_domain>/oauth2callback/<idp_provider_name>** URL.
- ⑧ Defines a secret of a client registered with the OpenID provider.
- ⑨ The **Issuer Identifier** described in the OpenID spec. You must use **https** without query or fragment component.
- ⑩ Defines a mapping method that controls how mappings are established between identities of this provider and **User** objects.

3. Save the file to apply the changes.

9.2. CONFIGURING THE OAUTH SERVER FOR A HOSTED CLUSTER BY USING THE WEB CONSOLE

You can configure the internal OAuth server for your hosted cluster by using the OpenShift Container Platform web console.

You can configure OAuth for the following supported identity providers:

- **oidc**
- **htpasswd**
- **keystone**
- **ldap**
- **basic-authentication**
- **request-header**
- **github**
- **gitlab**
- **google**

Adding any identity provider in the OAuth configuration removes the default **kubeadmin** user provider.



NOTE

When you configure identity providers, you must configure at least one **NodePool** replica in your hosted cluster in advance. Traffic for DNS resolution is sent through the worker nodes. You do not need to configure the **NodePool** replicas in advance for the **htpasswd** and **request-header** identity providers.

Prerequisites

- You logged in as a user with **cluster-admin** privileges.
- You created your hosted cluster.

Procedure

1. Navigate to **Home** → **API Explorer**.
2. Use the **Filter by kind** box to search for your **HostedCluster** resource.
3. Click the **HostedCluster** resource that you want to edit.
4. Click the **Instances** tab.
5. Click the Options menu  next to your hosted cluster name entry and click **Edit HostedCluster**.
6. Add the OAuth configuration in the YAML file:

```

spec:
  configuration:
    oauth:
      identityProviders:
        - openID: ①
          claims:
            email: ②
            - <email_address>
          name: ③
            - <display_name>
          preferredUsername: ④
            - <preferred_username>
          clientId: <client_id> ⑤
          clientSecret:
            name: <client_id_secret_name> ⑥
          issuer: https://example.com/identity ⑦
        mappingMethod: lookup ⑧
        name: IAM
        type: OpenID

```

- ① This provider name is prefixed to the value of the identity claim to form an identity name. The provider name is also used to build the redirect URL.
- ② Defines a list of attributes to use as the email address.
- ③ Defines a list of attributes to use as a display name.
- ④ Defines a list of attributes to use as a preferred user name.
- ⑤ Defines the ID of a client registered with the OpenID provider. You must allow the client to redirect to the https://oauth-openshift.apps.<cluster_name>. <cluster_domain>/oauth2callback/<idp_provider_name> URL.
- ⑥ Defines a secret of a client registered with the OpenID provider.
- ⑦ The [Issuer Identifier](#) described in the OpenID spec. You must use **https** without query or fragment component.
- ⑧ Defines a mapping method that controls how mappings are established between identities of this provider and **User** objects.

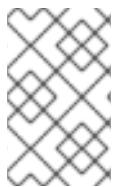
7. Click **Save**.

Additional resources

- To know more about supported identity providers, see "[Understanding identity provider configuration](#)" in *Authentication and authorization*.

9.3. ASSIGNING COMPONENTS IAM ROLES BY USING THE CCO IN A HOSTED CLUSTER ON AWS

You can assign components IAM roles that provide short-term, limited-privilege security credentials by using the Cloud Credential Operator (CCO) in hosted clusters on Amazon Web Services (AWS). By default, the CCO runs in a hosted control plane.



NOTE

The CCO supports a manual mode only for hosted clusters on AWS. By default, hosted clusters are configured in a manual mode. The management cluster might use modes other than manual.

9.4. VERIFYING THE CCO INSTALLATION IN A HOSTED CLUSTER ON AWS

You can verify that the Cloud Credential Operator (CCO) is running correctly in your hosted control plane.

Prerequisites

- You configured the hosted cluster on Amazon Web Services (AWS).

Procedure

1. Verify that the CCO is configured in a manual mode in your hosted cluster by running the following command:

```
$ oc get cloudcredentials <hosted_cluster_name> \  
-n <hosted_cluster_namespace> \  
-o=jsonpath={.spec.credentialsMode}
```

Expected output

```
Manual
```

2. Verify that the value for the **serviceAccountIssuer** resource is not empty by running the following command:

```
$ oc get authentication cluster --kubeconfig <hosted_cluster_name>.kubeconfig \  
-o jsonpath --template '{.spec.serviceAccountIssuer}'
```

Example output

```
https://aos-hypershift-ci-oidc-29999.s3.us-east-2.amazonaws.com/hypershift-ci-29999
```

9.5. ENABLING OPERATORS TO SUPPORT CCO-BASED WORKFLOWS WITH AWS STS

As an Operator author designing your project to run on Operator Lifecycle Manager (OLM), you can enable your Operator to authenticate against AWS on STS-enabled OpenShift Container Platform clusters by customizing your project to support the Cloud Credential Operator (CCO).

With this method, the Operator is responsible for and requires RBAC permissions for creating the **CredentialsRequest** object and reading the resulting **Secret** object.



NOTE

By default, pods related to the Operator deployment mount a **serviceAccountToken** volume so that the service account token can be referenced in the resulting **Secret** object.

Prerequisites

- OpenShift Container Platform 4.14 or later
- Cluster in STS mode
- OLM-based Operator project

Procedure

1. Update your Operator project's **ClusterServiceVersion** (CSV) object:

- a. Ensure your Operator has RBAC permission to create **CredentialsRequests** objects:

Example 9.1. Example clusterPermissions list

```
# ...
install:
spec:
  clusterPermissions:
    - rules:
        - apiGroups:
            - "cloudcredential.openshift.io"
        resources:
          - credentialsrequests
        verbs:
          - create
          - delete
          - get
          - list
          - patch
          - update
          - watch
```

- b. Add the following annotation to claim support for this method of CCO-based workflow with AWS STS:

```
# ...
metadata:
  annotations:
    features.operators.openshift.io/token-auth-aws: "true"
```

2. Update your Operator project code:

- a. Get the role ARN from the environment variable set on the pod by the **Subscription** object. For example:

```
// Get ENV var
roleARN := os.Getenv("ROLEARN")
setupLog.Info("getting role ARN", "role ARN = ", roleARN)
webIdentityTokenPath := "/var/run/secrets/openshift/serviceaccount/token"
```

- b. Ensure you have a **CredentialsRequest** object ready to be patched and applied. For example:

Example 9.2. Example **CredentialsRequest** object creation

```
import (
    minterv1 "github.com/openshift/cloud-credential-operator/pkg/apis/cloudcredential/v1"
    corev1 "k8s.io/api/core/v1"
    metav1 "k8s.io/apimachinery/pkg/apis/meta/v1"
)

var in = minterv1.AWSProviderSpec{
    StatementEntries: []minterv1.StatementEntry{
        {
            Action: []string{
                "s3:*",
            },
            Effect: "Allow",
            Resource: "arn:aws:s3:::*",
        },
    },
    STSIAMRoleARN: "<role_arn>",
}

var codec = minterv1.Codec
var ProviderSpec, _ = codec.EncodeProviderSpec(in.DeepCopyObject())

const (
    name     = "<credential_request_name>"
    namespace = "<namespace_name>"
)

var CredentialsRequestTemplate = &minterv1.CredentialsRequest{
    ObjectMeta: metav1.ObjectMeta{
        Name:     name,
        Namespace: "openshift-cloud-credential-operator",
    },
    Spec: minterv1.CredentialsRequestSpec{
        ProviderSpec: ProviderSpec,
        SecretRef: corev1.ObjectReference{
            Name:     "<secret_name>",
            Namespace: namespace,
        },
        ServiceAccountNames: []string{
            "<service_account_name>",
        },
    },
}
```

```

    CloudTokenPath: "",  

  },  

}

```

Alternatively, if you are starting from a **CredentialsRequest** object in YAML form (for example, as part of your Operator project code), you can handle it differently:

Example 9.3. Example **CredentialsRequest** object creation in YAML form

```

// CredentialsRequest is a struct that represents a request for credentials  

type CredentialsRequest struct {  

  APIVersion string `yaml:"apiVersion"  

  Kind string `yaml:"kind"  

  Metadata struct {  

    Name string `yaml:"name"  

    Namespace string `yaml:"namespace"  

  } `yaml:"metadata"  

  Spec struct {  

    SecretRef struct {  

      Name string `yaml:"name"  

      Namespace string `yaml:"namespace"  

    } `yaml:"secretRef"  

    ProviderSpec struct {  

      APIVersion string `yaml:"apiVersion"  

      Kind string `yaml:"kind"  

      StatementEntries []struct {  

        Effect string `yaml:"effect"  

        Action []string `yaml:"action"  

        Resource string `yaml:"resource"  

      } `yaml:"statementEntries"  

      STSIAMRoleARN string `yaml:"stsIAMRoleARN"  

    } `yaml:"providerSpec"  

// added new field  

    CloudTokenPath string `yaml:"cloudTokenPath"  

  } `yaml:"spec"  

}  

// ConsumeCredsRequestAddingTokenInfo is a function that takes a YAML filename  

and two strings as arguments  

// It unmarshals the YAML file to a CredentialsRequest object and adds the token  

information.  

func ConsumeCredsRequestAddingTokenInfo(fileName, tokenString, tokenPath  

string) (*CredentialsRequest, error) {  

  // open a file containing YAML form of a CredentialsRequest  

  file, err := os.Open(fileName)  

  if err != nil {  

    return nil, err  

  }  

  defer file.Close()  

  // create a new CredentialsRequest object  

  cr := &CredentialsRequest{}
```

```
// decode the yaml file to the object
decoder := yaml.NewDecoder(file)
err = decoder.Decode(cr)
if err != nil {
    return nil, err
}

// assign the string to the existing field in the object
cr.Spec.CloudTokenPath = tokenPath

// return the modified object
return cr, nil
}
```



NOTE

Adding a **CredentialsRequest** object to the Operator bundle is not currently supported.

- c. Add the role ARN and web identity token path to the credentials request and apply it during Operator initialization:

Example 9.4. Example applying **CredentialsRequest** object during Operator initialization

```
// apply CredentialsRequest on install
credReq := credreq.CredentialsRequestTemplate
credReq.Spec.CloudTokenPath = webIdentityTokenPath

c := mgr.GetClient()
if err := c.Create(context.TODO(), credReq); err != nil {
    if !errors.Exists(err) {
        setupLog.Error(err, "unable to create CredRequest")
        os.Exit(1)
    }
}
```

- d. Ensure your Operator can wait for a **Secret** object to show up from the CCO, as shown in the following example, which is called along with the other items you are reconciling in your Operator:

Example 9.5. Example wait for**Secret** object

```
// WaitForSecret is a function that takes a Kubernetes client, a namespace, and a v1
// "k8s.io/api/core/v1" name as arguments
// It waits until the secret object with the given name exists in the given namespace
// It returns the secret object or an error if the timeout is exceeded
func WaitForSecret(client kubernetes.Interface, namespace, name string)
(*v1.Secret, error) {
    // set a timeout of 10 minutes
    timeout := time.After(10 * time.Minute) ①

    // set a polling interval of 10 seconds
    ...
```

```

ticker := time.NewTicker(10 * time.Second)

// loop until the timeout or the secret is found
for {
    select {
    case <-timeout:
        // timeout is exceeded, return an error
        return nil, fmt.Errorf("timed out waiting for secret %s in namespace %s", name,
namespace)
        // add to this error with a pointer to instructions for following a manual path to a
Secret that will work on STS
    case <-ticker.C:
        // polling interval is reached, try to get the secret
        secret, err := client.CoreV1().Secrets(namespace).Get(context.Background(),
name, metav1.GetOptions{})
        if err != nil {
            if errors.NotFound(err) {
                // secret does not exist yet, continue waiting
                continue
            } else {
                // some other error occurred, return it
                return nil, err
            }
        } else {
            // secret is found, return it
            return secret, nil
        }
    }
}
}

```

1

The **timeout** value is based on an estimate of how fast the CCO might detect an added **CredentialsRequest** object and generate a **Secret** object. You might consider lowering the time or creating custom feedback for cluster administrators that could be wondering why the Operator is not yet accessing the cloud resources.

- Set up the AWS configuration by reading the secret created by the CCO from the credentials request and creating the AWS config file containing the data from that secret:

Example 9.6. Example AWS configuration creation

```

func SharedCredentialsFileFromSecret(secret *corev1.Secret) (string, error) {
    var data []byte
    switch {
    case len(secret.Data["credentials"]) > 0:
        data = secret.Data["credentials"]
    default:
        return "", errors.New("invalid secret for aws credentials")
    }

    f, err := ioutil.TempFile("", "aws-shared-credentials")
    if err != nil {
        return "", errors.Wrap(err, "failed to create file for shared credentials")
    }
}

```

```
    }
    defer f.Close()
    if _, err := f.Write(data); err != nil {
        return "", errors.Wrapf(err, "failed to write credentials to %s", f.Name())
    }
    return f.Name(), nil
}
```



IMPORTANT

The secret is assumed to exist, but your Operator code should wait and retry when using this secret to give time to the CCO to create the secret.

Additionally, the wait period should eventually time out and warn users that the OpenShift Container Platform cluster version, and therefore the CCO, might be an earlier version that does not support the **CredentialsRequest** object workflow with STS detection. In such cases, instruct users that they must add a secret by using another method.

- f. Configure the AWS SDK session, for example:

Example 9.7. Example AWS SDK session configuration

```
sharedCredentialsFile, err := SharedCredentialsFileFromSecret(secret)
if err != nil {
    // handle error
}
options := session.Options{
    SharedConfigState: session.SharedConfigEnable,
    SharedConfigFiles: []string{sharedCredentialsFile},
}
```

Additional resources

- [Cluster Operators reference page for the Cloud Credential Operator](#)

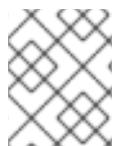
CHAPTER 10. HANDLING MACHINE CONFIGURATION FOR HOSTED CONTROL PLANES

In a standalone OpenShift Container Platform cluster, a machine config pool manages a set of nodes. You can handle a machine configuration by using the **MachineConfigPool** custom resource (CR).

TIP

You can reference any **machineconfiguration.openshift.io** resources in the **nodepool.spec.config** field of the **NodePool** CR.

In hosted control planes, the **MachineConfigPool** CR does not exist. A node pool contains a set of compute nodes. You can handle a machine configuration by using node pools.



NOTE

In OpenShift Container Platform 4.18 or later, the default container runtime for worker nodes is changed from runC to crun.

10.1. CONFIGURING NODE POOLS FOR HOSTED CONTROL PLANES

On hosted control planes, you can configure node pools by creating a **MachineConfig** object inside of a config map in the management cluster.

Procedure

- To create a **MachineConfig** object inside of a config map in the management cluster, enter the following information:

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: <configmap_name>
  namespace: clusters
data:
  config: |
    apiVersion: machineconfiguration.openshift.io/v1
    kind: MachineConfig
    metadata:
      labels:
        machineconfiguration.openshift.io/role: worker
      name: <machineconfig_name>
    spec:
      config:
      ignition:
        version: 3.2.0
      storage:
        files:
          - contents:
              source: data:...
              mode: 420
              overwrite: true
            path: ${PATH} ①
```

- 1 Sets the path on the node where the **MachineConfig** object is stored.
2. After you add the object to the config map, you can apply the config map to the node pool as follows:

```
$ oc edit nodepool <nodepool_name> --namespace <hosted_cluster_namespace>

apiVersion: hypershift.openshift.io/v1alpha1
kind: NodePool
metadata:
# ...
name: nodepool-1
namespace: clusters
# ...
spec:
config:
- name: <configmap_name> 1
# ...
```

- 1 Replace **<configmap_name>** with the name of your config map.

10.2. REFERENCING THE KUBELET CONFIGURATION IN NODE POOLS

To reference your kubelet configuration in node pools, you add the kubelet configuration in a config map and then apply the config map in the **NodePool** resource.

Procedure

1. Add the kubelet configuration inside of a config map in the management cluster by entering the following information:

Example ConfigMap object with the kubelet configuration

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: <configmap_name> 1
  namespace: clusters
data:
  config: |
    apiVersion: machineconfiguration.openshift.io/v1
    kind: KubeletConfig
    metadata:
      name: <kubeletconfig_name> 2
    spec:
      kubeletConfig:
        registerWithTaints:
        - key: "example.sh/unregistered"
          value: "true"
          effect: "NoExecute"
```

- 1 Replace <configmap_name> with the name of your config map.
 - 2 Replace <kubeletconfig_name> with the name of the **KubeletConfig** resource.
2. Apply the config map to the node pool by entering the following command:

```
$ oc edit nodepool <nodepool_name> --namespace clusters 1
```

- 1 Replace <nodepool_name> with the name of your node pool.

Example NodePool resource configuration

```
apiVersion: hypershift.openshift.io/v1alpha1
kind: NodePool
metadata:
# ...
  name: nodepool-1
  namespace: clusters
# ...
spec:
  config:
    - name: <configmap_name> 1
# ...
```

- 1 Replace <configmap_name> with the name of your config map.

10.3. CONFIGURING NODE TUNING IN A HOSTED CLUSTER

To set node-level tuning on the nodes in your hosted cluster, you can use the Node Tuning Operator. In hosted control planes, you can configure node tuning by creating config maps that contain **Tuned** objects and referencing those config maps in your node pools.

Procedure

- 1 Create a config map that contains a valid tuned manifest, and reference the manifest in a node pool. In the following example, a **Tuned** manifest defines a profile that sets **vm.dirty_ratio** to 55 on nodes that contain the **tuned-1-node-label** node label with any value. Save the following **ConfigMap** manifest in a file named **tuned-1.yaml**:

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: tuned-1
  namespace: clusters
data:
  tuning: |
    apiVersion: tuned.openshift.io/v1
    kind: Tuned
    metadata:
      name: tuned-1
      namespace: openshift-cluster-node-tuning-operator
```

```

spec:
  profile:
    - data: |
        [main]
        summary=Custom OpenShift profile
        include=openshift-node
        [sysctl]
        vm.dirty_ratio="55"
      name: tuned-1-profile
  recommend:
    - priority: 20
      profile: tuned-1-profile

```



NOTE

If you do not add any labels to an entry in the **spec.recommend** section of the Tuned spec, node-pool-based matching is assumed, so the highest priority profile in the **spec.recommend** section is applied to nodes in the pool. Although you can achieve more fine-grained node-label-based matching by setting a label value in the Tuned **.spec.recommend.match** section, node labels will not persist during an upgrade unless you set the **.spec.management.upgradeType** value of the node pool to **InPlace**.

2. Create the **ConfigMap** object in the management cluster:

```
$ oc --kubeconfig="$MGMT_KUBECONFIG" create -f tuned-1.yaml
```

3. Reference the **ConfigMap** object in the **spec.tuningConfig** field of the node pool, either by editing a node pool or creating one. In this example, assume that you have only one **NodePool**, named **nodepool-1**, which contains 2 nodes.

```

apiVersion: hypershift.openshift.io/v1alpha1
kind: NodePool
metadata:
  ...
  name: nodepool-1
  namespace: clusters
  ...
spec:
  ...
  tuningConfig:
    - name: tuned-1
status:
  ...

```



NOTE

You can reference the same config map in multiple node pools. In hosted control planes, the Node Tuning Operator appends a hash of the node pool name and namespace to the name of the Tuned CRs to distinguish them. Outside of this case, do not create multiple TuneD profiles of the same name in different Tuned CRs for the same hosted cluster.

Verification

Now that you have created the **ConfigMap** object that contains a **Tuned** manifest and referenced it in a **NodePool**, the Node Tuning Operator syncs the **Tuned** objects into the hosted cluster. You can verify which **Tuned** objects are defined and which TuneD profiles are applied to each node.

1. List the **Tuned** objects in the hosted cluster:

```
$ oc --kubeconfig="$HC_KUBECONFIG" get tuned.tuned.openshift.io \
-n openshift-cluster-node-tuning-operator
```

Example output

NAME	AGE
default	7m36s
rendered	7m36s
tuned-1	65s

2. List the **Profile** objects in the hosted cluster:

```
$ oc --kubeconfig="$HC_KUBECONFIG" get profile.tuned.openshift.io \
-n openshift-cluster-node-tuning-operator
```

Example output

NAME	TUNED	APPLIED	DEGRADED	AGE
nodepool-1-worker-1	tuned-1-profile	True	False	7m43s
nodepool-1-worker-2	tuned-1-profile	True	False	7m14s



NOTE

If no custom profiles are created, the **openshift-node** profile is applied by default.

3. To confirm that the tuning was applied correctly, start a debug shell on a node and check the sysctl values:

```
$ oc --kubeconfig="$HC_KUBECONFIG" \
debug node/nodepool-1-worker-1 -- chroot /host sysctl vm.dirty_ratio
```

Example output

```
vm.dirty_ratio = 55
```

10.4. DEPLOYING THE SR-IOV OPERATOR FOR HOSTED CONTROL PLANES

After you configure and deploy your hosting service cluster, you can create a subscription to the SR-IOV Operator on a hosted cluster. The SR-IOV pod runs on worker machines rather than the control plane.

Prerequisites

You must configure and deploy the hosted cluster on AWS.

Procedure

1. Create a namespace and an Operator group:

```
apiVersion: v1
kind: Namespace
metadata:
  name: openshift-sriov-network-operator
---
apiVersion: operators.coreos.com/v1
kind: OperatorGroup
metadata:
  name: sriov-network-operators
  namespace: openshift-sriov-network-operator
spec:
  targetNamespaces:
    - openshift-sriov-network-operator
```

2. Create a subscription to the SR-IOV Operator:

```
apiVersion: operators.coreos.com/v1alpha1
kind: Subscription
metadata:
  name: sriov-network-operator-subscription
  namespace: openshift-sriov-network-operator
spec:
  channel: stable
  name: sriov-network-operator
  config:
    nodeSelector:
      node-role.kubernetes.io/worker: ""
  source: redhat-operators
  sourceNamespace: openshift-marketplace
```

Verification

1. To verify that the SR-IOV Operator is ready, run the following command and view the resulting output:

```
$ oc get csv -n openshift-sriov-network-operator
```

Example output

NAME	DISPLAY	VERSION	REPLACES
PHASE			
sriov-network-operator.4.18.0-202211021237	SR-IOV Network Operator	4.18.0-	
202211021237	sriov-network-operator.4.18.0-202210290517	Succeeded	

2. To verify that the SR-IOV pods are deployed, run the following command:

```
$ oc get pods -n openshift-sriov-network-operator
```

10.5. CONFIGURING THE NTP SERVER FOR HOSTED CLUSTERS

You can configure the Network Time Protocol (NTP) server for your hosted clusters by using Butane.

Procedure

- 1 Create a Butane config file, **99-worker-chrony.bu**, that includes the contents of the **chrony.conf** file. For more information about Butane, see "Creating machine configs with Butane".

Example 99-worker-chrony.bu configuration

```
# ...
variant: openshift
version: 4.18.0
metadata:
  name: 99-worker-chrony
  labels:
    machineconfiguration.openshift.io/role: worker
storage:
  files:
    - path: /etc/chrony.conf
      mode: 0644 ①
      overwrite: true
    contents:
      inline: |
        pool 0.rhel.pool.ntp.org iburst ②
        driftfile /var/lib/chrony/drift
        makestep 1.0 3
        rtcsync
        logdir /var/log/chrony
# ...
```

- ① Specify an octal value mode for the **mode** field in the machine config file. After creating the file and applying the changes, the **mode** field is converted to a decimal value.
- ② Specify any valid, reachable time source, such as the one provided by your Dynamic Host Configuration Protocol (DHCP) server.



NOTE

For machine-to-machine communication, the NTP on the User Datagram Protocol (UDP) port is **123**. If you configured an external NTP time server, you must open UDP port **123**.

- 2 Use Butane to generate a **MachineConfig** object file, **99-worker-chrony.yaml**, that contains a configuration that Butane sends to the nodes. Run the following command:

```
$ butane 99-worker-chrony.bu -o 99-worker-chrony.yaml
```

Example 99-worker-chrony.yaml configuration

```
# Generated by Butane; do not edit
apiVersion: machineconfiguration.openshift.io/v1
kind: MachineConfig
metadata:
  labels:
    machineconfiguration.openshift.io/role: worker
  name: <machineconfig_name>
spec:
  config:
    ignition:
      version: 3.2.0
  storage:
    files:
      - contents:
          source: data:...
          mode: 420
          overwrite: true
          path: /example/path
```

- Add the contents of the **99-worker-chrony.yaml** file inside of a config map in the management cluster:

Example config map

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: <configmap_name>
  namespace: <namespace> ①
data:
  config: |
    apiVersion: machineconfiguration.openshift.io/v1
    kind: MachineConfig
    metadata:
      labels:
        machineconfiguration.openshift.io/role: worker
      name: <machineconfig_name>
    spec:
      config:
        ignition:
          version: 3.2.0
      storage:
        files:
          - contents:
              source: data:...
              mode: 420
              overwrite: true
              path: /example/path
# ...
```

- ① Replace **<namespace>** with the name of your namespace where you created the node pool, such as **clusters**.

- Apply the config map to your node pool by running the following command:

```
$ oc edit nodepool <nodepool_name> --namespace <hosted_cluster_namespace>
```

Example NodePool configuration

```
apiVersion: hypershift.openshift.io/v1alpha1
kind: NodePool
metadata:
# ...
  name: nodepool-1
  namespace: clusters
# ...
spec:
  config:
    - name: <configmap_name> ①
# ...
```

① Replace **<configmap_name>** with the name of your config map.

- Add the list of your NTP servers in the **infra-env.yaml** file, which defines the **InfraEnv** custom resource (CR):

Example infra-env.yaml file

```
apiVersion: agent-install.openshift.io/v1beta1
kind: InfraEnv
# ...
spec:
  additionalNTPSources:
    - <ntp_server> ①
    - <ntp_server1>
    - <ntp_server2>
# ...
```

① Replace **<ntp_server>** with the name of your NTP server. For more details about creating a host inventory and the **InfraEnv** CR, see "Creating a host inventory".

- Apply the **InfraEnv** CR by running the following command:

```
$ oc apply -f infra-env.yaml
```

Verification

- Check the following fields to know the status of your host inventory:
 - conditions**: The standard Kubernetes conditions indicating if the image was created successfully.
 - isoDownloadURL**: The URL to download the Discovery Image.
 - createdTime**: The time at which the image was last created. If you modify the **InfraEnv** CR, ensure that you have updated the timestamp before downloading a new image.
- Verify that your host inventory is created by running the following command:

```
$ oc describe infraenv <infraenv_resource_name> -n <infraenv_namespace>
```



NOTE

If you modify the **InfraEnv** CR, confirm that the **InfraEnv** CR has created a new Discovery Image by looking at the **createdTime** field. If you already booted hosts, boot them again with the latest Discovery Image.

Additional resources

- [Creating machine configs with Butane](#)
- [Creating a host inventory](#)

CHAPTER 11. USING FEATURE GATES IN A HOSTED CLUSTER

You can use feature gates in a hosted cluster to enable features that are not part of the default set of features. You can enable the **TechPreviewNoUpgrade** feature set by using feature gates in your hosted cluster.

11.1. ENABLING FEATURE SETS BY USING FEATURE GATES

You can enable the **TechPreviewNoUpgrade** feature set in a hosted cluster by editing the **HostedCluster** custom resource (CR) with the OpenShift CLI.

Prerequisites

- You installed the OpenShift CLI (**oc**).

Procedure

1. Open the **HostedCluster** CR for editing on the hosting cluster by running the following command:

```
$ oc edit <hosted_cluster_name> -n <hosted_cluster_namespace>
```

2. Define the feature set by entering a value in the **featureSet** field. For example:

```
apiVersion: hypershift.openshift.io/v1beta1
kind: HostedCluster
metadata:
  name: <hosted_cluster_name> ①
  namespace: <hosted_cluster_namespace> ②
spec:
  configuration:
    featureGate:
      featureSet: TechPreviewNoUpgrade ③
```

- ① Specifies your hosted cluster name.
- ② Specifies your hosted cluster namespace.
- ③ This feature set is a subset of the current Technology Preview features.



WARNING

Enabling the **TechPreviewNoUpgrade** feature set on your cluster cannot be undone and prevents minor version updates. This feature set allows you to enable these Technology Preview features on test clusters, where you can fully test them. Do not enable this feature set on production clusters.

3. Save the file to apply the changes.

Verification

- Verify that the **TechPreviewNoUpgrade** feature gate is enabled in your hosted cluster by running the following command:

```
$ oc get featuregate cluster -o yaml
```

Additional resources

- FeatureGate [config.openshift.io/v1]

CHAPTER 12. OBSERVABILITY FOR HOSTED CONTROL PLANES

You can gather metrics for hosted control planes by configuring metrics sets. The HyperShift Operator can create or delete monitoring dashboards in the management cluster for each hosted cluster that it manages.

12.1. CONFIGURING METRICS SETS FOR HOSTED CONTROL PLANES

Hosted control planes for Red Hat OpenShift Container Platform creates **ServiceMonitor** resources in each control plane namespace that allow a Prometheus stack to gather metrics from the control planes. The **ServiceMonitor** resources use metrics relabelings to define which metrics are included or excluded from a particular component, such as etcd or the Kubernetes API server. The number of metrics that are produced by control planes directly impacts the resource requirements of the monitoring stack that gathers them.

Instead of producing a fixed number of metrics that apply to all situations, you can configure a metrics set that identifies a set of metrics to produce for each control plane. The following metrics sets are supported:

- **Telemetry**: These metrics are needed for telemetry. This set is the default set and is the smallest set of metrics.
- **SRE**: This set includes the necessary metrics to produce alerts and allow the troubleshooting of control plane components.
- **All**: This set includes all of the metrics that are produced by standalone OpenShift Container Platform control plane components.

To configure a metrics set, set the **METRICS_SET** environment variable in the HyperShift Operator deployment by entering the following command:

```
$ oc set env -n hypershift deployment/operator METRICS_SET=All
```

12.1.1. Configuring the SRE metrics set

When you specify the **SRE** metrics set, the HyperShift Operator looks for a config map named **sre-metric-set** with a single key: **config**. The value of the **config** key must contain a set of **RelabelConfigs** that are organized by control plane component.

You can specify the following components:

- **etcd**
- **kubeAPIServer**
- **kubeControllerManager**
- **openshiftAPIServer**
- **openshiftControllerManager**
- **openshiftRouteControllerManager**
- **cvo**

- **olm**
- **catalogOperator**
- **registryOperator**
- **nodeTuningOperator**
- **controlPlaneOperator**
- **hostedClusterConfigOperator**

A configuration of the **SRE** metrics set is illustrated in the following example:

```

kubeAPIServer:
  - action: "drop"
    regex: "etcd_(debugging|disk|server).*"
    sourceLabels: ["__name__"]
  - action: "drop"
    regex: "apiserver_admission_controller_admission_latencies_seconds_.*"
    sourceLabels: ["__name__"]
  - action: "drop"
    regex: "apiserver_admission_step_admission_latencies_seconds_.*"
    sourceLabels: ["__name__"]
  - action: "drop"
    regex:
      "scheduler_(e2e_scheduling_latency_microseconds|scheduling_algorithm_predicate_evaluation|scheduling_algorithm_priority_evaluation|scheduling_algorithm_preemption_evaluation|scheduling_algorithm_latency_microseconds|binding_latency_microseconds|scheduling_latency_seconds)"
      sourceLabels: ["__name__"]
  - action: "drop"
    regex:
      "apiserver_(request_count|request_latencies|request_latencies_summary|dropped_requests|storage_data_key_generation_latencies_microseconds|storage_transformation_failures_total|storage_transformation_latencies_microseconds|proxy_tunnel_sync_latency_secs)"
      sourceLabels: ["__name__"]
  - action: "drop"
    regex:
      "docker_(operations|operations_latency_microseconds|operations_errors|operations_timeout)"
      sourceLabels: ["__name__"]
  - action: "drop"
    regex:
      "reflector_(items_per_list|items_per_watch|list_duration_seconds|lists_total|short_watches_total|watch_duration_seconds|watches_total)"
      sourceLabels: ["__name__"]
  - action: "drop"
    regex:
      "etcd_(helper_cache_hit_count|helper_cache_miss_count|helper_cache_entry_count|request_cache_get_latencies_summary|request_cache_add_latencies_summary|request_latencies_summary)"
      sourceLabels: ["__name__"]
  - action: "drop"
    regex: "transformation_(transformation_latencies_microseconds|failures_total)"
    sourceLabels: ["__name__"]
  - action: "drop"
    regex:
      "network_plugin_operations_latency_microseconds|sync_proxy_rules_latency_microseconds|rest_client"

```

```

_request_latency_seconds"
  sourceLabels: ["__name__"]
- action: "drop"
  regex: "apiserver_request_duration_seconds_bucket;
(0.15|0.25|0.3|0.35|0.4|0.45|0.6|0.7|0.8|0.9|1.25|1.5|1.75|2.5|3|3.5|4.5|6|7|8|9|15|25|30|50)"
  sourceLabels: ["__name__", "le"]
kubeControllerManager:
- action: "drop"
  regex: "etcd_(debugging|disk|request|server).*"
  sourceLabels: ["__name__"]
- action: "drop"
  regex: "rest_client_request_latency_seconds_(bucket|count|sum)"
  sourceLabels: ["__name__"]
- action: "drop"
  regex: "root_ca_cert_publisher_sync_duration_seconds_(bucket|count|sum)"
  sourceLabels: ["__name__"]
openshiftAPIServer:
- action: "drop"
  regex: "etcd_(debugging|disk|server).*"
  sourceLabels: ["__name__"]
- action: "drop"
  regex: "apiserver_admission_controller_admission_latencies_seconds_.*"
  sourceLabels: ["__name__"]
- action: "drop"
  regex: "apiserver_admission_step_admission_latencies_seconds_.*"
  sourceLabels: ["__name__"]
- action: "drop"
  regex: "apiserver_request_duration_seconds_bucket;
(0.15|0.25|0.3|0.35|0.4|0.45|0.6|0.7|0.8|0.9|1.25|1.5|1.75|2.5|3|3.5|4.5|6|7|8|9|15|25|30|50)"
  sourceLabels: ["__name__", "le"]
openshiftControllerManager:
- action: "drop"
  regex: "etcd_(debugging|disk|request|server).*"
  sourceLabels: ["__name__"]
openshiftRouteControllerManager:
- action: "drop"
  regex: "etcd_(debugging|disk|request|server).*"
  sourceLabels: ["__name__"]
olm:
- action: "drop"
  regex: "etcd_(debugging|disk|server).*"
  sourceLabels: ["__name__"]
catalogOperator:
- action: "drop"
  regex: "etcd_(debugging|disk|server).*"
  sourceLabels: ["__name__"]
cvo:
- action: drop
  regex: "etcd_(debugging|disk|server).*"
  sourceLabels: ["__name__"]

```

12.2. ENABLING MONITORING DASHBOARDS IN A HOSTED CLUSTER

You can enable monitoring dashboards in a hosted cluster by creating a config map.

Procedure

- Create the **hypershift-operator-install-flags** config map in the **local-cluster** namespace. See the following example configuration:

```
kind: ConfigMap
apiVersion: v1
metadata:
  name: hypershift-operator-install-flags
  namespace: local-cluster
data:
  installFlagsToAdd: "--monitoring-dashboards --metrics-set=All" ①
  installFlagsToRemove: ""
```

- ① The **--monitoring-dashboards --metrics-set>All** flag adds the monitoring dashboard for all metrics.

- Wait a couple of minutes for the HyperShift Operator deployment in the **hypershift** namespace to be updated to include the following environment variable:

```
- name: MONITORING_DASHBOARDS
  value: "1"
```

When monitoring dashboards are enabled, for each hosted cluster that the HyperShift Operator manages, the Operator creates a config map named **cp-<hosted_cluster_namespace>-<hosted_cluster_name>** in the **openshift-config-managed** namespace, where **<hosted_cluster_namespace>** is the namespace of the hosted cluster and **<hosted_cluster_name>** is the name of the hosted cluster. As a result, a new dashboard is added in the administrative console of the management cluster.

- To view the dashboard, log in to the management cluster's console and go to the dashboard for the hosted cluster by clicking **Observe → Dashboards**.
- Optional: To disable monitoring dashboards in a hosted cluster, remove the **--monitoring-dashboards --metrics-set>All** flag from the **hypershift-operator-install-flags** config map. When you delete a hosted cluster, its corresponding dashboard is also deleted.

12.2.1. Dashboard customization

To generate dashboards for each hosted cluster, the HyperShift Operator uses a template that is stored in the **monitoring-dashboard-template** config map in the Operator namespace (**hypershift**). This template contains a set of Grafana panels that contain the metrics for the dashboard. You can edit the content of the config map to customize the dashboards.

When a dashboard is generated, the following strings are replaced with values that correspond to a specific hosted cluster:

Name	Description
NAME	The name of the hosted cluster
NAMESPACE	The namespace of the hosted cluster

Name	Description
<code>__CONTROL_PLANE_NAMESPACE__</code>	The namespace where the control plane pods of the hosted cluster are placed
<code>__CLUSTER_ID__</code>	The UUID of the hosted cluster, which matches the <code>_id</code> label of the hosted cluster metrics

CHAPTER 13. NETWORKING FOR HOSTED CONTROL PLANES

For standalone OpenShift Container Platform, proxy support is mainly about ensuring that workloads in the cluster are configured to use the HTTP or HTTPS proxy to access external services, honoring the **NO_PROXY** setting if one is configured, and accepting any trust bundle that is configured for the proxy.

For hosted control planes, proxy support involves the following additional use cases.

13.1. CONTROL PLANE WORKLOADS THAT NEED TO ACCESS EXTERNAL SERVICES

Operators that run in the control plane need to access external services through the proxy that is configured for the hosted cluster. The proxy is usually accessible only through the data plane. The control plane workloads are as follows:

- The Control Plane Operator needs to validate and obtain endpoints from certain identity providers when it creates the OAuth server configuration.
- The OAuth server needs non-LDAP identity provider access.
- The OpenShift API server handles image registry metadata import.
- The Ingress Operator needs access to validate external canary routes.

In a hosted cluster, you must send traffic that originates from the Control Plane Operator, Ingress Operator, OAuth server, and OpenShift API server pods through the data plane to the configured proxy and then to its final destination.



NOTE

Some operations are not possible when a hosted cluster is reduced to zero compute nodes; for example, when you import OpenShift image streams from a registry that requires proxy access.

13.2. COMPUTE NODES THAT NEED TO ACCESS AN IGNITION ENDPOINT

When compute nodes need a proxy to access the ignition endpoint, you must configure the proxy in the user-data stub that is configured on the compute node when it is created. For cases where machines need a proxy to access the ignition URL, the proxy configuration is included in the stub.

The stub resembles the following example:

```
---
{"ignition":{"config":{"merge":[{"httpHeaders":[{"name":"Authorization","value":"Bearer ..."}, {"name":"TargetConfigVersionHash","value":"a4c1b0dd"}],"source":"https://ignition.controlplanehost.example.com/ignition","verification":{} }], "replace":{"verification":{}}, "proxy":{"httpProxy":"http://proxy.example.org:3128", "httpsProxy":"https://proxy.example.org:3129", "noProxy":"host.example.org"}, "security":{"tls":{"certificateAuthorities":[{"source": "...", "verification": {}}]}}, "timeouts":{}, "version":"3.2.0"}, "passwd":{}, "storage":{}, "systemd":{}}
---
```

13.3. COMPUTE NODES THAT NEED TO ACCESS THE API SERVER

This use case is relevant to self-managed hosted control planes, not to Red Hat OpenShift Service on AWS with hosted control planes.

For communication with the control plane, hosted control planes uses a local proxy in every compute node that listens on IP address 172.20.0.1 and forwards traffic to the API server. If an external proxy is required to access the API server, that local proxy needs to use the external proxy to send traffic out. When a proxy is not needed, hosted control planes uses **haproxy** for the local proxy, which only forwards packets via TCP. When a proxy is needed, hosted control planes uses a custom proxy, **control-plane-operator-kubernetes-default-proxy**, to send traffic through the external proxy.

13.4. MANAGEMENT CLUSTERS THAT NEED EXTERNAL ACCESS

The HyperShift Operator has a controller that monitors the OpenShift global proxy configuration of the management cluster and sets the proxy environment variables on its own deployment. Control plane deployments that need external access are configured with the proxy environment variables of the management cluster.

13.5. MANAGEMENT CLUSTER THAT USES A PROXY AND A HOSTED CLUSTER WITH A SECONDARY NETWORK AND NO DEFAULT POD NETWORK

If a management cluster uses a proxy configuration and you are configuring a hosted cluster with a secondary network but are not attaching the default pod network, add the CIDR of the secondary network to the proxy configuration. Specifically, you need to add the CIDR of the secondary network to the **noProxy** section of the proxy configuration for the management cluster. Otherwise, the Kubernetes API server will route some API requests through the proxy. In the hosted cluster configuration, the CIDR of the secondary network is automatically added to the **noProxy** section.

13.6. ADDITIONAL RESOURCES

- [Configuring the cluster-wide proxy](#)

CHAPTER 14. TROUBLESHOOTING HOSTED CONTROL PLANES

If you encounter issues with hosted control planes, see the following information to guide you through troubleshooting.

14.1. GATHERING INFORMATION TO TROUBLESHOOT HOSTED CONTROL PLANES

When you need to troubleshoot an issue with hosted clusters, you can gather information by running the **must-gather** command. The command generates output for the management cluster and the hosted cluster.

The output for the management cluster contains the following content:

- **Cluster-scoped resources:** These resources are node definitions of the management cluster.
- **The hypershift-dump compressed file:** This file is useful if you need to share the content with other people.
- **Namespaced resources:** These resources include all of the objects from the relevant namespaces, such as config maps, services, events, and logs.
- **Network logs:** These logs include the OVN northbound and southbound databases and the status for each one.
- **Hosted clusters:** This level of output involves all of the resources inside of the hosted cluster.

The output for the hosted cluster contains the following content:

- **Cluster-scoped resources:** These resources include all of the cluster-wide objects, such as nodes and CRDs.
- **Namespaced resources:** These resources include all of the objects from the relevant namespaces, such as config maps, services, events, and logs.

Although the output does not contain any secret objects from the cluster, it can contain references to the names of secrets.

Prerequisites

- You must have **cluster-admin** access to the management cluster.
- You need the **name** value for the **HostedCluster** resource and the namespace where the CR is deployed.
- You must have the **hcp** command-line interface installed. For more information, see "Installing the hosted control planes command-line interface".
- You must have the OpenShift CLI (**oc**) installed.
- You must ensure that the **kubeconfig** file is loaded and is pointing to the management cluster.

Procedure

- To gather the output for troubleshooting, enter the following command:

```
$ oc adm must-gather \
--image=registry.redhat.io/multicluster-engine/must-gather-rhel9:v<mce_version> \
/usr/bin/gather hosted-cluster-namespace=HOSTEDCLUSTERNAMESPACE \
hosted-cluster-name=HOSTEDCLUSTERNAME \
--dest-dir=NAME ; tar -cvzf NAME.tgz NAME
```

where:

- You replace **<mce_version>** with the version of multicluster engine Operator that you are using; for example, **2.6**.
- The **hosted-cluster-namespace=HOSTEDCLUSTERNAMESPACE** parameter is optional. If you do not include it, the command runs as though the hosted cluster is in the default namespace, which is **clusters**.
- If you want to save the results of the command to a compressed file, specify the **--dest-dir=NAME** parameter and replace **NAME** with the name of the directory where you want to save the results.

Additional resources

- [Installing the hosted control planes command-line interface](#)

14.2. GATHERING OPENSHIFT CONTAINER PLATFORM DATA FOR A HOSTED CLUSTER

You can gather OpenShift Container Platform debugging information for a hosted cluster by using the multicluster engine Operator web console or by using the CLI.

14.2.1. Gathering data for a hosted cluster by using the CLI

You can gather OpenShift Container Platform debugging information for a hosted cluster by using the CLI.

Prerequisites

- You must have **cluster-admin** access to the management cluster.
- You need the **name** value for the **HostedCluster** resource and the namespace where the CR is deployed.
- You must have the **hcp** command-line interface installed. For more information, see "Installing the hosted control planes command-line interface".
- You must have the OpenShift CLI (**oc**) installed.
- You must ensure that the **kubeconfig** file is loaded and is pointing to the management cluster.

Procedure

1. Generate the **kubeconfig** file by entering the following command:

```
$ hcp create kubeconfig --namespace <hosted_cluster_namespace> \  
--name <hosted_cluster_name> > <hosted_cluster_name>.kubeconfig
```

2. After you save the **kubeconfig** file, you can access the hosted cluster by entering the following example command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get nodes
```

3. Collect the must-gather information by entering the following command:

```
$ oc adm must-gather
```

14.2.2. Gathering data for a hosted cluster by using the web console

You can gather OpenShift Container Platform debugging information for a hosted cluster by using the multicloud engine Operator web console.

Prerequisites

- You must have **cluster-admin** access to the management cluster.
- You need the **name** value for the **HostedCluster** resource and the namespace where the CR is deployed.
- You must have the **hcp** command-line interface installed. For more information, see "Installing the hosted control planes command-line interface".
- You must have the OpenShift CLI (**oc**) installed.
- You must ensure that the **kubeconfig** file is loaded and is pointing to the management cluster.

Procedure

1. In the web console, select **All Clusters** and select the cluster you want to troubleshoot.
2. In the upper-right corner, select **Download kubeconfig**.
3. Export the downloaded **kubeconfig** file.
4. Collect the must-gather information by entering the following command:

```
$ oc adm must-gather
```

14.3. ENTERING THE MUST-GATHER COMMAND IN A DISCONNECTED ENVIRONMENT

Complete the following steps to run the **must-gather** command in a disconnected environment.

Procedure

1. In a disconnected environment, mirror the Red Hat operator catalog images into their mirror registry. For more information, see *Install on disconnected networks*.

- Run the following command to extract logs, which reference the image from their mirror registry:

```
REGISTRY=registry.example.com:5000
IMAGE=$REGISTRY/multicluster-engine/must-gather-
rhel8@sha256:ff9f37eb400dc1f7d07a9b6f2da9064992934b69847d17f59e385783c071b9d8

$ oc adm must-gather \
--image=$IMAGE /usr/bin/gather \
hosted-cluster-namespace=HOSTEDCLUSTERNAMESPACE \
hosted-cluster-name=HOSTEDCLUSTERNAME \
--dest-dir=./data
```

Additional resources

- [Install on disconnected networks](#)

14.4. TROUBLESHOOTING HOSTED CLUSTERS ON OPENSHIFT VIRTUALIZATION

When you troubleshoot a hosted cluster on OpenShift Virtualization, start with the top-level **HostedCluster** and **NodePool** resources and then work down the stack until you find the root cause. The following steps can help you discover the root cause of common issues.

14.4.1. HostedCluster resource is stuck in a partial state

If a hosted control plane is not coming fully online because a **HostedCluster** resource is pending, identify the problem by checking prerequisites, resource conditions, and node and Operator status.

Procedure

- Ensure that you meet all of the prerequisites for a hosted cluster on OpenShift Virtualization.
- View the conditions on the **HostedCluster** and **NodePool** resources for validation errors that prevent progress.
- By using the **kubeconfig** file of the hosted cluster, inspect the status of the hosted cluster:
 - View the output of the **oc get clusteroperators** command to see which cluster Operators are pending.
 - View the output of the **oc get nodes** command to ensure that worker nodes are ready.

14.4.2. No worker nodes are registered

If a hosted control plane is not coming fully online because the hosted control plane has no worker nodes registered, identify the problem by checking the status of various parts of the hosted control plane.

Procedure

- View the **HostedCluster** and **NodePool** conditions for failures that indicate what the problem might be.

- Enter the following command to view the KubeVirt worker node virtual machine (VM) status for the **NodePool** resource:

```
$ oc get vm -n <namespace>
```

- If the VMs are stuck in the provisioning state, enter the following command to view the CDI import pods within the VM namespace for clues about why the importer pods have not completed:

```
$ oc get pods -n <namespace> | grep "import"
```

- If the VMs are stuck in the starting state, enter the following command to view the status of the virt-launcher pods:

```
$ oc get pods -n <namespace> -l kubevirt.io=virt-launcher
```

If the virt-launcher pods are in a pending state, investigate why the pods are not being scheduled. For example, not enough resources might exist to run the virt-launcher pods.

- If the VMs are running but they are not registered as worker nodes, use the web console to gain VNC access to one of the affected VMs. The VNC output indicates whether the ignition configuration was applied. If a VM cannot access the hosted control plane ignition server on startup, the VM cannot be provisioned correctly.
- If the ignition configuration was applied but the VM is still not registering as a node, see [Identifying the problem: Access the VM console logs](#) to learn how to access the VM console logs during startup.

Additional resources

- [Identifying the problem: Access the VM console logs](#)

14.4.3. Worker nodes are stuck in the NotReady state

During cluster creation, nodes enter the **NotReady** state temporarily while the networking stack is rolled out. This part of the process is normal. However, if this part of the process takes longer than 15 minutes, identify the problem by investigating the node object and pods.

Procedure

- Enter the following command to view the conditions on the node object and determine why the node is not ready:

```
$ oc get nodes -o yaml
```

- Enter the following command to look for failing pods within the cluster:

```
$ oc get pods -A --field-selector=status.phase!=Running,status.phase!=Succeeded
```

14.4.4. Ingress and console cluster operators are not coming online

If a hosted control plane is not coming fully online because the Ingress and console cluster Operators are not online, check the wildcard DNS routes and load balancer.

Procedure

- If the cluster uses the default Ingress behavior, enter the following command to ensure that wildcard DNS routes are enabled on the OpenShift Container Platform cluster that the virtual machines (VMs) are hosted on:

```
$ oc patch ingresscontroller -n openshift-ingress-operator \
  default --type=json -p \
  '[{"op": "add", "path": "/spec/routeAdmission", "value": {"wildcardPolicy": "WildcardsAllowed"}}]'
```

- If you use a custom base domain for the hosted control plane, complete the following steps:
 - Ensure that the load balancer is targeting the VM pods correctly.
 - Ensure that the wildcard DNS entry is targeting the load balancer IP address.

14.4.5. Load balancer services for the hosted cluster are not available

If a hosted control plane is not coming fully online because the load balancer services are not becoming available, check events, details, and the Kubernetes Cluster Configuration Manager (KCCM) pod.

Procedure

- Look for events and details that are associated with the load balancer service within the hosted cluster.
- By default, load balancers for the hosted cluster are handled by the kubevirt-cloud-controller-manager within the hosted control plane namespace. Ensure that the KCCM pod is online and view its logs for errors or warnings. To identify the KCCM pod in the hosted control plane namespace, enter the following command:

```
$ oc get pods -n <hosted_control_plane_namespace> \
  -l app=cloud-controller-manager
```

14.4.6. Hosted cluster PVCs are not available

If a hosted control plane is not coming fully online because the persistent volume claims (PVCs) for a hosted cluster are not available, check the PVC events and details, and component logs.

Procedure

- Look for events and details that are associated with the PVC to understand which errors are occurring.
- If a PVC is failing to attach to a pod, view the logs for the kubevirt-csi-node **daemonset** component within the hosted cluster to further investigate the problem. To identify the kubevirt-csi-node pods for each node, enter the following command:

```
$ oc get pods -n openshift-cluster-csi-drivers -o wide \
  -l app=kubevirt-csi-driver
```

- If a PVC cannot bind to a persistent volume (PV), view the logs of the kubevirt-csi-controller component within the hosted control plane namespace. To identify the kubevirt-csi-controller pod within the hosted control plane namespace, enter the following command:

```
$ oc get pods -n <hcp namespace> -l app=kubevirt-csi-driver
```

14.4.7. VM nodes are not correctly joining the cluster

If a hosted control plane is not coming fully online because the VM nodes are not correctly joining the cluster, access the VM console logs.

Procedure

- To access the VM console logs, complete the steps in [How to get serial console logs for VMs part of OpenShift Virtualization Hosted Control Plane clusters](#).

14.4.8. RHCOS image mirroring fails

For hosted control planes on OpenShift Virtualization in a disconnected environment, **oc-mirror** fails to automatically mirror the Red Hat Enterprise Linux CoreOS (RHCOS) image to the internal registry. When you create your first hosted cluster, the Kubevirt virtual machine does not boot, because the boot image is not available in the internal registry.

To resolve this issue, manually mirror the RHCOS image to the internal registry.

Procedure

1. Get the internal registry name by running the following command:

```
$ oc get imagecontentsourcepolicy -o json \
| jq -r '.items[].spec.repositoryDigestMirrors[0].mirrors[0]'
```

2. Get a payload image by running the following command:

```
$ oc get clusterversion version -ojsonpath='{.status.desired.image}'
```

3. Extract the **0000_50_installer_coreos-bootimages.yaml** file that contains boot images from your payload image on the hosted cluster. Replace **<payload_image>** with the name of your payload image. Run the following command:

```
$ oc image extract \
--file /release-manifests/0000_50_installer_coreos-bootimages.yaml \
<payload_image> --confirm
```

4. Get the RHCOS image by running the following command:

```
$ cat 0000_50_installer_coreos-bootimages.yaml | jq -r .data.stream \
| jq -r '.architectures.x86_64.images.kubevirt."digest-ref"'
```

5. Mirror the RHCOS image to your internal registry. Replace **<rhcoss_image>** with your RHCOS image; for example, **quay.io/openshift-release-dev/ocp-v4.0-art-dev@sha256:d9643ead36b1c026be664c9c65c11433c6cdf71bfd93ba229141d134a4a6dd94**.

Replace `<internal_registry>` with the name of your internal registry; for example, `virthost.ostest.test.metalkube.org:5000/localimages/ocp-v4.0-art-dev`. Run the following command:

```
$ oc image mirror <rhcoss_image> <internal_registry>
```

6. Create a YAML file named **rhcoss-boot-kubevirt.yaml** that defines the **ImageDigestMirrorSet** object. See the following example configuration:

```
apiVersion: config.openshift.io/v1
kind: ImageDigestMirrorSet
metadata:
  name: rhcos-boot-kubevirt
spec:
  repositoryDigestMirrors:
    - mirrors:
      - virthost.ostest.test.metalkube.org:5000/localimages/ocp-v4.0-art-dev ①
      source: quay.io/openshift-release-dev/ocp-v4.0-art-dev ②
```

- ① Specify the name of your internal registry, for example, `virthost.ostest.test.metalkube.org:5000/localimages/ocp-v4.0-art-dev`.
- ② Specify your RHCOS image without its digest, for example, `quay.io/openshift-release-dev/ocp-v4.0-art-dev`.

7. Apply the **rhcoss-boot-kubevirt.yaml** file to create the **ImageDigestMirrorSet** object by running the following command:

```
$ oc apply -f rhcos-boot-kubevirt.yaml
```

14.4.9. Return non-bare-metal clusters to the late binding pool

If you are using late binding managed clusters without **BareMetalHosts**, you must complete additional manual steps to delete a late binding cluster and return the nodes back to the Discovery ISO.

For late binding managed clusters without **BareMetalHosts**, removing cluster information does not automatically return all nodes to the Discovery ISO.

Procedure

To unbind the non-bare-metal nodes with late binding, complete the following steps:

1. Remove the cluster information. For more information, see *Removing a cluster from management*.
2. Clean the root disks.
3. Reboot manually with the Discovery ISO.

Additional resources

- [Removing a cluster from management](#)

14.5. TROUBLESHOOTING HOSTED CLUSTERS ON BARE METAL

The following information applies to troubleshooting hosted control planes on bare metal.

14.5.1. Nodes fail to be added to hosted control planes on bare metal

When you scale up a hosted control planes cluster with nodes that were provisioned by using Assisted Installer, the host fails to pull the ignition with a URL that contains port 22642. That URL is invalid for hosted control planes and indicates that an issue exists with the cluster.

Procedure

1. To determine the issue, review the assisted-service logs:

```
$ oc logs -n multicluster-engine <assisted_service_pod_name> ①
```

- 1 Specify the Assisted Service pod name.

2. In the logs, find errors that resemble these examples:

```
error="failed to get pull secret for update: invalid pull secret data in secret pull-secret"
```

```
pull secret must contain auth for \"registry.redhat.io\"
```

3. To fix this issue, see "Add the pull secret to the namespace" in the multicluster engine for Kubernetes Operator documentation.



NOTE

To use hosted control planes, you must have multicluster engine Operator installed, either as a standalone operator or as part of Red Hat Advanced Cluster Management. Because the operator has a close association with Red Hat Advanced Cluster Management, the documentation for the operator is published within that product's documentation. Even if you do not use Red Hat Advanced Cluster Management, the parts of its documentation that cover multicluster engine Operator are relevant to hosted control planes.

Additional resources

- [Add the pull secret to the namespace](#)

14.6. RESTARTING HOSTED CONTROL PLANE COMPONENTS

If you are an administrator for hosted control planes, you can use the `hypershift.openshift.io/restart-date` annotation to restart all control plane components for a particular `HostedCluster` resource. For example, you might need to restart control plane components for certificate rotation.

Procedure

- To restart a control plane, annotate the `HostedCluster` resource by entering the following command:

```
$ oc annotate hostedcluster \  
-n <hosted_cluster_namespace> \  
<hosted_cluster_name> \  
hypershift.openshift.io/restart-date=$(date --iso-8601=seconds) 1
```

- 1 The control plane is restarted whenever the value of the annotation changes. The **date** command serves as the source of a unique string. The annotation is treated as a string, not a timestamp.

Verification

After you restart a control plane, the following hosted control planes components are typically restarted:



NOTE

You might see some additional components restarting as a side effect of changes implemented by the other components.

- catalog-operator
- certified-operators-catalog
- cluster-api
- cluster-autoscaler
- cluster-policy-controller
- cluster-version-operator
- community-operators-catalog
- control-plane-operator
- hosted-cluster-config-operator
- ignition-server
- ingress-operator
- konnectivity-agent
- konnectivity-server
- kube-apiserver
- kube-controller-manager
- kube-scheduler
- machine-approvers
- oauth-openshift
- olm-operator

- openshift-apiserver
- openshift-controller-manager
- openshift-oauth-apiserver
- packageserver
- redhat-marketplace-catalog
- redhat-operators-catalog

14.7. PAUSING THE RECONCILIATION OF A HOSTED CLUSTER AND HOSTED CONTROL PLANE

If you are a cluster instance administrator, you can pause the reconciliation of a hosted cluster and hosted control plane. You might want to pause reconciliation when you back up and restore an etcd database or when you need to debug problems with a hosted cluster or hosted control plane.

Procedure

1. To pause reconciliation for a hosted cluster and hosted control plane, populate the **pausedUntil** field of the **HostedCluster** resource.

- To pause the reconciliation until a specific time, enter the following command:

```
$ oc patch -n <hosted_cluster_namespace> \
  hostedclusters/<hosted_cluster_name> \
  -p '{"spec":{"pausedUntil":"<timestamp>"}' \
  --type=merge ①
```

- ① Specify a timestamp in the RFC339 format, for example, **2024-03-03T03:28:48Z**. The reconciliation is paused until the specified time is passed.

- To pause the reconciliation indefinitely, enter the following command:

```
$ oc patch -n <hosted_cluster_namespace> \
  hostedclusters/<hosted_cluster_name> \
  -p '{"spec":{"pausedUntil":"true"}' \
  --type=merge
```

The reconciliation is paused until you remove the field from the **HostedCluster** resource.

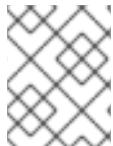
When the pause reconciliation field is populated for the **HostedCluster** resource, the field is automatically added to the associated **HostedControlPlane** resource.

2. To remove the **pausedUntil** field, enter the following patch command:

```
$ oc patch -n <hosted_cluster_namespace> \
  hostedclusters/<hosted_cluster_name> \
  -p '{"spec":{"pausedUntil":null}}' \
  --type=merge
```

14.8. SCALING DOWN THE DATA PLANE TO ZERO

If you are not using the hosted control plane, to save the resources and cost you can scale down a data plane to zero.



NOTE

Ensure you are prepared to scale down the data plane to zero. Because the workload from the worker nodes disappears after scaling down.

Procedure

- Set the **kubeconfig** file to access the hosted cluster by running the following command:

```
$ export KUBECONFIG=<install_directory>/auth/kubeconfig
```

- Get the name of the **NodePool** resource associated to your hosted cluster by running the following command:

```
$ oc get nodepool --namespace <hosted_cluster_namespace>
```

- Optional: To prevent the pods from draining, add the **nodeDrainTimeout** field in the **NodePool** resource by running the following command:

```
$ oc edit nodepool <nodepool_name> --namespace <hosted_cluster_namespace>
```

Example output

```
apiVersion: hypershift.openshift.io/v1alpha1
kind: NodePool
metadata:
# ...
  name: nodepool-1
  namespace: clusters
# ...
spec:
  arch: amd64
  clusterName: clustername 1
  management:
    autoRepair: false
    replace:
      rollingUpdate:
        maxSurge: 1
        maxUnavailable: 0
        strategy: RollingUpdate
      upgradeType: Replace
    nodeDrainTimeout: 0s 2
# ...
```

- 1** Defines the name of your hosted cluster.
- 2** Specifies the total amount of time that the controller spends to drain a node. By default, the **nodeDrainTimeout: 0s** setting blocks the node draining process.

**NOTE**

To allow the node draining process to continue for a certain period of time, you can set the value of the **nodeDrainTimeout** field accordingly, for example, **nodeDrainTimeout: 1m**.

4. Scale down the **NodePool** resource associated to your hosted cluster by running the following command:

```
$ oc scale nodepool/<nodepool_name> --namespace <hosted_cluster_namespace> \
--replicas=0
```

**NOTE**

After scaling down the data plan to zero, some pods in the control plane stay in the **Pending** status and the hosted control plane stays up and running. If necessary, you can scale up the **NodePool** resource.

5. Optional: Scale up the **NodePool** resource associated to your hosted cluster by running the following command:

```
$ oc scale nodepool/<nodepool_name> --namespace <hosted_cluster_namespace> --
replicas=1
```

After rescaling the **NodePool** resource, wait for couple of minutes for the **NodePool** resource to become available in a **Ready** state.

Verification

- Verify that the value for the **nodeDrainTimeout** field is greater than **0s** by running the following command:

```
$ oc get nodepool -n <hosted_cluster_namespace> <nodepool_name> -
jsonpath='{.spec.nodeDrainTimeout}'
```

CHAPTER 15. DESTROYING A HOSTED CLUSTER

15.1. DESTROYING A HOSTED CLUSTER ON AWS

You can destroy a hosted cluster and its managed cluster resource on Amazon Web Services (AWS) by using the command-line interface (CLI).

15.1.1. Destroying a hosted cluster on AWS by using the CLI

You can use the command-line interface (CLI) to destroy a hosted cluster on Amazon Web Services (AWS).

Procedure

1. Delete the managed cluster resource on multicluster engine Operator by running the following command:

```
$ oc delete managedcluster <hosted_cluster_name> ①
```

- 1 Replace **<hosted_cluster_name>** with the name of your cluster.

2. Delete the hosted cluster and its backend resources by running the following command:

```
$ hcp destroy cluster aws \
--name <hosted_cluster_name> ①
--infra-id <infra_id> ②
--role-arn <arn_role> ③
--sts-creds <path_to_sts_credential_file> ④
--base-domain <basedomain> ⑤
```

- 1 Specify the name of your hosted cluster, for instance, **example**.
- 2 Specify the infrastructure name for your hosted cluster.
- 3 Specify the Amazon Resource Name (ARN), for example, **arn:aws:iam::820196288204:role/myrole**.
- 4 Specify the path to your AWS Security Token Service (STS) credentials file, for example, **/home/user/sts-creds/sts-creds.json**.
- 5 Specify your base domain, for example, **example.com**.



IMPORTANT

If your session token for AWS Security Token Service (STS) is expired, retrieve the STS credentials in a JSON file named **sts-creds.json** by running the following command:

```
$ aws sts get-session-token --output json > sts-creds.json
```

15.2. DESTROYING A HOSTED CLUSTER ON BARE METAL

You can destroy hosted clusters on bare metal by using the command-line interface (CLI) or the multicluster engine Operator web console.

15.2.1. Destroying a hosted cluster on bare metal by using the CLI

You can use the **hcp** command-line interface (CLI) to destroy a hosted cluster on bare metal.

Procedure

- Delete the hosted cluster and its backend resources by running the following command:

```
$ hcp destroy cluster agent --name <hosted_cluster_name> ①
```

- Specify the name of your hosted cluster.

15.2.2. Destroying a hosted cluster on bare metal by using the web console

You can use the multicluster engine Operator web console to destroy a hosted cluster on bare metal.

Procedure

- In the console, click **Infrastructure → Clusters**.
- On the **Clusters** page, select the cluster that you want to destroy.
- In the **Actions** menu, select **Destroy clusters** to remove the cluster.

15.3. DESTROYING A HOSTED CLUSTER ON OPENSHIFT VIRTUALIZATION

You can destroy a hosted cluster and its managed cluster resource on OpenShift Virtualization by using the command-line interface (CLI).

15.3.1. Destroying a hosted cluster on OpenShift Virtualization by using the CLI

You can use the command-line interface (CLI) to destroy a hosted cluster and its managed cluster resource on OpenShift Virtualization.

Procedure

- Delete the managed cluster resource on multicluster engine Operator by running the following command:

```
$ oc delete managedcluster <hosted_cluster_name>
```

- Delete the hosted cluster and its backend resources by running the following command:

```
$ hcp destroy cluster kubevirt --name <hosted_cluster_name>
```

15.4. DESTROYING A HOSTED CLUSTER ON IBM Z

You can destroy a hosted cluster on **x86** bare metal with IBM Z compute nodes and its managed cluster resource by using the command-line interface (CLI).

15.4.1. Destroying a hosted cluster on x86 bare metal with IBM Z compute nodes

To destroy a hosted cluster and its managed cluster on **x86** bare metal with IBM Z® compute nodes, you can use the command-line interface (CLI).

Procedure

1. Scale the **NodePool** object to **0** nodes by running the following command:

```
$ oc -n <hosted_cluster_namespace> scale nodepool <nodepool_name> \
--replicas 0
```

After the **NodePool** object is scaled to **0**, the compute nodes are detached from the hosted cluster. In OpenShift Container Platform version 4.17, this function is applicable only for IBM Z with KVM. For z/VM and LPAR, you must delete the compute nodes manually.

If you want to re-attach compute nodes to the cluster, you can scale up the **NodePool** object with the number of compute nodes that you want. For z/VM and LPAR to reuse the agents, you must re-create them by using the **Discovery** image.



IMPORTANT

If the compute nodes are not detached from the hosted cluster or are stuck in the **Notready** state, delete the compute nodes manually by running the following command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig delete \
node <compute_node_name>
```

If you are using an OSA network device in Processor Resource/Systems Manager (PR/SM) mode, auto scaling is not supported. You must delete the old agent manually and scale up the node pool because the new agent joins during the scale down process.

2. Verify the status of the compute nodes by entering the following command:

```
$ oc --kubeconfig <hosted_cluster_name>.kubeconfig get nodes
```

After the compute nodes are detached from the hosted cluster, the status of the agents is changed to **auto-assign**.

3. Delete the agents from the cluster by running the following command:

```
$ oc -n <hosted_control_plane_namespace> delete agent <agent_name>
```



NOTE

You can delete the virtual machines that you created as agents after you delete the agents from the cluster.

4. Destroy the hosted cluster by running the following command:

```
$ hcp destroy cluster agent --name <hosted_cluster_name> \  
--namespace <hosted_cluster_namespace>
```

15.5. DESTROYING A HOSTED CLUSTER ON IBM POWER

You can destroy a hosted cluster on IBM Power by using the command-line interface (CLI).

15.5.1. Destroying a hosted cluster on IBM Power by using the CLI

To destroy a hosted cluster on IBM Power, you can use the `hcp` command-line interface (CLI).

Procedure

- Delete the hosted cluster by running the following command:

```
$ hcp destroy cluster agent  
--name=<hosted_cluster_name> \①  
--namespace=<hosted_cluster_namespace> \②  
--cluster-grace-period <duration> ③
```

- ① Replace `<hosted_cluster_name>` with the name of your hosted cluster.
- ② Replace `<hosted_cluster_namespace>` with the name of your hosted cluster namespace.
- ③ Specifies the duration to destroy the hosted cluster completely, for example, **20m0s**.

15.6. DESTROYING A HOSTED CLUSTER ON NON-BARE-METAL AGENT MACHINES

You can destroy hosted clusters on non-bare-metal agent machines by using the command-line interface (CLI) or the multicluster engine Operator web console.

15.6.1. Destroying a hosted cluster on non-bare-metal agent machines

You can use the `hcp` command-line interface (CLI) to destroy a hosted cluster on non-bare-metal agent machines.

Procedure

- Delete the hosted cluster and its backend resources by running the following command:

```
$ hcp destroy cluster agent --name <hosted_cluster_name> ①
```

- ① Replace `<hosted_cluster_name>` with the name of your hosted cluster.

15.6.2. Destroying a hosted cluster on non-bare-metal agent machines by using the web console

You can use the multicluster engine Operator web console to destroy a hosted cluster on non-bare-metal agent machines.

Procedure

1. In the console, click **Infrastructure → Clusters**.
2. On the **Clusters** page, select the cluster that you want to destroy.
3. In the **Actions** menu, select **Destroy clusters** to remove the cluster.

CHAPTER 16. MANUALLY IMPORTING A HOSTED CLUSTER

Hosted clusters are automatically imported into multicluster engine Operator after the hosted control plane becomes available.

16.1. LIMITATIONS OF MANAGING IMPORTED HOSTED CLUSTERS

Hosted clusters are automatically imported into the local multicluster engine for Kubernetes Operator, unlike a standalone OpenShift Container Platform or third party clusters. Hosted clusters run some of their agents in the *hosted mode* so that the agents do not use the resources of your cluster.

If you choose to automatically import hosted clusters, you can update node pools and the control plane in hosted clusters by using the **HostedCluster** resource on the management cluster. To update node pools and a control plane, see "Updating node pools in a hosted cluster" and "Updating a control plane in a hosted cluster".

You can import hosted clusters into a location other than the local multicluster engine Operator by using the Red Hat Advanced Cluster Management (RHACM). For more information, see "Discovering multicluster engine for Kubernetes Operator hosted clusters in Red Hat Advanced Cluster Management".

In this topology, you must update your hosted clusters by using the command-line interface or the console of the local multicluster engine for Kubernetes Operator where the cluster is hosted. You cannot update the hosted clusters through the RHACM hub cluster.

16.2. ADDITIONAL RESOURCES

- [Updating node pools in a hosted cluster](#)
- [Updating a control plane in a hosted cluster](#)
- [Discovering multicluster engine for Kubernetes Operator hosted clusters in Red Hat Advanced Cluster Management](#)

16.3. MANUALLY IMPORTING HOSTED CLUSTERS

If you want to import hosted clusters manually, complete the following steps.

Procedure

1. In the console, click **Infrastructure → Clusters** and select the hosted cluster that you want to import.
2. Click **Import hosted cluster**.



NOTE

For your *discovered* hosted cluster, you can also import from the console, but the cluster must be in an upgradable state. Import on your cluster is disabled if the hosted cluster is not in an upgradable state because the hosted control plane is not available. Click **Import** to begin the process. The status is **Importing** while the cluster receives updates and then changes to **Ready**.

16.4. MANUALLY IMPORTING A HOSTED CLUSTER ON AWS

You can also import a hosted cluster on Amazon Web Services (AWS) with the command-line interface.

Procedure

- Create your **ManagedCluster** resource by using the following sample YAML file:

```
apiVersion: cluster.open-cluster-management.io/v1
kind: ManagedCluster
metadata:
  annotations:
    import.open-cluster-management.io/hosting-cluster-name: local-cluster
    import.open-cluster-management.io/klusterlet-deploy-mode: Hosted
    open-cluster-management/created-via: hypershift
  labels:
    cloud: auto-detect
    cluster.open-cluster-management.io/clusterset: default
    name: <hosted_cluster_name> ①
    vendor: OpenShift
    name: <hosted_cluster_name>
spec:
  hubAcceptsClient: true
  leaseDurationSeconds: 60
```

- Replace **<hosted_cluster_name>** with the name of your hosted cluster.

- Run the following command to apply the resource:

```
$ oc apply -f <file_name> ①
```

- Replace **<file_name>** with the YAML file name you created in the previous step.

- If you have Red Hat Advanced Cluster Management installed, create your **KlusterletAddonConfig** resource by using the following sample YAML file. If you have installed multicloud engine Operator only, skip this step:

```
apiVersion: agent.open-cluster-management.io/v1
kind: KlusterletAddonConfig
metadata:
  name: <hosted_cluster_name> ①
  namespace: <hosted_cluster_namespace> ②
spec:
  clusterName: <hosted_cluster_name>
  clusterNamespace: <hosted_cluster_namespace>
  clusterLabels:
    cloud: auto-detect
    vendor: auto-detect
  applicationManager:
    enabled: true
  certPolicyController:
    enabled: true
  iamPolicyController:
```

```

enabled: true
policyController:
  enabled: true
searchCollector:
  enabled: false

```

- 1 Replace <hosted_cluster_name> with the name of your hosted cluster.
- 2 Replace <hosted_cluster_namespace> with the name of your hosted cluster namespace.

4. Run the following command to apply the resource:

```
$ oc apply -f <file_name> 1
```

- 1 Replace <file_name> with the YAML file name you created in the previous step.

5. After the import process is complete, your hosted cluster becomes visible in the console. You can also check the status of your hosted cluster by running the following command:

```
$ oc get managedcluster <hosted_cluster_name>
```

16.5. DISABLING THE AUTOMATIC IMPORT OF HOSTED CLUSTERS INTO MULTICLUSTER ENGINE OPERATOR

Hosted clusters are automatically imported into multicluster engine Operator after the control plane becomes available. If needed, you can disable the automatic import of hosted clusters.

Any hosted clusters that were previously imported are not affected, even if you disable automatic import. When you upgrade to multicluster engine Operator 2.5 and automatic import is enabled, all hosted clusters that are not imported are automatically imported if their control planes are available.



NOTE

If Red Hat Advanced Cluster Management is installed, all Red Hat Advanced Cluster Management add-ons are also enabled.

When automatic import is disabled, only newly created hosted clusters are not automatically imported. Hosted clusters that were already imported are not affected. You can still manually import clusters by using the console or by creating the **ManagedCluster** and **KlusterletAddonConfig** custom resources.

Procedure

To disable the automatic import of hosted clusters, complete the following steps:

1. On the hub cluster, open the **hypershift-addon-deploy-config** specification that is in the **AddonDeploymentConfig** resource in the namespace where multicluster engine Operator is installed by entering the following command:

```
$ oc edit addondeploymentconfig hypershift-addon-deploy-config \
-n multicluster-engine
```

2. In the **spec.customizedVariables** section, add the **autoImportDisabled** variable with value of **"true"**, as shown in the following example:

```
apiVersion: addon.open-cluster-management.io/v1alpha1
kind: AddOnDeploymentConfig
metadata:
  name: hypershift-addon-deploy-config
  namespace: multicluster-engine
spec:
  customizedVariables:
    - name: hcMaxNumber
      value: "80"
    - name: hcThresholdNumber
      value: "60"
    - name: autoImportDisabled
      value: "true"
```

3. To re-enable automatic import, set the value of the **autoImportDisabled** variable to **"false"** or remove the variable from the **AddonDeploymentConfig** resource.